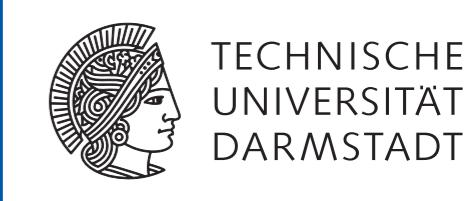


# Computer Vision I

View-Based Modeling - 15.05.2013





# Announcements

- ◆ **No class next week!**
  - ◆ Project review meeting
- ◆ Homework assignment 1:
  - ◆ Is out...
  - ◆ Signup for the "Testat" current read only -- will go online again soon
- ◆ **Exam**
  - ◆ 21. August 2013, 10:00 - 12:00



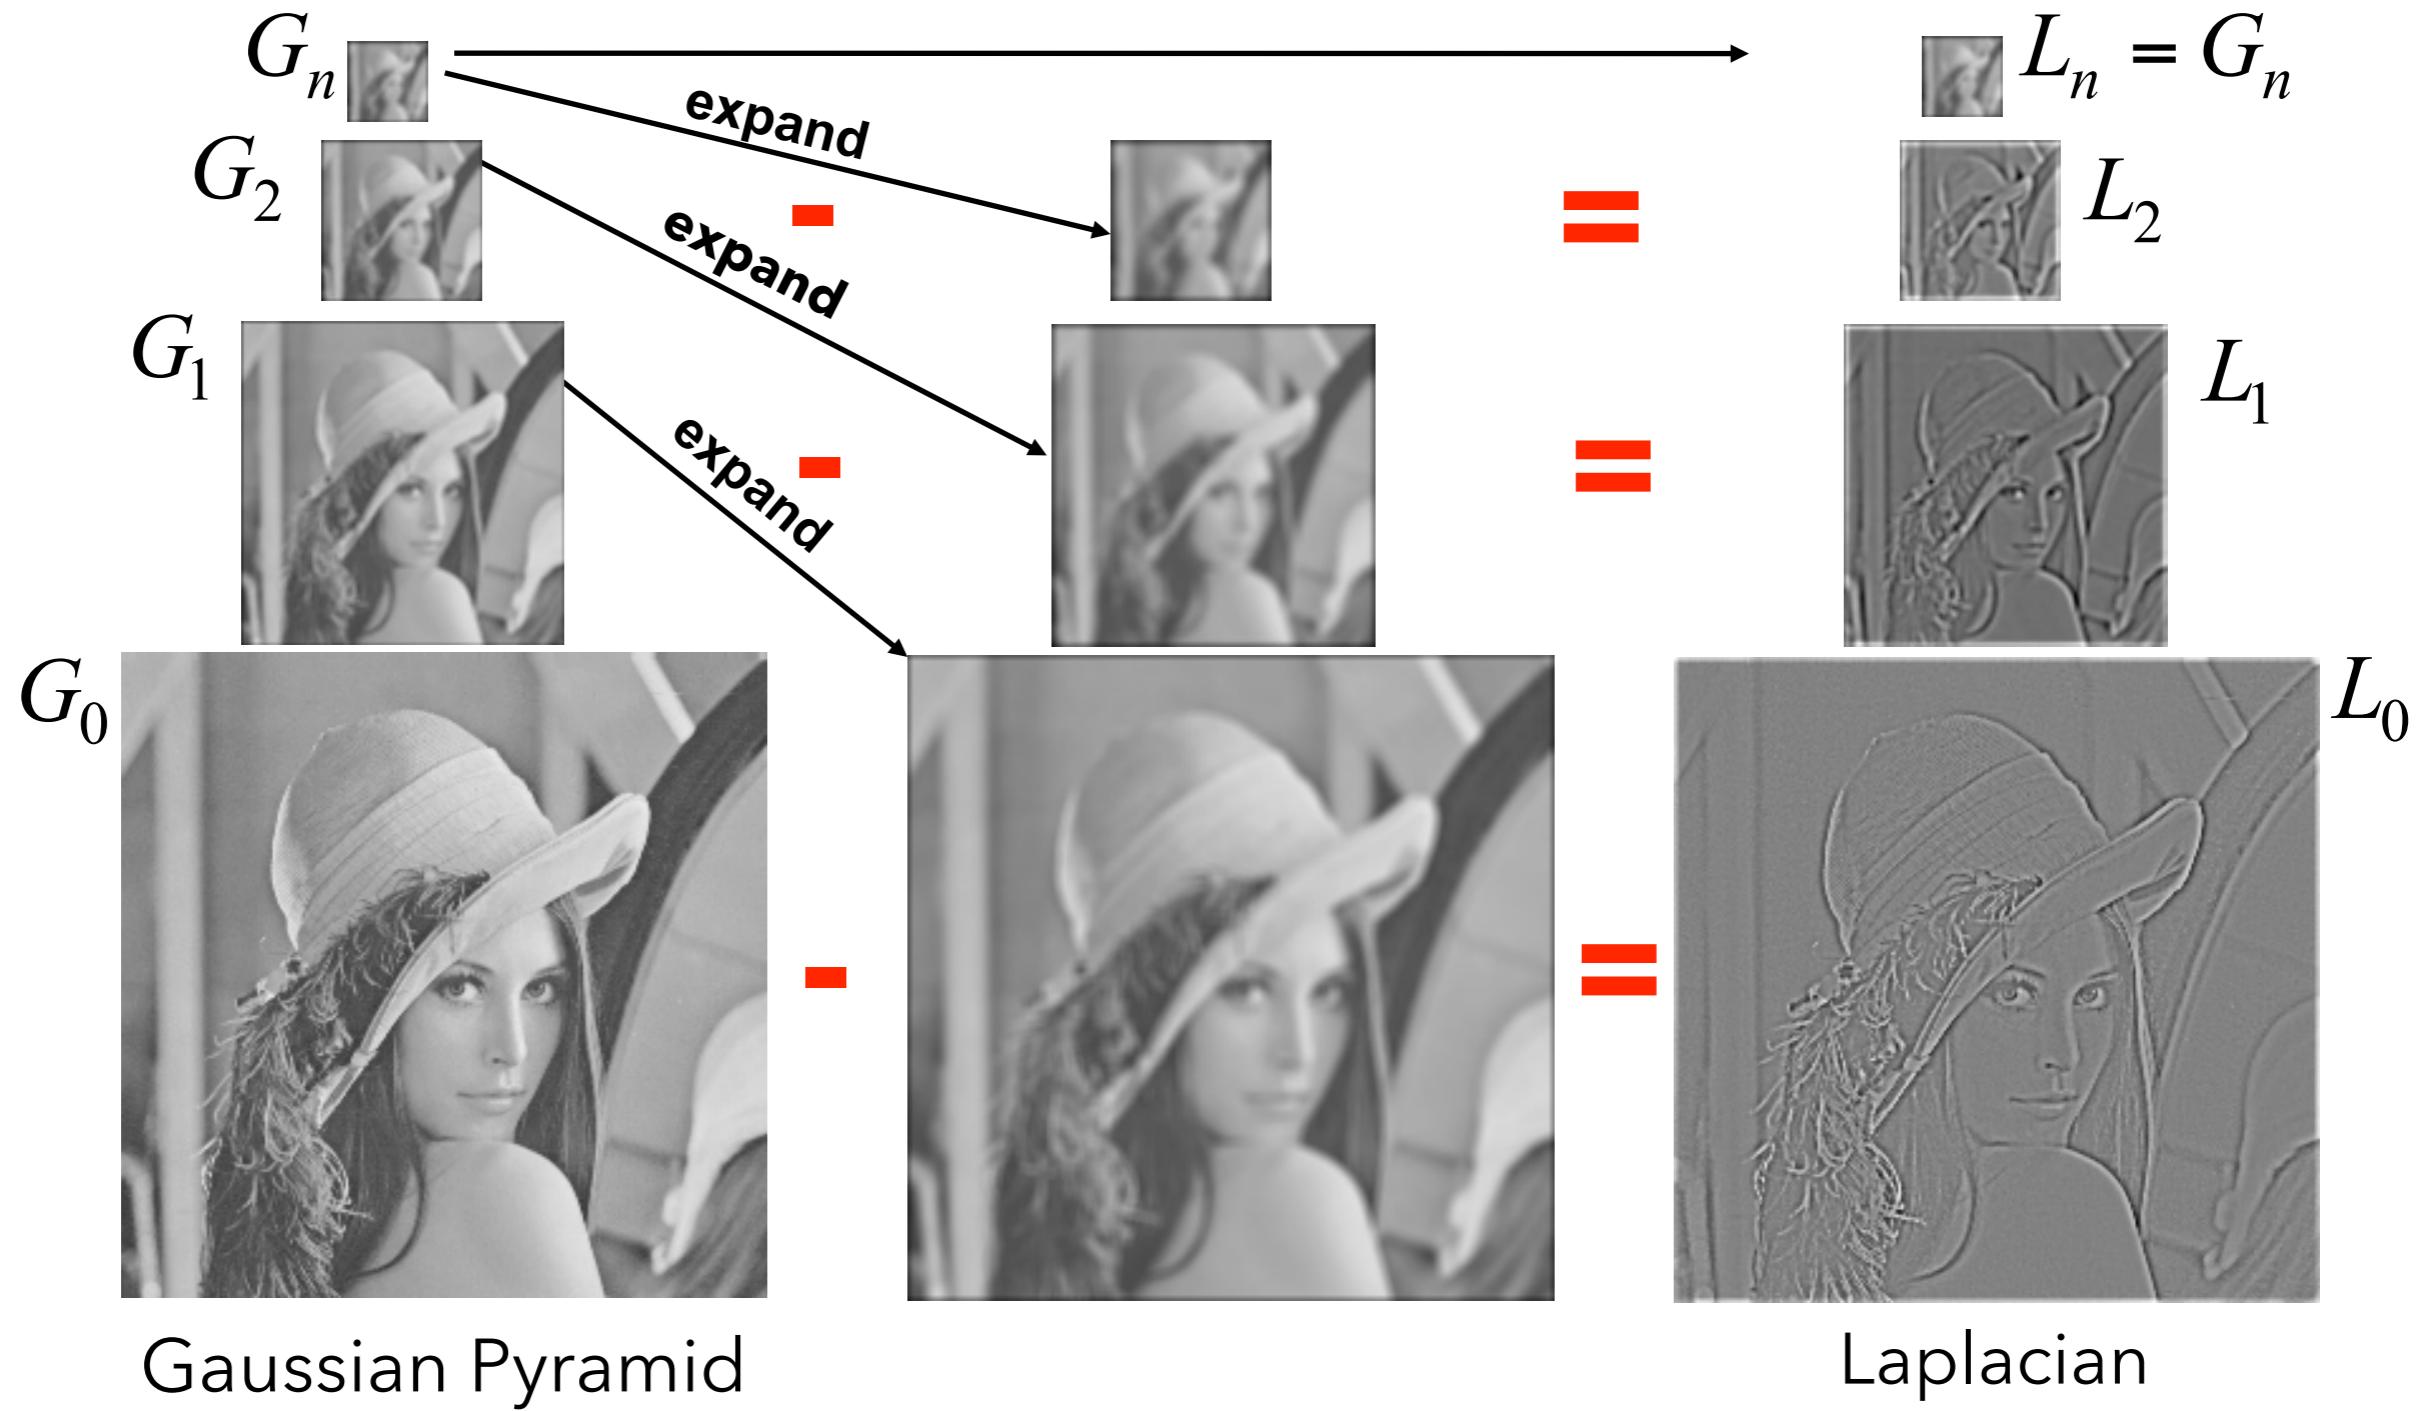
# Wrap up - Basics of Digital Image Processing

- ◆ Images
- ◆ Filtering
  - ◆ Linear filtering
  - ◆ Non-linear filtering & morphology
- ◆ Multi-scale image representation
  - ◆ Gaussian pyramid
  - ◆ Laplacian pyramid
- ◆ Edge detection
  - ◆ 'Recognition using line drawings'
  - ◆ Image derivatives (1st and 2nd order)

# Laplacian pyramid

$$L_i = G_i - \text{expand}(G_{i+1})$$

$$G_i = L_i + \text{expand}(G_{i+1})$$

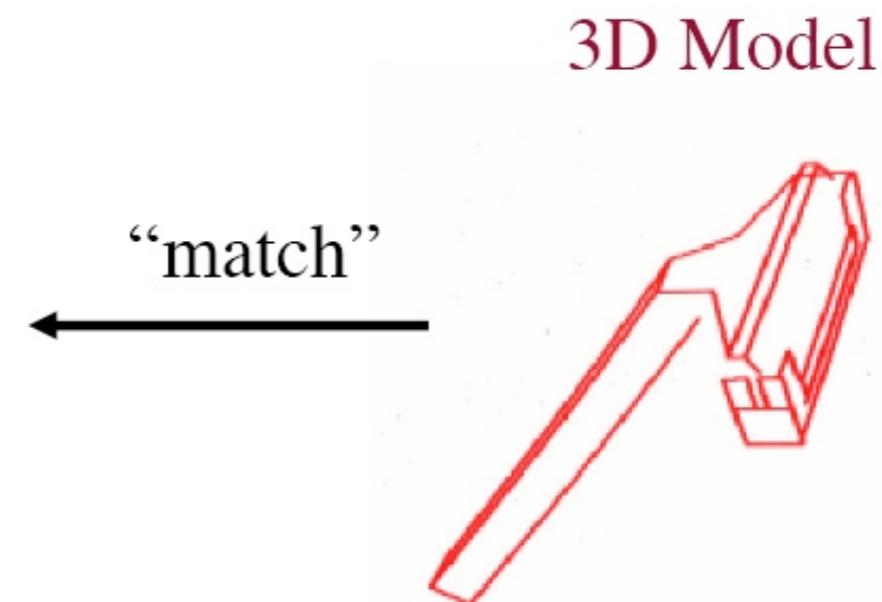
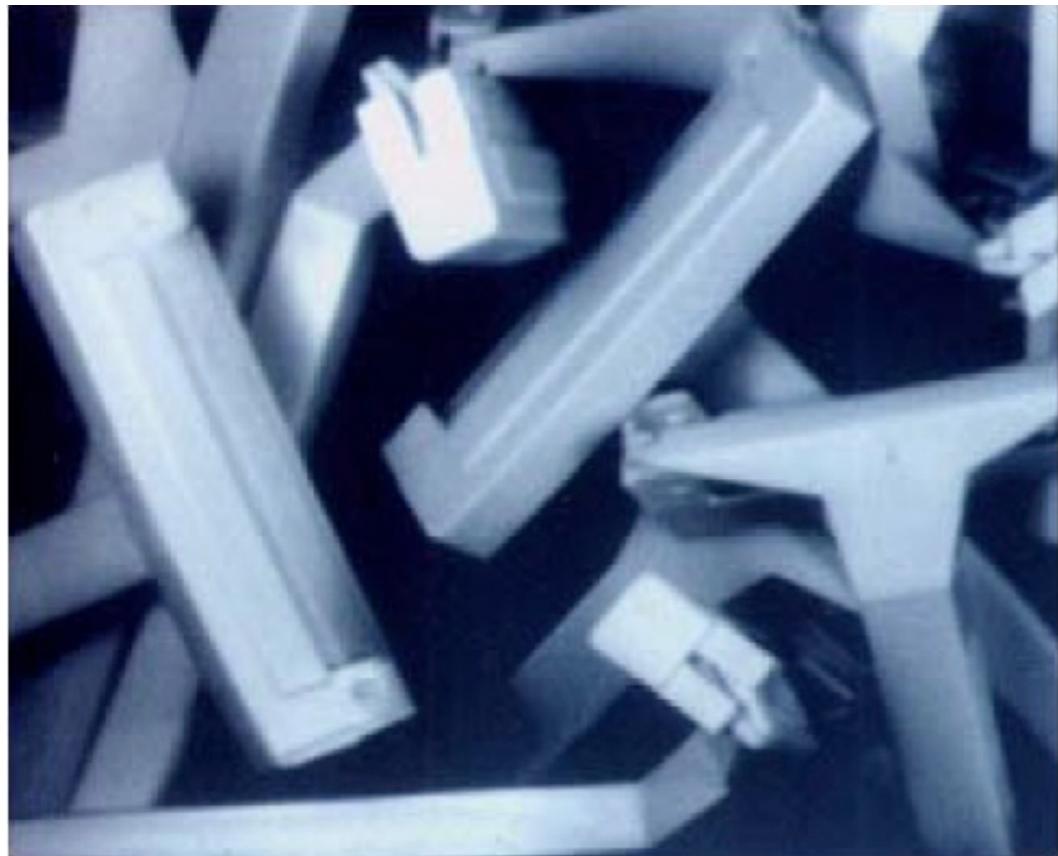


# Laplacian pyramid

- ◆ The Laplacian pyramid provides a simple frequency decomposition into subbands.
- ◆ Each level / subband contains only image structure of a particular range of spatial frequencies:
  - ◆ E.g., the finest level contains all the high-frequency detail.
- ◆ Observation: We can get back the original image by reversing the decomposition.

# Example of Recognition & Localization

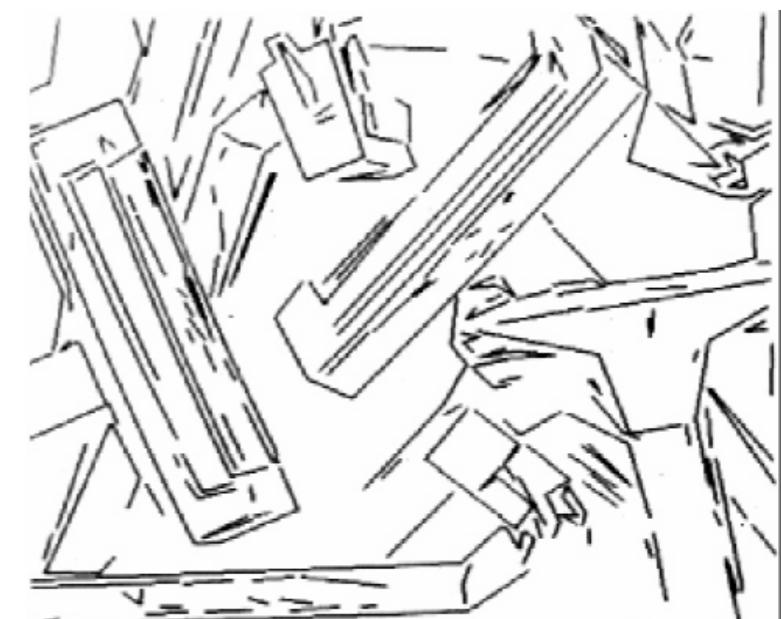
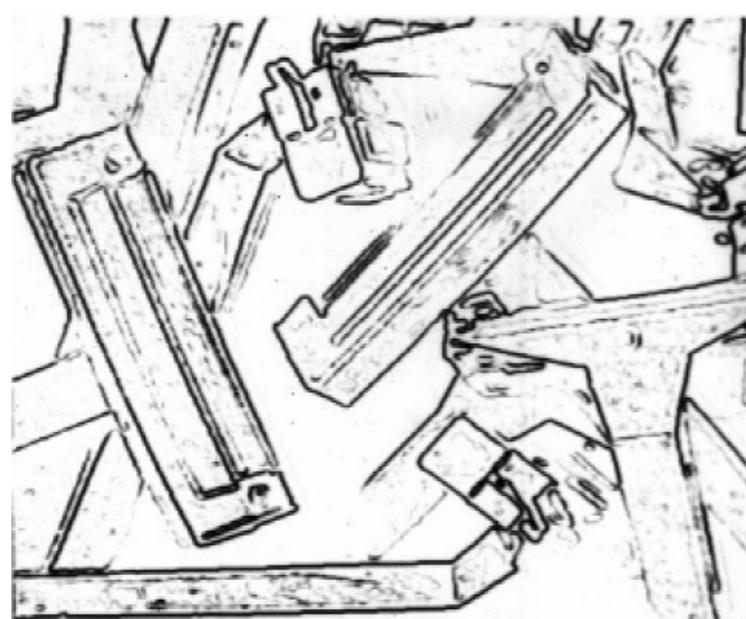
- ◆ Simple approach [David Lowe]



Parameters: 3D position  
and orientation

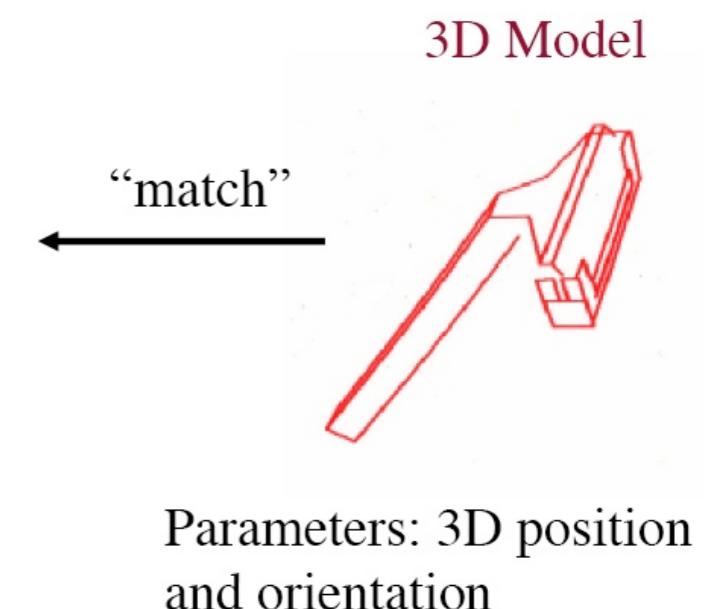
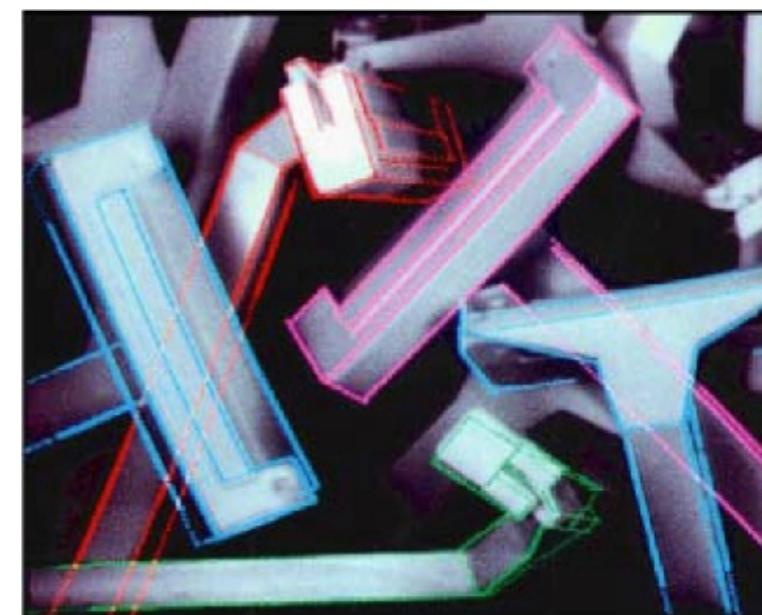
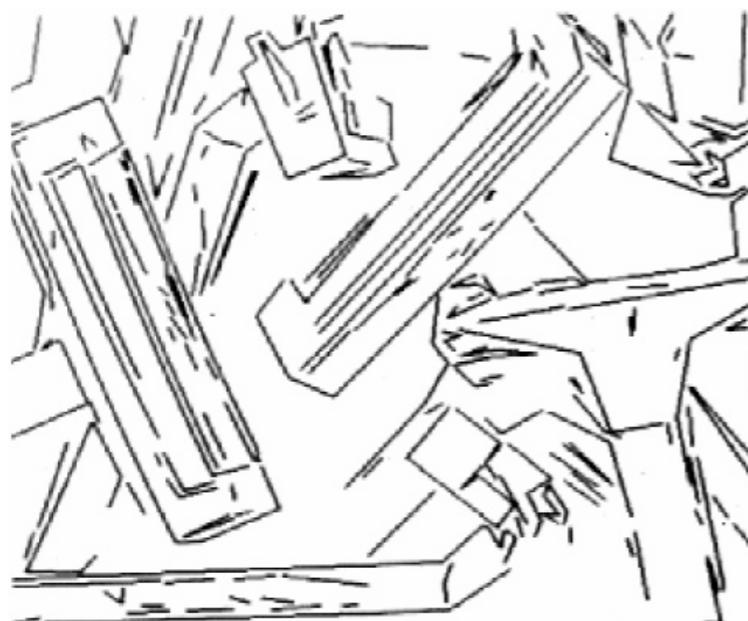
# Example of Recognition & Localization

- ◆ Simple approach [David Lowe]
  - ◆ 1. "Filter" image to **find brightness changes**
  - ◆ 2. **"Fit"** **lines** to the raw measurements



# Example of Recognition & Localization

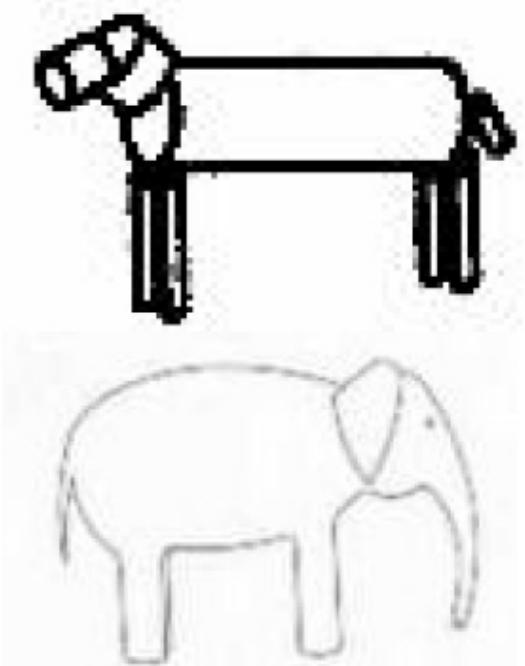
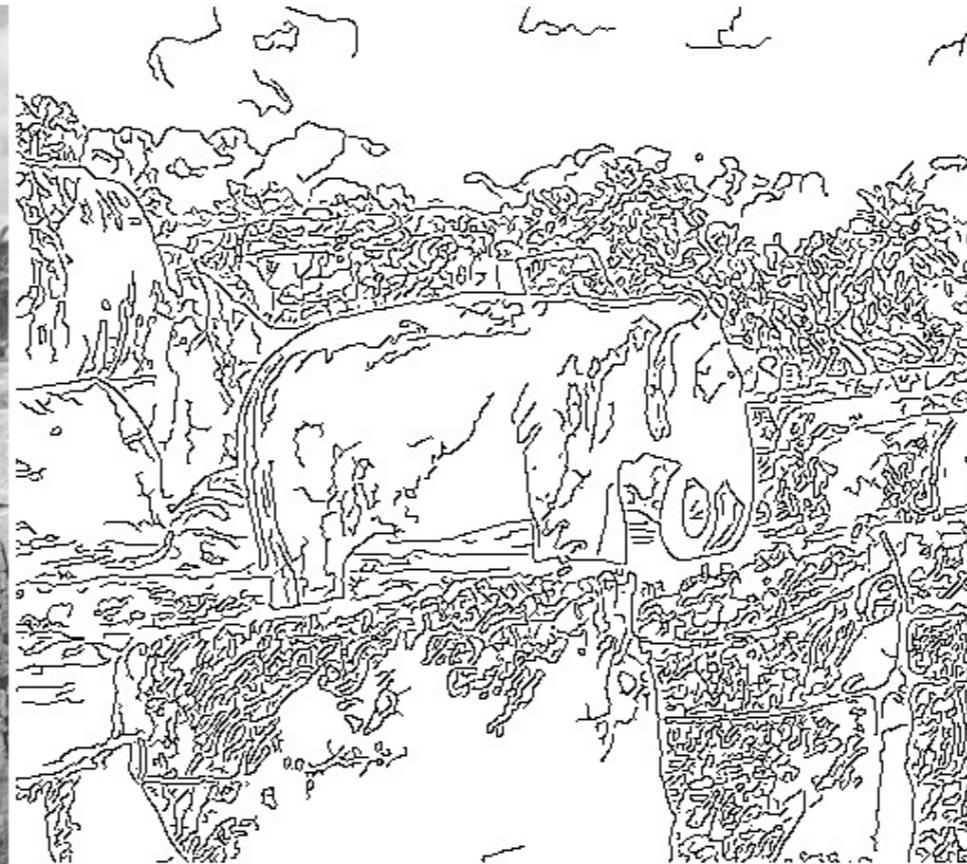
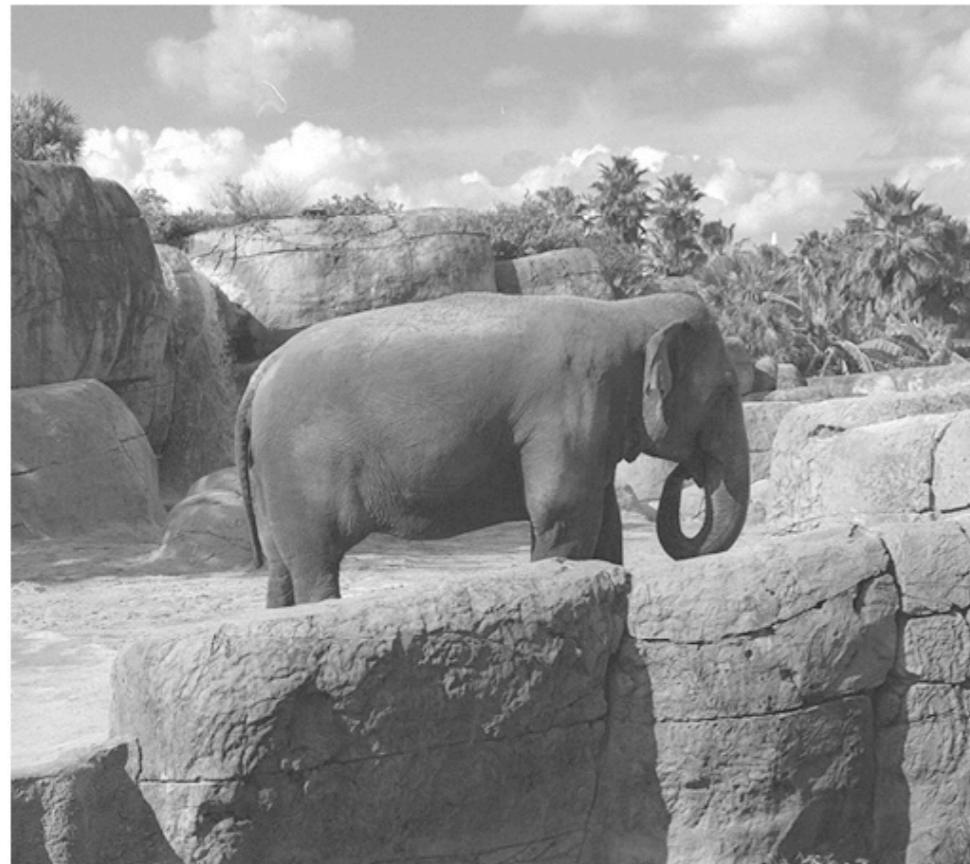
- ◆ Simple approach [David Lowe]
  - ◆ 3. “Project” model into the image and **“match” to lines** (solving for 3D pose)



# Line drawing vs. edge detection



# Models vs. edge detection



Match “model” to measurements?

University of South Florida

# Search and Recognition



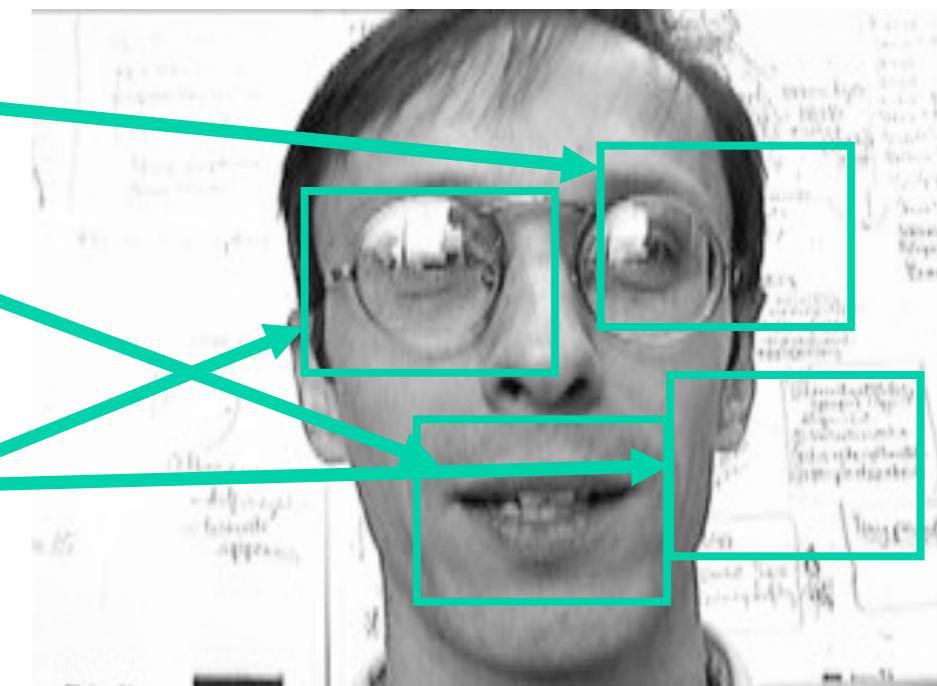
- ◆ What if we do not have a line-based object model?
  - ◆ Edges may not help us.



- ◆ How can we find the mouth?
- ◆ How can we recognize the “expression”?

# Naïve View-Based Approach

Database of mouth “templates”

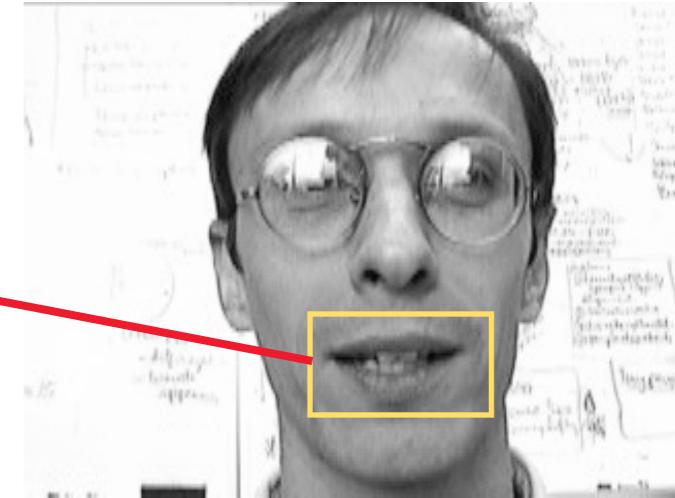


- Search every image region (at every scale).
- Compare each template; chose the best match.

# View-Based Methods

- ◆ Idea: Represent objects by their appearance in an ensemble of images, including different poses, illuminations, configurations of shape, ...
- ◆ Approach covered here:
  - ◆ Subspace methods (also called "Eigen methods")

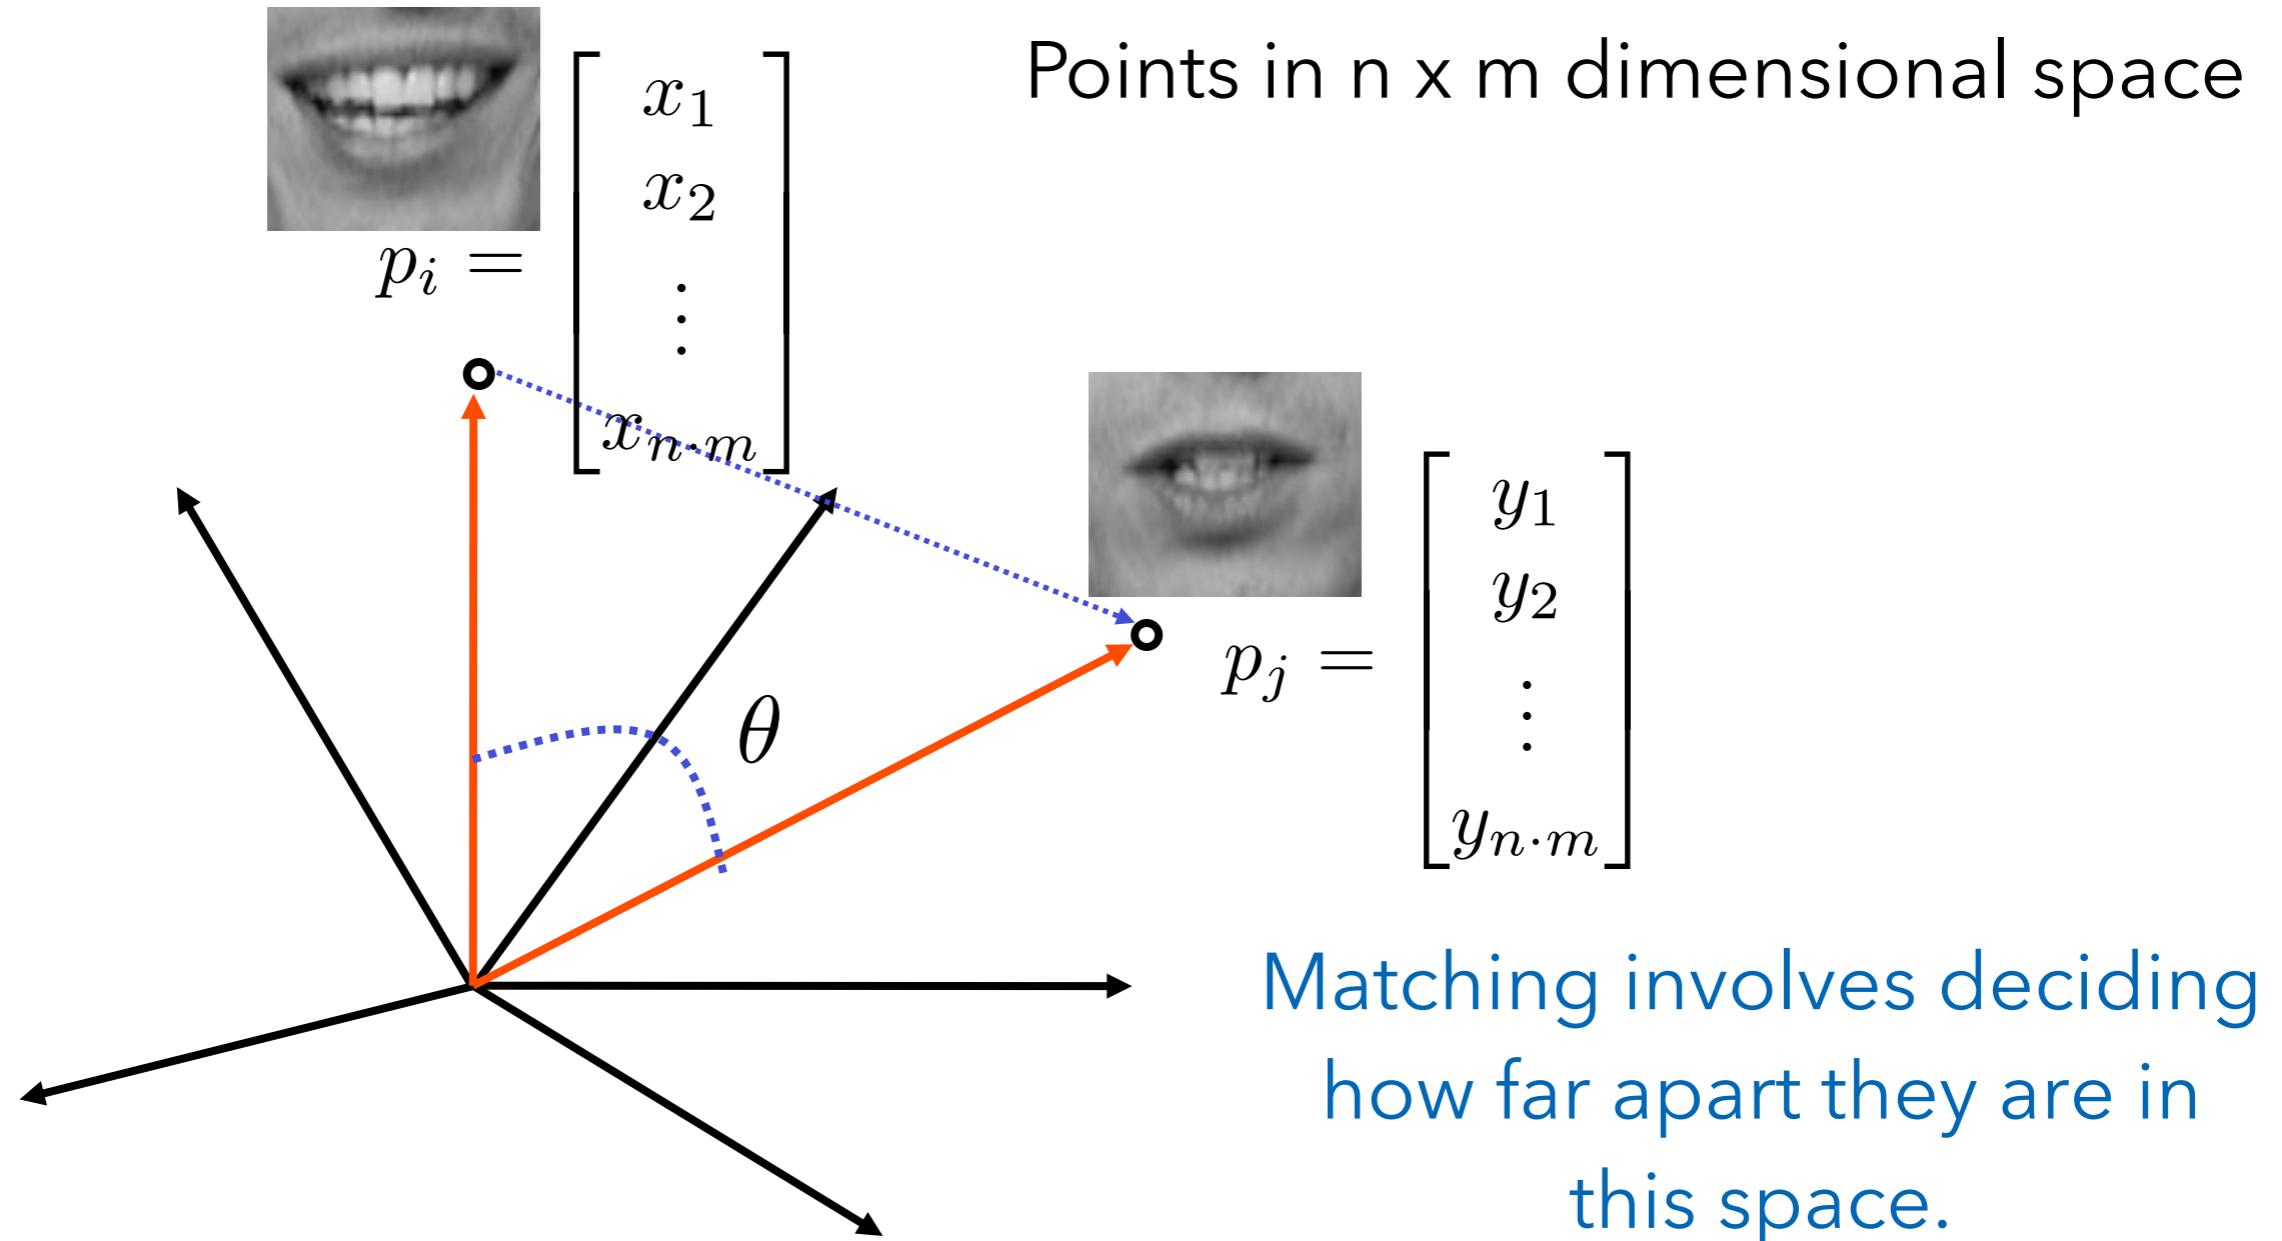
# Images as Vectors



$$= \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_{n \cdot m} \end{bmatrix} \quad = \quad \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_{n \cdot m} \end{bmatrix}$$

e.g. standard lexicographic ordering

# Images as Points



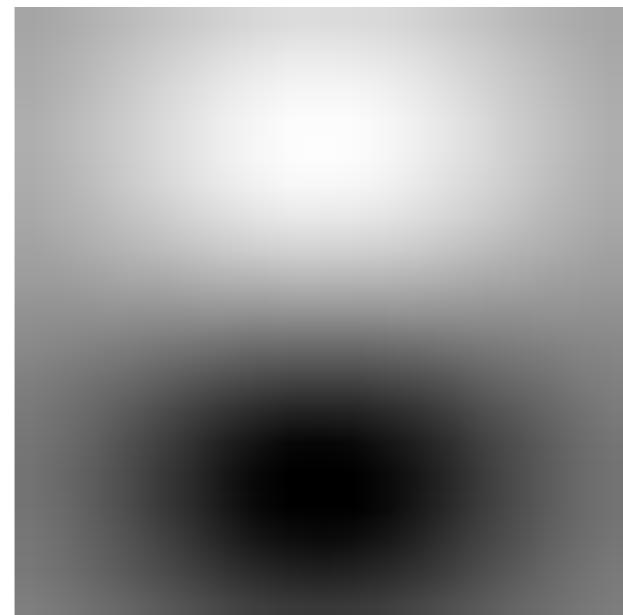
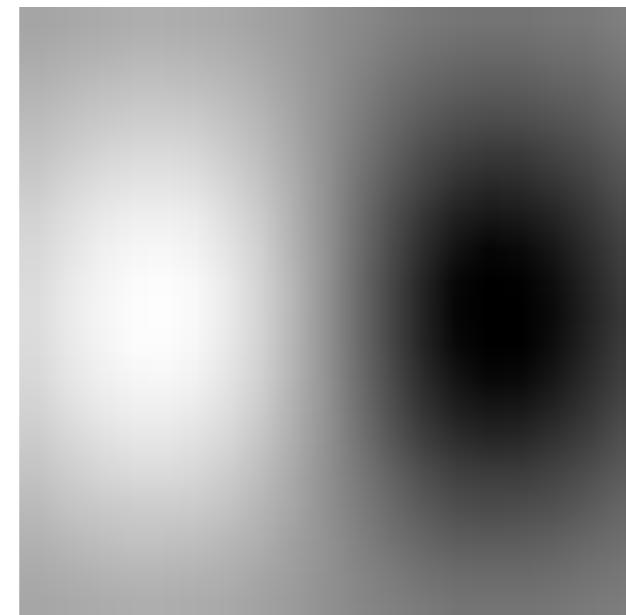
# Filters are Templates

- ◆ Remember image filtering from last lecture.
- ◆ Applying a filter at some pixel can be seen as taking a **dot-product** between the image and some vector.
- ◆ Filtering the image is a set of dot products:

$$f[m, n] = (I * g)[m, n] = \sum_{k,l} I[m - k, n - l]g[k, l]$$

# Filters as Templates

- ◆ Insight:
  - ◆ Filters look like the effects they are intended to find.
  - ◆ Filters find effects they look like.
  - ◆ We can use filters to match a template against an image.



# Template Matching

Correlation  
(sum of  
product of  
"signals")

$$\mathbf{a}^T \mathbf{b} = |\mathbf{a}| |\mathbf{b}| \cos \theta$$

Angle  
between the  
vectors

Template or filter as a  
vector

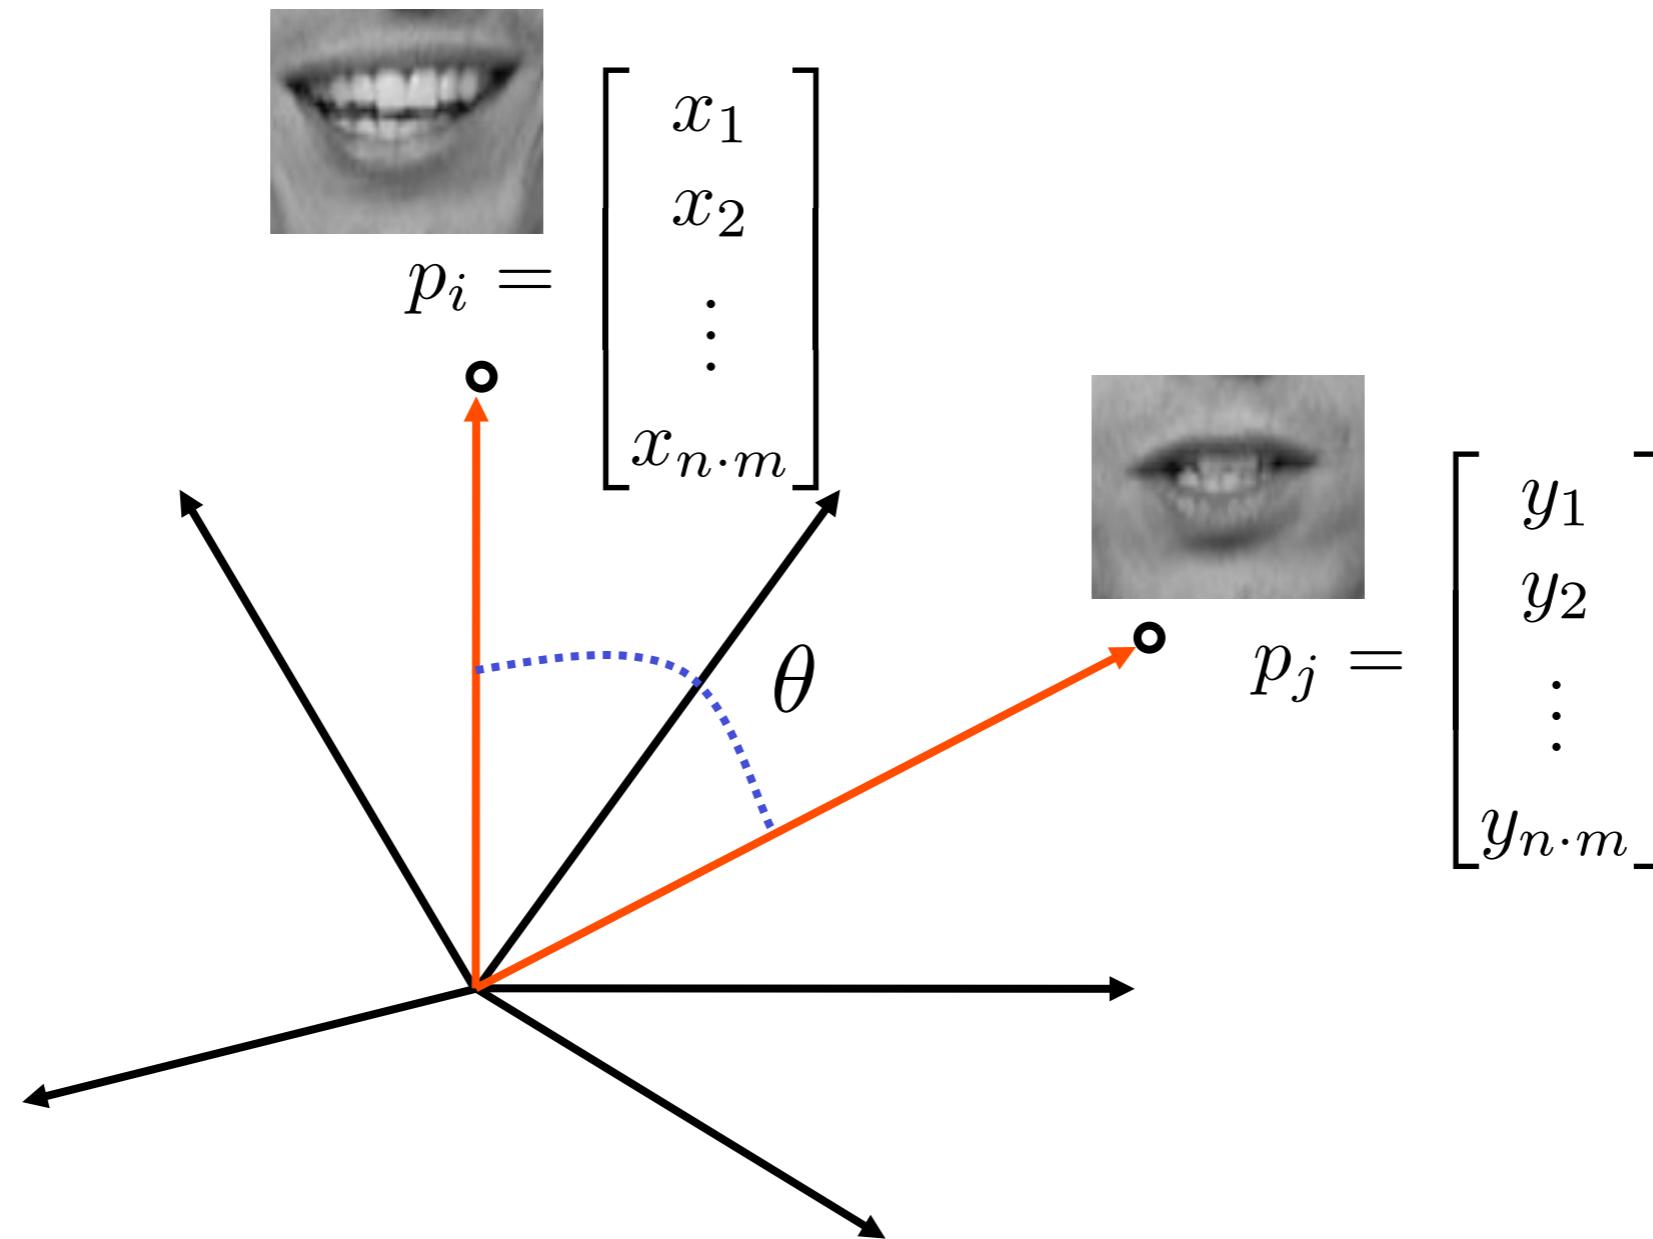
Image patch as a  
vector

Normalized correlation:

$$\frac{\mathbf{a}^T \mathbf{b}}{|\mathbf{a}| |\mathbf{b}|} = \cos \theta$$

# Dot Product

- ◆ When we filter, we measure the angle (cosine of it, really) between the filter template and the image patch.

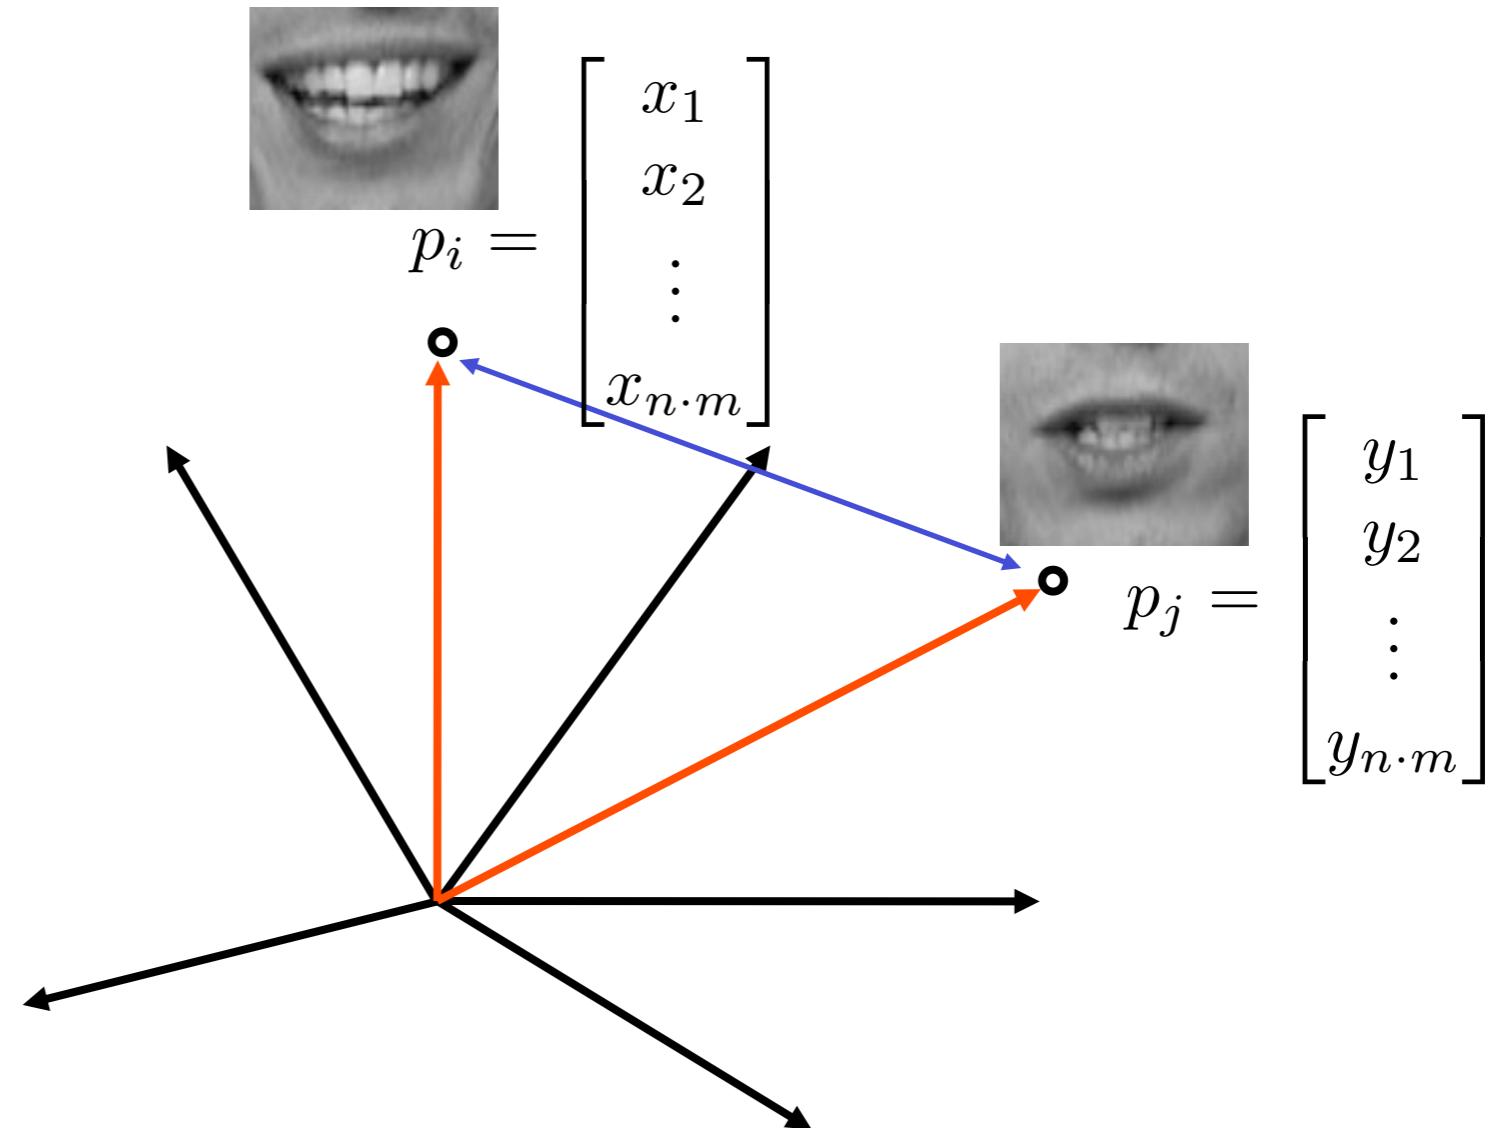


# SSD Matching

- ◆ An alternative is to minimize the **Sum of Squared Differences (SSD)**:

$$E(I, T) = \sum_{i,j} (I(i, j) - T(i, j))^2$$

- ◆ Distance metric



# Template Methods

- ◆ **Image templates** (simplest view-based method - straw man):
  - ◆ Keep an image of every object from different viewing directions, lighting conditions, etc.
  - ◆ Nearest neighbor cross-correlation matching with images in model database (or robust matching for clutter & occlusion).
- ◆ **Obvious problems:**
  - ◆ Storage and computation costs become unreasonable as the number of objects increases.
  - ◆ May require very large ensemble of 'training' images.

Fleet & Szeliski

# Subspace Methods

- ◆ How can we find **more efficient representations** for the ensemble of views, and more efficient methods for matching?
- ◆ Idea: images are not random... especially images of the same object (category) have similar appearance

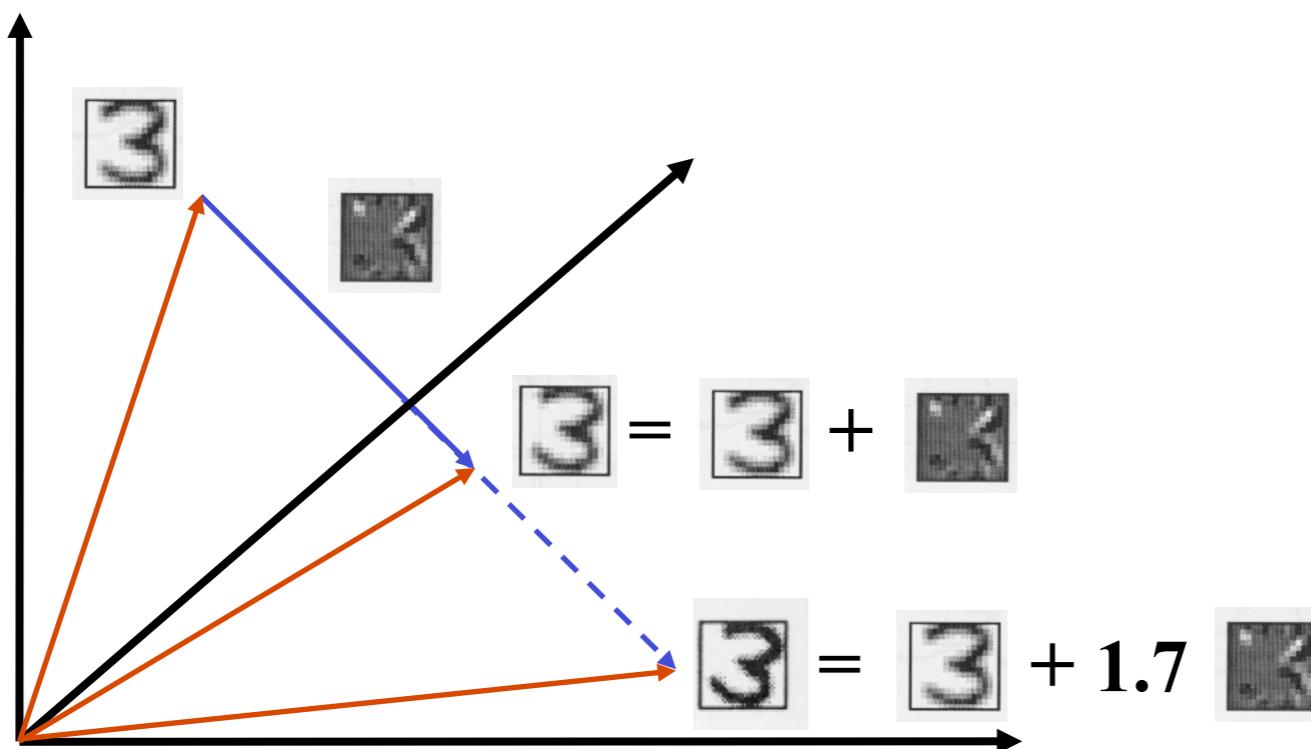


E.g., let images be represented as points in a high-dimensional space (e.g., one dimension per pixel)

Fleet & Szeliski

# Linear Dimensionality Reduction

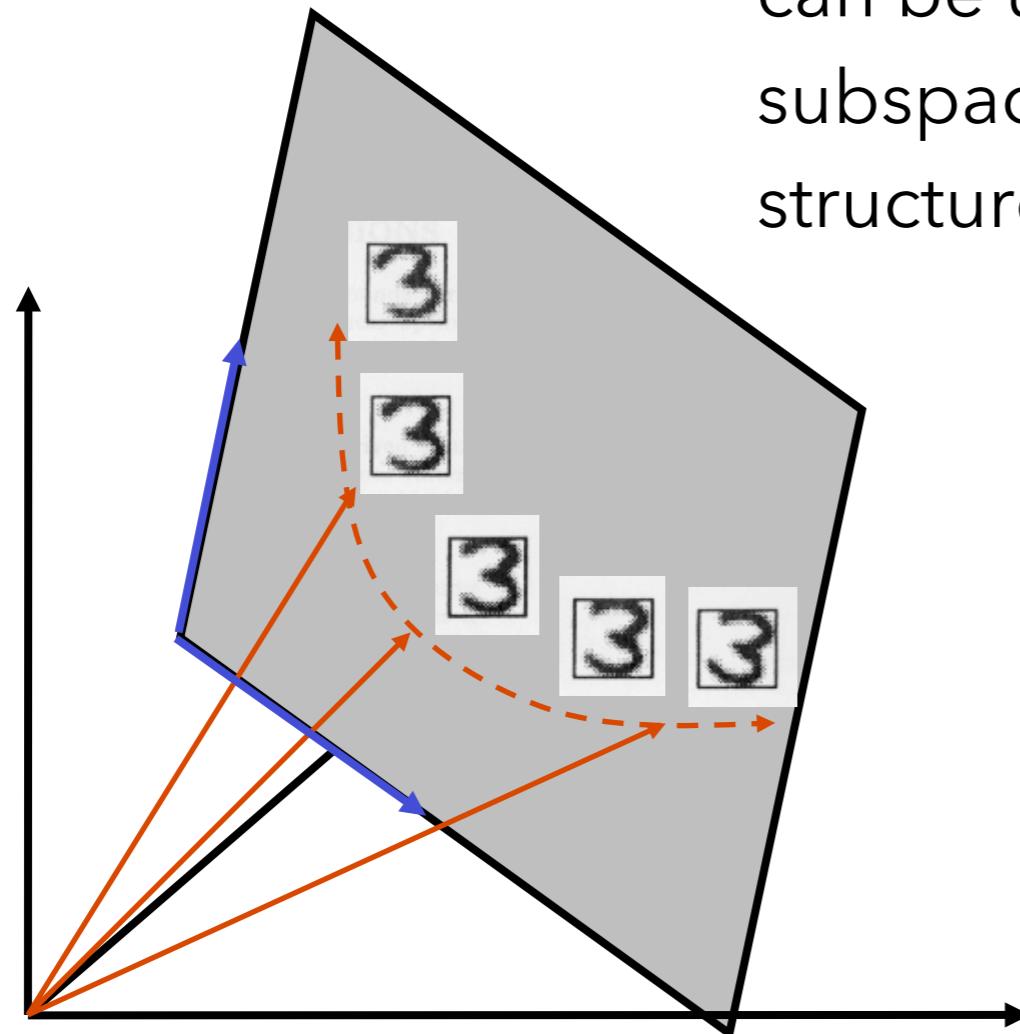
- Given that differences are structured, we can use '**basis images**' to transform images into other images in the same space.



Fleet & Szeliski

# Linear Dimensionality Reduction

What linear transformations of the images can be used to define a lower-dimensional subspace that captures most of the structure in the image ensemble?



Fleet & Szeliski

# Approach

- ◆ Find a lower dimensional representation that captures the variability in the data.
- ◆ Search using this low dimensional model.

Data point  $n$ :

$$\mathbf{x}^n \in \mathbb{R}^M$$

Low-dimensional representation:

$$\mathbf{a}^n \in \mathbb{R}^D \quad D \ll M$$

Mapping:  $\mathbf{x}^n \rightarrow \mathbf{a}^n$

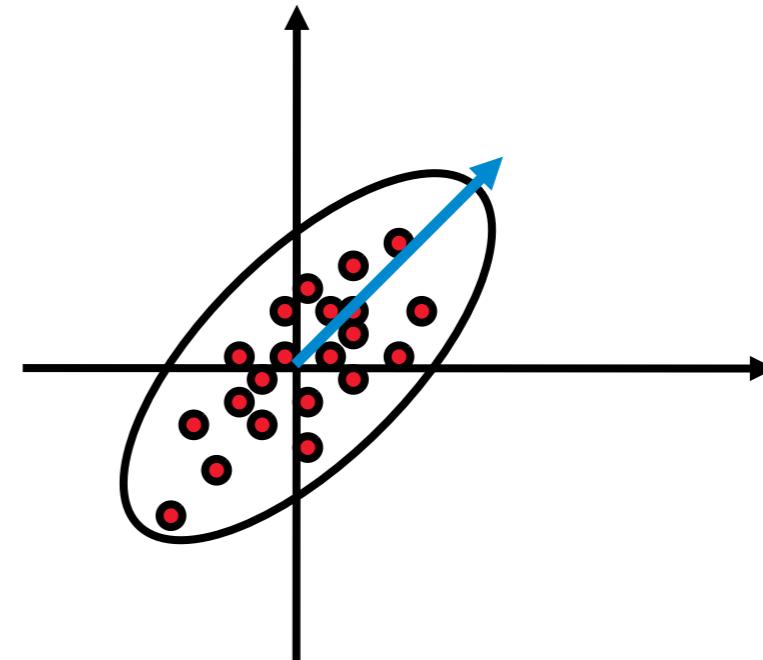
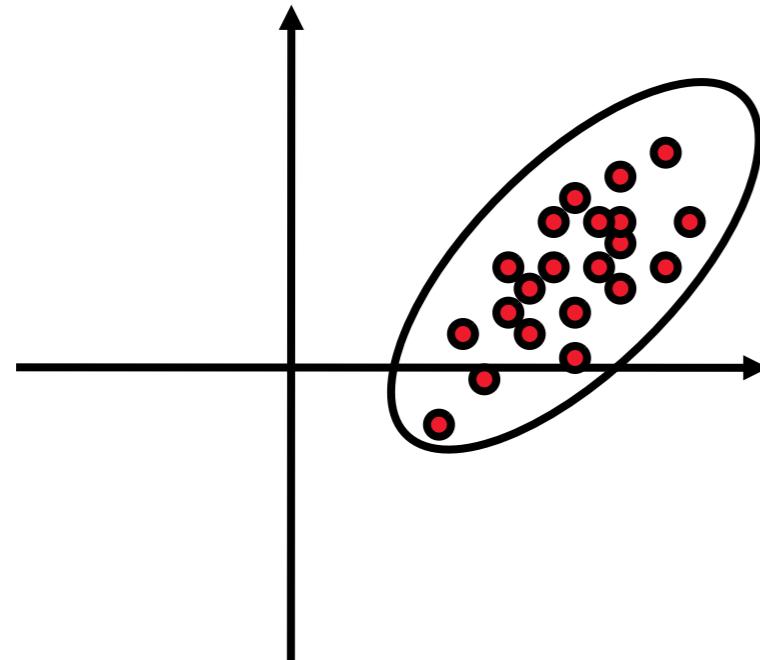
◆ Restrict this mapping to be a linear function:

$$\mathbf{a}^n = \mathbf{B}\mathbf{x}^n, \quad \mathbf{B} \in \mathbb{R}^{D \times M}$$

◆ Note:

- ◆ From now on, I will use boldface symbols to denote vectors, e.g.:  
 $\mathbf{x}, \mathbf{y}, \dots$

# Intuition



If I give you the mean and one vector to represent the data, what vector would you choose?

Why?

# Recall: Basis Representations

We can always write a vector as:

$$\mathbf{x} = \sum_{i=1}^M a_i \mathbf{u}_i$$

where

$$\mathbf{u}_i^T \mathbf{u}_j = \delta_{ij}$$

Kronecker delta = 1 if  $i = j$ ,  
0 otherwise.

orthonormal

Example:

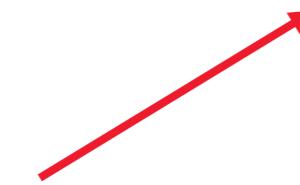
$$\begin{bmatrix} 3 \\ 7 \end{bmatrix} = 3 \begin{bmatrix} 1 \\ 0 \end{bmatrix} + 7 \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

# Recall: Basis Representations

$$\begin{aligned}
 \mathbf{x} &= a_1 \mathbf{u}_1 + a_2 \mathbf{u}_2 \\
 a_1 &= \mathbf{u}_1^T (\mathbf{x} - a_2 \mathbf{u}_2) \\
 &= \mathbf{u}_1^T \mathbf{x} - a_2 \mathbf{u}_1^T \mathbf{u}_2 \\
 &= \mathbf{u}_1^T \mathbf{x}
 \end{aligned}$$

More generally:

$$a_i = \mathbf{u}_i^T \mathbf{x}$$

  
 Scalar coefficient      projection

# Decomposition

$$\mathbf{x}^n = \sum_{i=1}^D a_i \mathbf{u}_i + \sum_{j=D+1}^M b_j \mathbf{u}_j$$

Approximation  $\tilde{\mathbf{x}}^n$       Error

Want the  $D$  bases that minimize the mean squared error over the training data

$$\arg \min E(\mathbf{u}_1, \dots, \mathbf{u}_D) = \sum_{n=1}^N \|\mathbf{x}^n - \tilde{\mathbf{x}}^n\|^2$$

# Minimizing the Error

- ◆ Rewrite the error (assuming a single basis vector):

$$\begin{aligned}
 E(\mathbf{u}) &= \sum_{n=1}^N \|\mathbf{x}^n - \tilde{\mathbf{x}}^n\|^2 \\
 &= \sum_{n=1}^N \|\mathbf{x}^n - (\mathbf{u}^T \mathbf{x}^n) \mathbf{u}\|^2 \\
 &= \sum_{n=1}^N \|\mathbf{x}^n\|^2 - 2(\mathbf{u}^T \mathbf{x}^n)^2 + (\mathbf{u}^T \mathbf{x}^n)^2 \cdot \mathbf{u}^T \mathbf{u} \\
 &= \sum_{n=1}^N \|\mathbf{x}^n\|^2 - (\mathbf{u}^T \mathbf{x}^n)^2 = \sum_{n=1}^N \|\mathbf{x}^n\|^2 - (a_n)^2
 \end{aligned}$$

# Minimizing the Error

- ◆ Rewrite the error (assuming a single basis vector):

$$E(\mathbf{u}) = \sum_{n=1}^N \|\mathbf{x}^n\|^2 - (\mathbf{u}^T \mathbf{x}^n)^2 = \sum_{n=1}^N \|\mathbf{x}^n\|^2 - (a_n)^2$$

- ◆ Minimizing the error is equivalent to maximizing the variance of the projection:

$$\frac{1}{N} \sum_{n=1}^N a_n^2 \quad \text{(assuming mean 0)}$$

- ◆ We can ensure a zero mean projection by subtracting the mean from every data point:

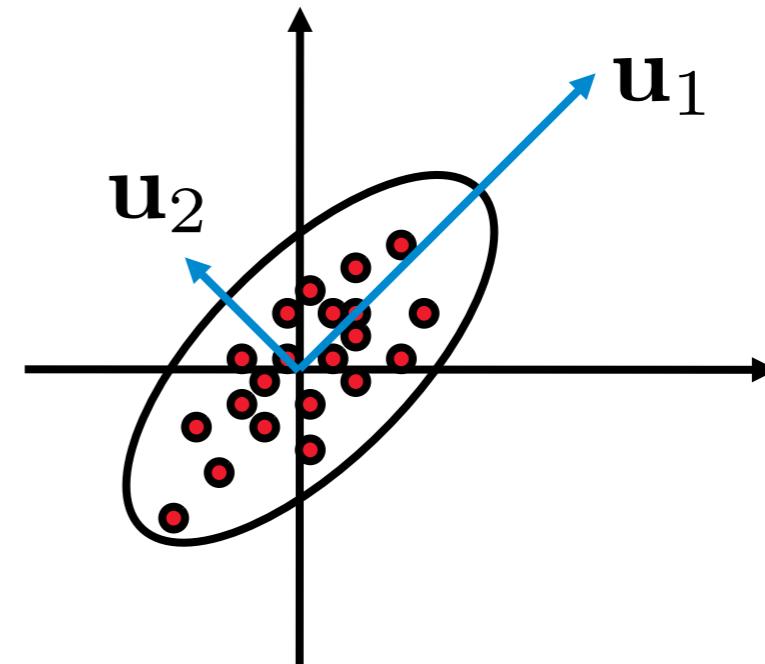
$$\mathbf{x}^n - \bar{\mathbf{x}}$$

# Intuition



$$\mathbf{x}^n - \bar{\mathbf{x}} = \sum_{i=1}^D a_i \mathbf{u}_i + \sum_{j=D+1}^M b_j \mathbf{u}_j$$

$$\tilde{\mathbf{x}}^n = \sum_{i=1}^D a_i \mathbf{u}_i + \bar{\mathbf{x}}$$



Projecting onto  $\mathbf{u}_1$  captures the majority of the variance and hence projecting onto it minimizes the error

How do we find the axis of largest variance?

# Statistics Review: Univariate

Sample mean:

(univariate case)

$$\bar{x} = \langle x \rangle = E[x] = \frac{1}{N} \sum_{n=1}^N x^n$$

Sample variance:

$$\begin{aligned}\hat{\sigma}^2 &= \langle (x - \bar{x})^2 \rangle = \text{Var}[x] = \frac{1}{N} \sum_{n=1}^N (x^n - \bar{x})^2 \\ &= \langle x^2 \rangle - \langle 2\bar{x}x \rangle + \langle \bar{x}^2 \rangle = \langle x^2 \rangle - \bar{x}^2\end{aligned}$$

# Statistics Review: Multivariate



Sample mean:

(multivariate case)

$$\bar{\mathbf{x}} = \langle \mathbf{x} \rangle = E[\mathbf{x}] = \frac{1}{N} \sum_{n=1}^N \mathbf{x}^n$$

Sample covariance:

$$\begin{aligned}\text{cov}(x_i, x_j) = \sigma_{ij} &= \langle (x_i - \bar{x}_i)(x_j - \bar{x}_j) \rangle \\ &= \langle x_i x_j \rangle - \langle x_i \bar{x}_j \rangle - \langle \bar{x}_i x_j \rangle + \langle \bar{x}_i \bar{x}_j \rangle \\ &= \langle x_i x_j \rangle - \bar{x}_i \bar{x}_j \\ \text{cov}(x_i, x_i) = \sigma_i^2 &= \langle x_i^2 \rangle - \bar{x}_i^2\end{aligned}$$

# Statistical Correlation

- ◆ The **covariance** of two variates  $X_i$  and  $X_j$  of a random vector  $\mathbf{X}$  provides a measure of how strongly correlated these variables are, and the derived quantity is:

$$\text{cor}(x_i, x_j) = \frac{\text{cov}(x_i, x_j)}{\sigma_i \sigma_j}$$

# Covariance Matrix

- ◆ The **covariance matrix** describes the covariance between all components of a random vector, e.g.:

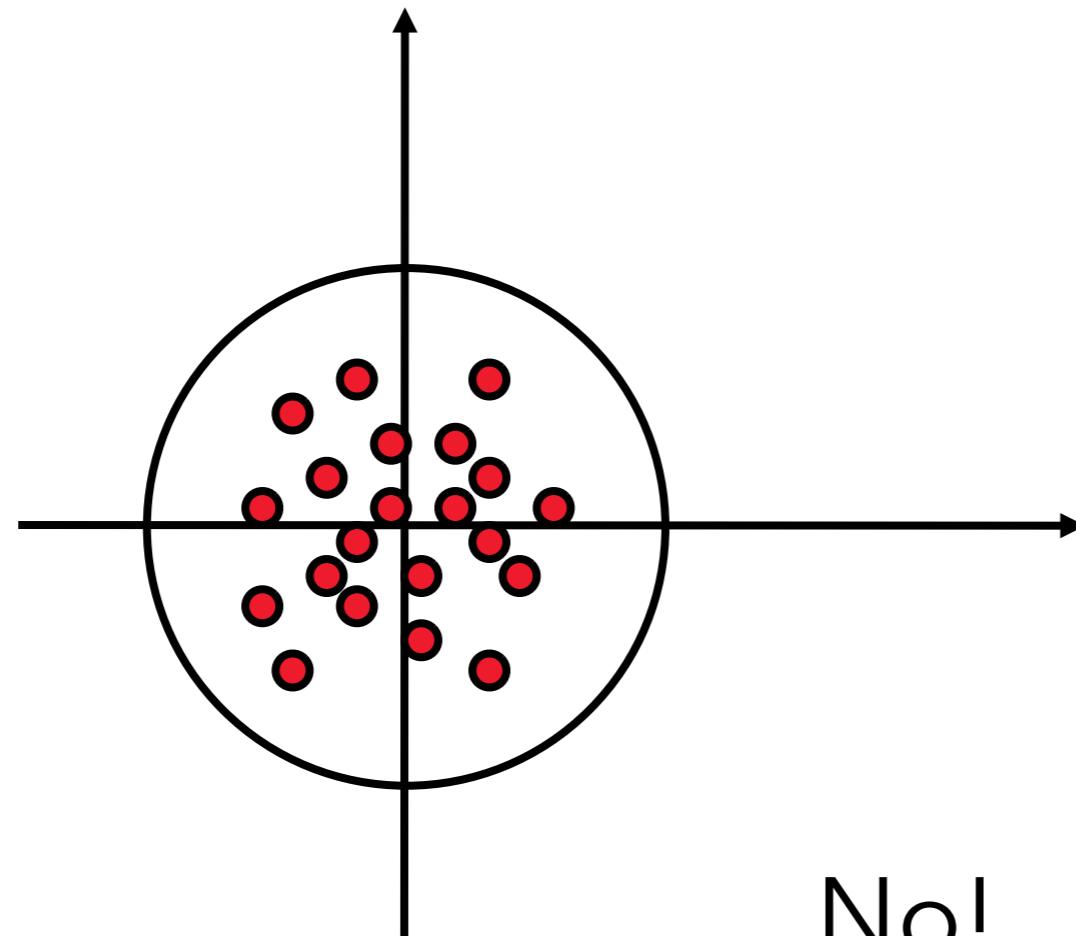
$$\mathbf{C} = \begin{bmatrix} \sigma_{11} & \sigma_{12} \\ \sigma_{21} & \sigma_{22} \end{bmatrix}$$

$$\hat{\mathbf{C}} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}^n - \bar{\mathbf{x}})(\mathbf{x}^n - \bar{\mathbf{x}})^T$$

# Recall: Outer Product

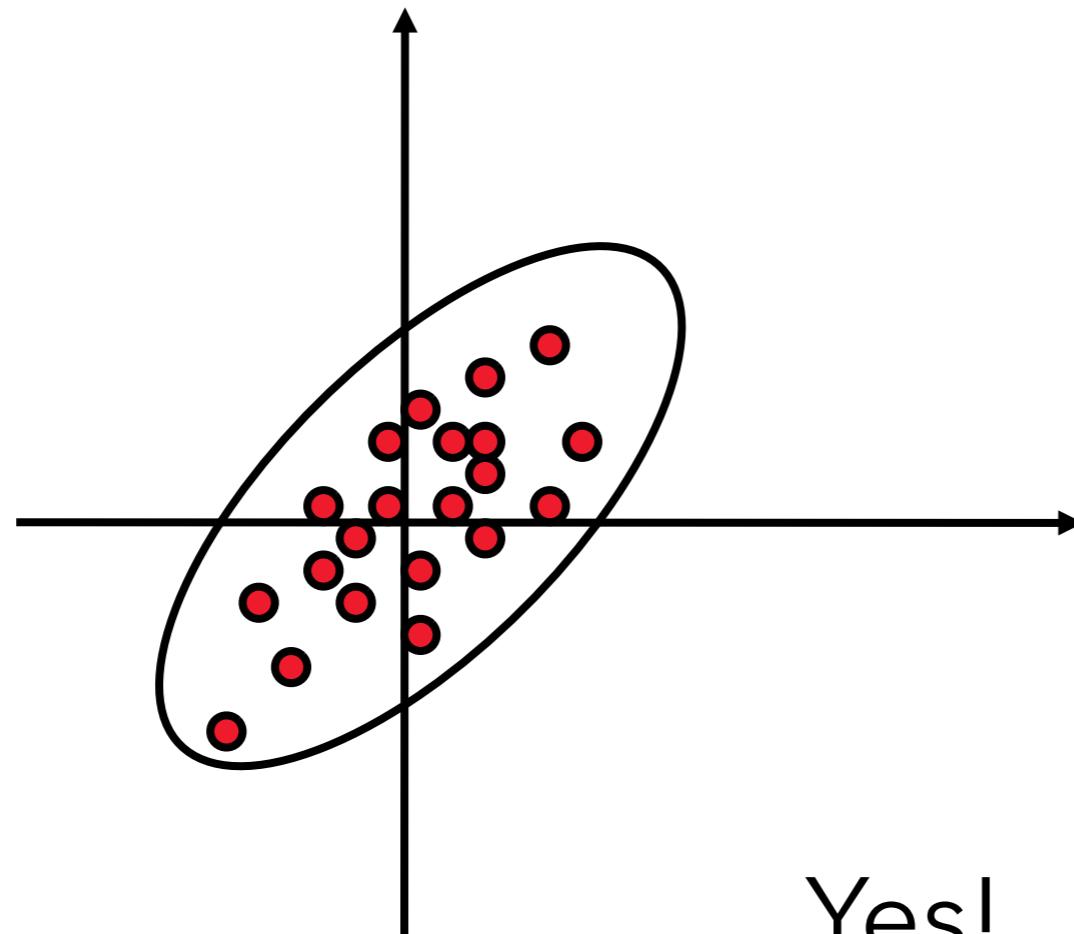
$$\mathbf{x}\mathbf{x}^T = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \begin{bmatrix} x_1 & x_2 \end{bmatrix} = \begin{bmatrix} x_1x_1 & x_1x_2 \\ x_2x_1 & x_2x_2 \end{bmatrix}$$

# Correlated?



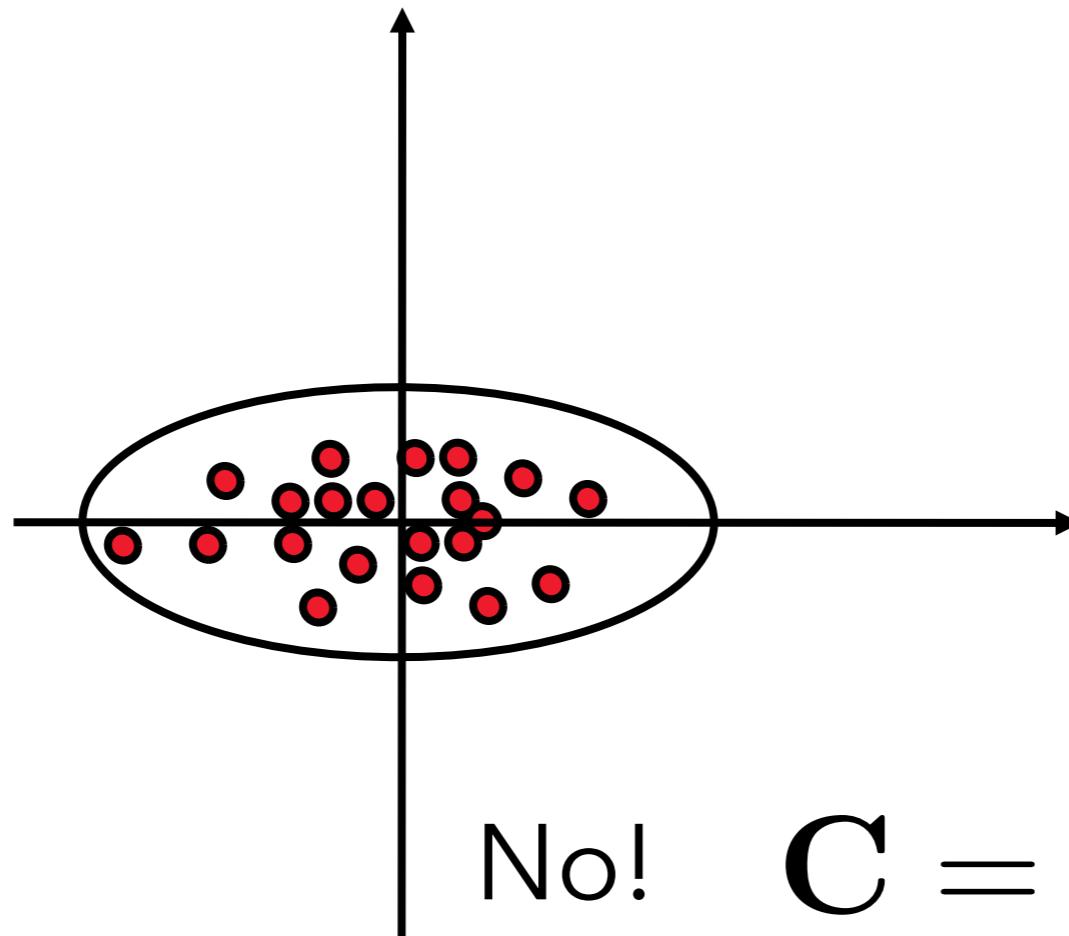
No!  $C = \begin{bmatrix} c & 0 \\ 0 & c \end{bmatrix}$

# Correlated?



Yes!  $C = \begin{bmatrix} a & b \\ b & c \end{bmatrix}$

# Correlated?

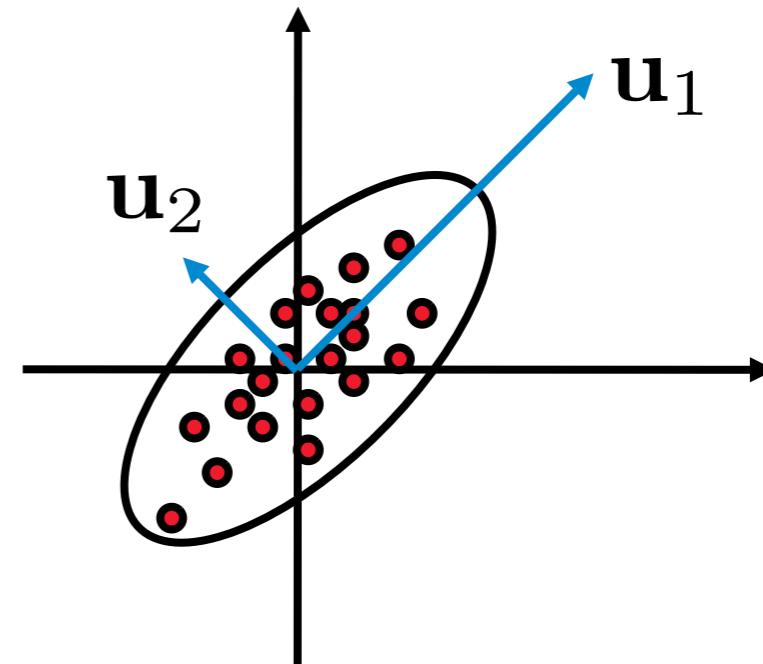


$$\mathbf{C} = \begin{bmatrix} c & 0 \\ 0 & d \end{bmatrix}, \quad c > d$$

# Intuition

$$\mathbf{x}^n - \bar{\mathbf{x}} = \sum_{i=1}^D a_i \mathbf{u}_i + \sum_{j=D+1}^M b_j \mathbf{u}_j$$

$$\tilde{\mathbf{x}}^n = \sum_{i=1}^D a_i \mathbf{u}_i + \bar{\mathbf{x}}$$



Data is correlated in its original coordinate system!

Note that axes of the projection are orthogonal and **decorrelate** the data; i.e. in the coordinate frame of these axes, the data is **uncorrelated**.  
(side note: this only works for Gaussians).

# Formulation of problem

- ◆ Let  $\mathbf{X} = [\mathbf{x}^1, \dots, \mathbf{x}^N] \in \mathbb{R}^{M \times N}$  be a matrix of  $N$  vectors in  $M$ -dimensional input space.
  - ◆ Let  $\mathbf{u} \in \mathbb{R}^M$  be a direction (a vector of length 1) in the input space.
  - ◆ The projection of the  $n$ -th vector  $\mathbf{x}^n$  onto the vector  $\mathbf{u}$  can be calculated as:
- $$a_n = \mathbf{u}^T \mathbf{x}^n = \sum_{i=1}^M x_i^n u_i$$
- ◆ The goal is to find a direction  $\mathbf{u}$  that maximizes the variance of the projections of all input vectors  $\mathbf{x}^n$ ,  $n = 1 \dots N$

# Mean & Variance

- ◆ The mean of the projection

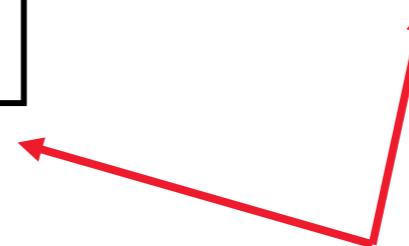
$$\bar{a} = \frac{1}{N} \sum_{n=1}^N a_n = \frac{1}{N} \sum_{n=1}^N \mathbf{u}^T \mathbf{x}^n = \frac{1}{N} \mathbf{u}^T \left[ \sum_{n=1}^N \mathbf{x}^n \right] = \mathbf{u}^T \bar{\mathbf{x}}$$

- ◆ ... is the projection of the mean.

# Mean & Variance

- ◆ Variance of the projection:

$$\begin{aligned}
 \sigma^2 &= \frac{1}{N} \sum_{n=1}^N (a_n - \bar{a})^2 = \frac{1}{N} \sum_{n=1}^N (\mathbf{u}^T \mathbf{x}^n - \mathbf{u}^T \bar{\mathbf{x}})^2 \\
 &= \frac{1}{N} \sum_{n=1}^N (\mathbf{u}^T (\mathbf{x}^n - \bar{\mathbf{x}}))^2 = \frac{1}{N} \sum_{n=1}^N \mathbf{u}^T (\mathbf{x}^n - \bar{\mathbf{x}})(\mathbf{x}^n - \bar{\mathbf{x}})^T \mathbf{u} \\
 &= \mathbf{u}^T \left[ \frac{1}{N} \sum_{n=1}^N (\mathbf{x}^n - \bar{\mathbf{x}})(\mathbf{x}^n - \bar{\mathbf{x}})^T \right] \mathbf{u} = \mathbf{u}^T \mathbf{C} \mathbf{u}
 \end{aligned}$$


covariance matrix

# Maximizing variance

- ◆ Maximize  $\sigma^2$  under the constraint that  $\|\mathbf{u}\| = 1$ .
- ◆ Maximize the following function (Langrangian):

$$F(\mathbf{u}; \lambda) = \mathbf{u}^T \mathbf{C} \mathbf{u} - \lambda(\|\mathbf{u}\|^2 - 1) = \mathbf{u}^T \mathbf{C} \mathbf{u} - \lambda(\mathbf{u}^T \mathbf{u} - 1)$$

$$\frac{\partial F(\mathbf{u}; \lambda)}{\partial \mathbf{u}} = 2\mathbf{C}\mathbf{u} - 2\lambda\mathbf{u} = 0$$

$\Leftrightarrow$

$\mathbf{C}\mathbf{u} = \lambda\mathbf{u}$

with

$$\|\mathbf{u}\| = 1$$

- ◆ Compute eigenvectors and eigenvalues of  $\mathbf{C}$ :
- ◆ The largest eigenvalue = maximal variance.
- ◆ Corresponding eigenvector = direction with the max. variance

# Eigendecomposition

- ◆ Let the eigenvectors and eigenvalues of  $\mathbf{C}$  be  $\mathbf{u}_k$  and  $\lambda_k$  for  $k \leq M$ , i.e.,  $\mathbf{C}\mathbf{u}_k = \lambda_k\mathbf{u}_k$  with  $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_M$
- ◆ In matrix form:  $\mathbf{C}\mathbf{U} = \mathbf{U}\Lambda$ , where  $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_M)$  and  $\mathbf{U} = [\mathbf{u}_1, \dots, \mathbf{u}_M]$ .
- ◆ Because  $\mathbf{U}$  is orthonormal (we usually assume that the eigenvectors have unit norm) we know that  $\mathbf{U}\mathbf{U}^T = \mathbf{I}$ .
- ◆ This means that we can decompose  $\mathbf{C}$  as  $\mathbf{C} = \mathbf{U}\Lambda\mathbf{U}^T$

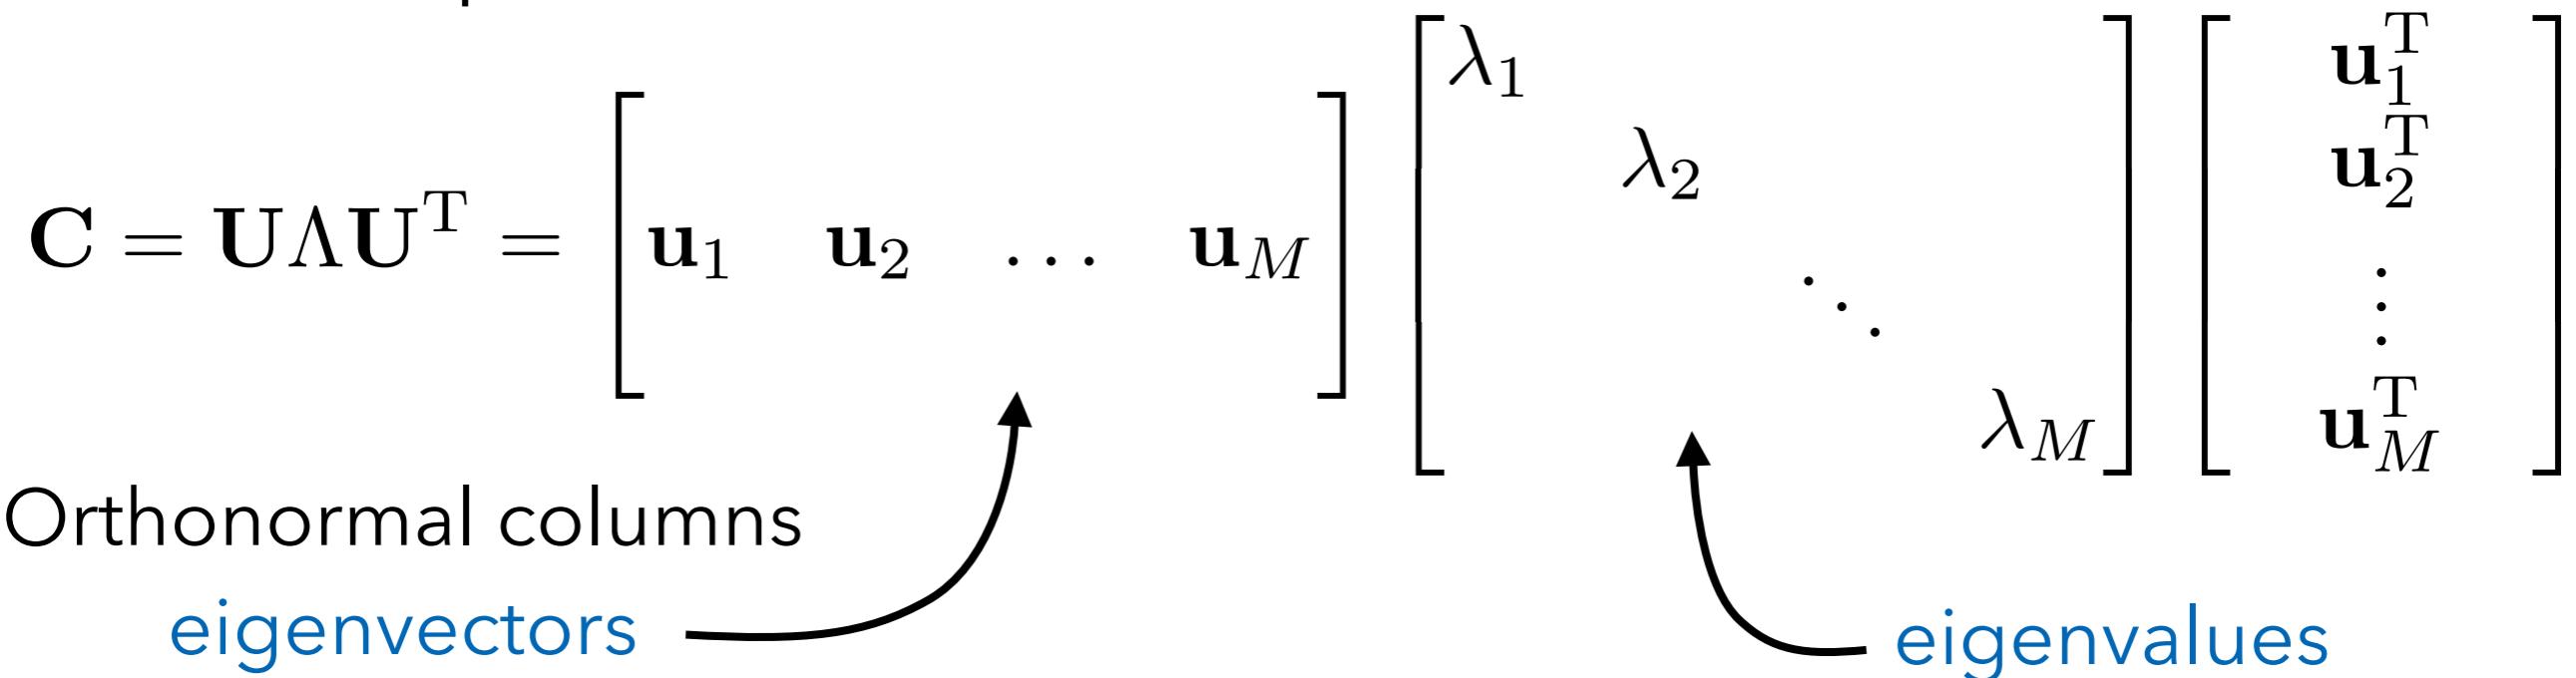
# Principal Component Analysis

- ◆  $\mathbf{C}$  is real, symmetric, and positive-definite. Thus we can decompose it as:

$$\mathbf{C} = \mathbf{U}\Lambda\mathbf{U}^T = \begin{bmatrix} \mathbf{u}_1 & \mathbf{u}_2 & \dots & \mathbf{u}_M \end{bmatrix} \begin{bmatrix} \lambda_1 & & & \\ & \lambda_2 & & \\ & & \ddots & \\ & & & \lambda_M \end{bmatrix} \begin{bmatrix} \mathbf{u}_1^T \\ \mathbf{u}_2^T \\ \vdots \\ \mathbf{u}_M^T \end{bmatrix}$$

Orthonormal columns  
eigenvectors

eigenvalues



- ◆ This is the **eigendecomposition** of  $\mathbf{C}$

# Principal Component Analysis

- ◆ **Observation:** If  $\lambda_k \approx 0$  for  $k > D$  for some  $D \ll M$  then we can use the subset of the first  $D$  eigenvectors to define a basis for approximating the data vectors.

$$\mathbf{x}^n - \bar{\mathbf{x}} = \sum_{i=1}^D a_i \mathbf{u}_i + \sum_{j=D+1}^M b_j \mathbf{u}_j$$

$$\mathbf{x}^n \approx \tilde{\mathbf{x}}^n = \bar{\mathbf{x}} + \sum_{i=1}^D a_i \mathbf{u}_i \quad \text{where} \quad a_i = \mathbf{u}_i^T (\mathbf{x}^n - \bar{\mathbf{x}})$$

- ◆ This representation has the **minimal mean squared error** (MSE) of all linear representations of dimension  $D$ :

$$\min E(\mathbf{u}_1, \dots, \mathbf{u}_D) = \sum_{n=1}^N \|\mathbf{x}^n - \tilde{\mathbf{x}}^n\|^2$$

# Principal Component Analysis

- ◆ Now we know how we can represent our data in a lower dimensional space in a principled fashion.
  - ◆ We compute the mean of the data, and subtract it.
  - ◆ We compute the covariance matrix, decompose it, and choose the first  $D$  eigenvalues.
  - ◆ This gives us an (eigen)basis for representing the data:

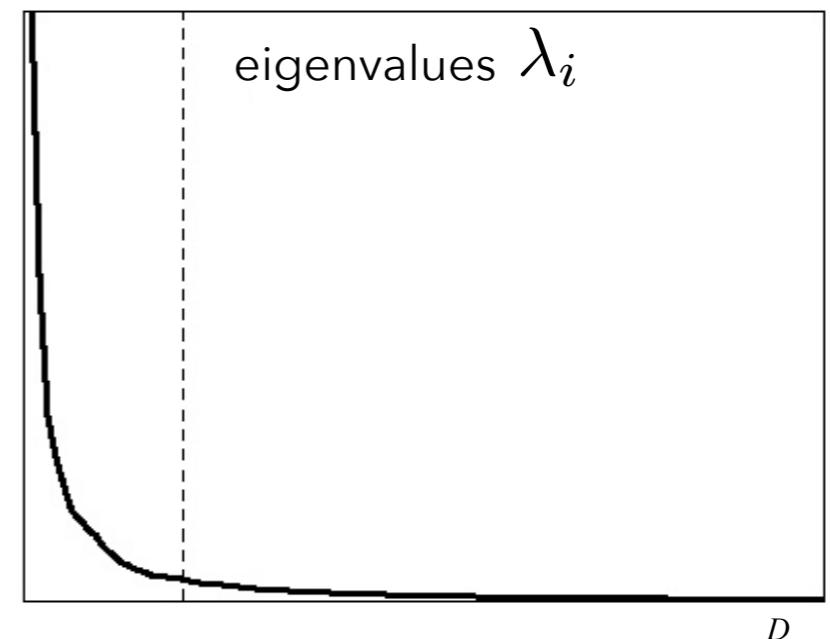
$$\mathbf{a}^n = \mathbf{B}^T(\mathbf{x}^n - \bar{\mathbf{x}}) \quad \text{where} \quad \mathbf{B} = [\mathbf{u}_1, \dots, \mathbf{u}_D]$$
$$\tilde{\mathbf{x}}^n = \bar{\mathbf{x}} + \mathbf{B}\mathbf{a}^n$$

- ◆ It is common to also normalize the variance of each dimension.
- ◆ Question:
  - ◆ How do we choose  $D$ ?

# Choosing $D$

- ◆ Larger  $D$  leads to better approximation.
- ◆ There are at least 2 good possibilities for choosing  $D$ :
  - ◆ First, choose  $D$  based on application performance, i.e. choose the smallest  $D$  that makes the application work well enough.
  - ◆ Second, choose  $D$  so that the eigenbasis captures some fraction of the variance, say  $\eta = 0.9$ 
    - ◆ Remember: Eigenvalue  $\lambda_i$  describes the marginal variance captured by  $\mathbf{u}_i$
    - ◆ Choose  $D$  using

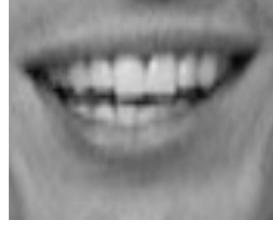
$$\sum_{i=1}^D \lambda_i \geq \eta \sum_{i=1}^M \lambda_i$$



# PCA in Practice



- ◆ Suppose we want to find a low-dimensional representation for mouth images.
- ◆ Let's assume the images are of moderate size:  $128 \times 128$
- ◆ How "big" is the covariance matrix?
  - ◆ The data vector has  $2^7 \cdot 2^7 = 2^{14}$  dimensions.
  - ◆ So the covariance matrix has  $2^{28}$  entries.
  - ◆ Takes up  $2^{31}$  Bytes, i.e. 2GB.
  - ◆ Too large to handle efficiently.
- ◆ What can we do to reduce this?


$$p_i = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_{n \times m} \end{bmatrix}$$

# More Efficiency

- ◆ Observation 1:
  - ◆ While we may work in spaces that have 100000s of dimensions, we often only have 100s or 1000s of examples.
- ◆ Observation 2:
  - ◆ We do not need to compute all eigenvectors and eigenvalues. Computing a few dozen of the largest ones usually suffices.
- ◆ But how do we exploit this?

# Rewrite PCA

- ◆ Data vectors in matrix form:  $\mathbf{X} = [\mathbf{x}^1, \dots, \mathbf{x}^N]$

- ◆ First subtract the mean:

$$\hat{\mathbf{X}} = \mathbf{X} - [\bar{\mathbf{x}}, \dots, \bar{\mathbf{x}}]$$

$$\bar{\mathbf{x}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}^n$$

- ◆ Express the covariance matrix differently:

$$\hat{\mathbf{C}} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}^n - \bar{\mathbf{x}})(\mathbf{x}^n - \bar{\mathbf{x}})^T = \frac{1}{N} \hat{\mathbf{X}} \hat{\mathbf{X}}^T$$

# Singular Value Decomposition



- ◆ SVD is a generalization of the eigendecomposition to rectangular matrices  $\hat{\mathbf{X}} \in \mathbb{R}^{M \times N}$

$$\hat{\mathbf{X}} = \mathbf{U} \mathbf{S} \mathbf{V}^T = \underbrace{\begin{bmatrix} \mathbf{u}_1 & \mathbf{u}_2 & \dots & \mathbf{u}_M \end{bmatrix}}_{M \times M} \underbrace{\begin{bmatrix} s_1 & & & \\ & s_2 & & \\ & & \ddots & \\ & & & s_N \end{bmatrix}}_{M \times N} \underbrace{\begin{bmatrix} \mathbf{v}_1^T \\ \mathbf{v}_2^T \\ \vdots \\ \mathbf{v}_N^T \end{bmatrix}}_{N \times N}$$

Orthonormal columns  
left-singular vectors

singular values  
(non-negative)

Orthonormal rows  
right-singular vectors

# How to use SVD for PCA

- ◆ Decompose the data matrix  $\hat{\mathbf{X}}$  with SVD:

$$\begin{aligned}
 \frac{1}{N} \hat{\mathbf{X}} \hat{\mathbf{X}}^T &= \frac{1}{N} \mathbf{U} \mathbf{S} \mathbf{V}^T (\mathbf{U} \mathbf{S} \mathbf{V}^T)^T \\
 &= \frac{1}{N} \mathbf{U} \mathbf{S} \mathbf{V}^T \mathbf{V} \mathbf{S}^T \mathbf{U}^T \\
 &= \frac{1}{N} \mathbf{U} \mathbf{S} \mathbf{S}^T \mathbf{U}^T \\
 &= \mathbf{U} \left( \frac{1}{N} \mathbf{S}^2 \right) \mathbf{U}^T \\
 &= \mathbf{U} \Lambda \mathbf{U}^T
 \end{aligned}$$

# How to use SVD for PCA

- ◆ If we perform SVD on the data matrix (after subtracting the mean) then:
  - ◆ The left-singular vectors give us the eigenvectors of the covariance matrix.
  - ◆ The singular values let us easily compute the eigenvalues of the covariance matrix:

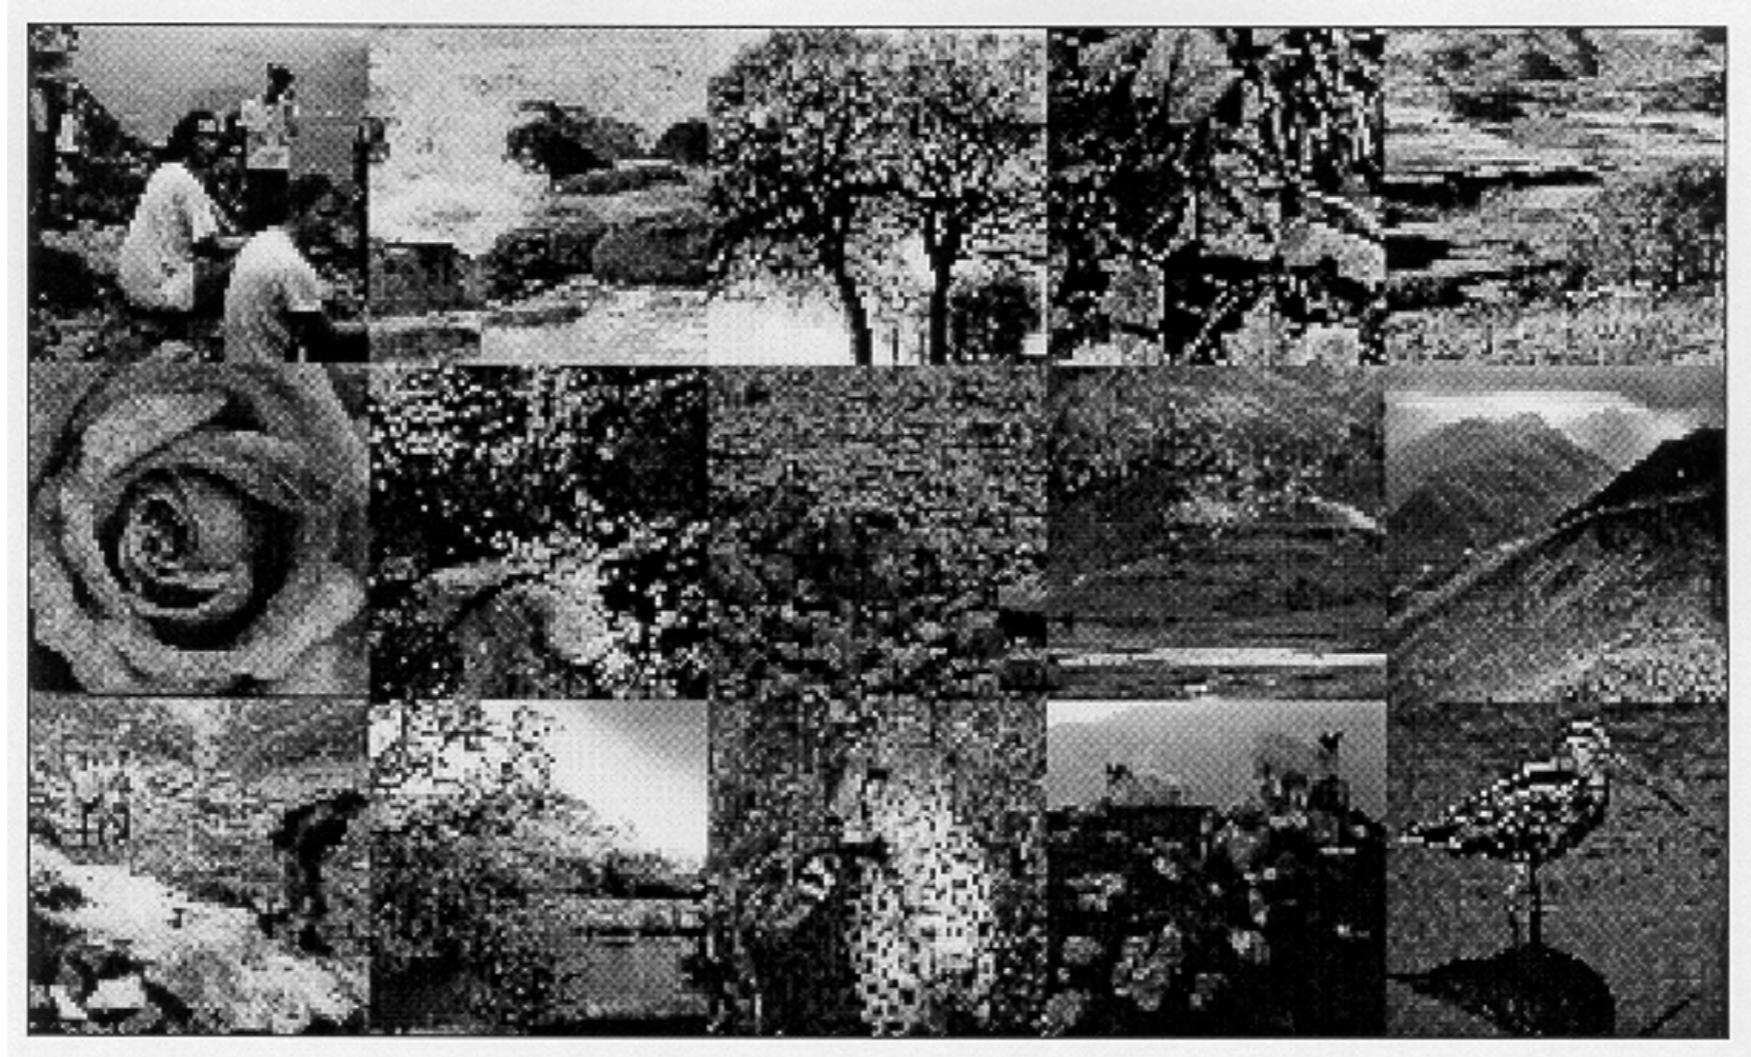
$$\lambda_i = \frac{1}{N} s_i^2$$

- ◆ **Advantage:** We never have to explicitly build and store the covariance matrix!

# How does this help us?

- ◆  $\mathbf{S} \in \mathbb{R}^{M \times N}$  is of manageable size now, and so is  $\mathbf{V} \in \mathbb{R}^{N \times N}$
- ◆ But  $\mathbf{U} \in \mathbb{R}^{M \times M}$  is still much too large.
- ◆ **Observation 1:** We only need the first  $N$  left-singular vectors, because only  $N$  singular values can be non-zero.
  - ◆ “Economy” decomposition, makes it all manageable memory-wise and computationally.
- ◆ **Observation 2:** We will typically need even fewer than  $N$  singular values and left-singular vectors.
  - ◆ Special variants that only compute a partial decomposition.

# Generic Image Ensembles



Fleet & Szeliski

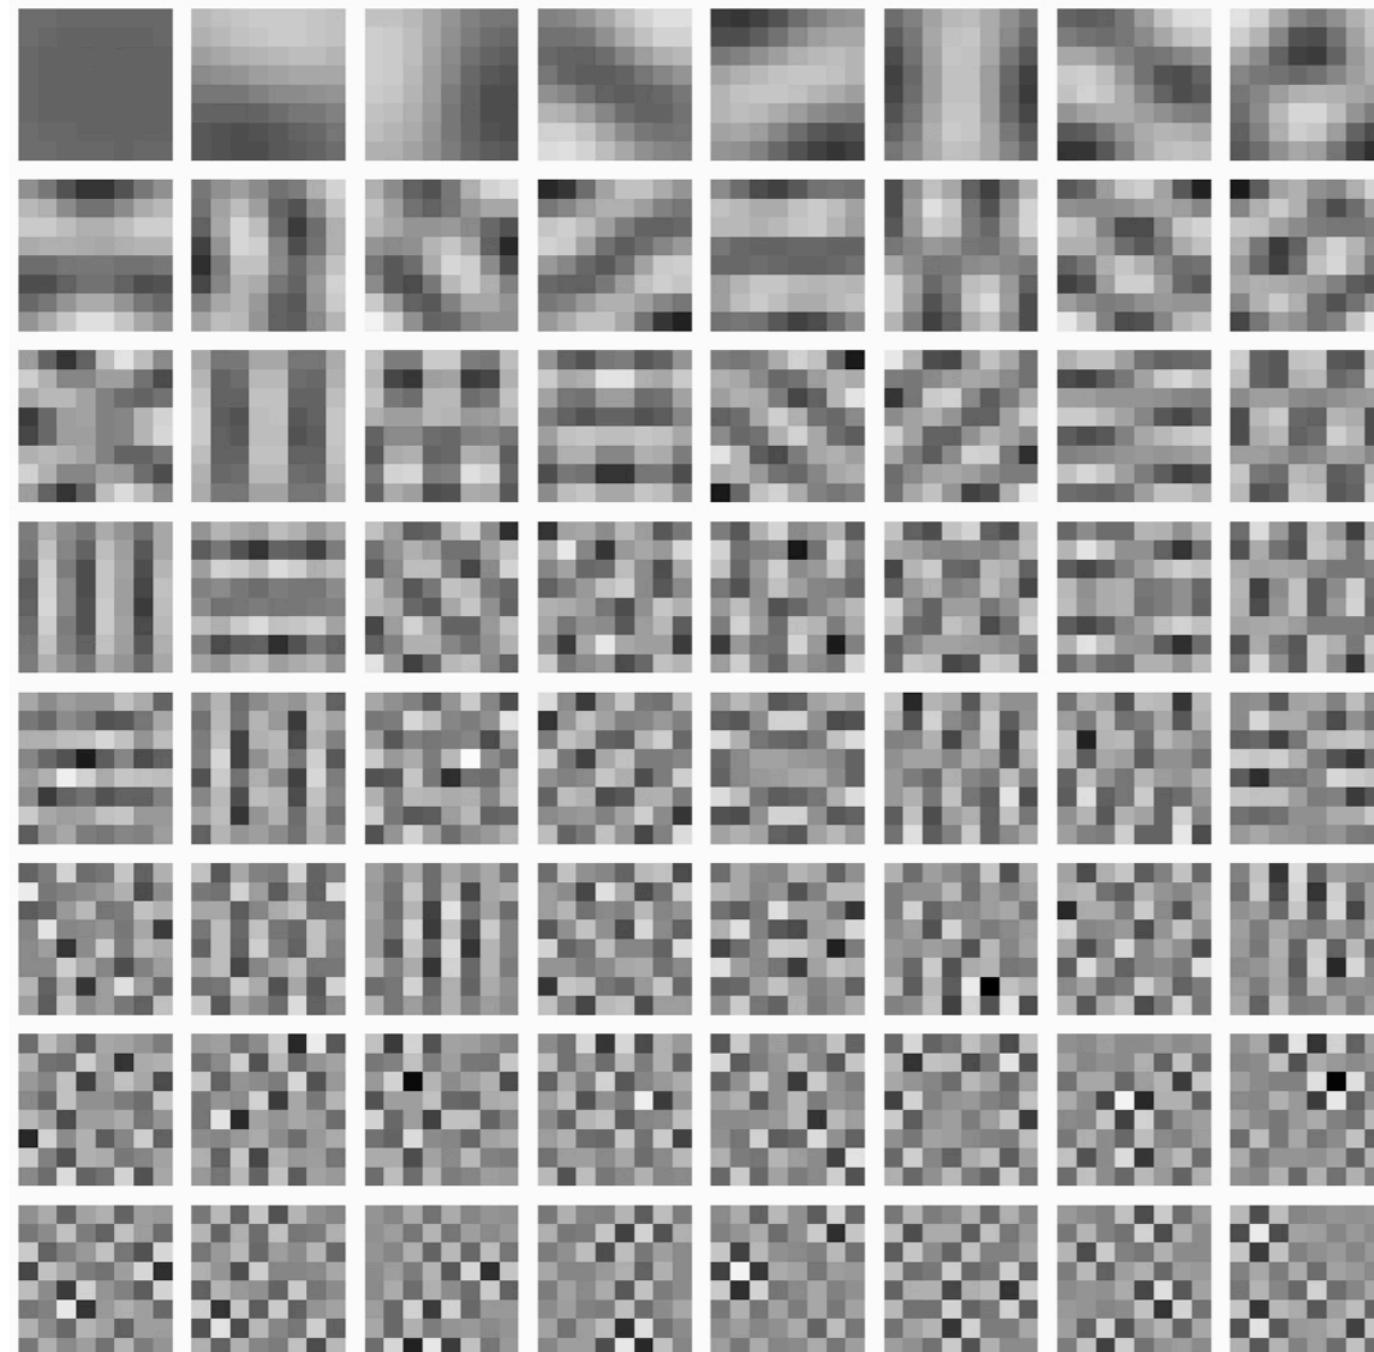
# Generic Image Ensembles



Is there a low dimensional model describing natural images?

Fleet & Szeliski

# PCA of Natural Image Patches



What do these bases look like?

8X8 image patches

Fleet & Szeliski

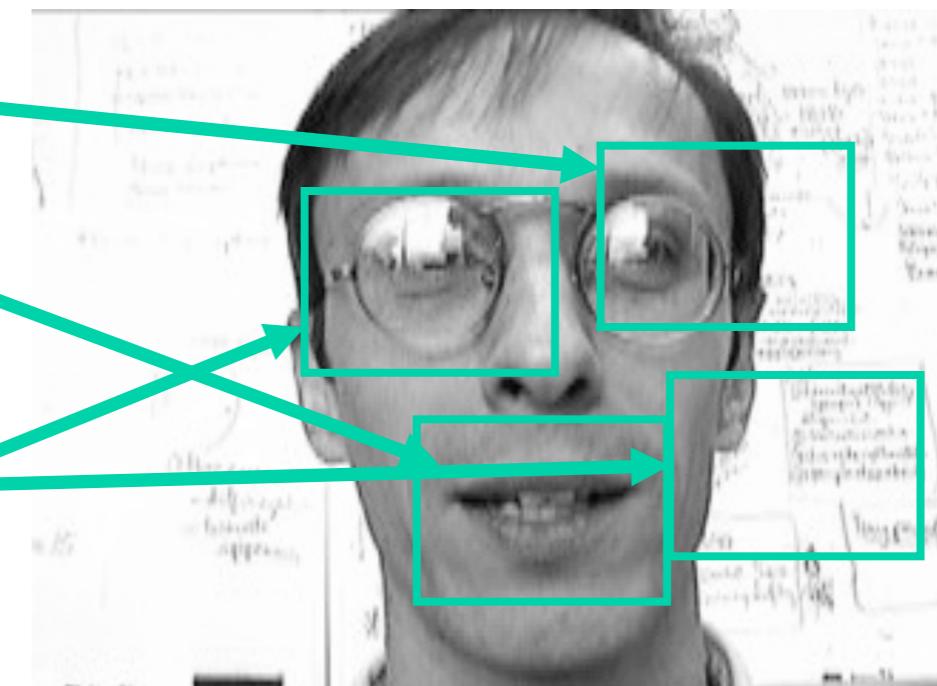
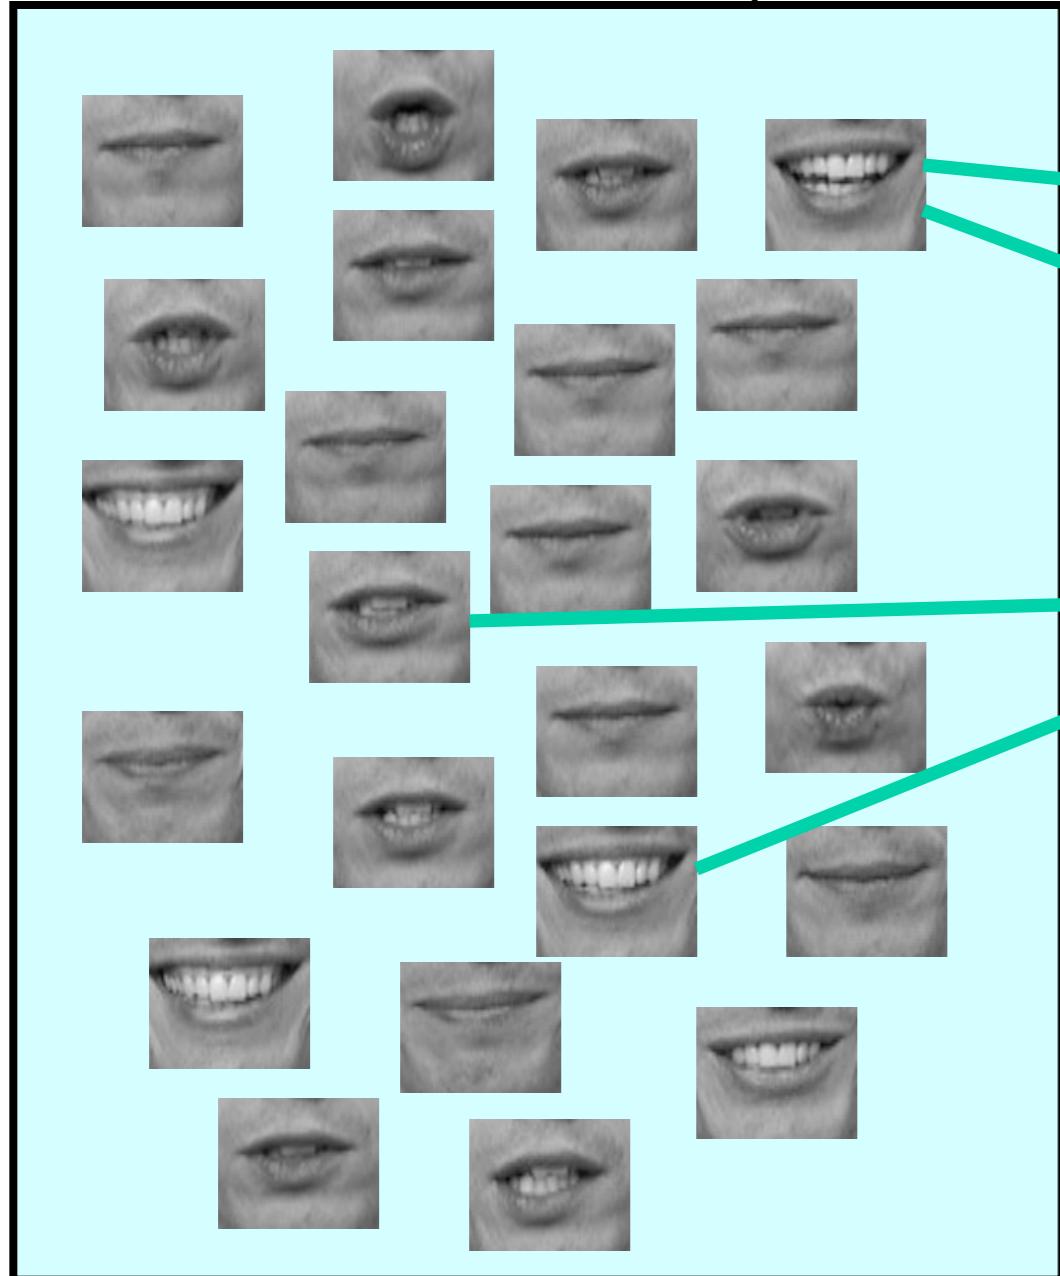
# Search and Recognition



- ◆ How can we find the mouth?
- ◆ How can we recognize the “expression”?

# Naïve View-Based Approach

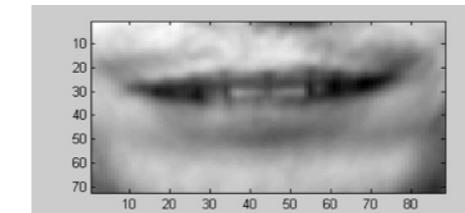
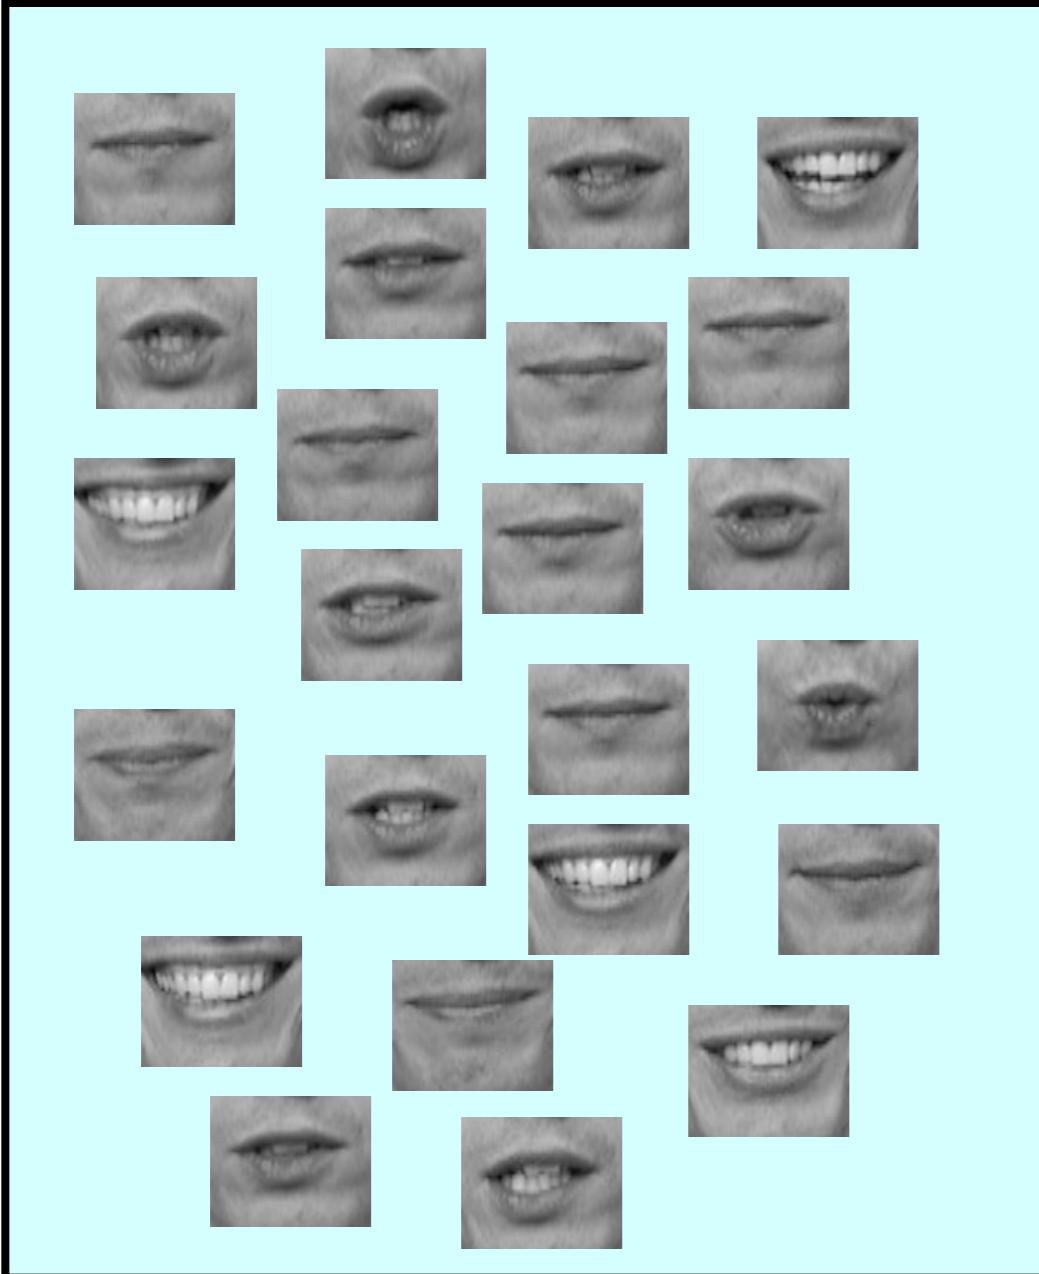
Database of mouth “templates”



- Search every image region (at every scale).
- Compare each template; chose the best match.

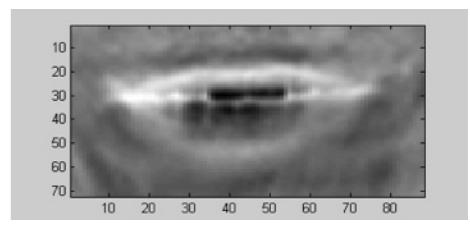
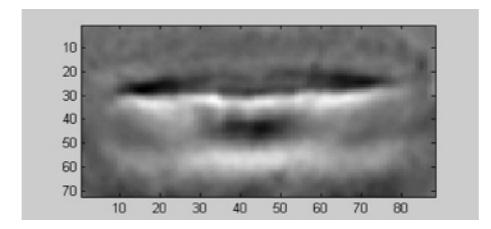
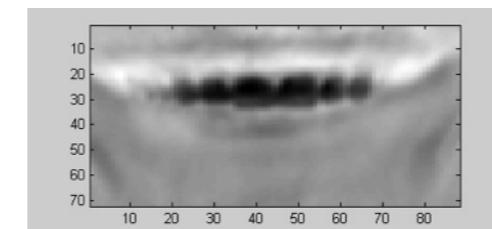
# View-Based Approach

Database of mouth “templates”



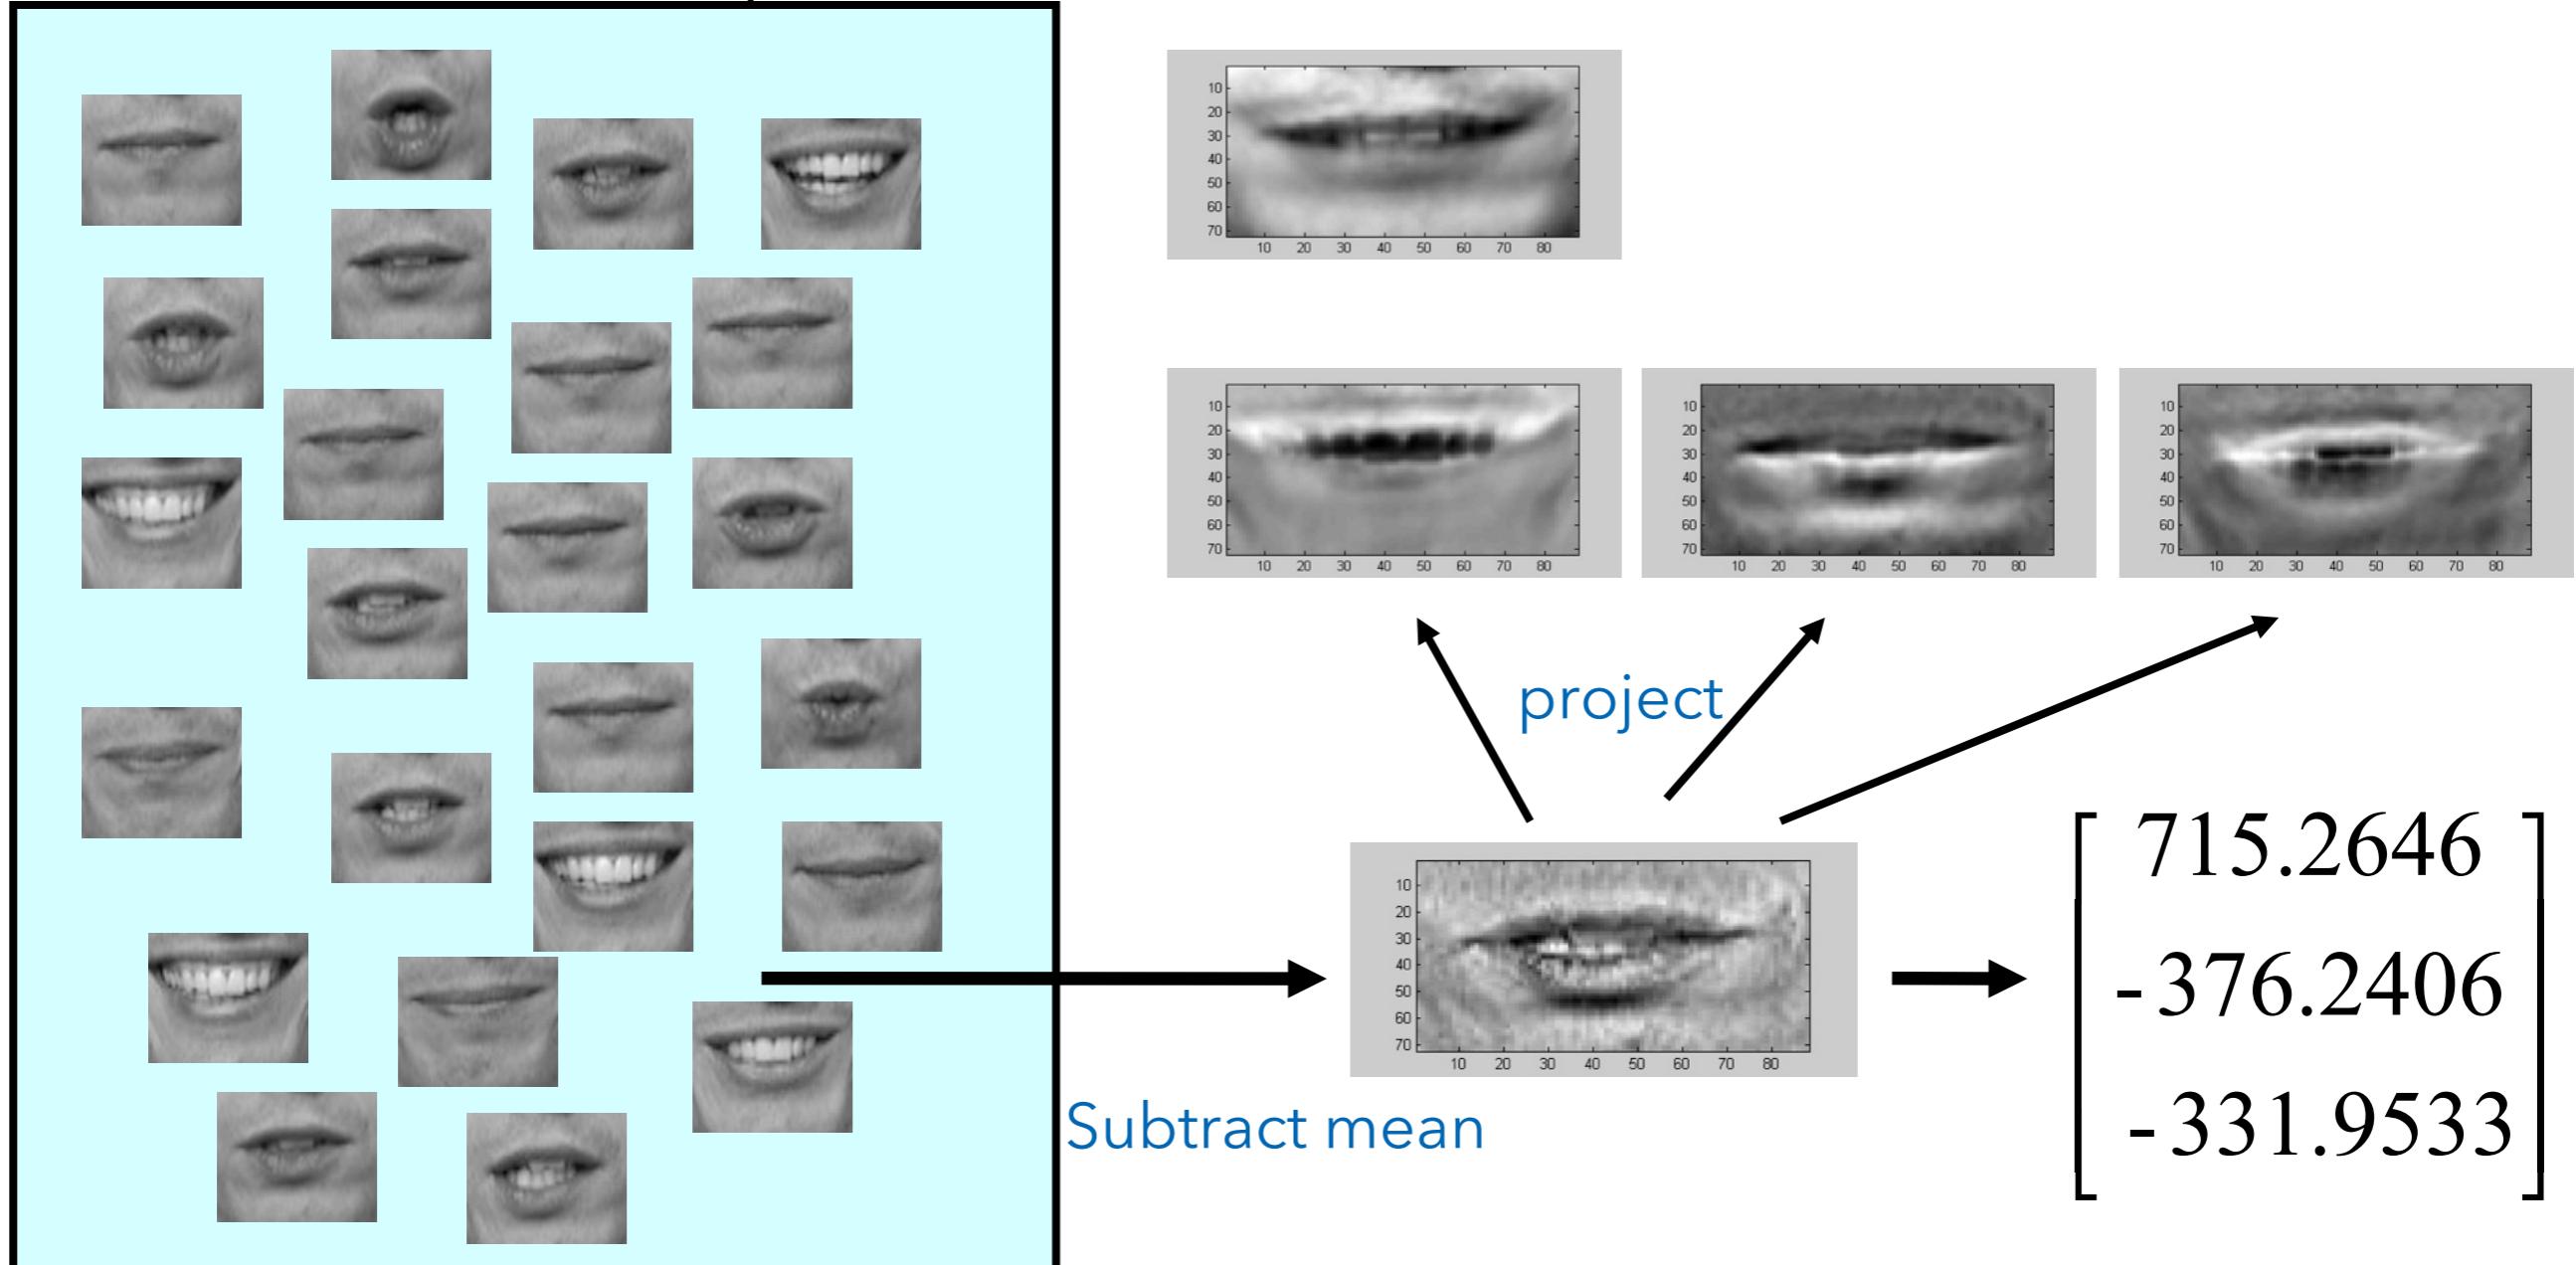
Mean

First three eigenvectors:

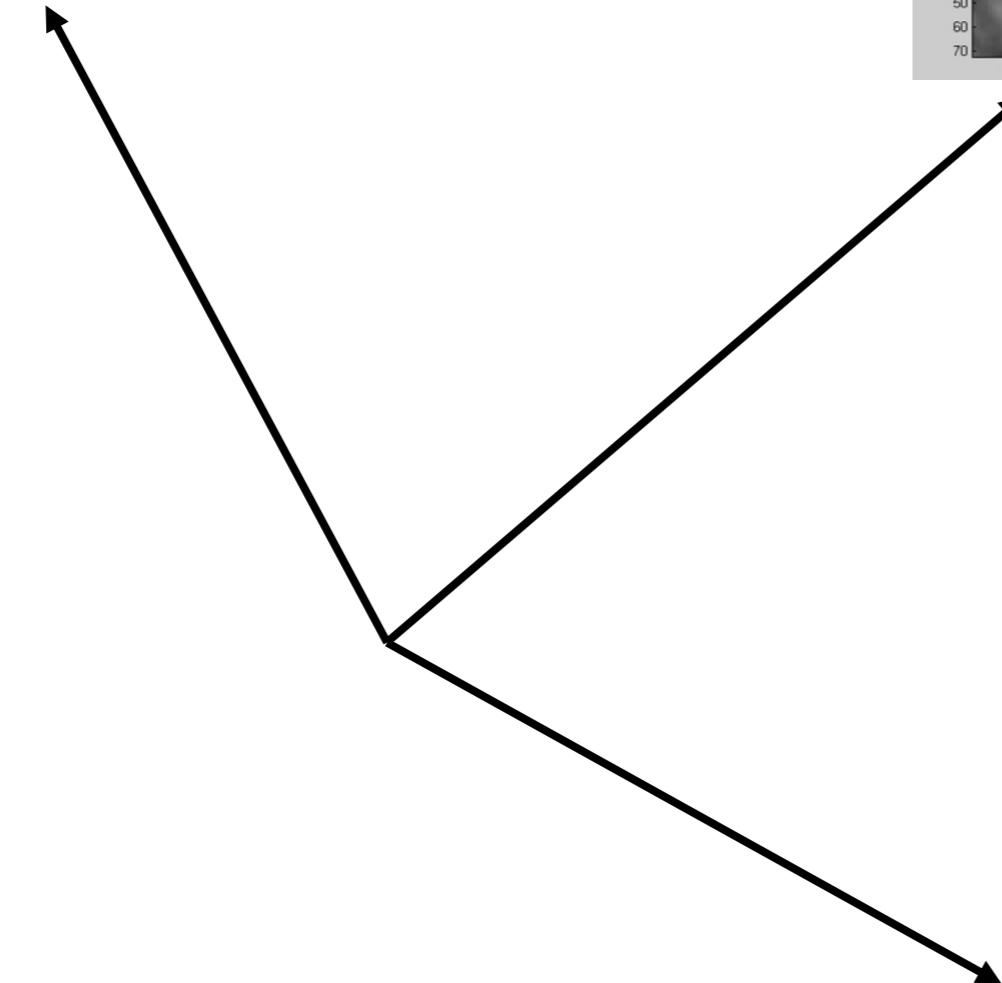
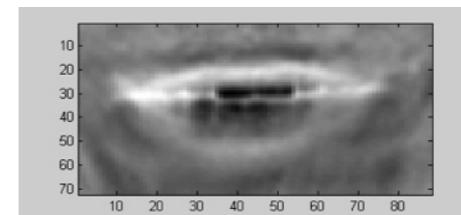
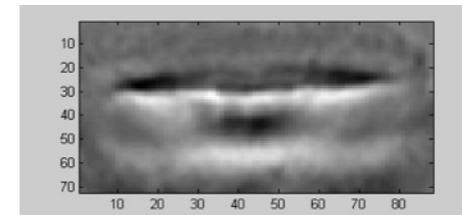
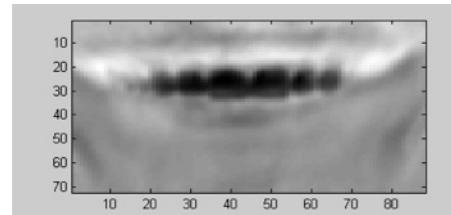


# View-Based Approach

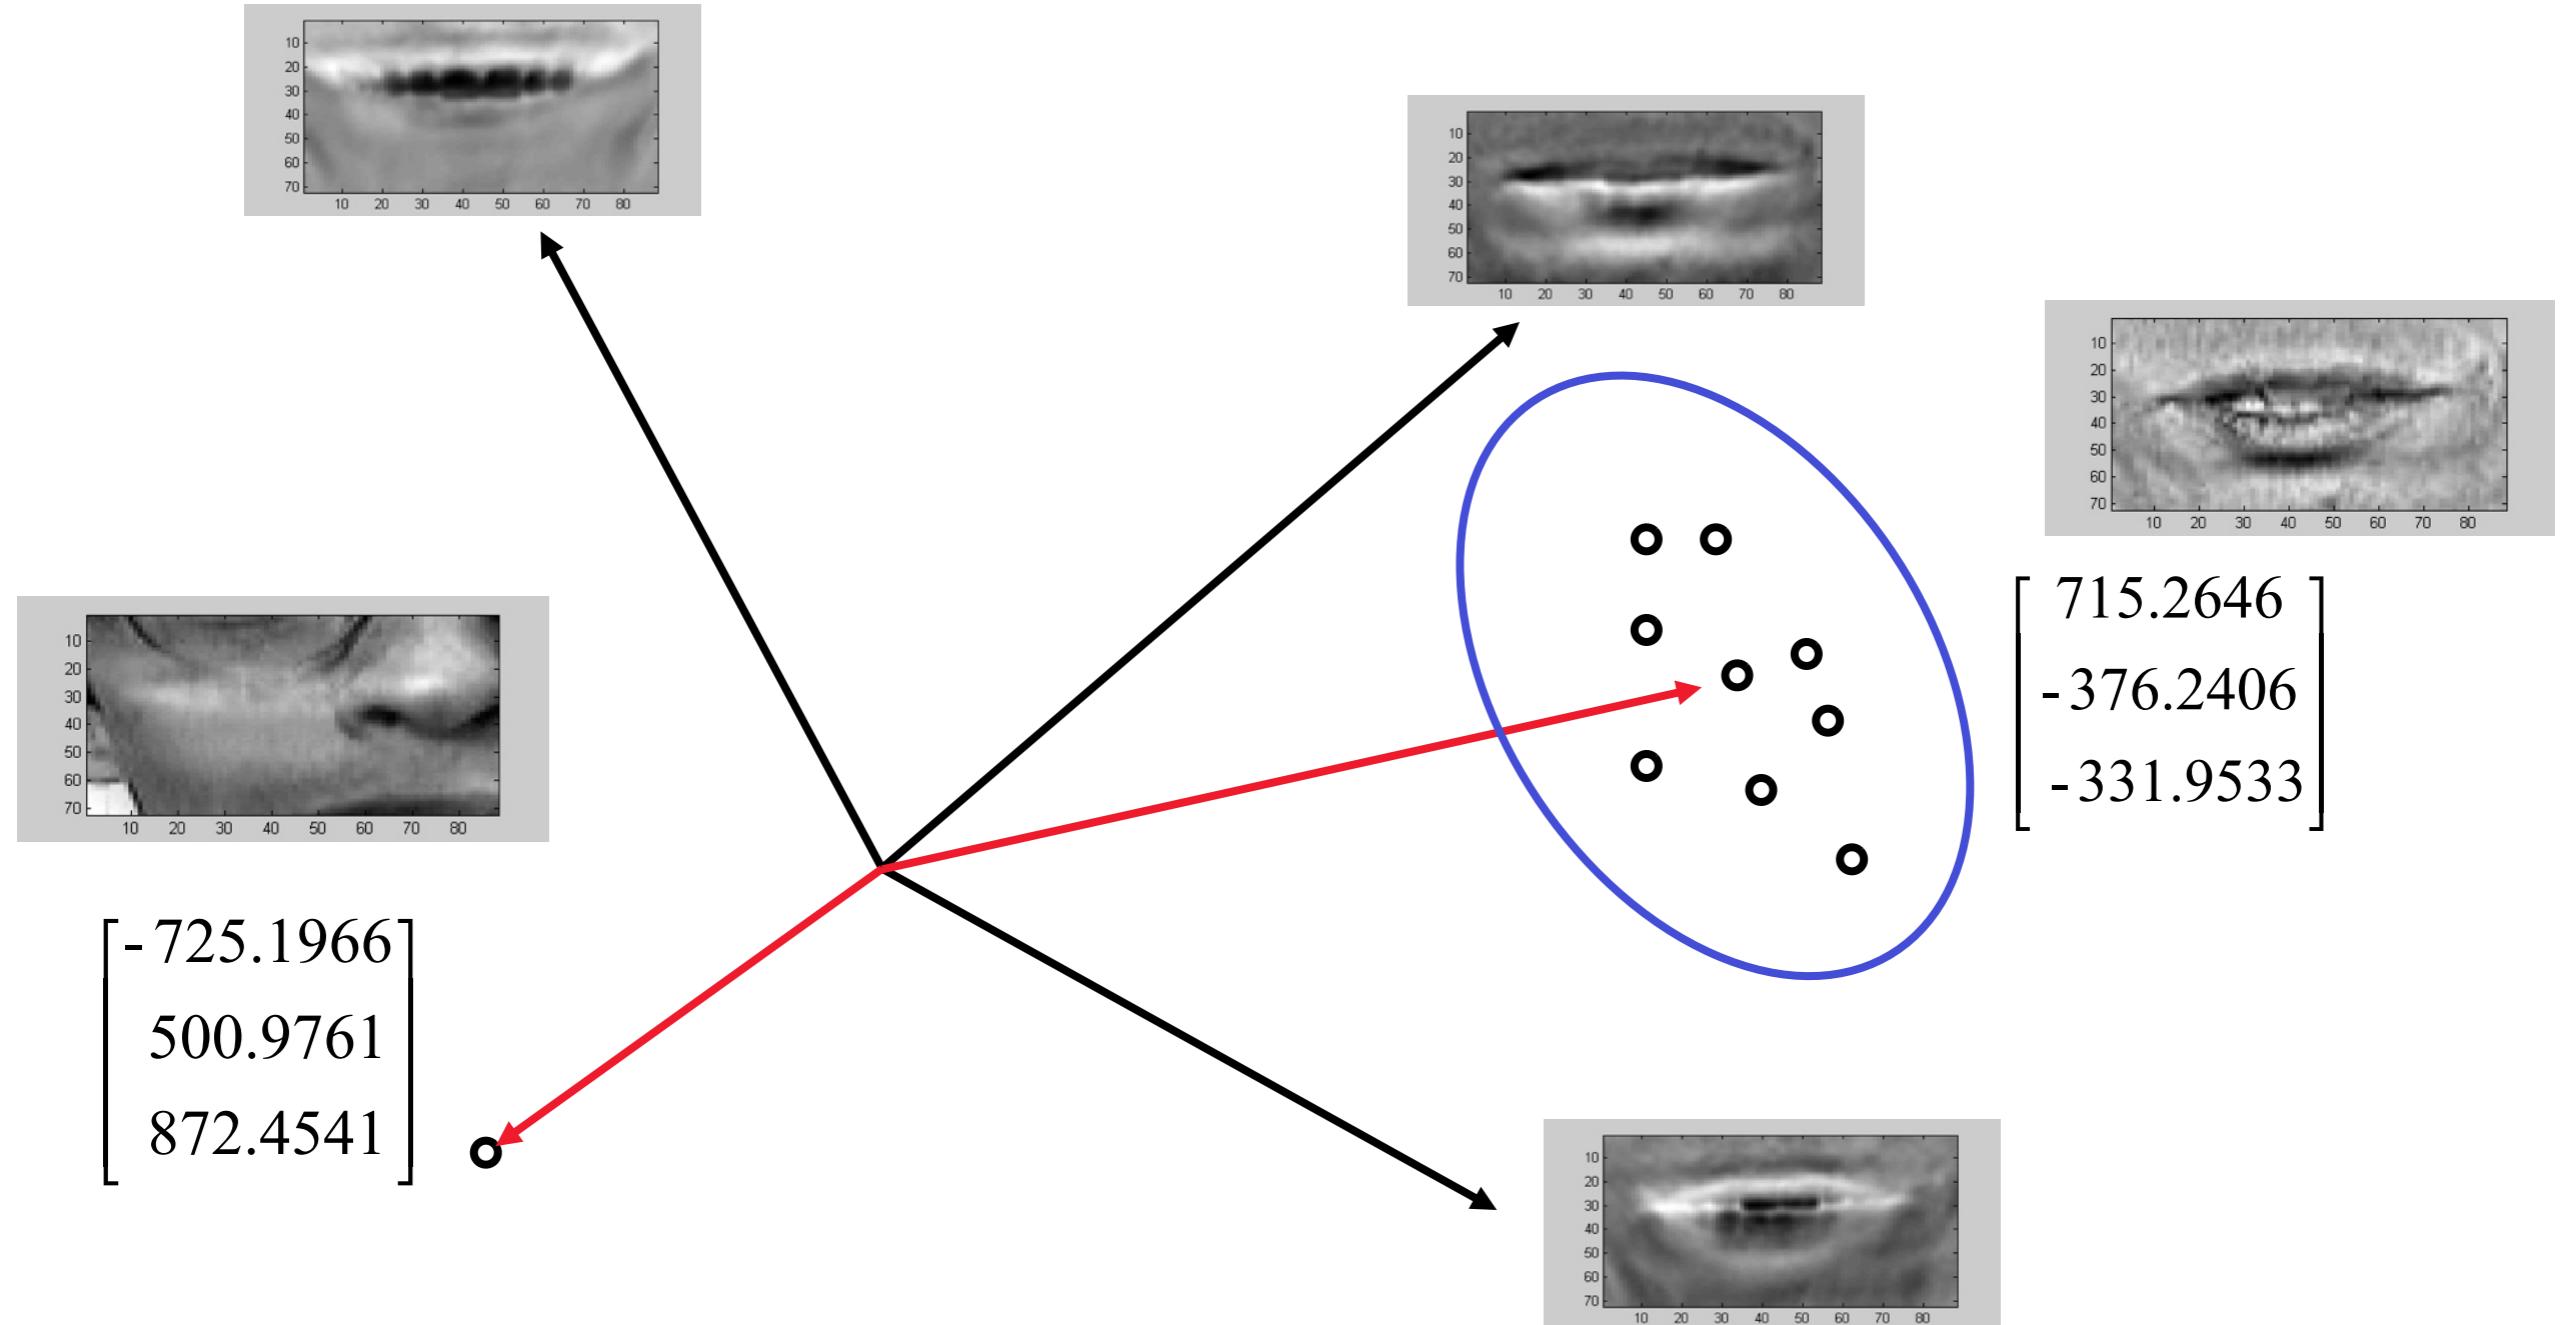
Database of mouth “templates”



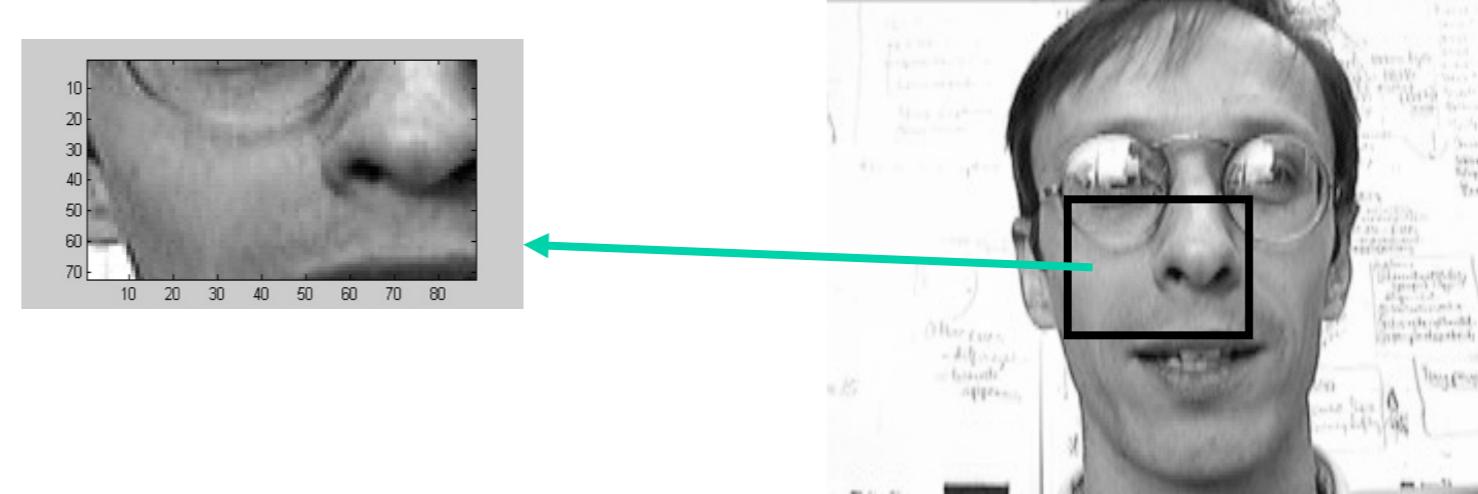
# Mouth Space



# Mouth Space



# Simple Search Strategy



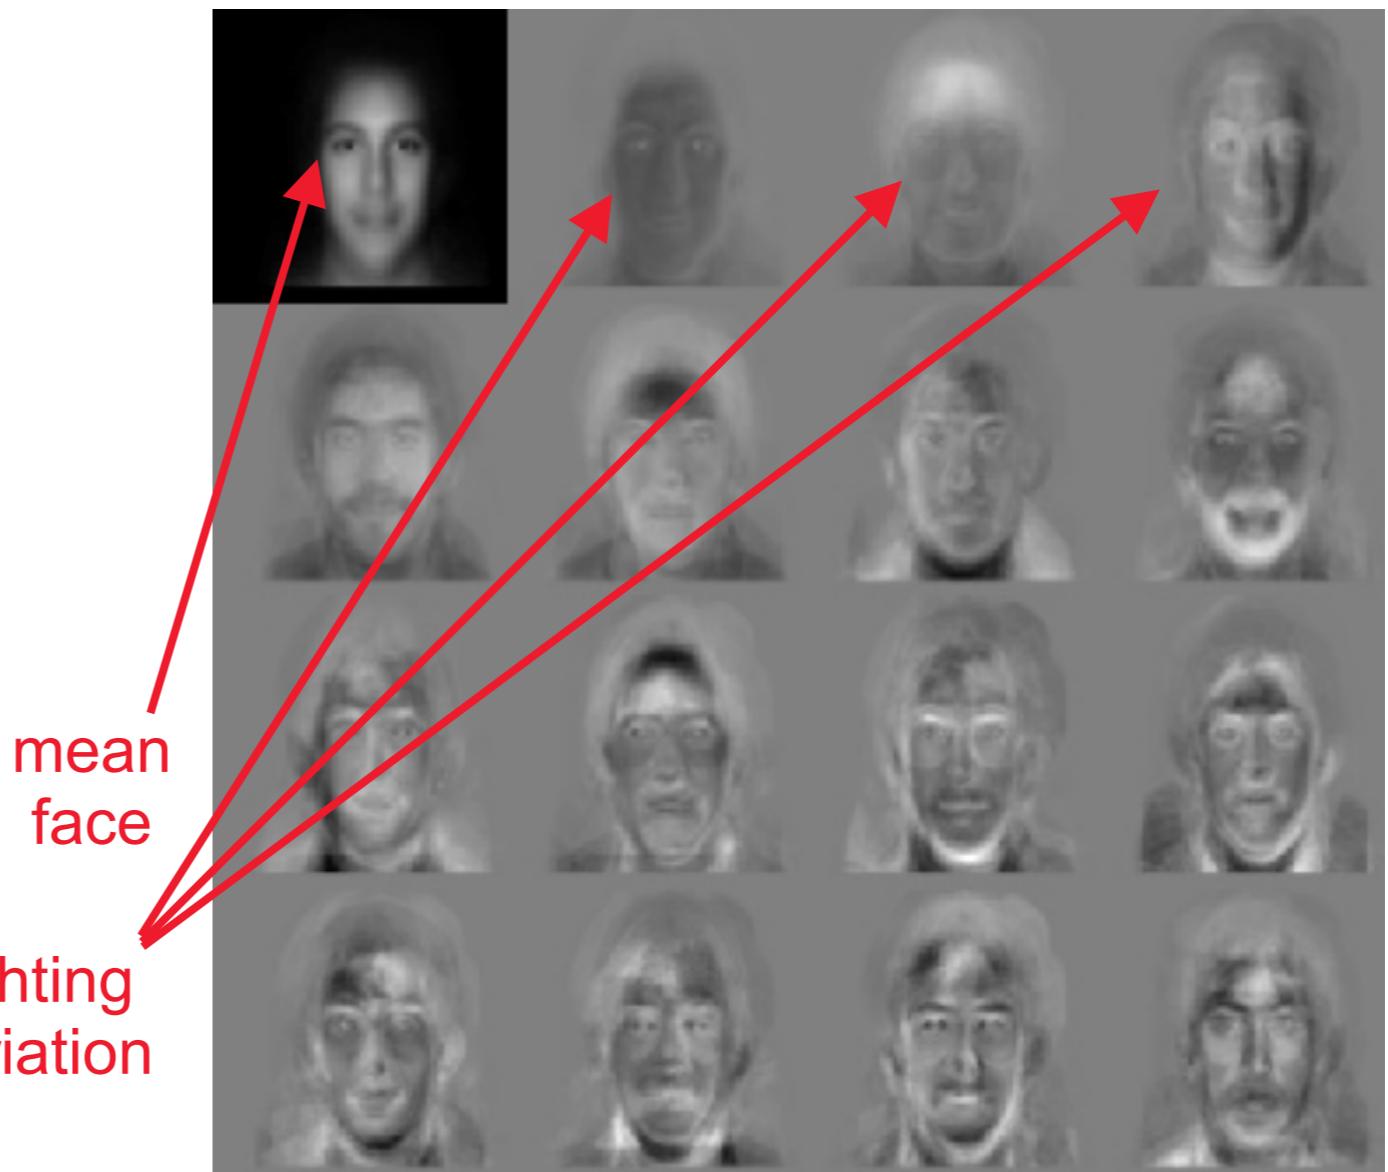
- ◆ Project each training image onto the low-dimensional subspace. Store the vectors of coefficients.
- ◆ For each image region:
  - ◆ Project it onto the low-dimensional subspace.
  - ◆ Compare this to each stored coefficient vector (cheap).
  - ◆ If the smallest distance is less than some threshold, then it is a mouth.

# EigenFaces

- ◆ First popular use of PCA for object recognition was for the detection and recognition of faces

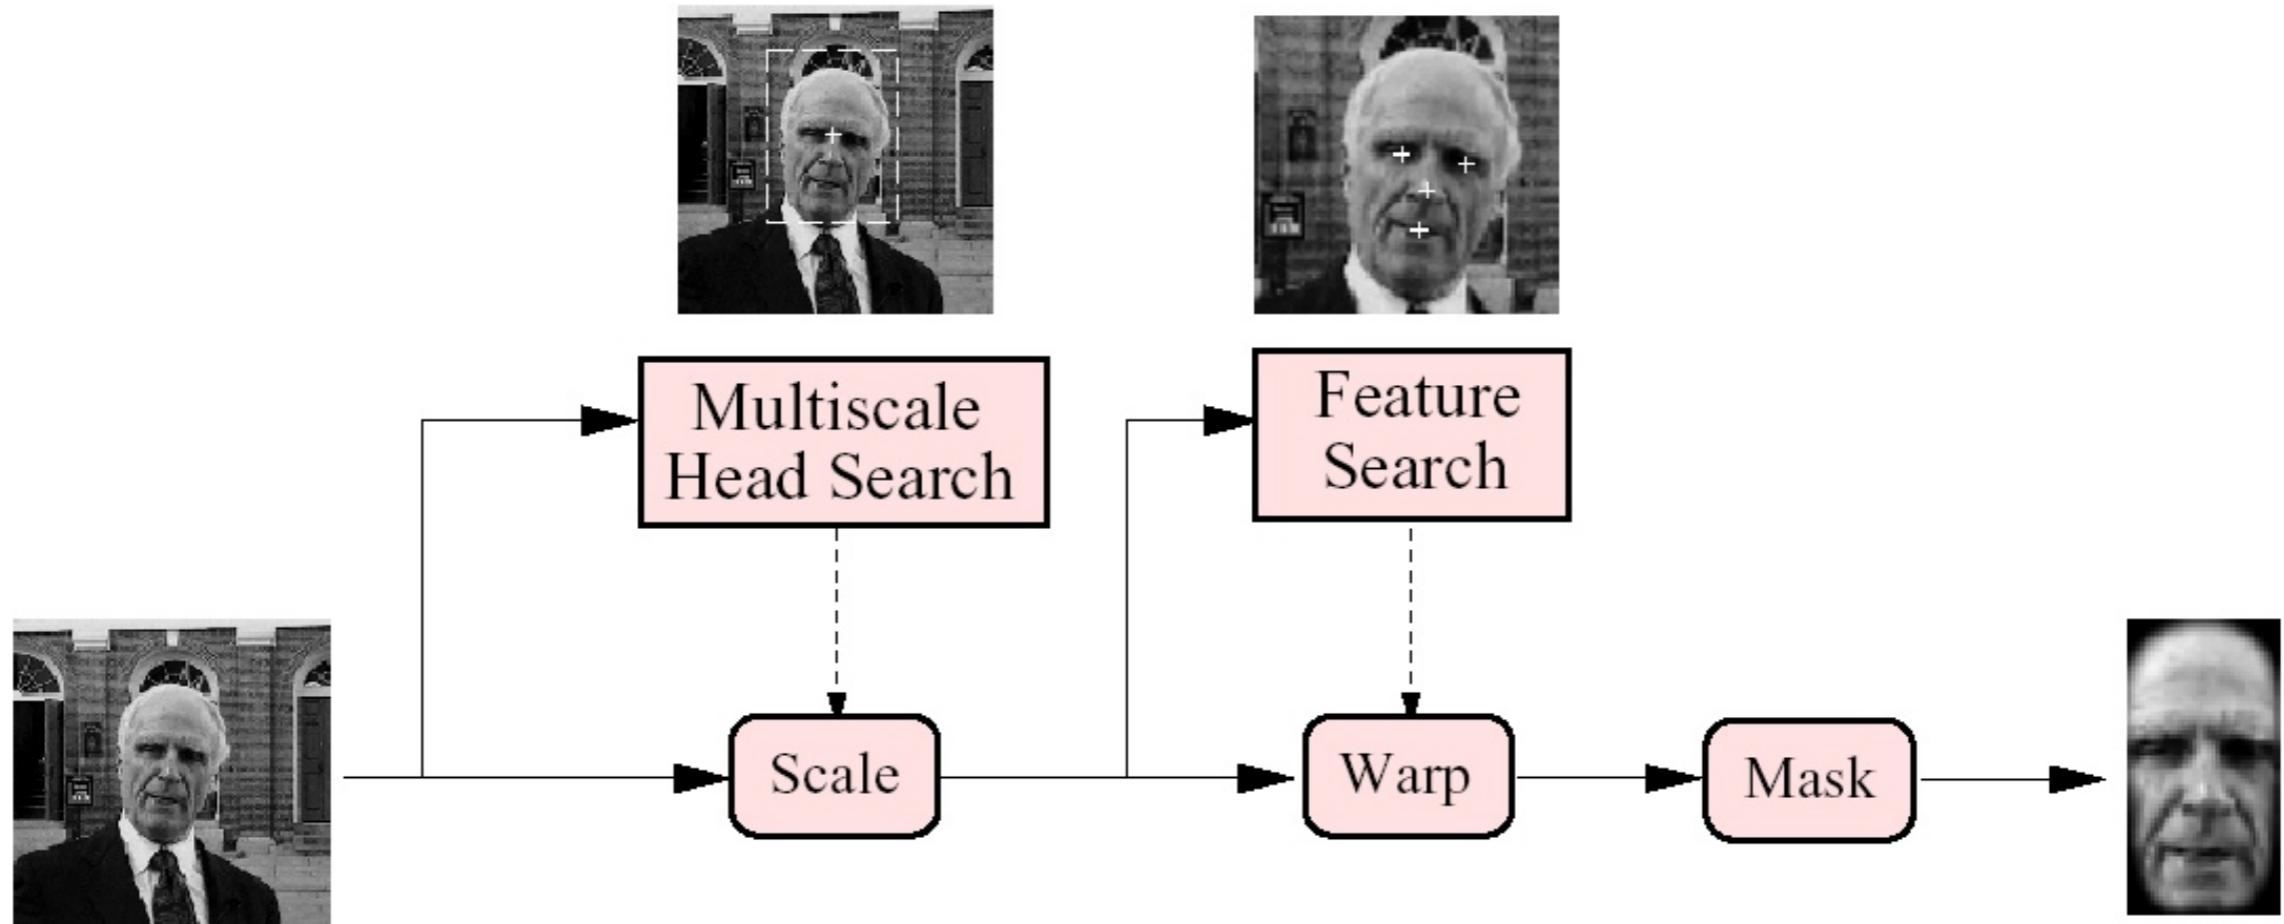
[Turk and Pentland, 1991]

- ▶ Collect a face ensemble
- ▶ Normalize for contrast, scale, & orientation.
- ▶ Remove backgrounds
- ▶ Apply PCA & choose the first D eigen-images that account for most of the variance of the data.



Fleet & Szeliski

# Face Recognition

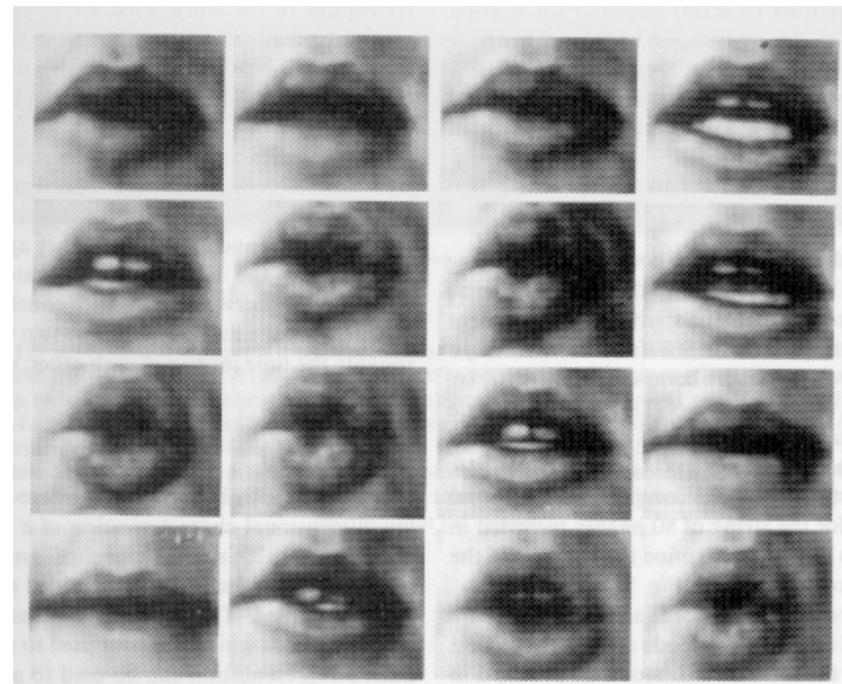


[Moghaddam, Jebara & Pentland, 2000]

Fleet & Szeliski

# EigenFeatures

Data:



PCA

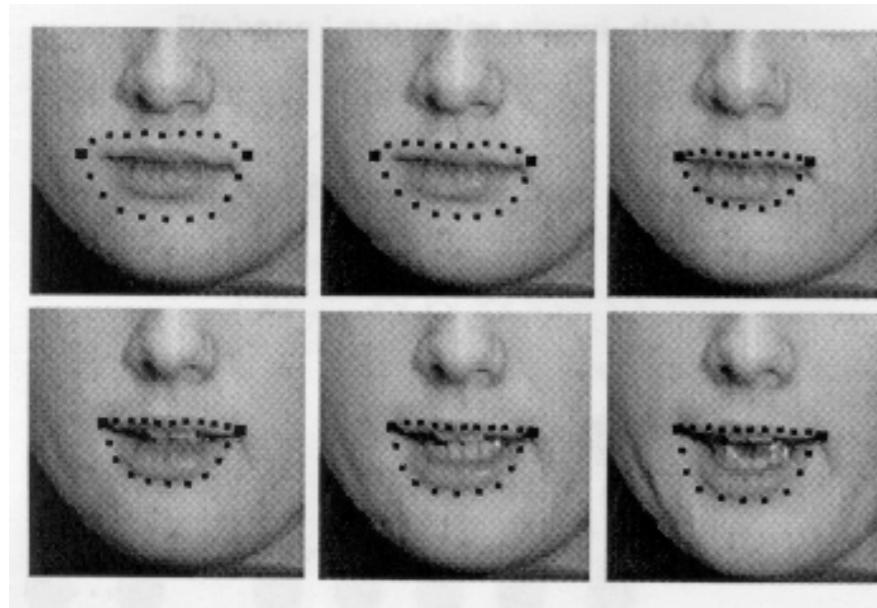
New Basis Vectors



Fleet & Szeliski

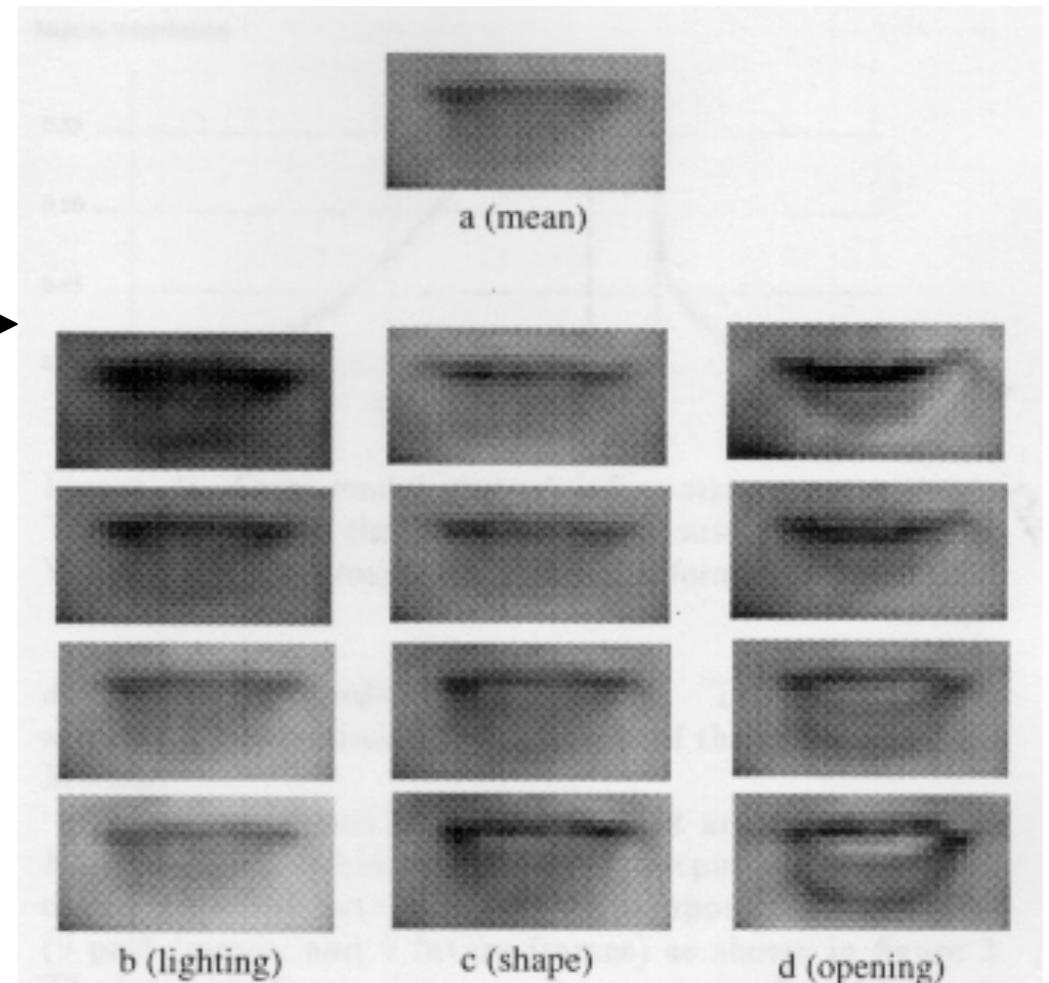
# EigenFeatures

Data:



PCA

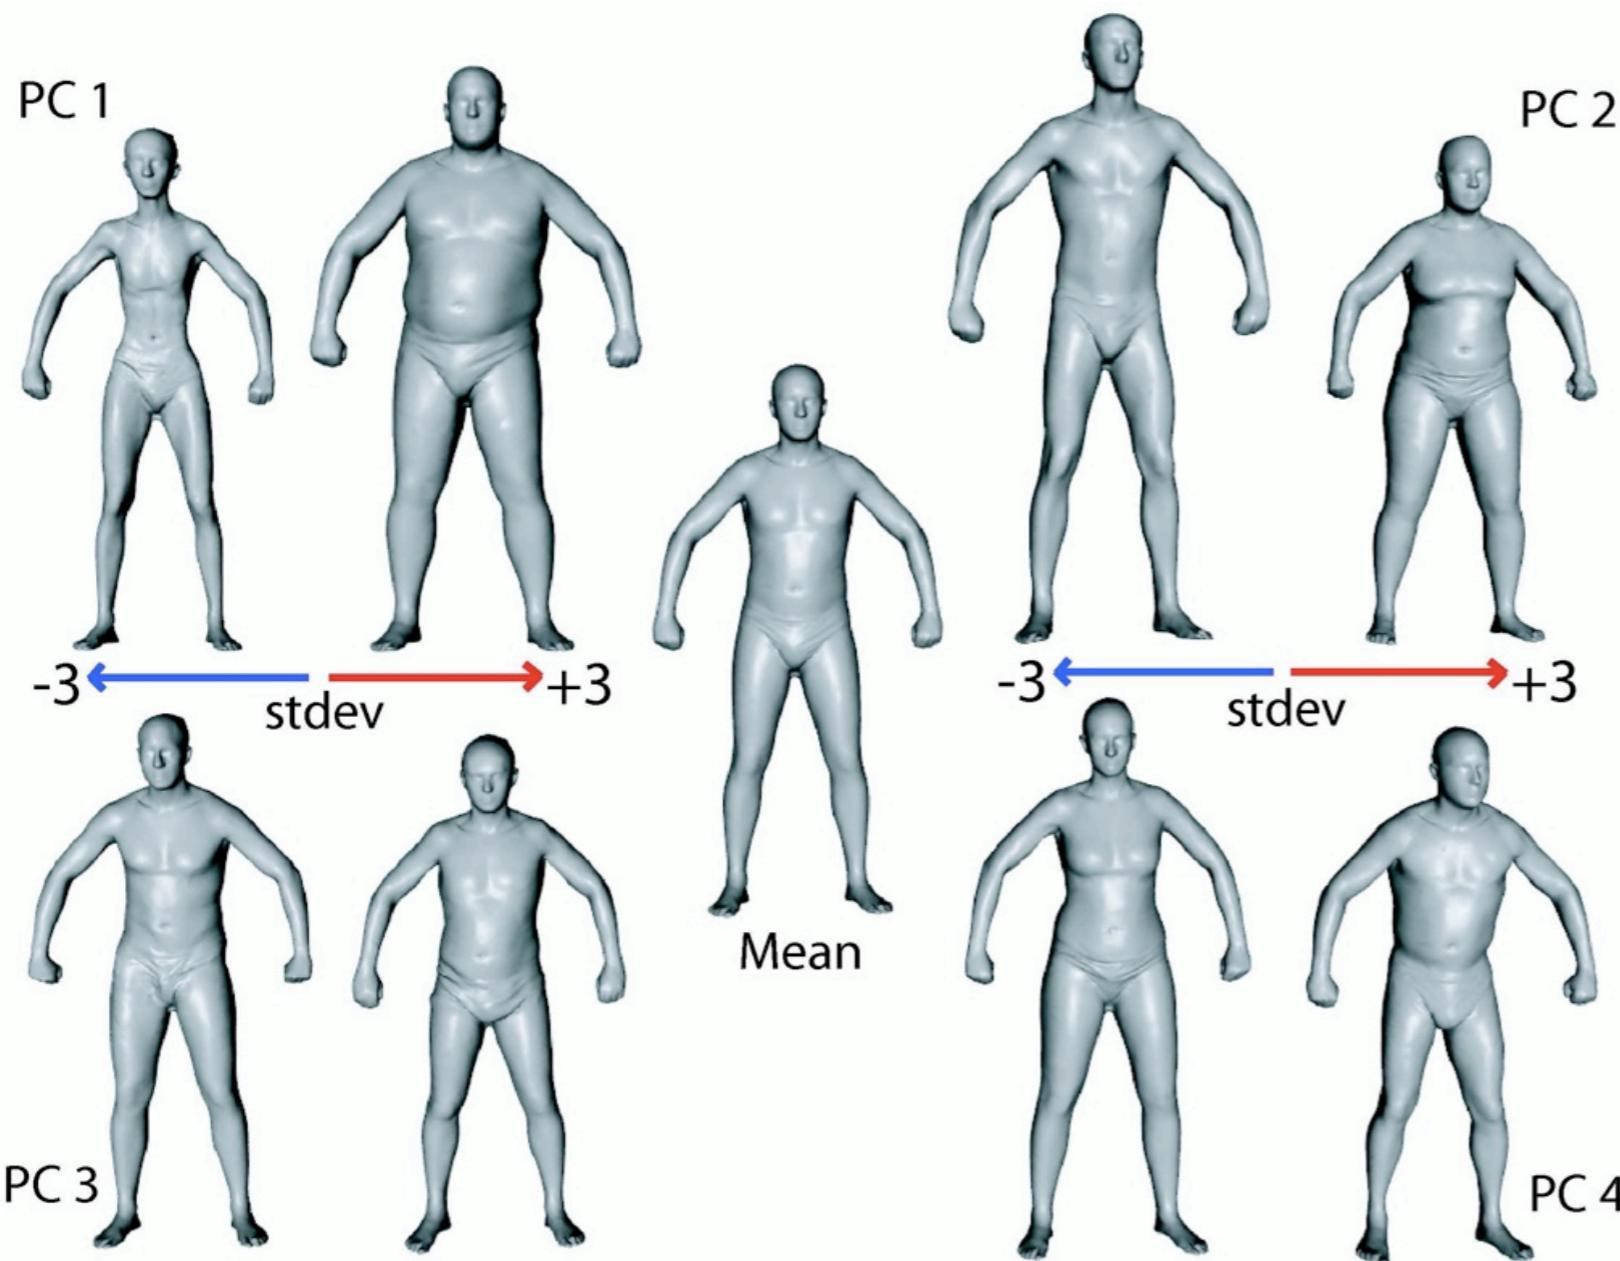
EigenLips



Fleet & Szeliski

# PCA Model of Body Shapes

- ◆ PCA on a detailed triangle model of human bodies:



[Anguelov et al. 05]



# Readings

- ◆ Edge detection, sec. 4.2
- ◆ PCA and face recognition, sec. 14.2
- ◆ For next time:
  - ◆ Feature points, sec. 4.1
  - ◆ Bag of word models, sec. 14.4.1



# Announcements

- ◆ **No class next week!**
  - ◆ Project review meeting
- ◆ Homework assignment 1:
  - ◆ Is out...
  - ◆ Signup for the "Testat" current read only -- will go online again soon
- ◆ **Exam**
  - ◆ 21. August 2013, 10:00 - 12:00