

# Computer Vision I

Bags of Visual Words - 05.06.2013



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

with slides from:

Bernt Schiele

Li Fei-Fei

Rick Szeliski

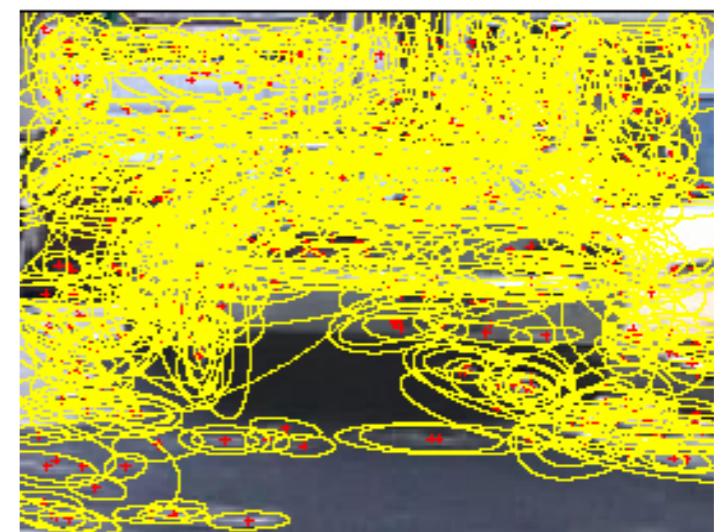
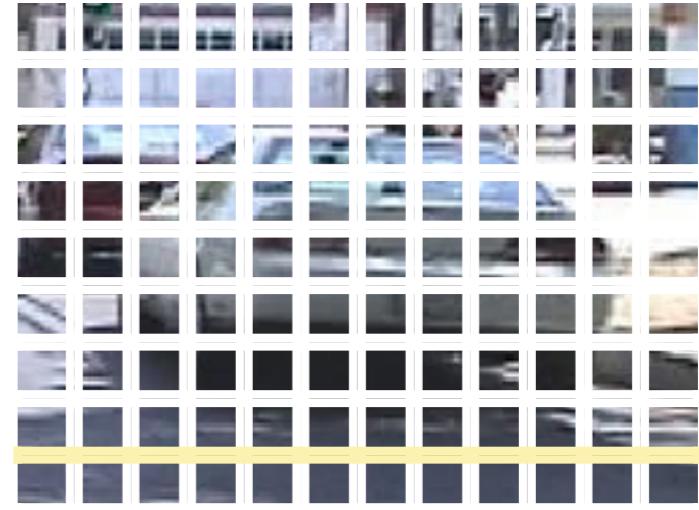
Steve Seitz

Thomas Hofmann



# Feature representation

- ◆ Dense representation (regular grid):
  - ◆ Color histogram approach [Swain & Ballard '91]
  - ◆ Multidimensional receptive field histograms [Schiele & Crowley '96-'00]
  
- ◆ Sparse feature representations:
  - ◆ Find key points with interest point detector
    - ◆ e.g. scale- or affine-invariant
  - ◆ Represent key points with feature vectors
    - ◆ e.g. SIFT, ShapeContext

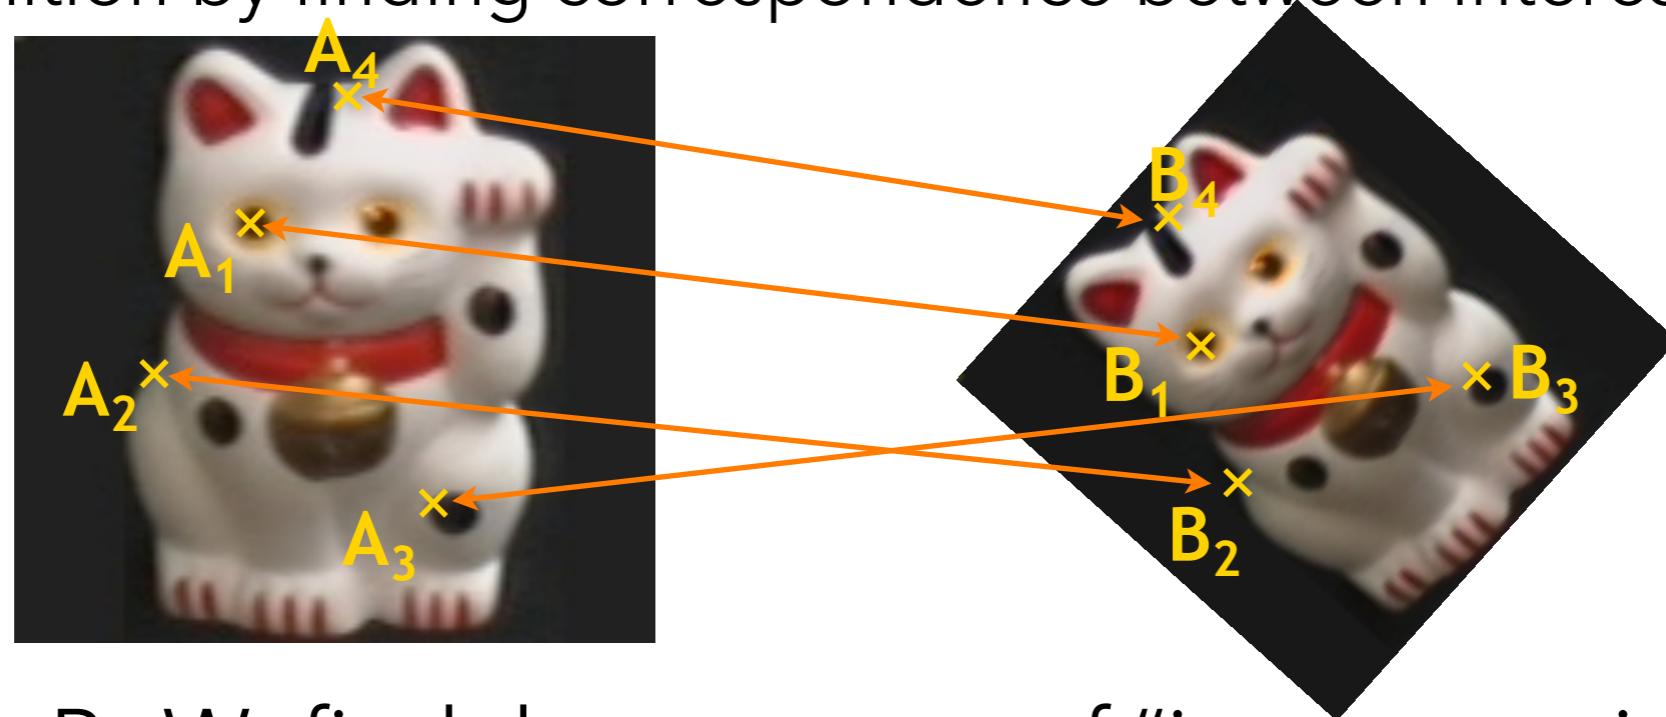


[Fei-Fei]

# Overview

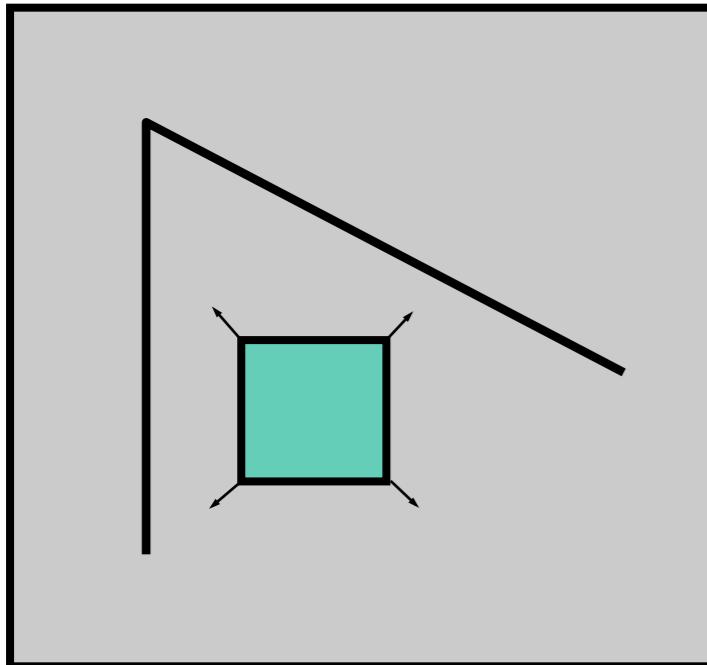
- ◆ **1 - Local Interest Point Detection**
  - ◆ Finding discriminative points (Harris, Hessian)
  - ◆ Scale invariant interest point detection (Harris-Laplace)
- ◆ 2 - Local Descriptors (= Features)
- ◆ 3 - Bag-of-Words Model (BoW)
  - ◆ BoW for Object Categorization

# Motivation for Local Interest Point Detection

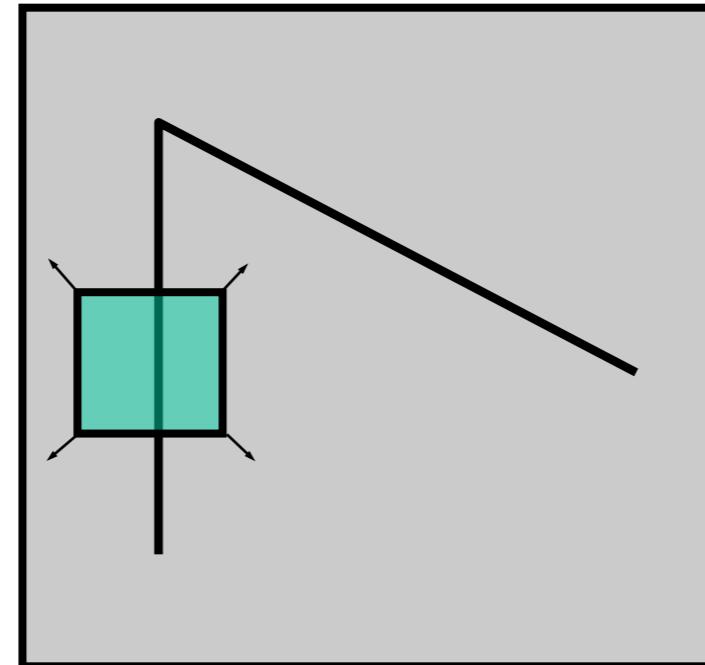
- ◆ General idea of interest point detection:
  - ◆ Recognition by finding correspondence between interest points
- ◆ Goal for BoW: find the same set of “interest points” that are
  - ◆ discriminative for object recognition
  - ◆ can be always found - even in the presence of arbitrary geometric and photometric transformations

# Feature detection

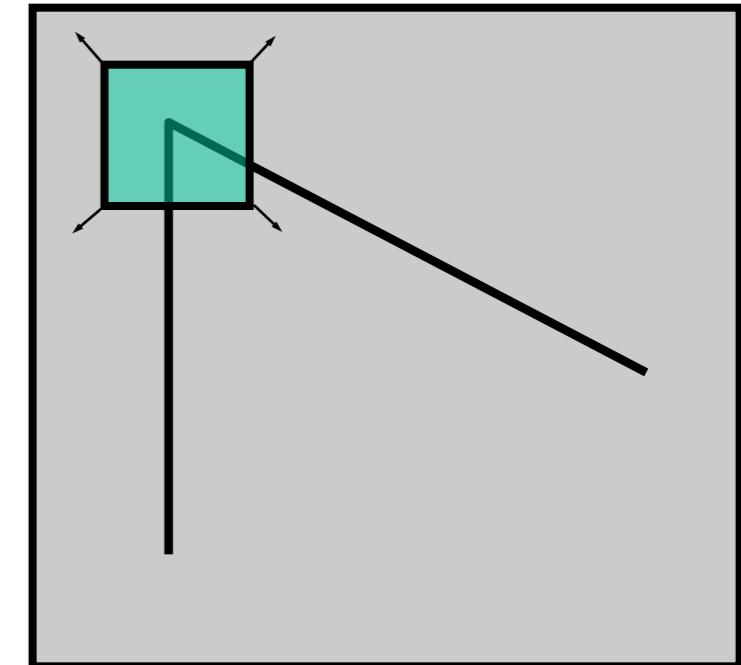
- ◆ Local measure of **feature uniqueness**
  - ◆ How does the window change when you shift it?
  - ◆ Shifting the window in any direction causes a big change



**"flat"** region:  
no change in all  
directions



**"edge"**:  
no change along the  
edge direction

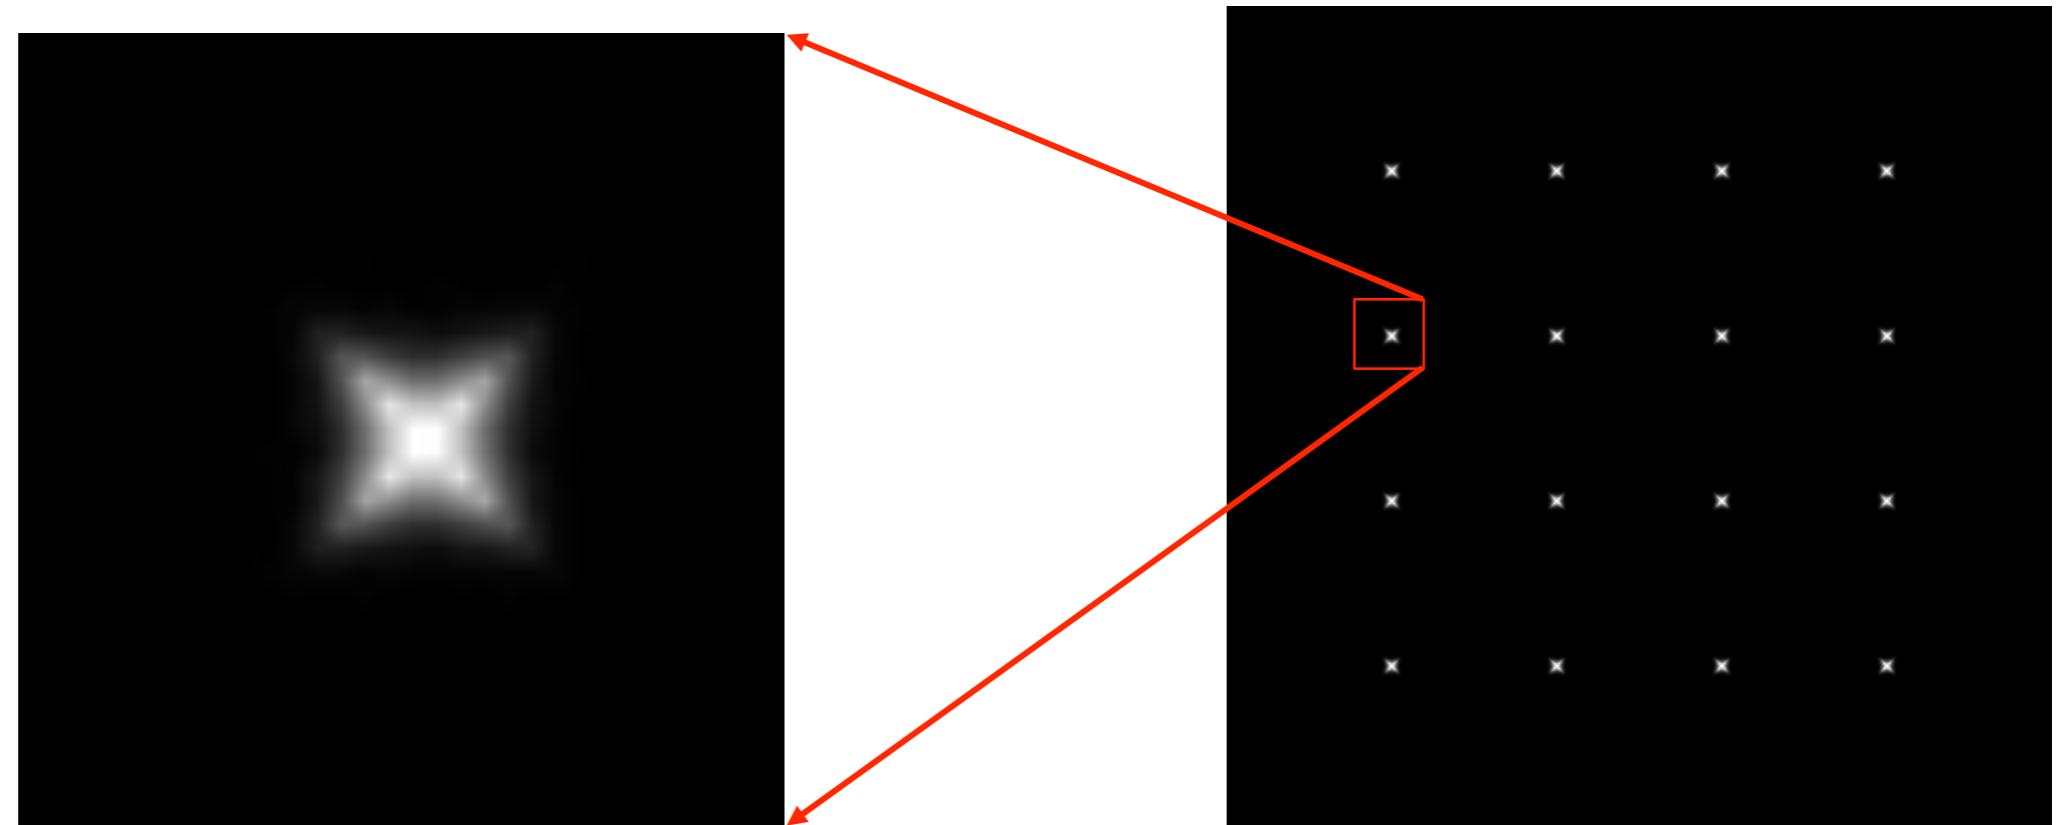


**"corner"**:  
significant change in  
all directions

[Szeliski & Seitz]

# Feature detection summary

- ◆ Here's what you do
  - ◆ Compute the gradient at each point in the image
  - ◆ Create the H matrix from the entries in the gradient
  - ◆ Compute the eigenvalues.
  - ◆ Find points with large response ( $\lambda_>$  threshold)
  - ◆ Choose those points where  $\lambda_>$  is a local maximum as features



# The Harris operator

- ◆  $\lambda_-$  is a variant of the “Harris operator” for feature detection

$$\begin{aligned} f &= \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2 \\ &= \det(H) - \alpha \cdot \text{trace}(A)^2 \end{aligned}$$

- ◆ The trace is the sum of the diagonals, i.e.,  $\text{trace}(H) = h_{11} + h_{22}$
- ◆ Very similar to  $\lambda_-$  but less expensive (no square root)
- ◆ Called the “Harris Corner Detector” or “Harris Operator”
- ◆ Lots of other detectors, this is one of the most popular

[Szeliski & Seitz]

# Harris features (in red)



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT



[Szeliski & Seitz]

# Discussion

- ◆ Interest-point detection so far:
  - ◆ Based on Harris or Hessian detector
  - ◆ Finds interesting, i.e. discriminative points  
(Harris detector was the “de-facto” standard for a long time)
  - ◆ Used for recognition, correspondence for stereo, sparse optical flow/motion, etc.
- ◆ Remember the goals of interest point detection:
  - ◆ Distinctiveness vs. invariance to transformation
  - ◆ Harris & Hessian find distinctive points - but they are **not** invariant to scale, affine and projective transformations

# Geometric Transformations



- ◆ Example of different geometric transformations:
  - ◆ (1) original
  - ◆ (2) similarity transformation (translation, image plane rotation, scaling)
  - ◆ (3) projective transformation



(1)



(2)



(3)

- ◆ Scale-Invariant Interest Point Detection
  - ◆ Matching images of different scales
  - ◆ Automatic scale selection
  - ◆ Scale invariant methods for feature extraction
    - ◆ Example: Harris-Laplace



# Matching image patches

- ◆ of different scales



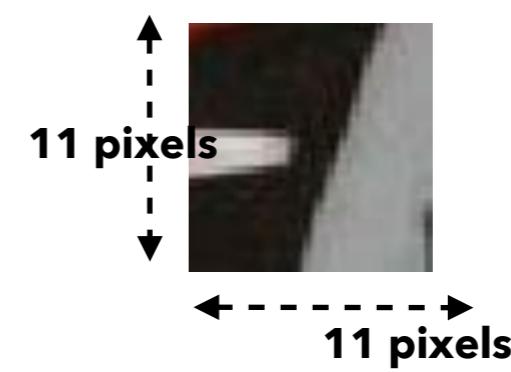
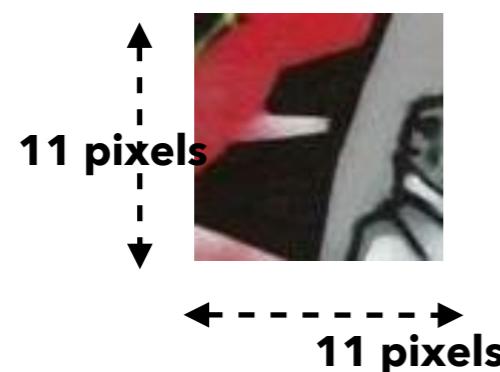
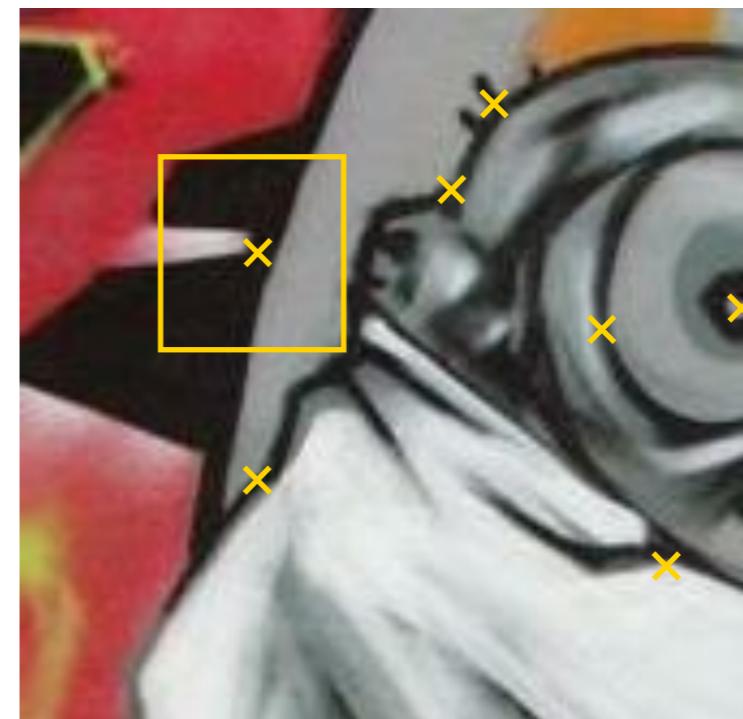
# Matching image patches

- ◆ Detecting interest points



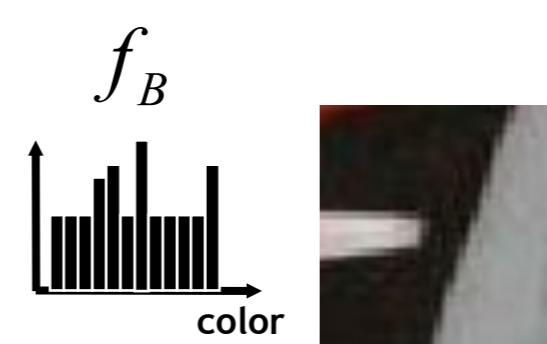
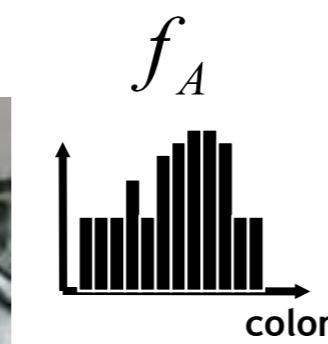
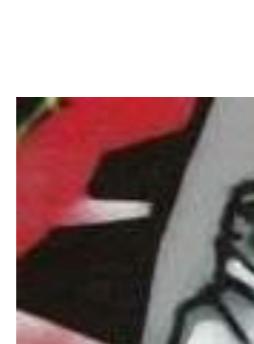
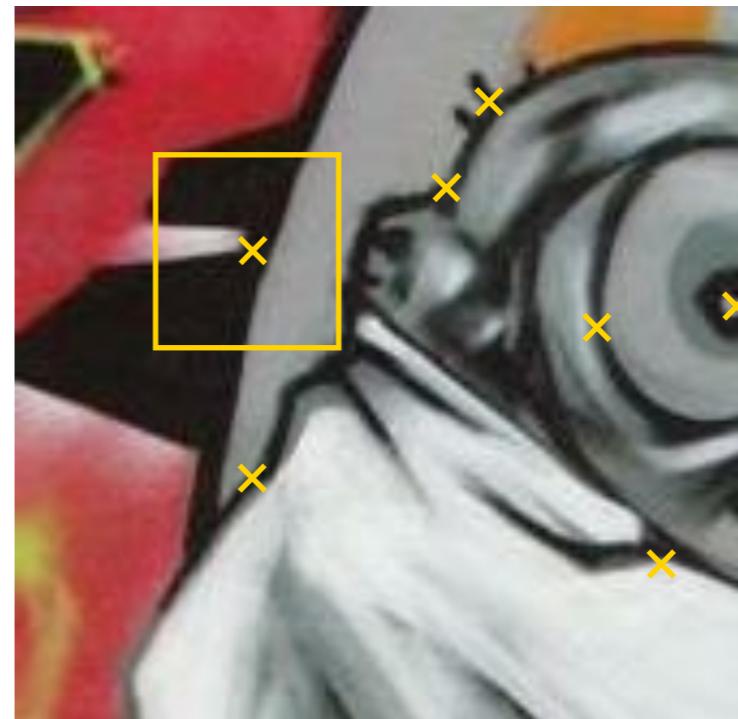
# Matching image patches

- ◆ Extracting patches



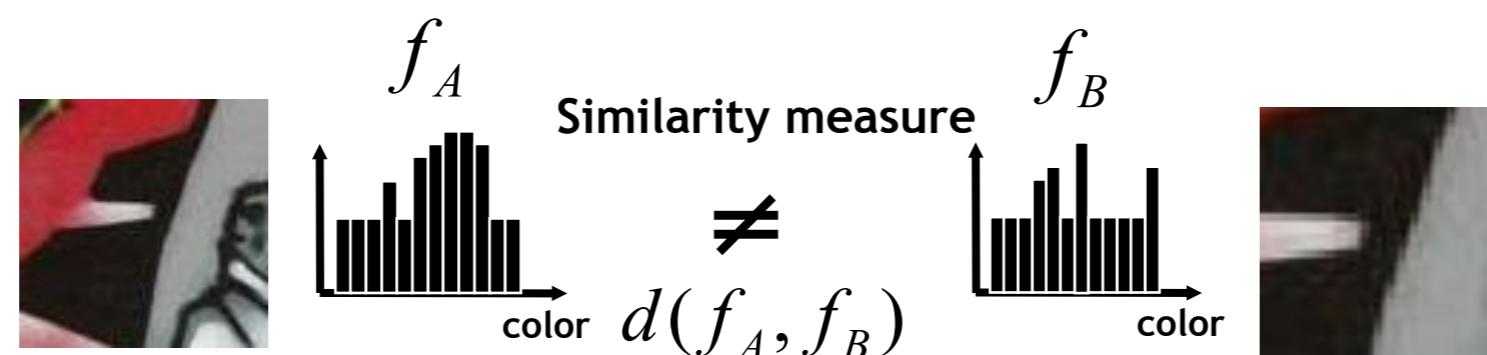
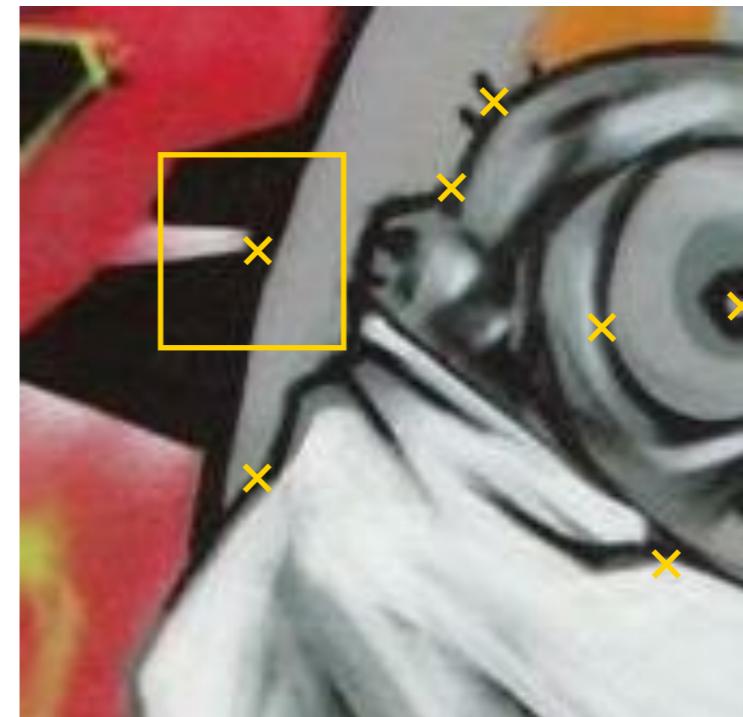
# Matching image patches

- ◆ Computing descriptors



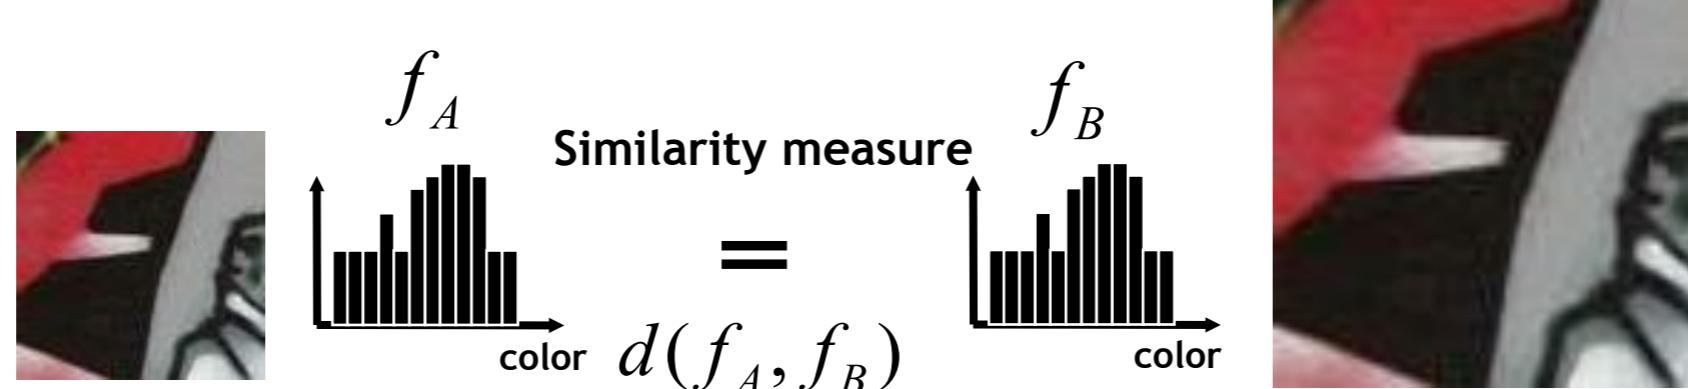
# Matching image patches

- ◆ Comparing descriptors
  - ◆ Impossible to match - different histograms due to different patch content



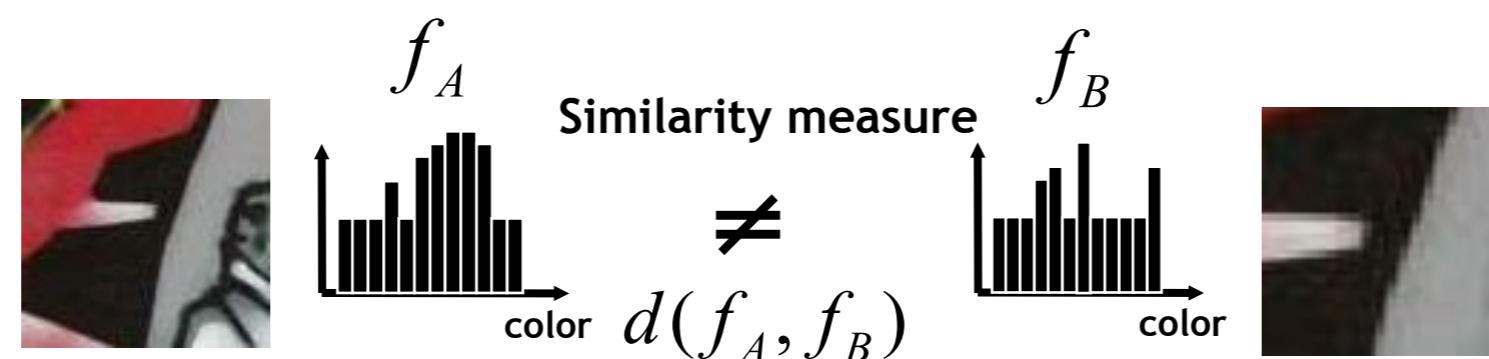
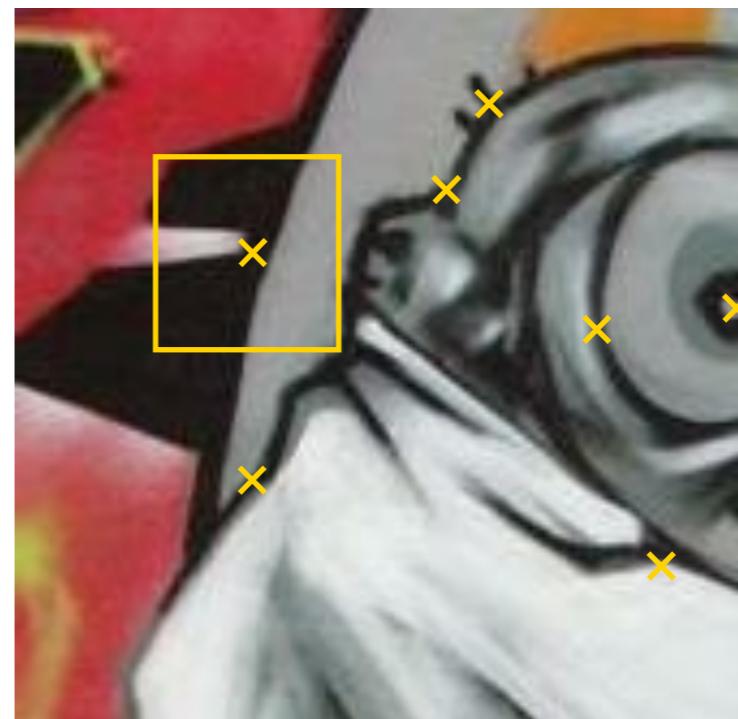
# Matching image patches

- ◆ Comparing descriptors
  - ◆ The patch should contain the same image - how to find the correct size?



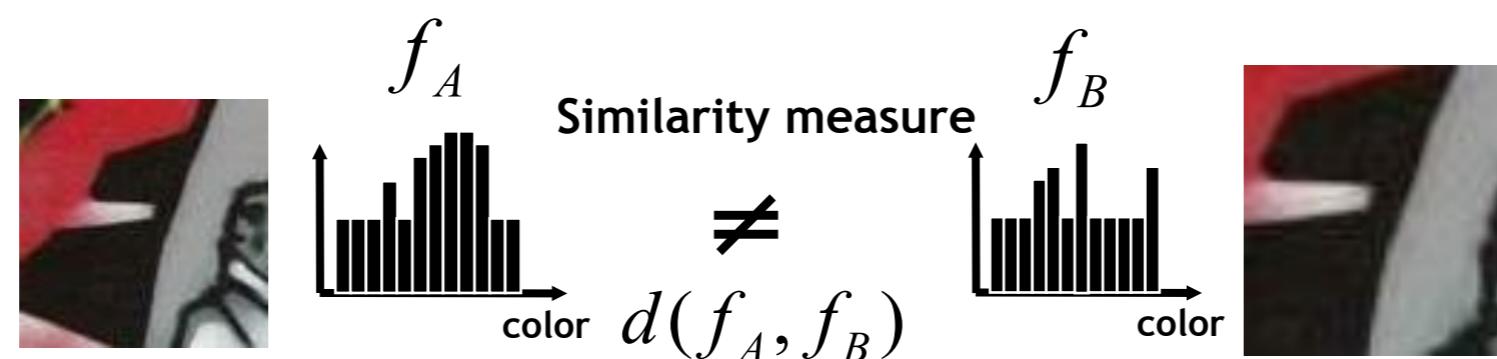
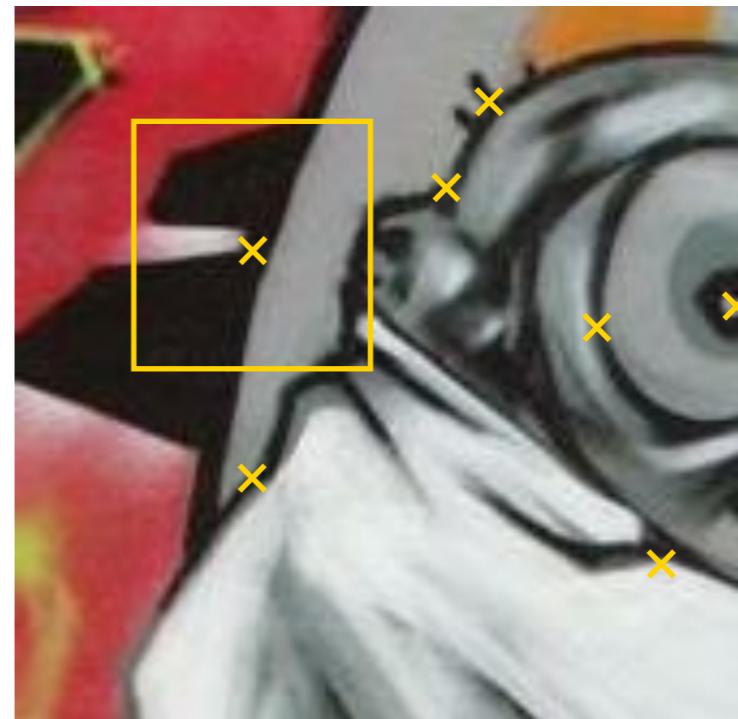
# Matching image patches

- ◆ Comparing descriptors while varying the patch size



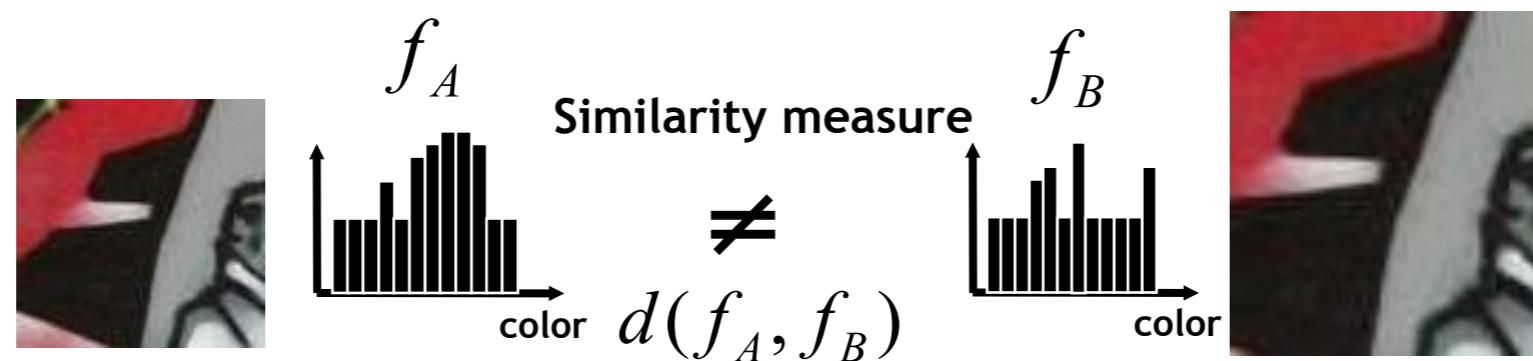
# Matching image patches

- ◆ Comparing descriptors while varying the patch size



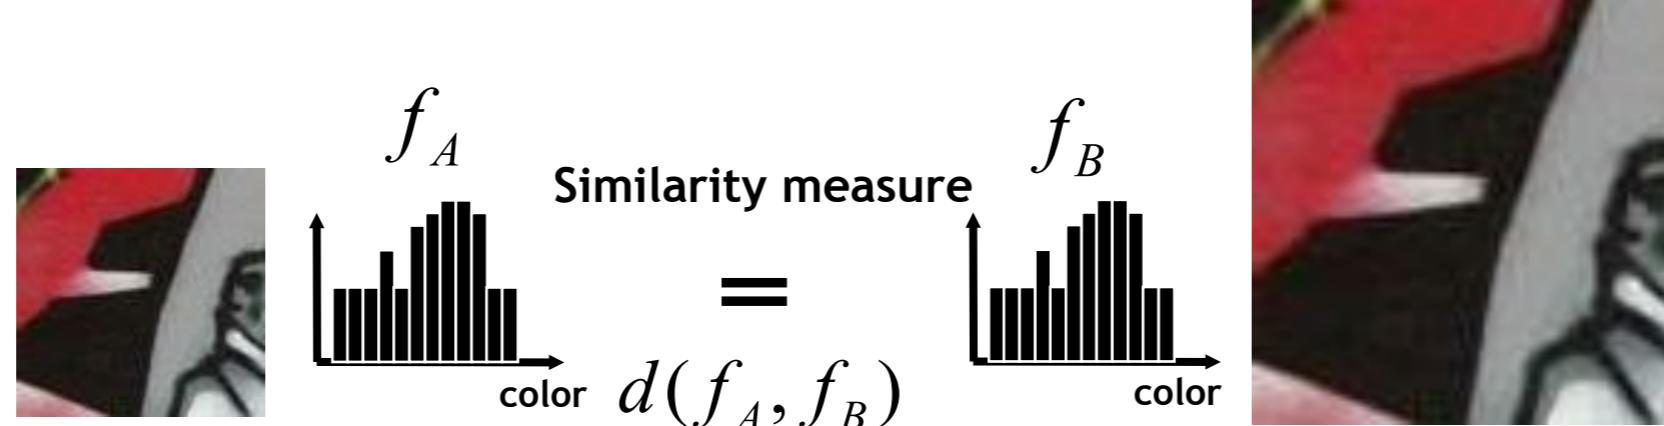
# Matching image patches

- ◆ Comparing descriptors while varying the patch size



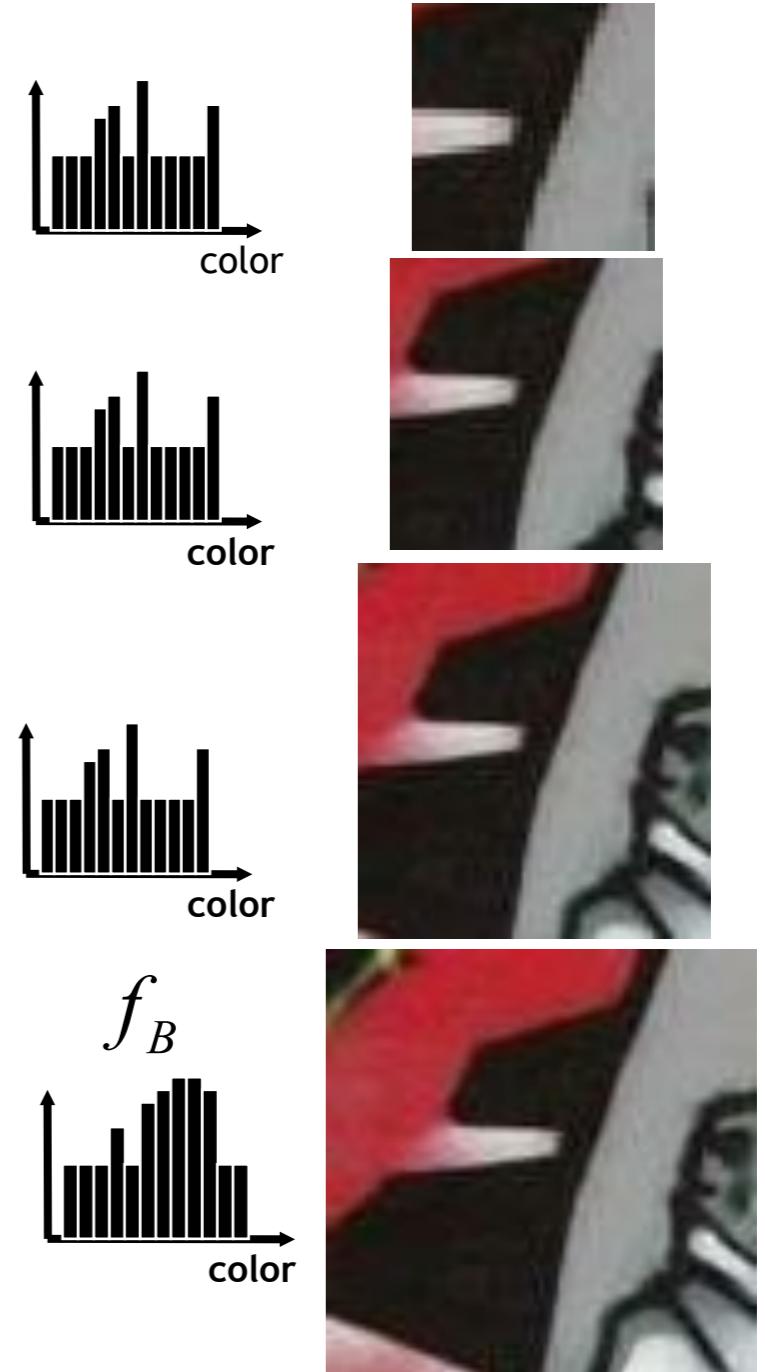
# Matching image patches

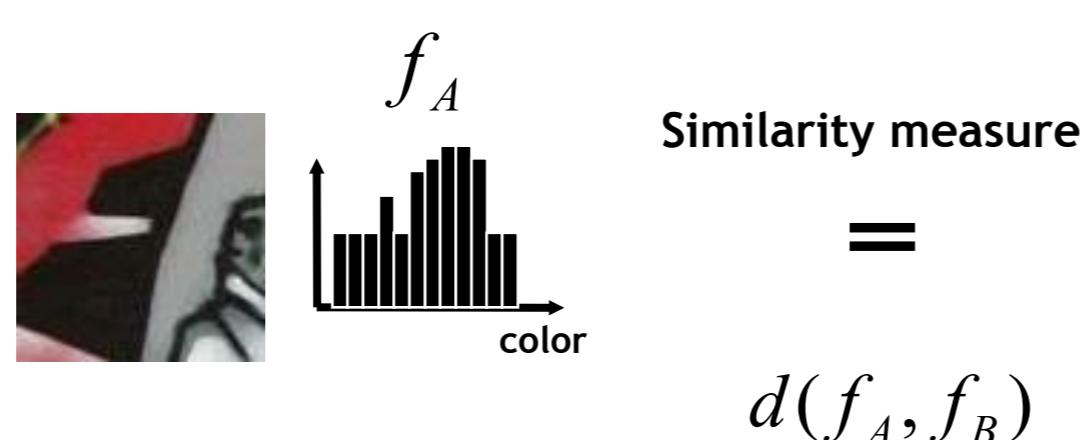
- ◆ Comparing descriptors while varying the patch size



# Matching image patches

- ◆ Comparing descriptors while varying the patch size
  - ◆ Computationally inefficient/prohibitive
  - ◆ Inefficient, but possible for matching
  - ◆ Prohibitive for retrieval in large databases
  - ◆ Prohibitive for recognition

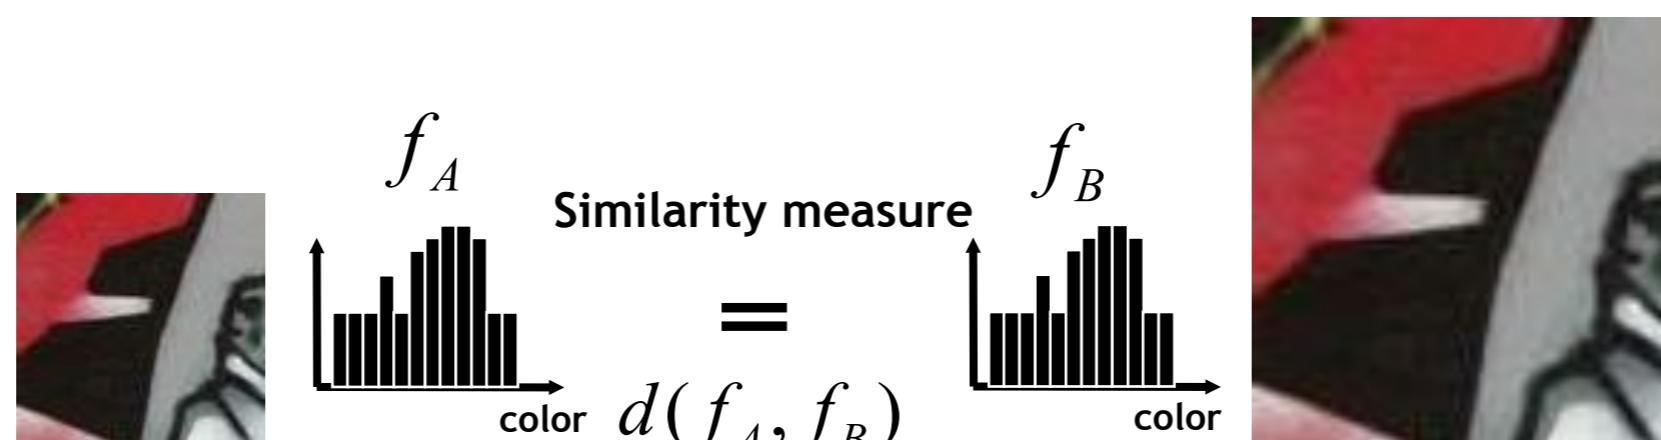




The diagram illustrates the similarity measure between two color histograms  $f_A$  and  $f_B$ . It shows two histograms side-by-side, each with an upward-pointing arrow and the word "color" at the bottom. Between the histograms is the text "Similarity measure =". Below the histograms is the formula  $d(f_A, f_B)$ .

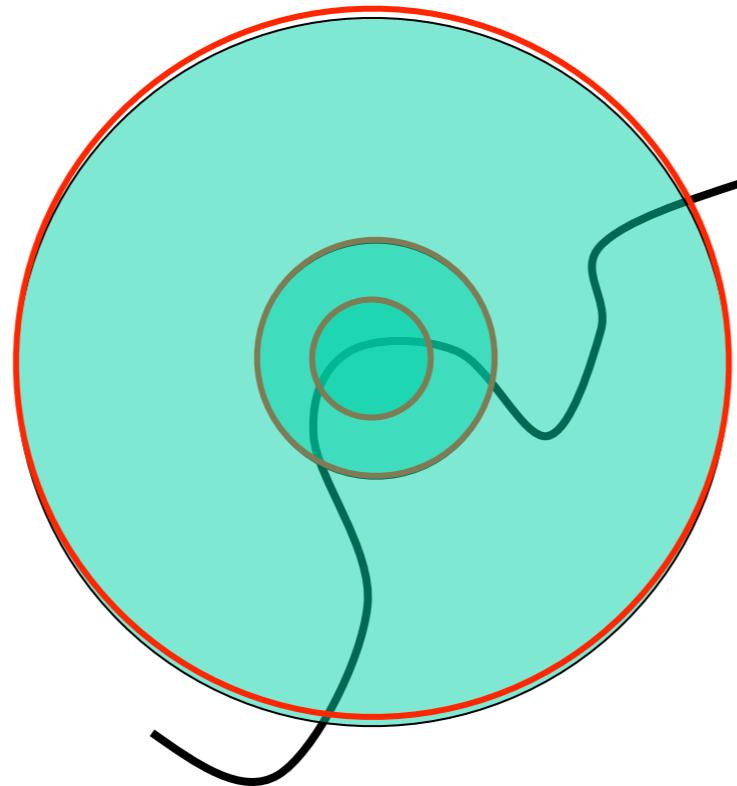
# Matching image patches

- ◆ Detector finds location and scale of interest points
  - ◆ In both images: independent automatic scale detection
  - ◆ by finding “characteristic” scale of an interest point



# Automatic scale selection

- ◆ Key idea:

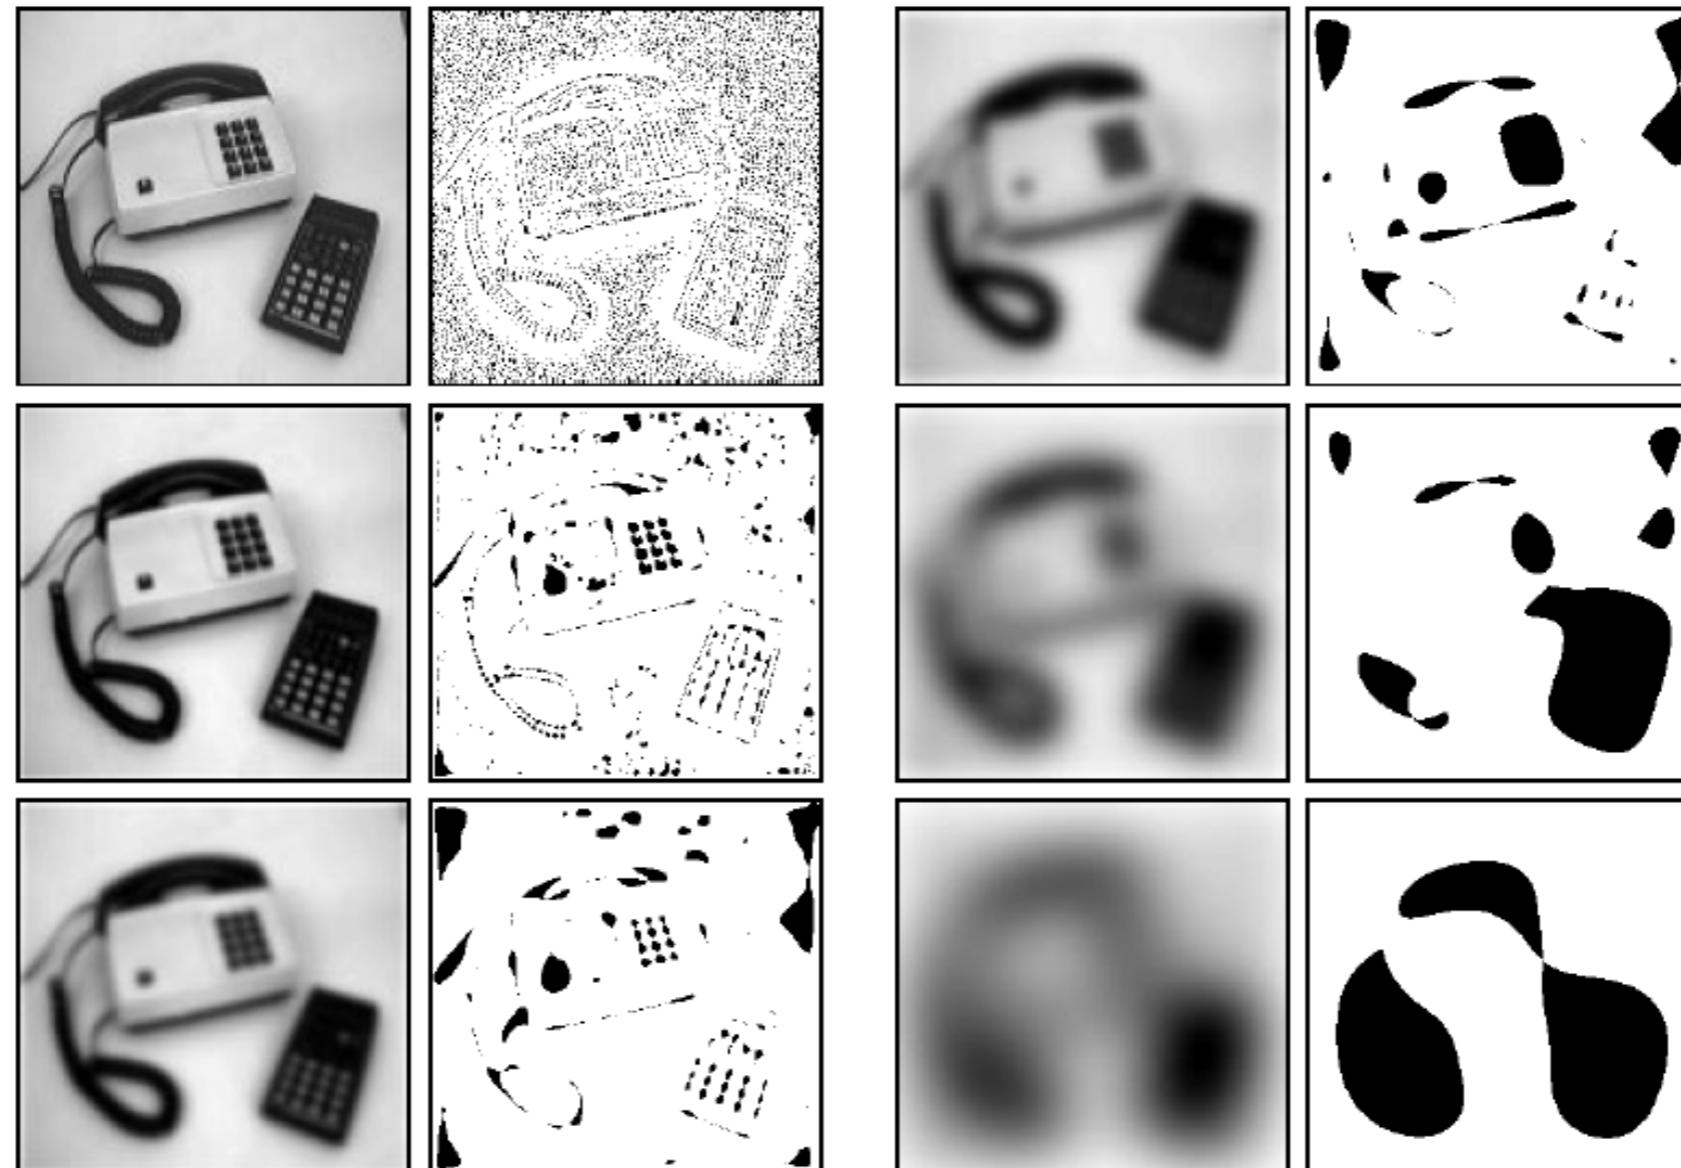


- ◆ Find scale that gives local maximum of some criterion.
- ◆ Which criterion?

[Seitz & Szeliski]

# Scale Space

- ◆ Level of details decreases monotonically
  - ◆ as the scale of Gaussian smoothing is increased



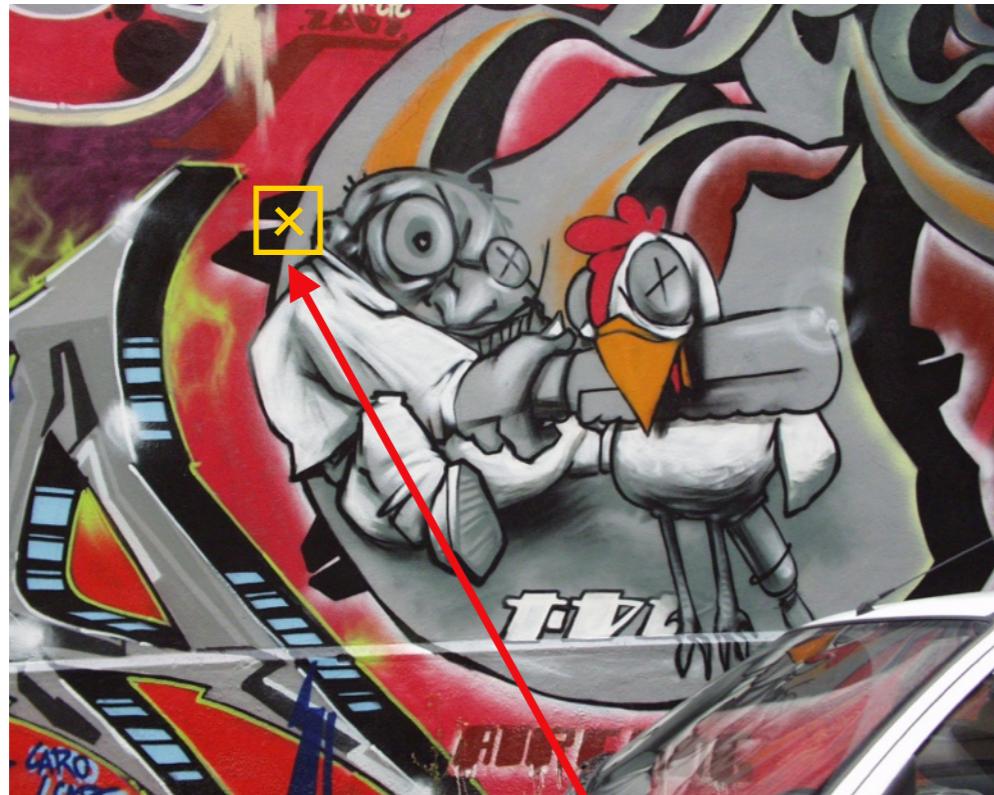
# Automatic scale selection



$$I_{r_1 \dots r_m}(x, \sigma) = I_{r_1 \dots r_m}(x', \sigma')$$

The **same derivative responses** if the patch contains the same image up to scale factor

# Automatic scale selection



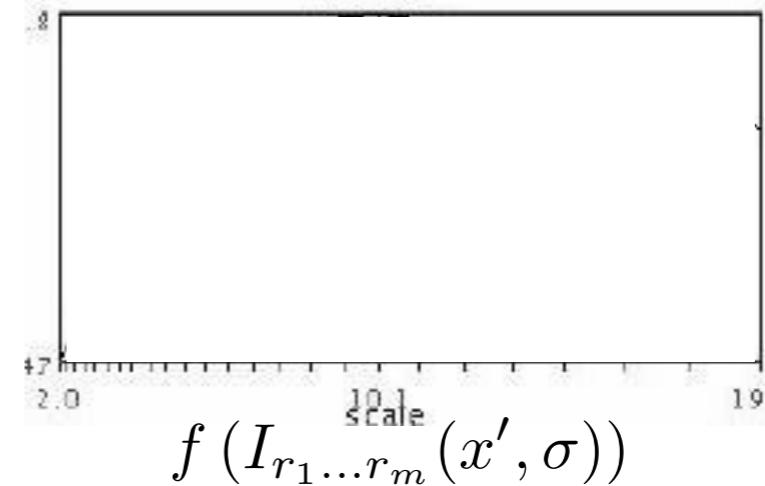
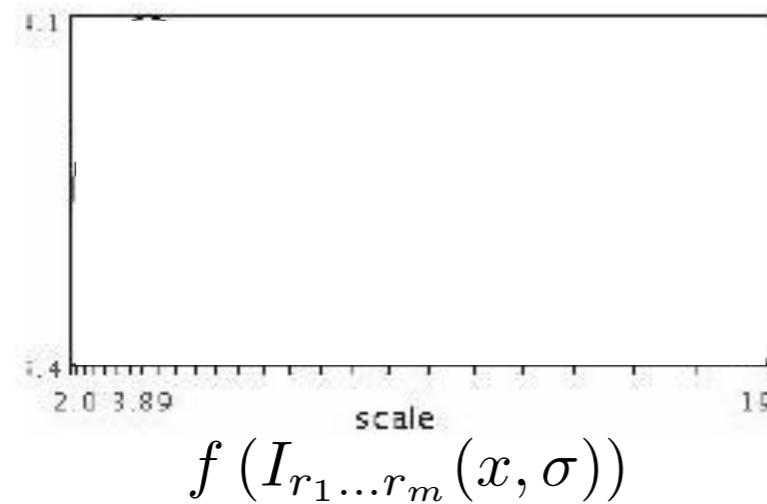
$$f(I_{r_1 \dots r_m}(x, \sigma)) = f(I_{r_1 \dots r_m}(x', \sigma'))$$

The **same operator responses** if the patch contains the same image up to scale factor

How to find corresponding patch sizes?

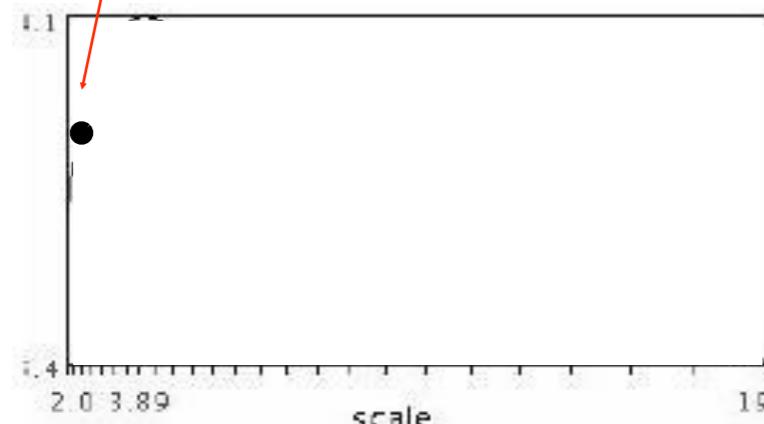
# Automatic scale selection

Function responses for increasing scale  
Scale trace (signature)

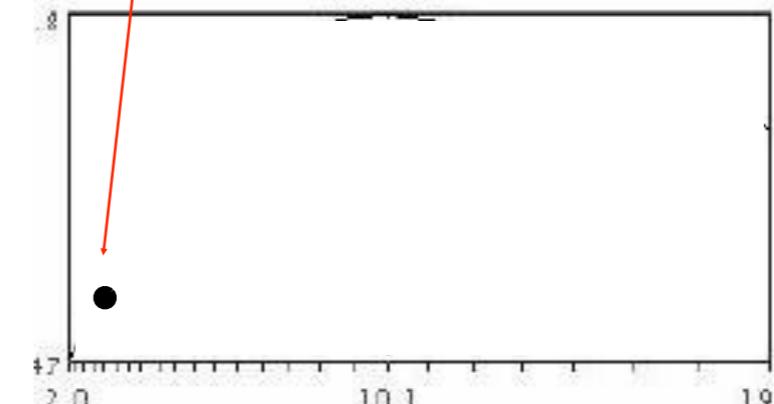


# Automatic scale selection

Function responses for increasing scale  
Scale trace (signature)



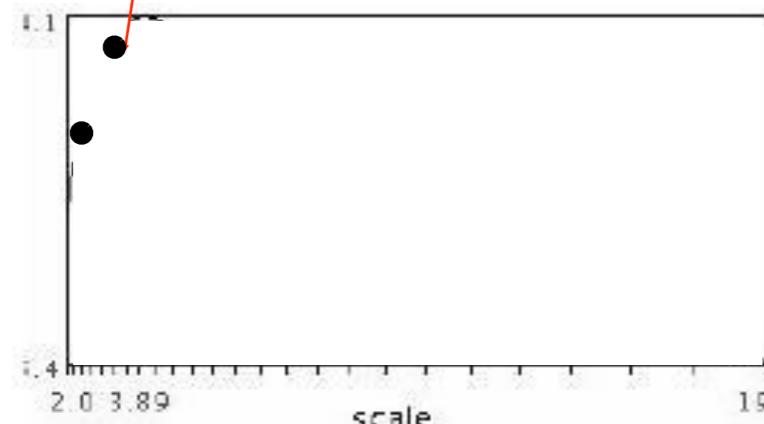
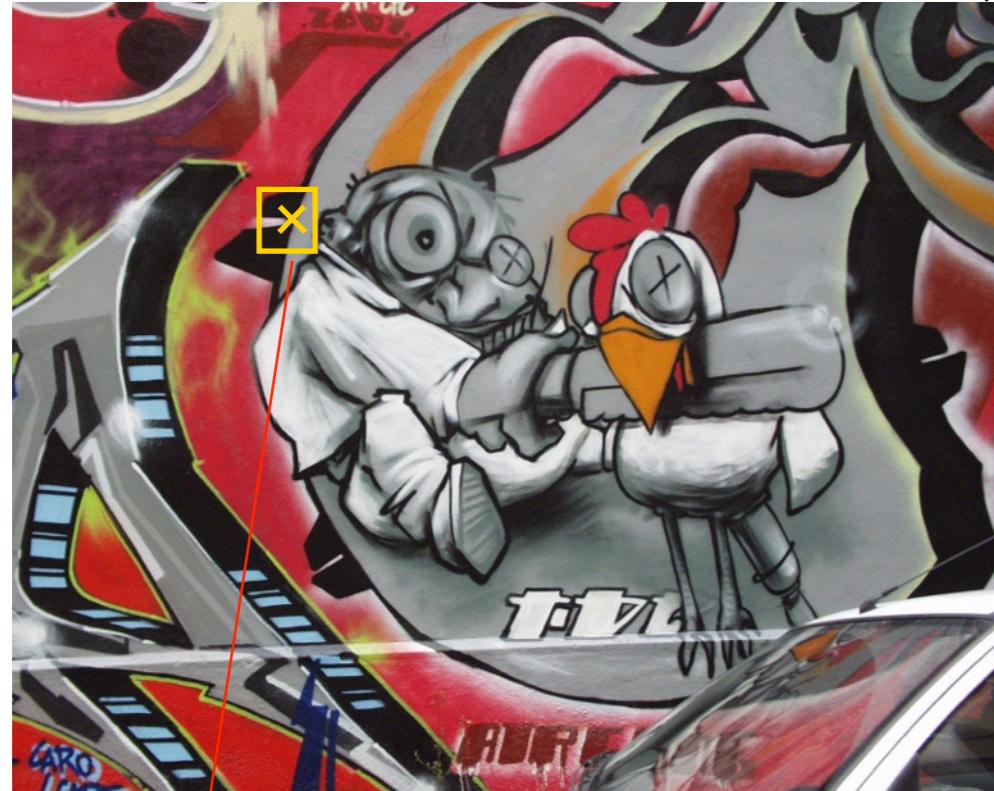
$$f(I_{r_1 \dots r_m}(x, \sigma))$$



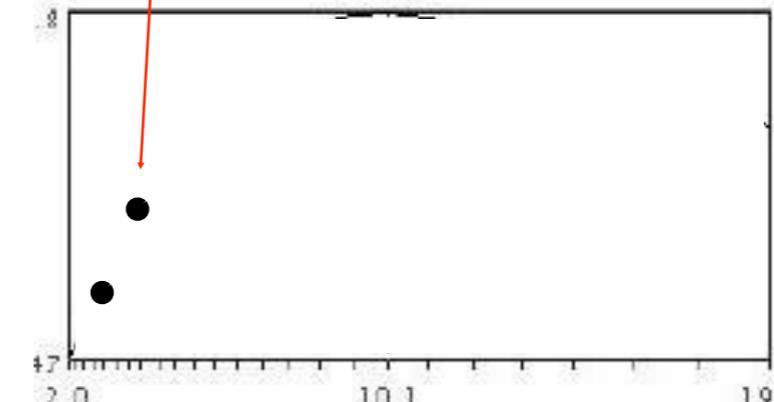
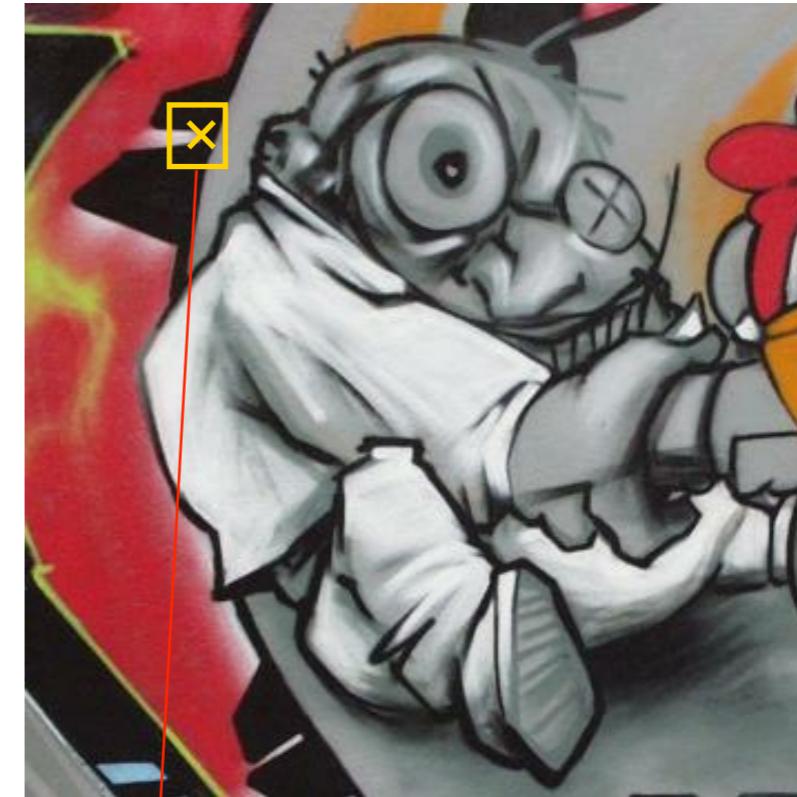
$$f(I_{r_1 \dots r_m}(x', \sigma))$$

# Automatic scale selection

Function responses for increasing scale  
Scale trace (signature)



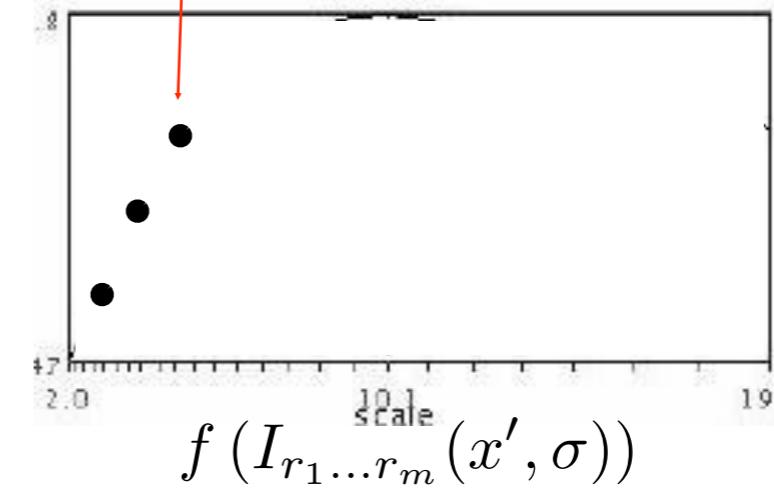
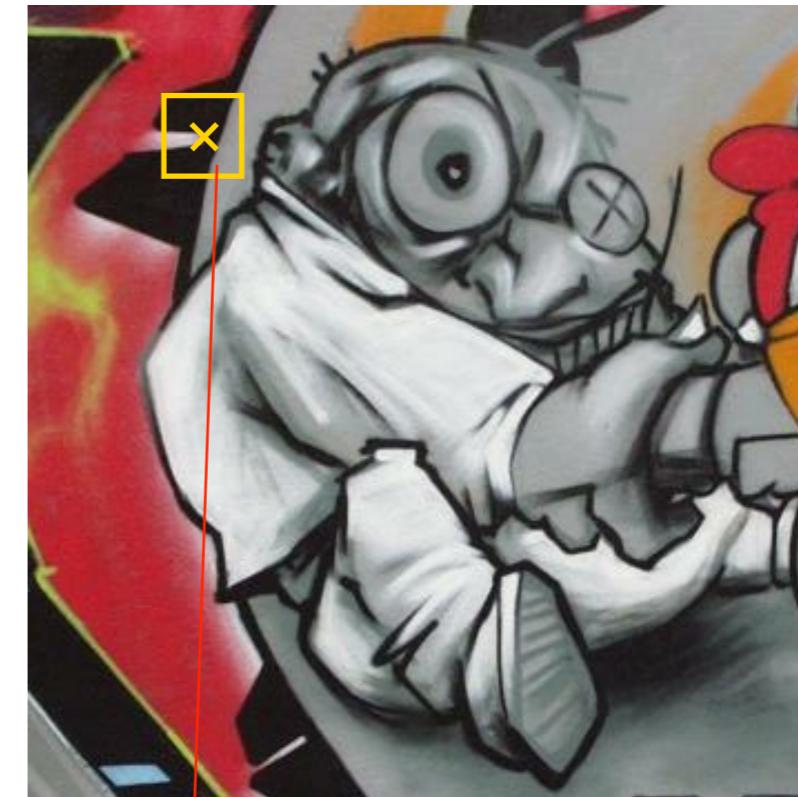
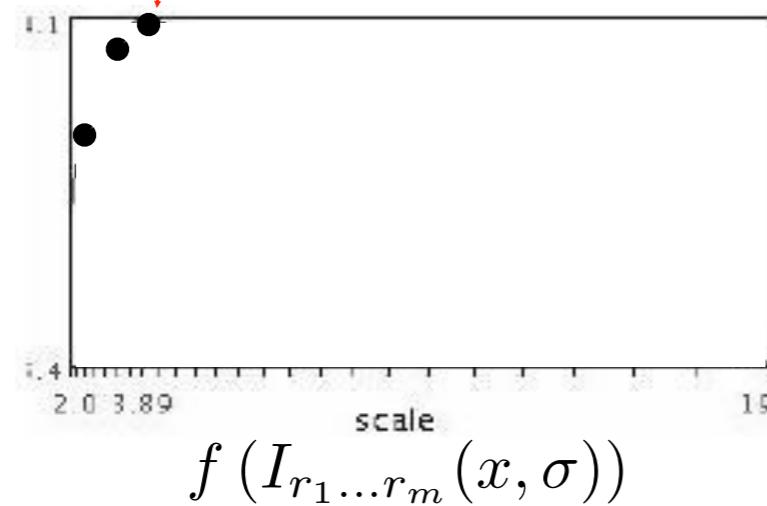
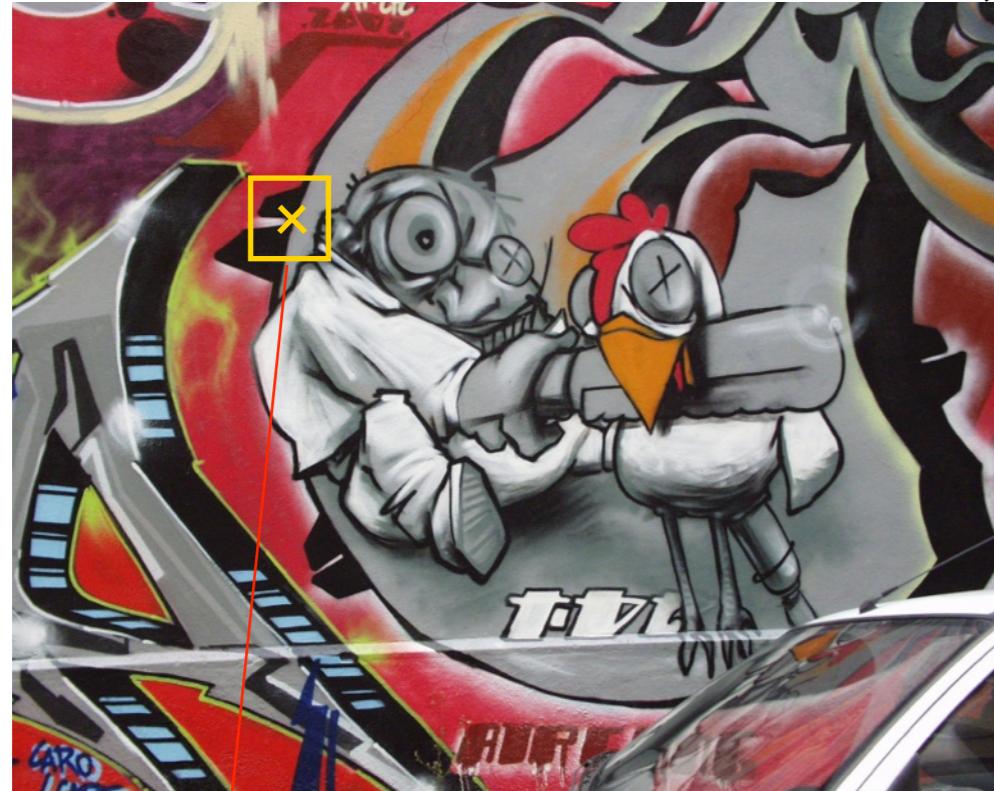
$$f(I_{r_1 \dots r_m}(x, \sigma))$$



$$f(I_{r_1 \dots r_m}(x', \sigma))$$

# Automatic scale selection

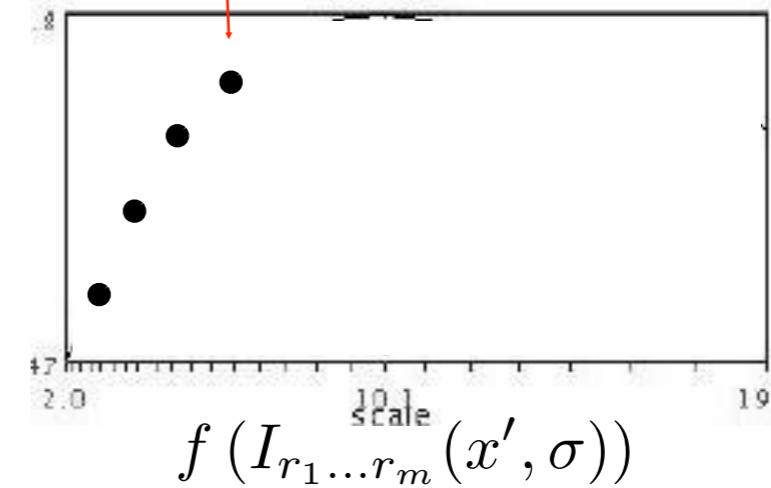
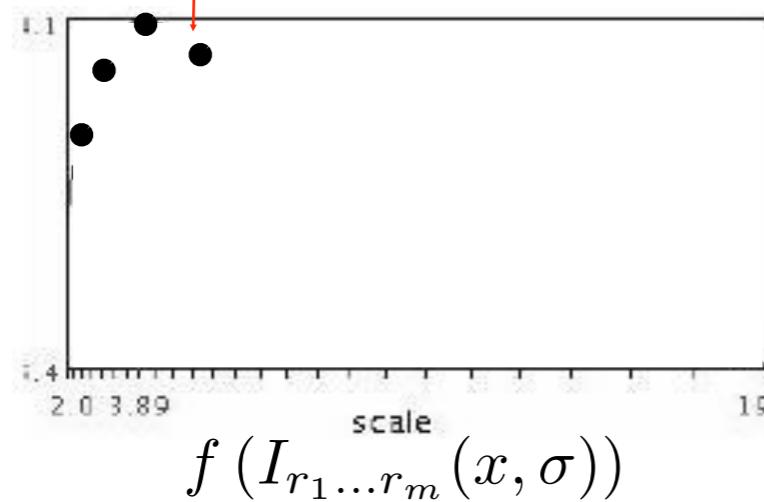
Function responses for increasing scale  
Scale trace (signature)



# Automatic scale selection

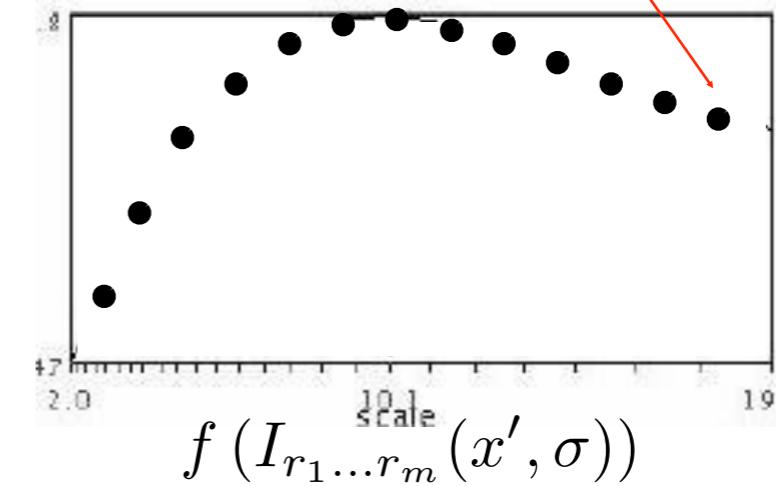
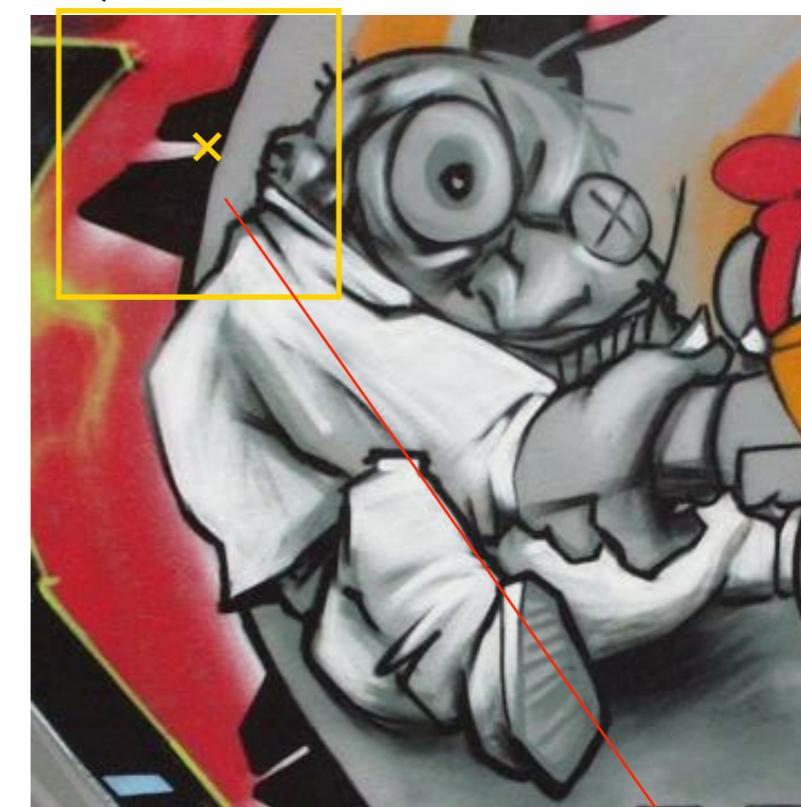
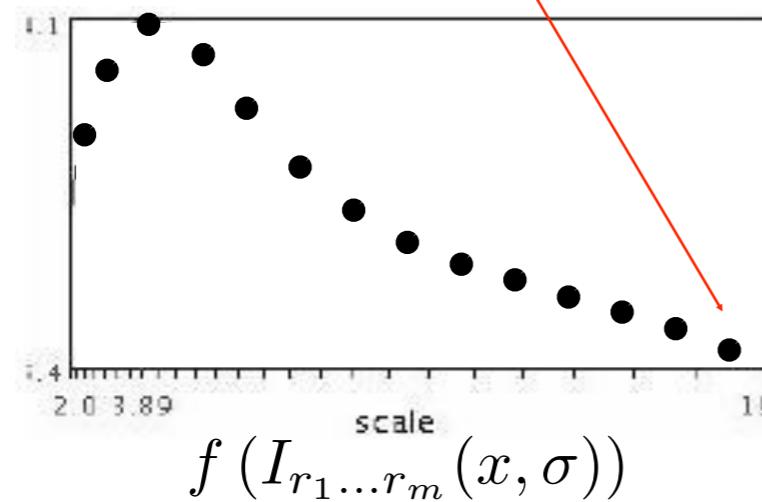


Function responses for increasing scale  
Scale trace (signature)



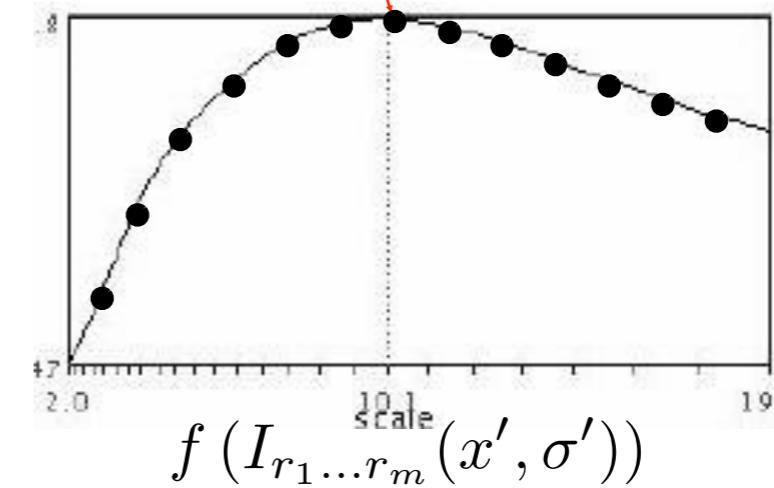
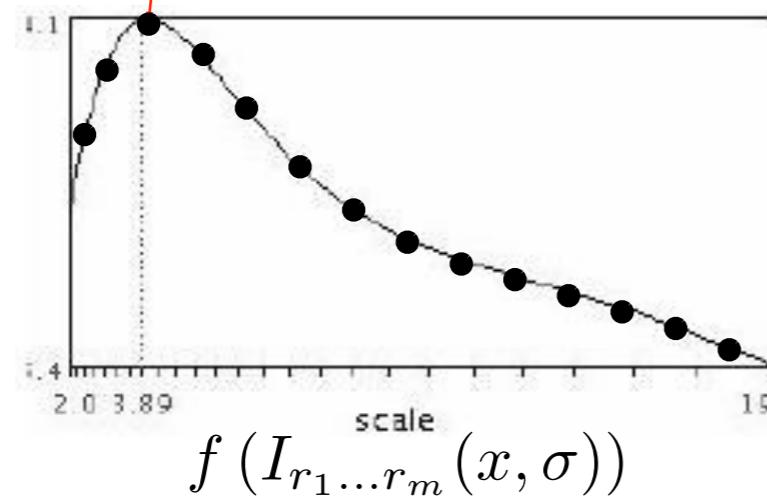
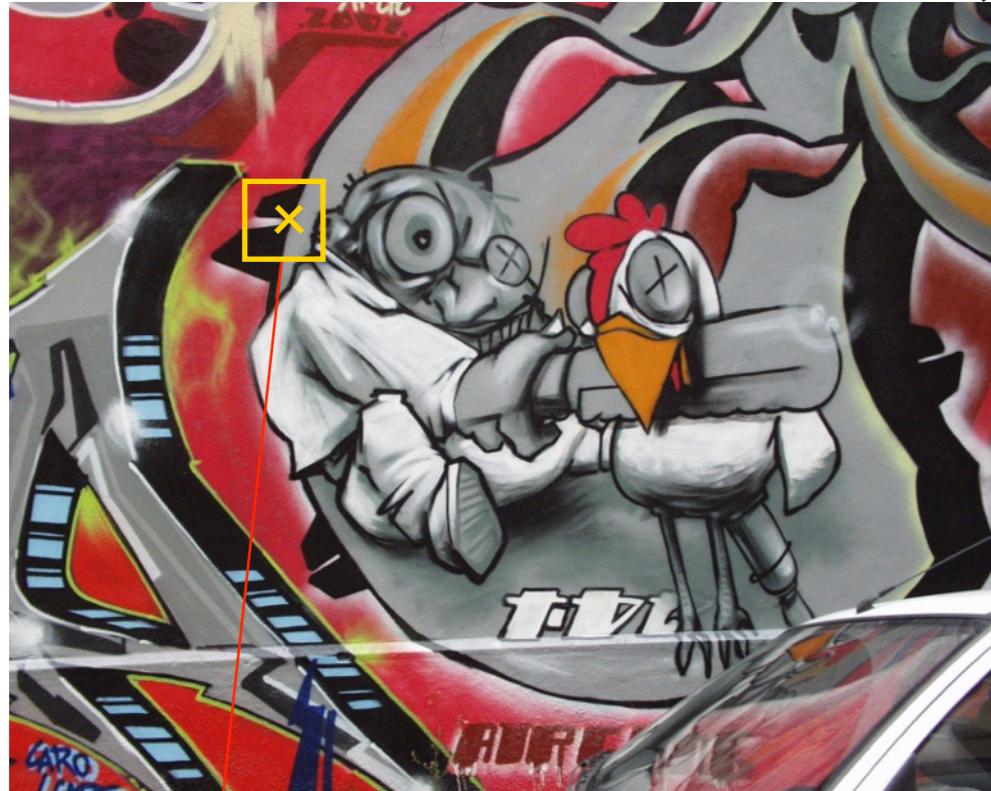
# Automatic scale selection

Function responses for increasing scale  
Scale trace (signature)



# Automatic scale selection

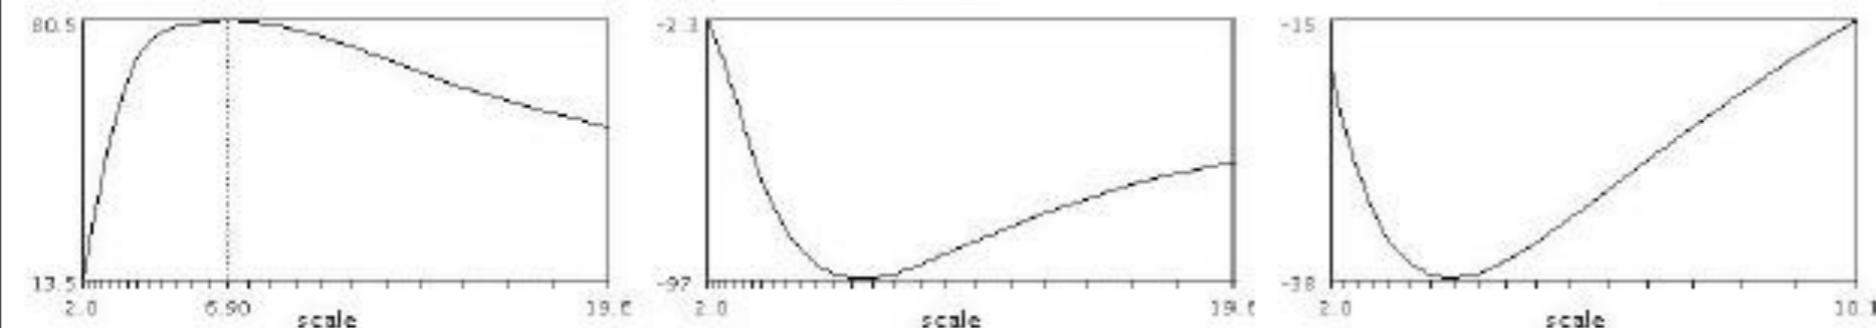
Function responses for increasing scale  
Scale trace (signature)



# Automatic scale selection



Laplacian =  $\text{trace}(\mathcal{H})$

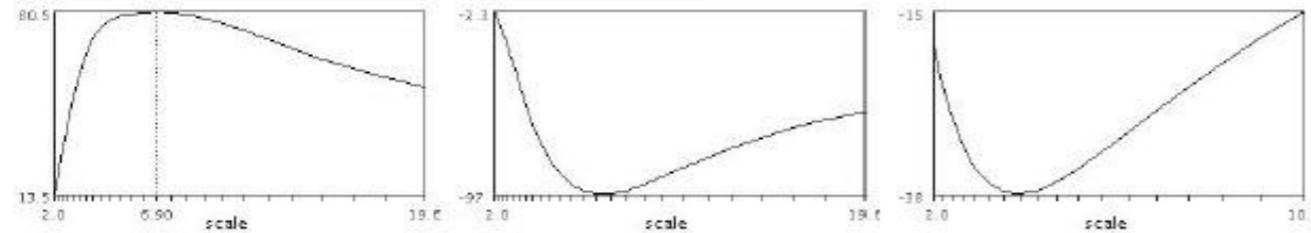


$$I_{xx} + I_{yy}$$

# Automatic scale selection

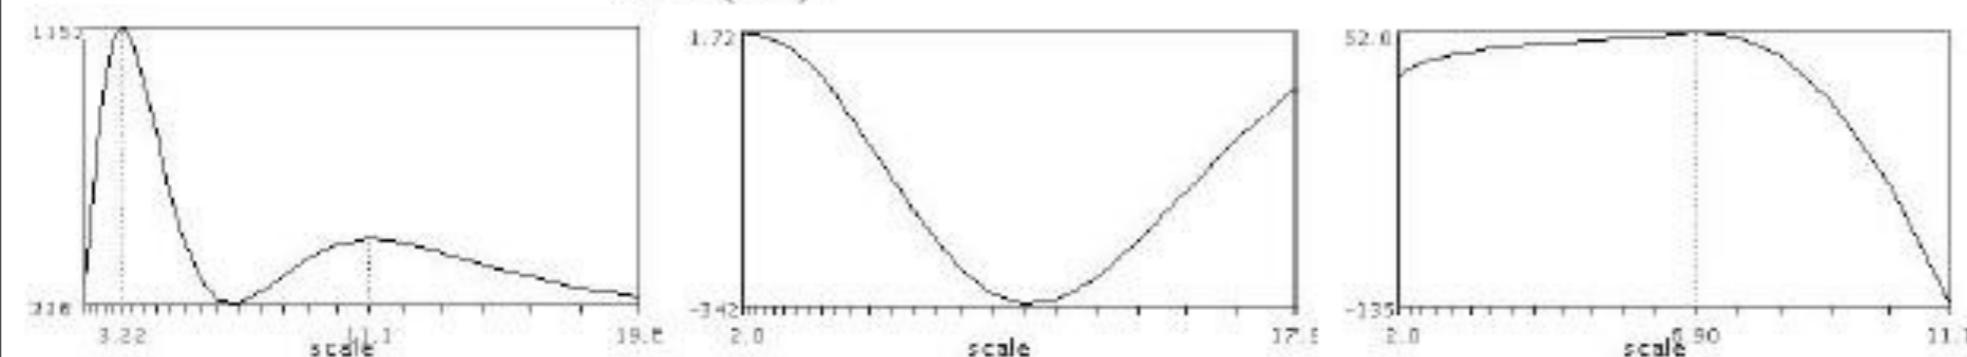


$\text{Laplacian} = \text{trace}(\mathcal{H})$



$$I_{xx} + I_{yy}$$

$\det(\mathcal{H})$

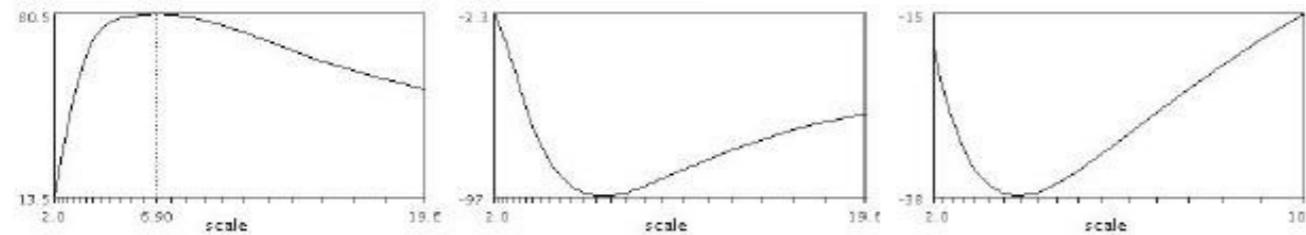


$$I_{xx}I_{yy} - I_{xy}^2$$

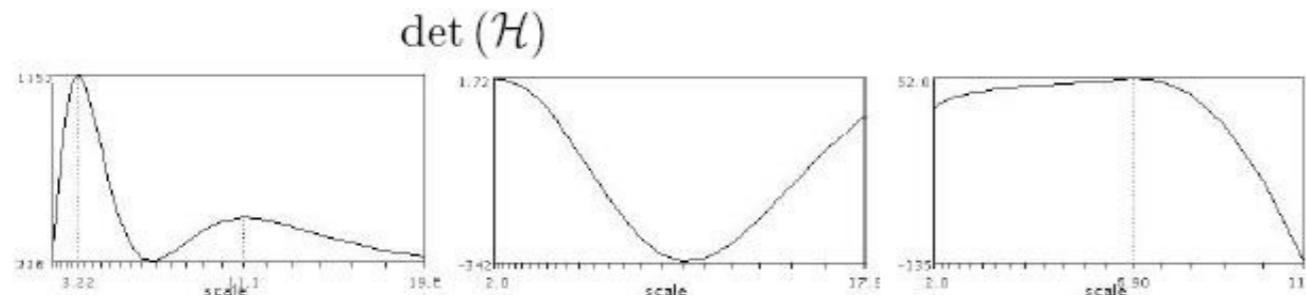
# Automatic scale selection



Laplacian =  $\text{trace}(\mathcal{H})$

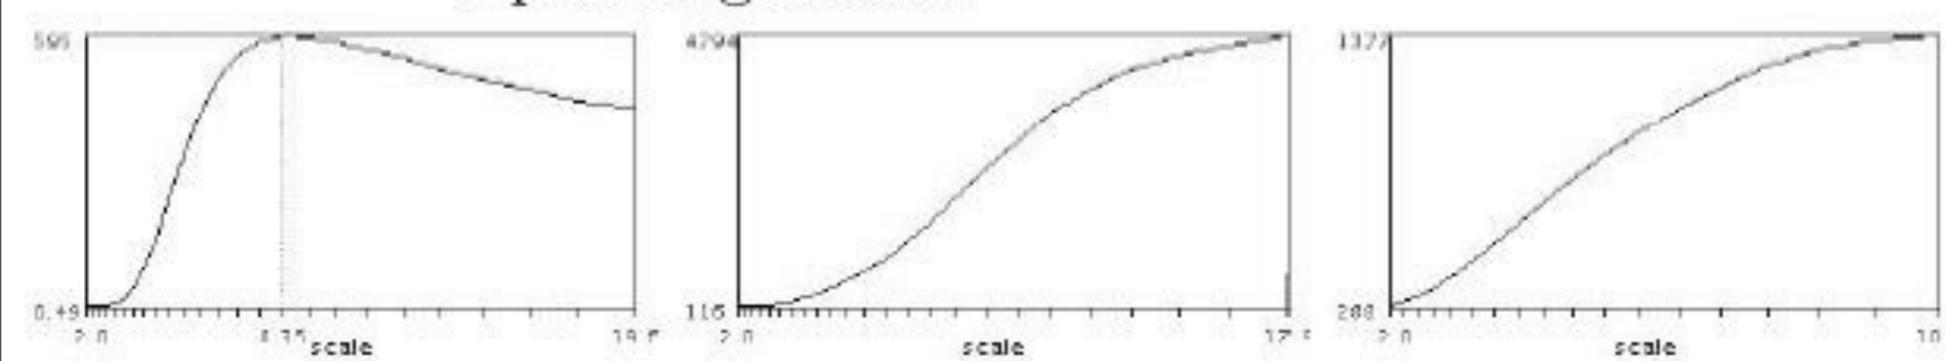


$$I_{xx} + I_{yy}$$



$$I_{xx}I_{yy} - I_{xy}^2$$

Squared gradient

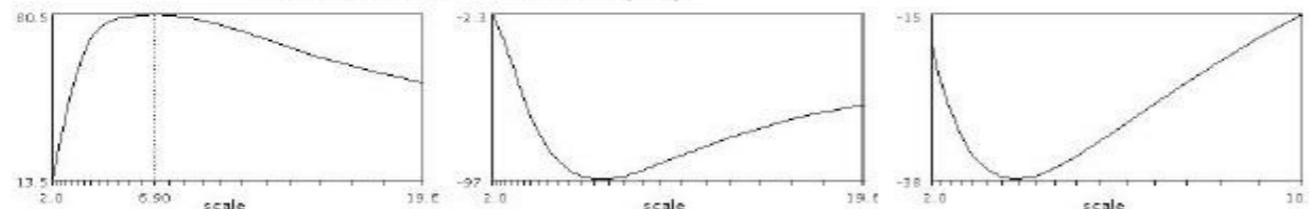


$$I_x^2 + I_y^2$$

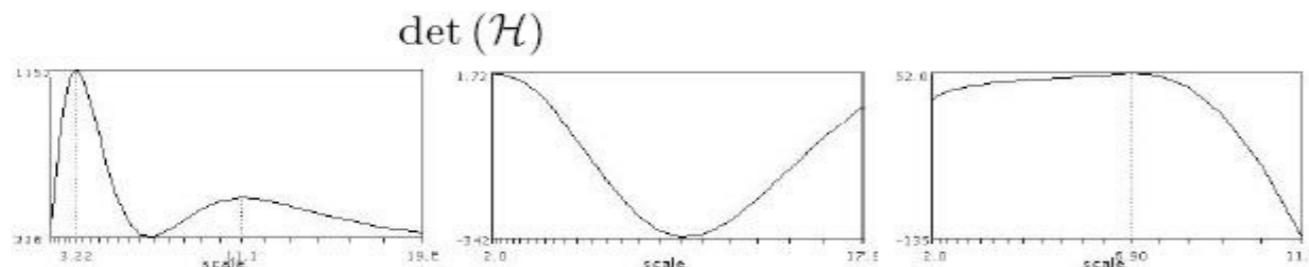
# Automatic scale selection



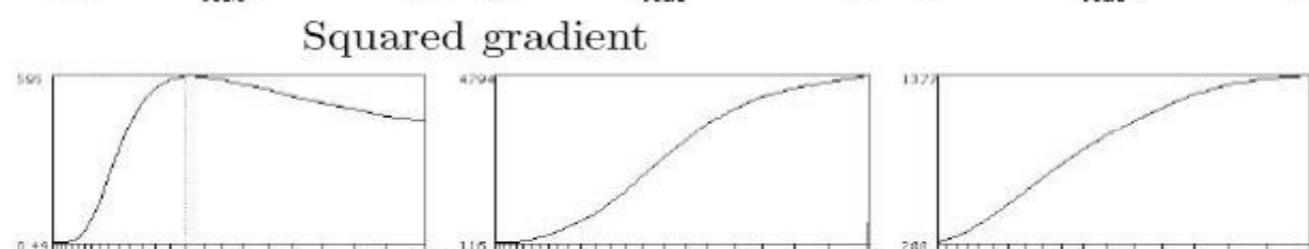
Laplacian =  $\text{trace}(\mathcal{H})$



$$I_{xx} + I_{yy}$$

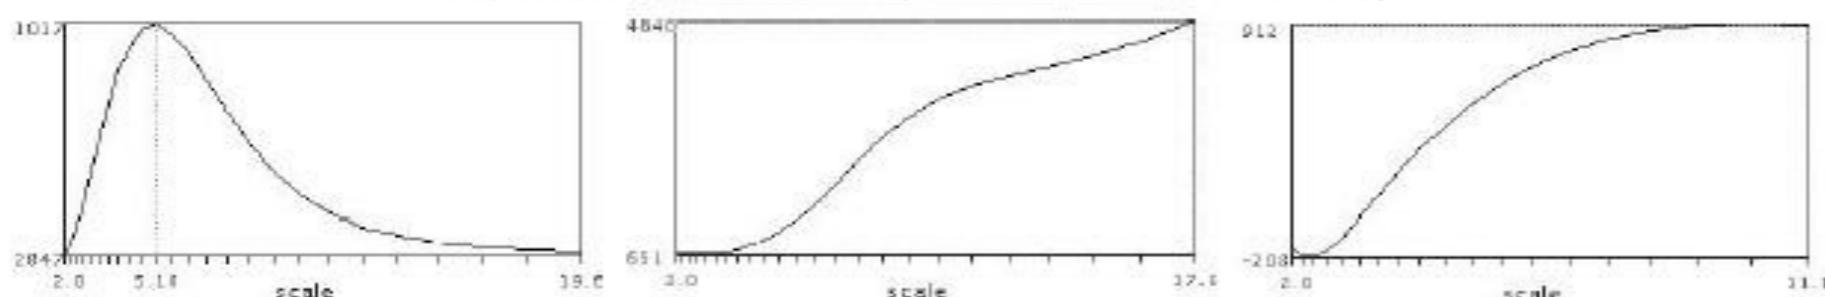


$$I_{xx}I_{yy} - I_{xy}^2$$



$$I_x^2 + I_y^2$$

Harris measure

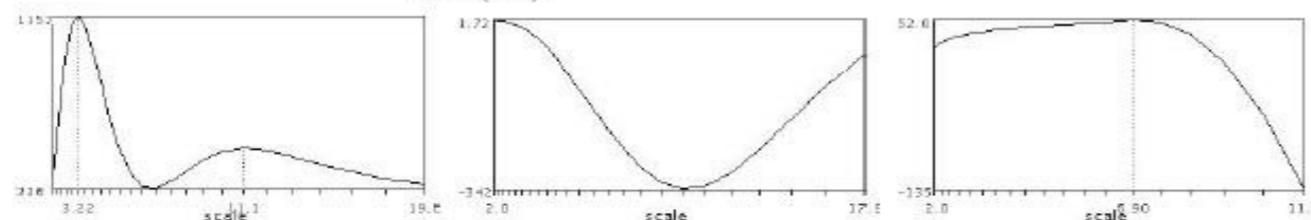
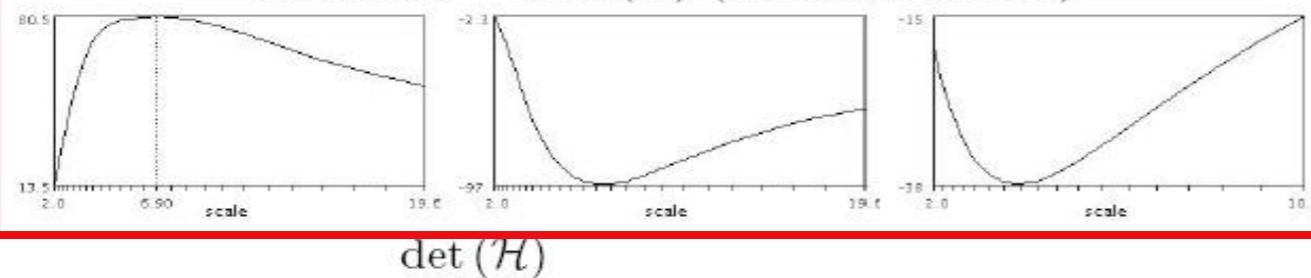


$$\det(H) - \alpha^2 \text{trace}(H)$$

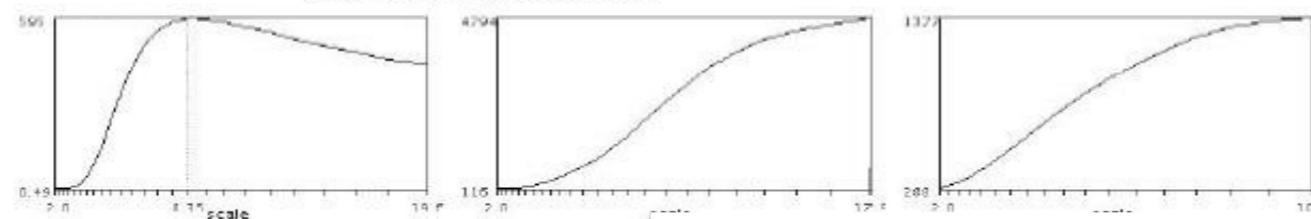
# Automatic scale selection



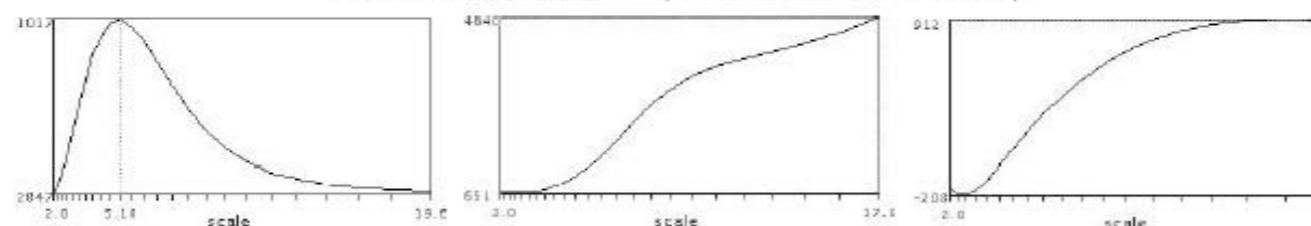
Laplacian =  $\text{trace}(\mathcal{H})$  (cf. equation 3.18)



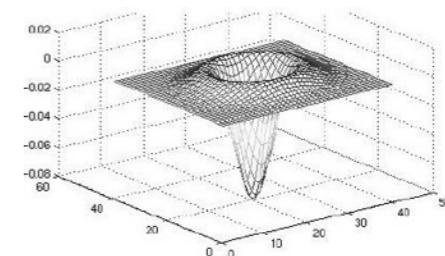
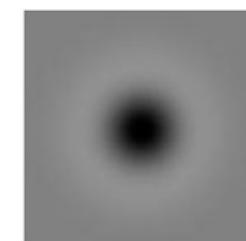
Squared gradient



Harris measure



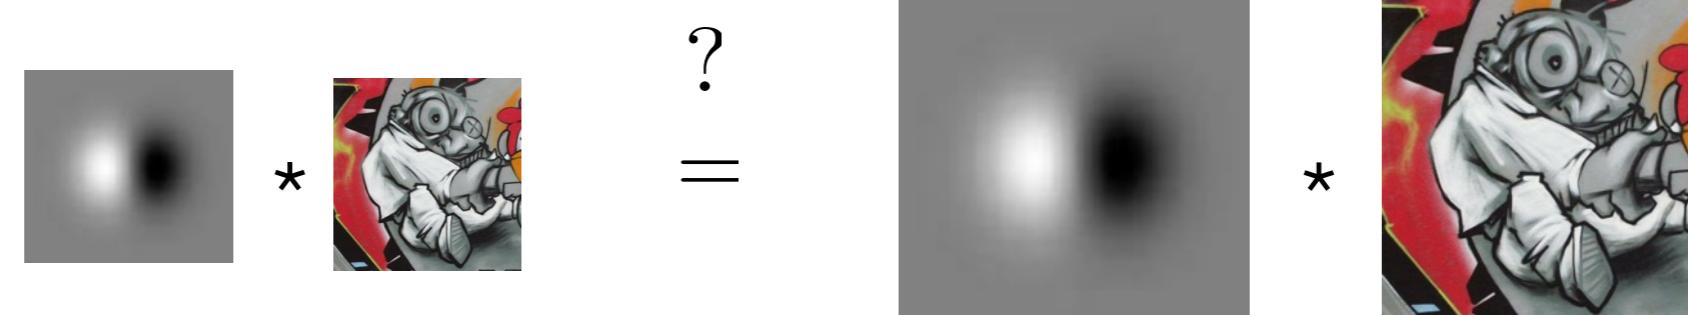
$$I_{xx} + I_{yy}$$



# Scale invariance - normalization



- ◆ The amplitude of spatial derivatives, in general, decreases with scale, due to the response being smoother at a larger scale
- ◆ In the case of the structures present at a large range of scales, like a corner or a step-edge, we would like to have the derivative constant over scale.



$$x' = sx \quad \sigma \quad I(x) \quad ? \quad = \quad \sigma' = s\sigma \quad I(x') = I(sx)$$

- ◆ then:

$$\frac{\partial}{\partial x} I(x') = \frac{\partial}{\partial x} I(sx) = \boxed{s} \frac{\partial}{\partial x} I(x)$$

# Scale invariance - normalization

- ◆ The amplitude of spatial derivatives, in general, decreases with scale, due to the response being smoother at a larger scale
- ◆ In the case of the structures present at a large range of scales, like a corner or a step-edge, we would like to have the derivative constant over scale.

$$\begin{array}{ccccc}
 \text{?} & & & & \\
 \sigma & * & I(x) & = & \sigma' = s\sigma \\
 & & & & * \\
 & & & & I(x') = I(sx)
 \end{array}$$

- ◆
- ◆  $I_{r_1 \dots r_m}(x') = \boxed{s^m} I_{r_1 \dots r_m}(sx)$
- ◆ In order to maintain the property of scale invariance the derivative function must be normalized with respect to the scale of derivation

$$I_{r_1 \dots r_m}(x') = \boxed{s^m} g_{r_1 \dots r_m}(\boxed{s\sigma}) * I(sx) = s^m g_{r_1 \dots r_m}(\sigma') * I(x')$$

# Laplacian of Gaussian (LoG)

- ◆ Local optima in scale space of Laplacian of Gaussian



$$I_{xx}(\sigma) + I_{yy}(\sigma)$$

# Laplacian of Gaussian (LoG)



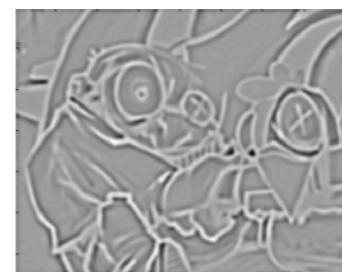
- ◆ Local optima in scale space of Laplacian of Gaussian



$$I_{xx}(\sigma) + I_{yy}(\sigma)$$



$\sigma$



# Laplacian of Gaussian (LoG)



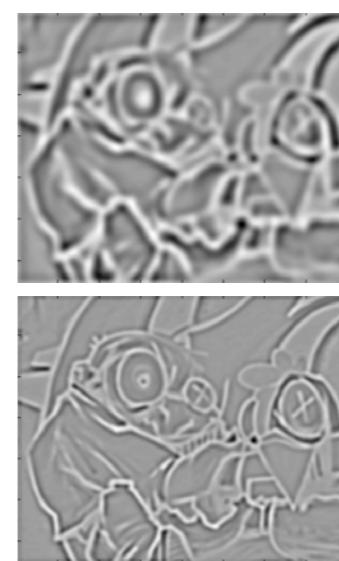
- ◆ Local optima in scale space of Laplacian of Gaussian



$$I_{xx}(\sigma) + I_{yy}(\sigma)$$

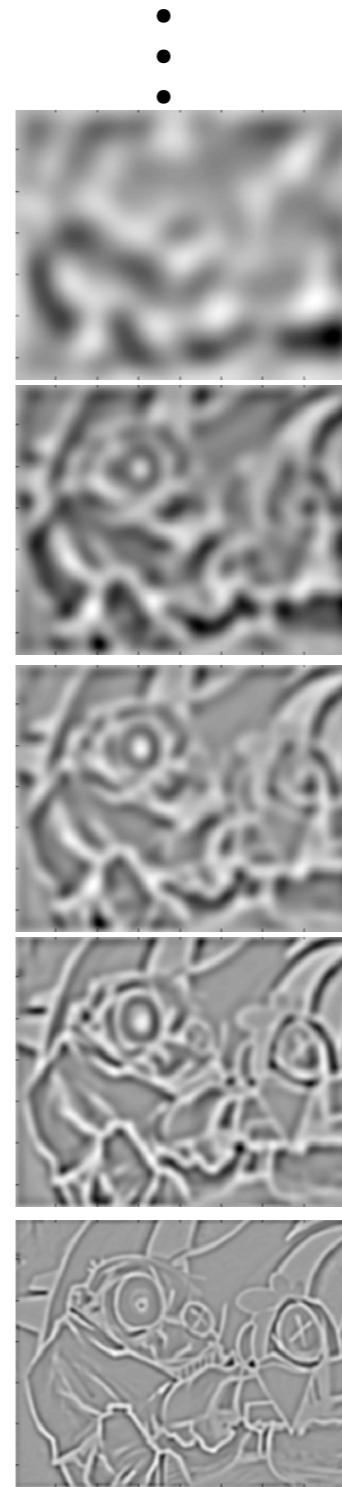
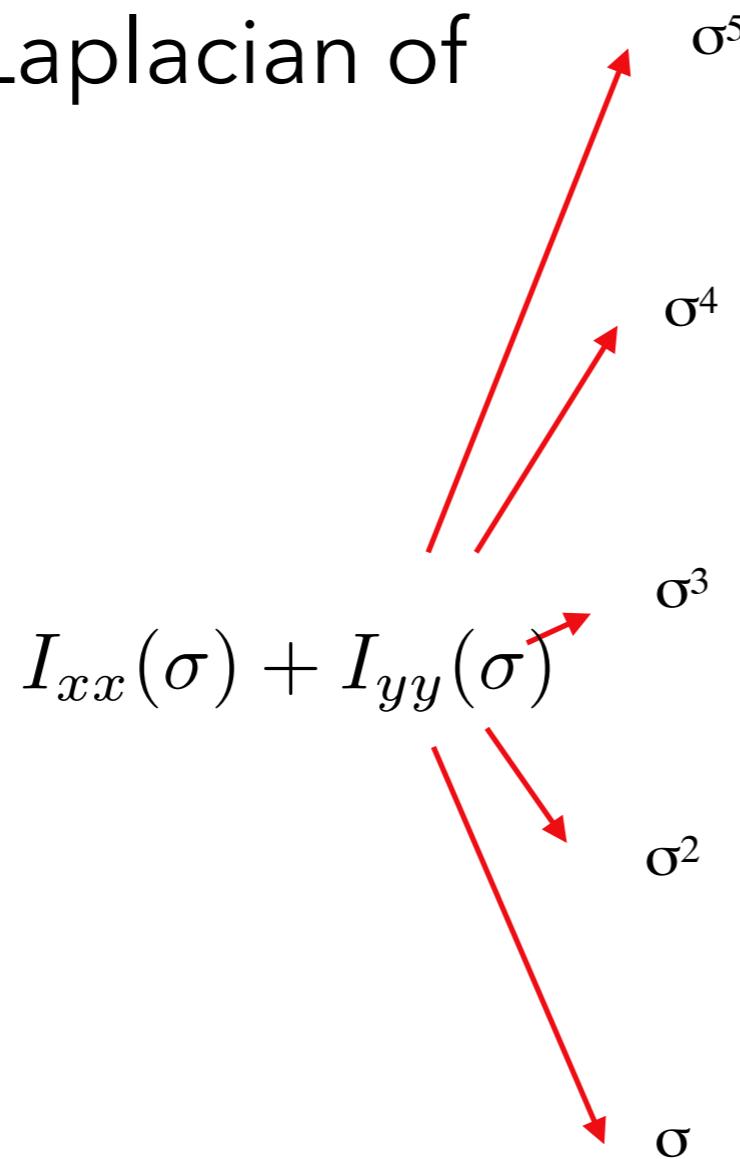
$\sigma^2$

$\sigma$



# Laplacian of Gaussian (LoG)

- ◆ Local optima in scale space of Laplacian of Gaussian

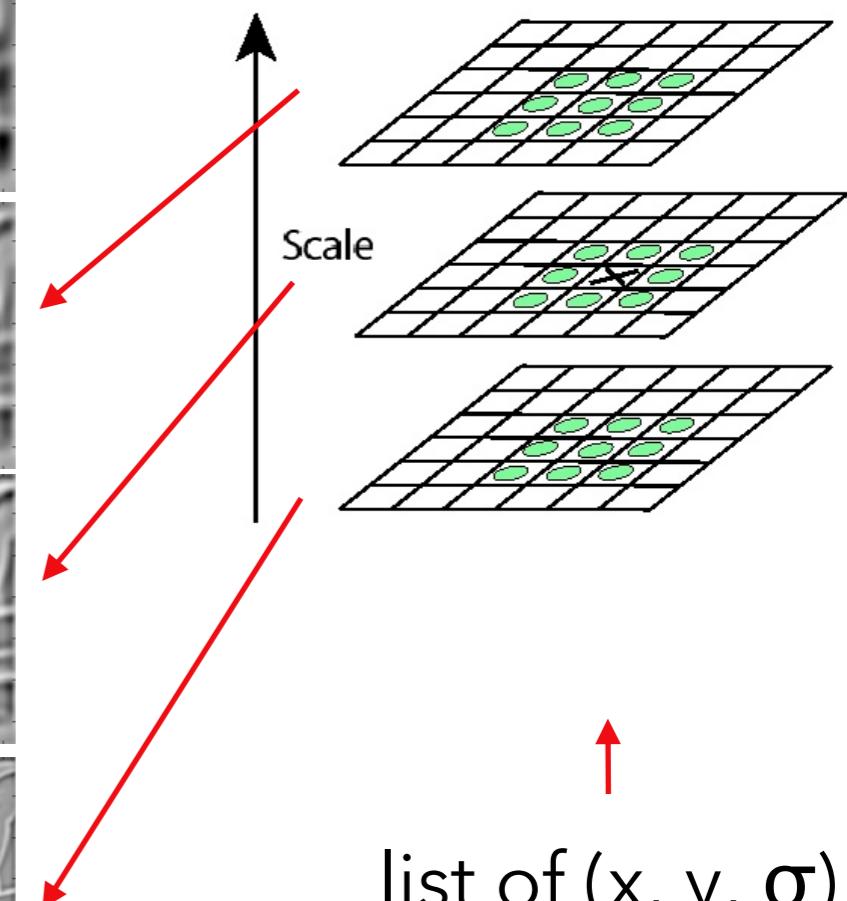
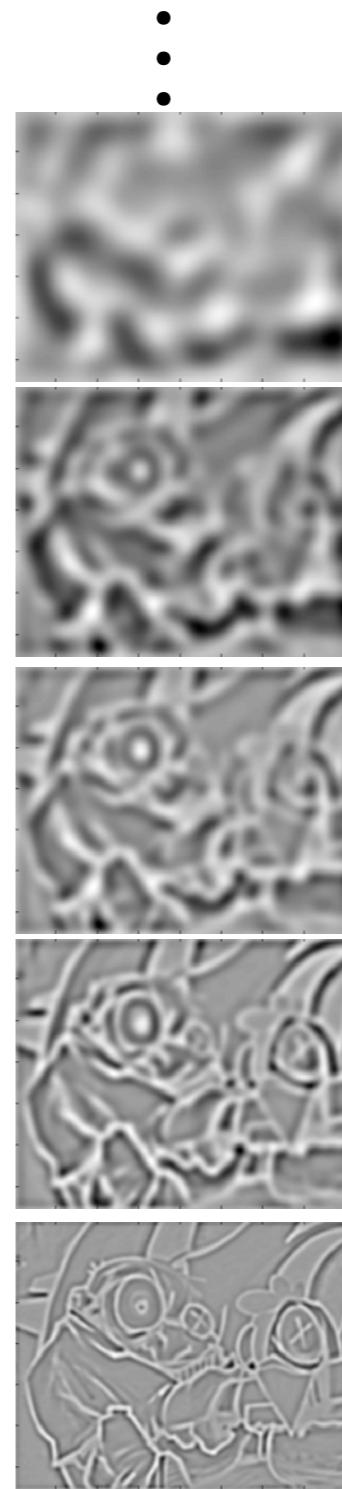


# Laplacian of Gaussian (LoG)

- ◆ Local optima in scale space of Laplacian of Gaussian



$$I_{xx}(\sigma) + I_{yy}(\sigma)$$

 $\sigma^5$  $\sigma^4$  $\sigma^3$  $\sigma^2$  $\sigma$ 

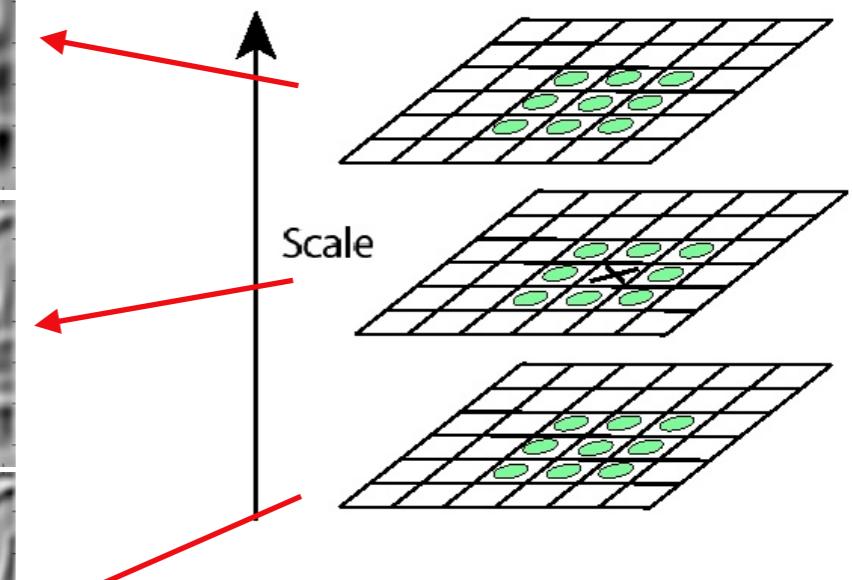
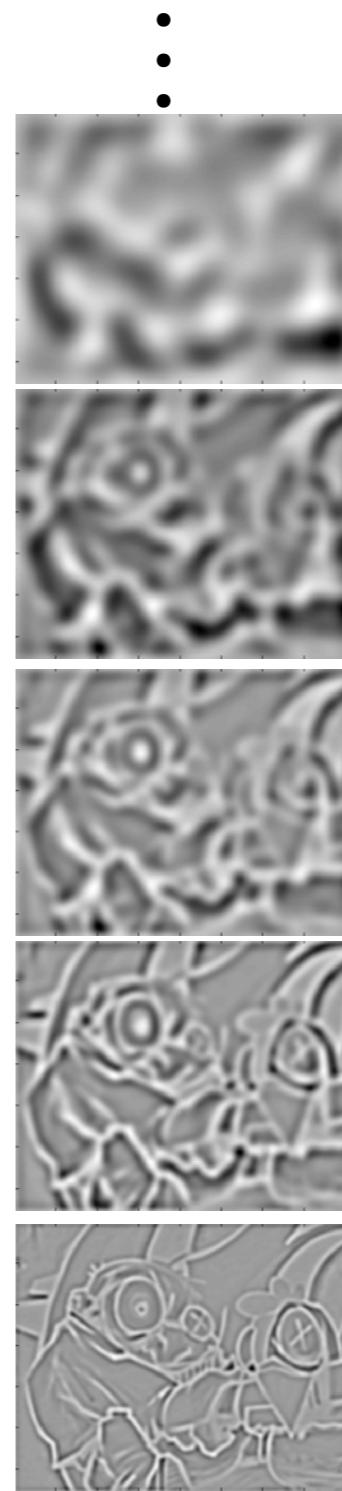
# Laplacian of Gaussian (LoG)

- ◆ Local optima in scale space of Laplacian of Gaussian



$$I_{xx}(\sigma) + I_{yy}(\sigma)$$

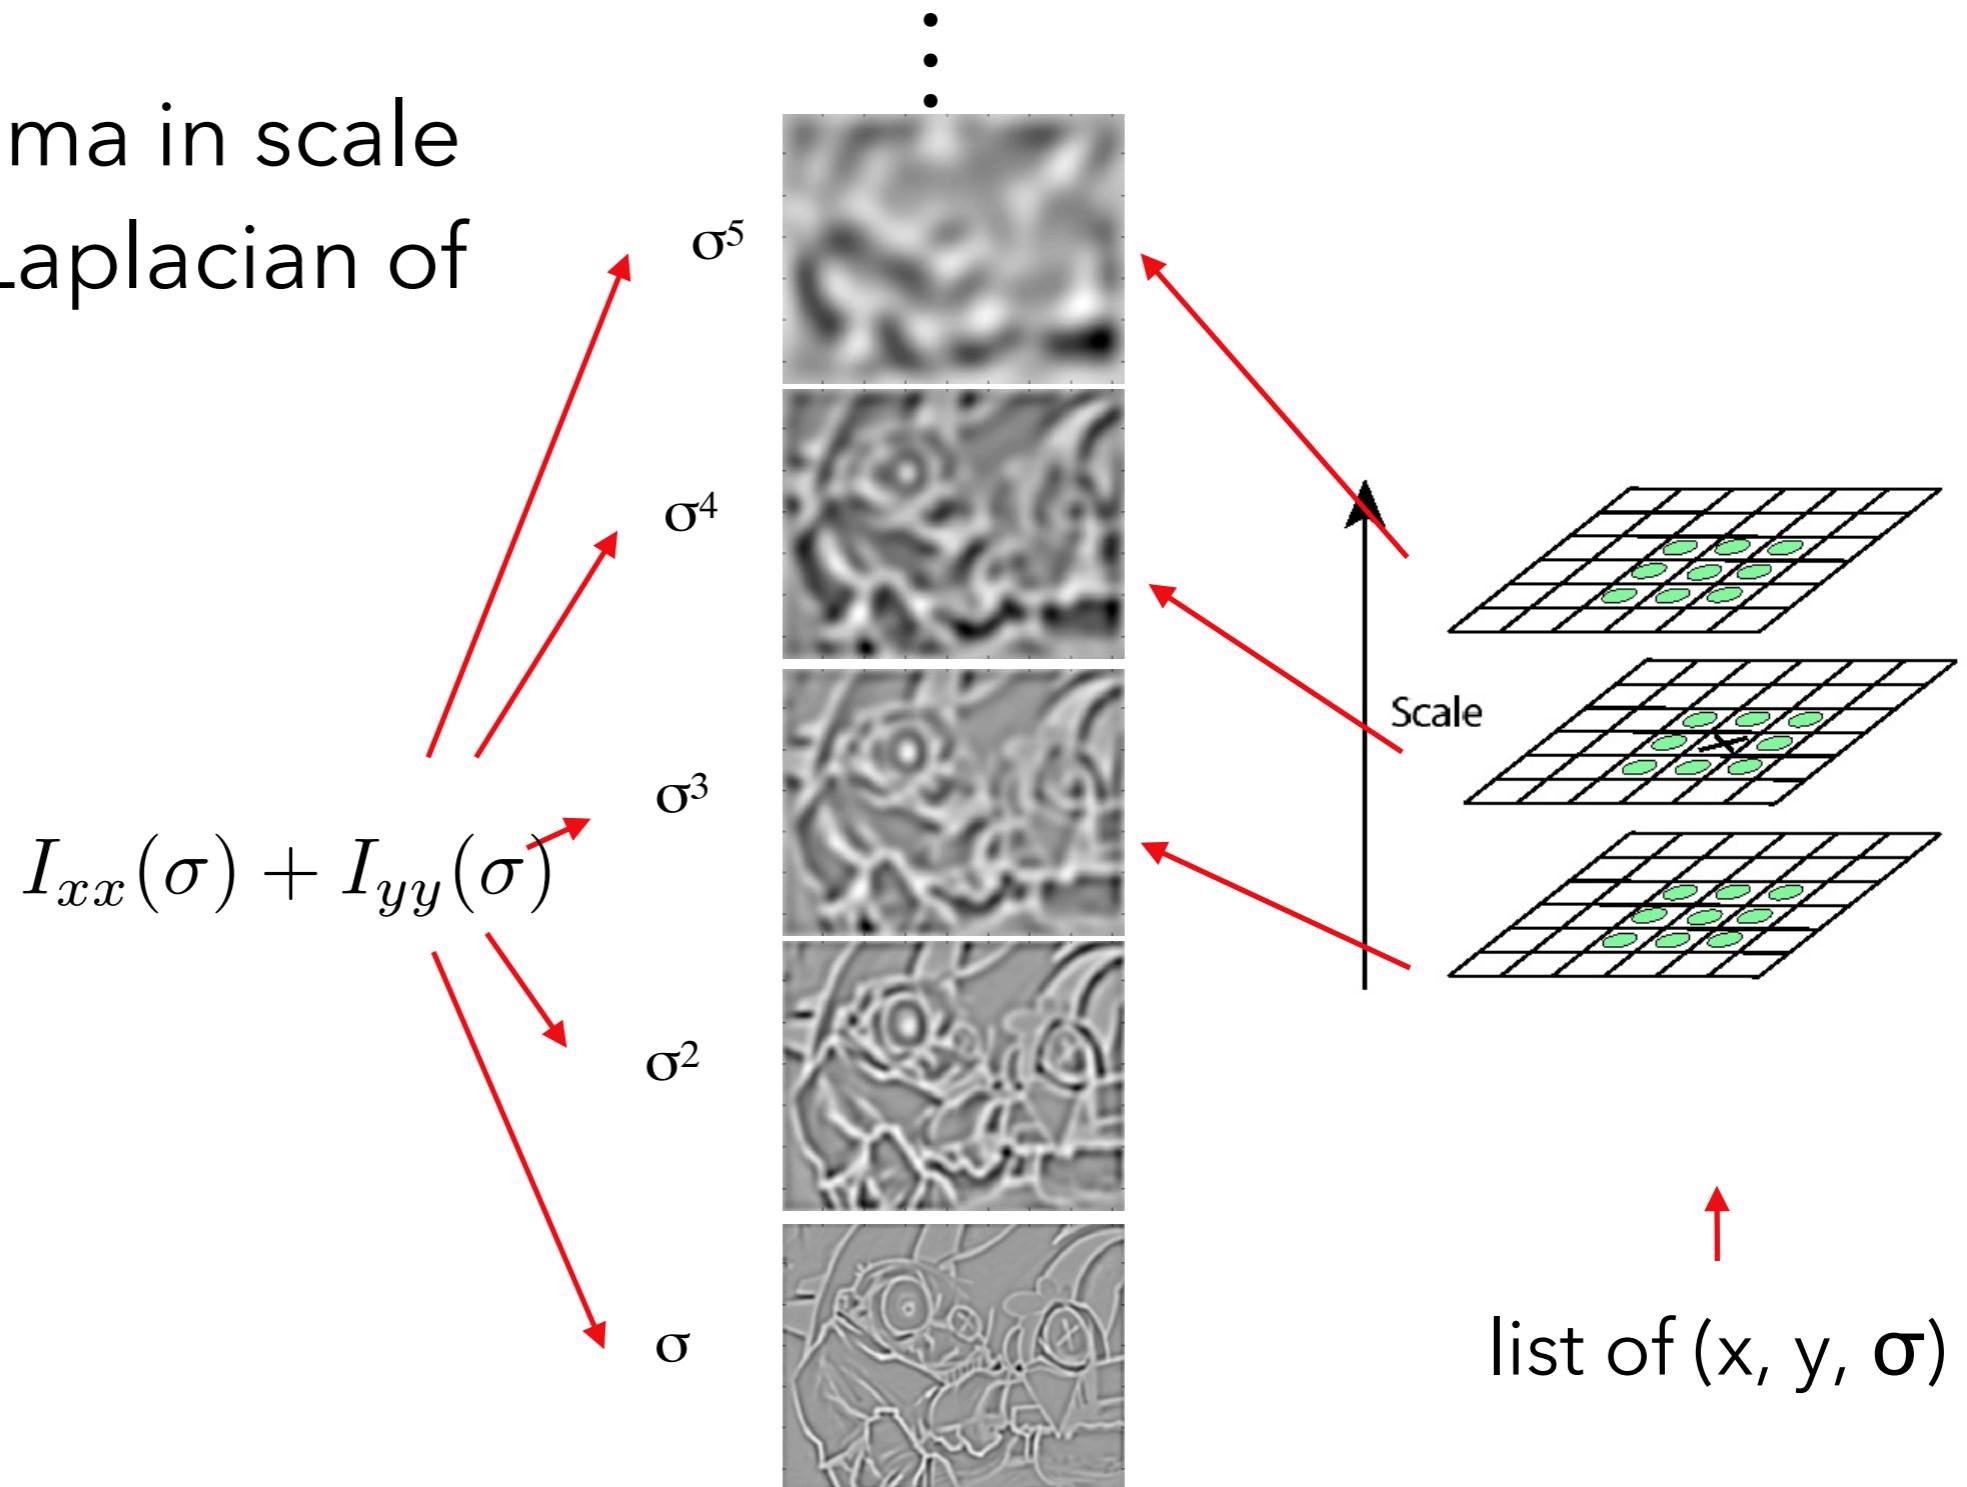
$\sigma^5$   
 $\sigma^4$   
 $\sigma^3$   
 $\sigma^2$   
 $\sigma$



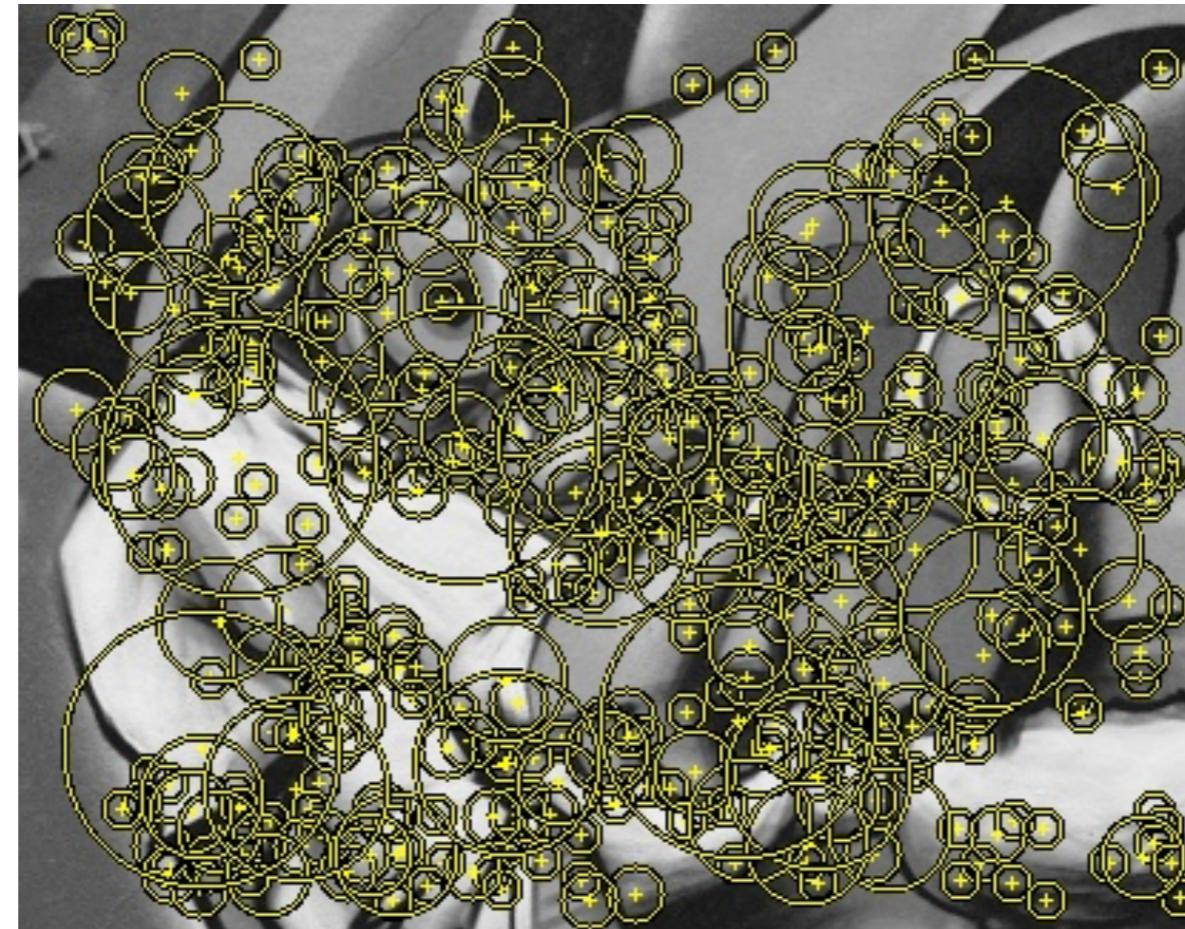
list of  $(x, y, \sigma)$

# Laplacian of Gaussian (LoG)

- ◆ Local optima in scale space of Laplacian of Gaussian



# Laplacian of Gaussian (LoG)



# Harris-Laplace (HarLap)

- ◆ Detecting multiscale Harris points



$$H = \sum_{(x,y) \in R} w(x,y) \nabla I(x,y) \nabla I(x,y)^T$$

$$f = \det(H) - \alpha^2 \text{trace}(H)$$

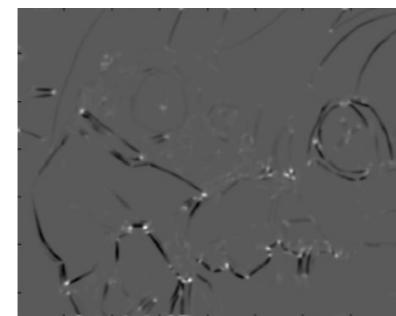
# Harris-Laplace (HarLap)

- ◆ Detecting multiscale Harris points



$$H = \sum_{(x,y) \in R} w(x,y) \nabla I(x,y) \nabla I(x,y)^T$$

$$f = \det(H) - \alpha^2 \text{trace}(H)$$



Computing Harris function

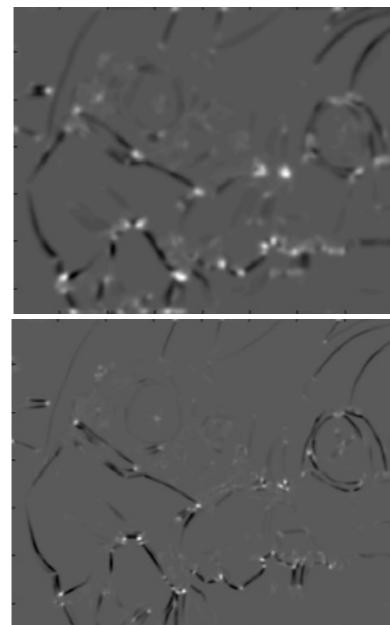
# Harris-Laplace (HarLap)



- ◆ Detecting multiscale Harris points



$$H = \sum_{(x,y) \in R} w(x,y) \nabla I(x,y) \nabla I(x,y)^T$$
$$f = \det(H) - \alpha^2 \text{trace}(H)$$



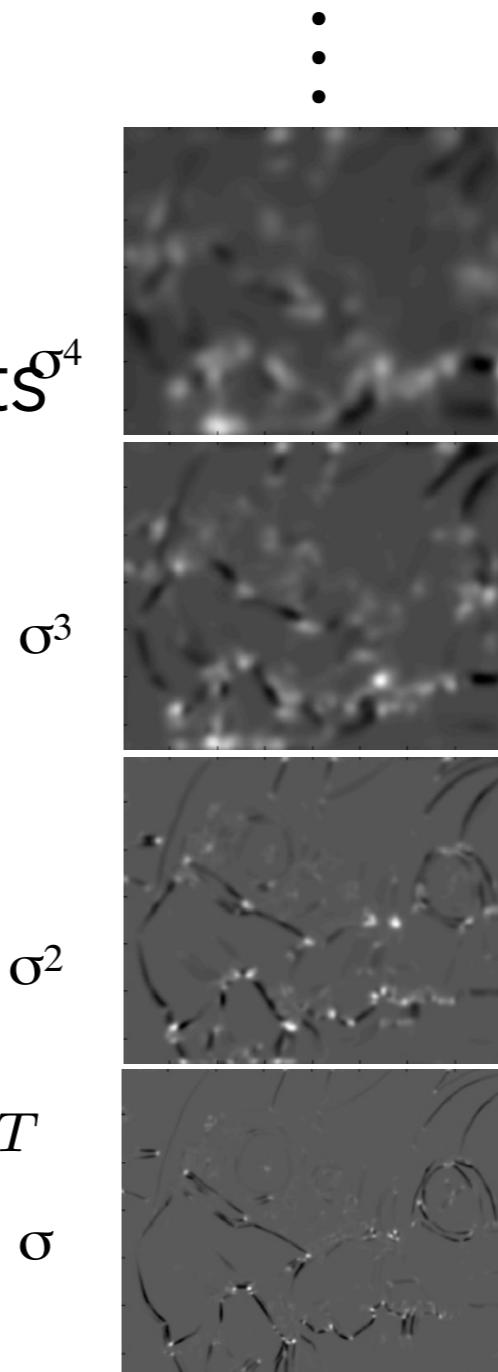
Computing Harris function

# Harris-Laplace (HarLap)

- ◆ Detecting multiscale Harris points - thousands of interest points<sup>σ<sup>4</sup></sup>



$$H = \sum_{(x,y) \in R} w(x,y) \nabla I(x,y) \nabla I(x,y)^T$$
$$f = \det(H) - \alpha^2 \text{trace}(H)$$



Computing Harris function

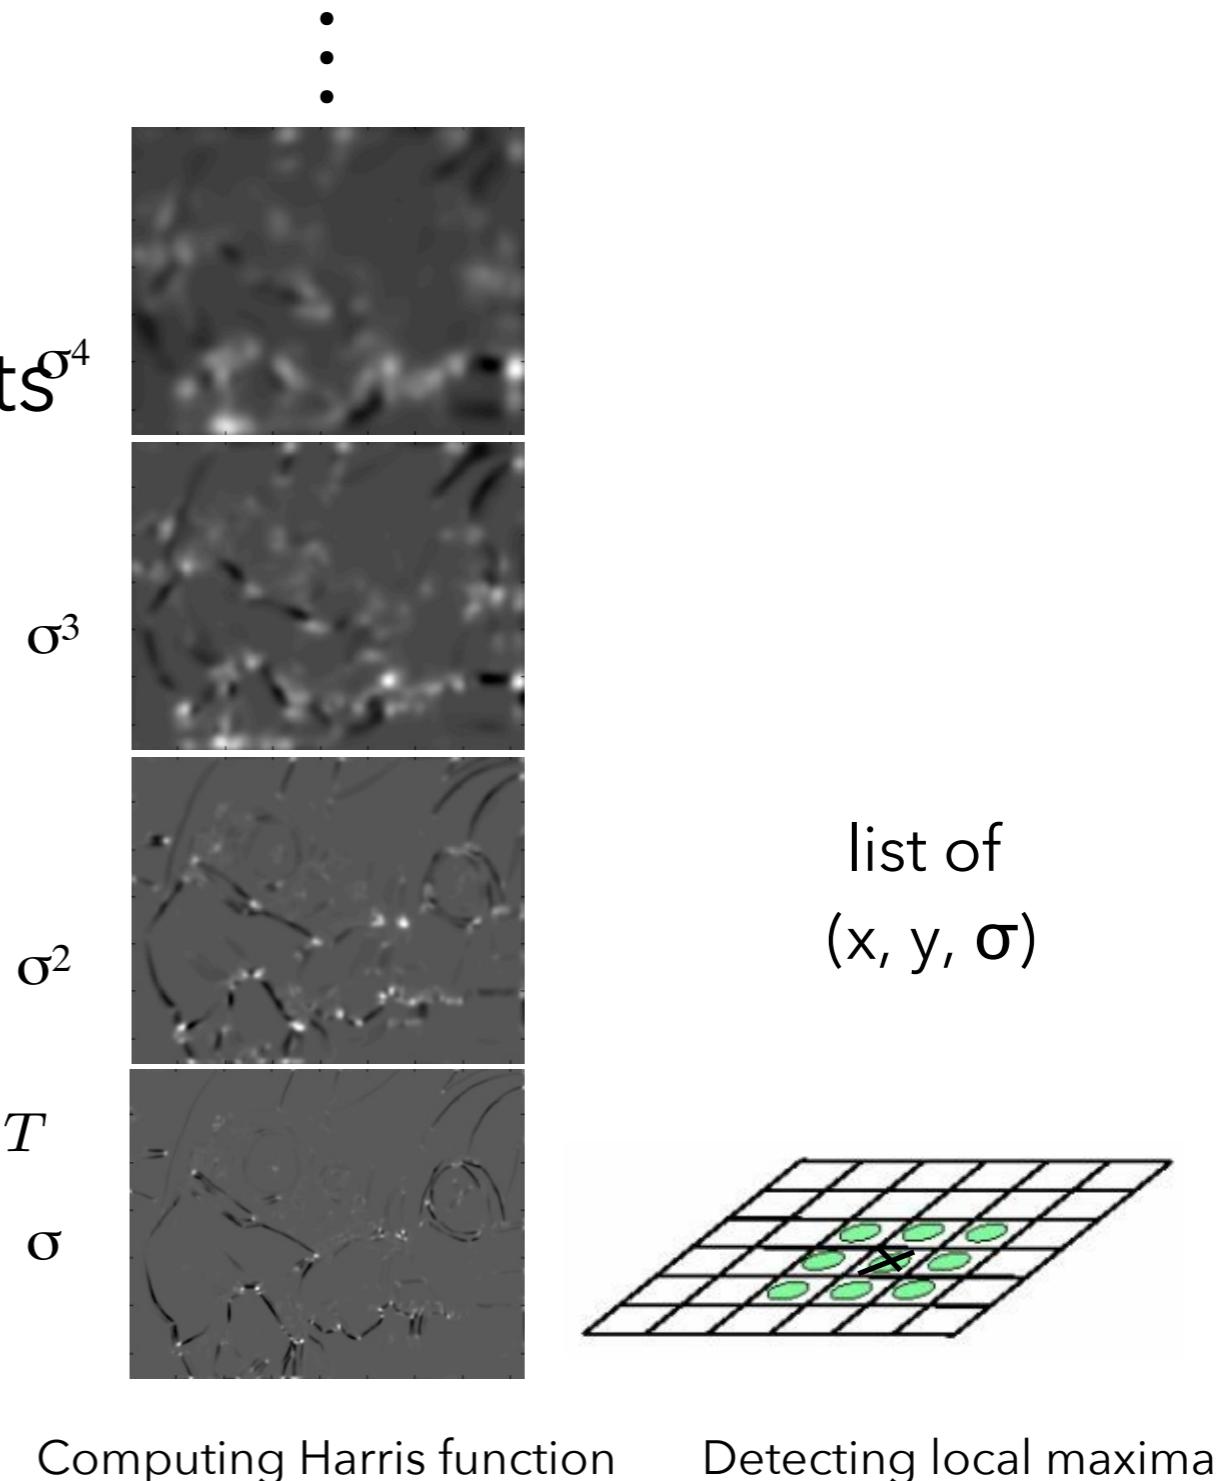
# Harris-Laplace (HarLap)

- ◆ Detecting multiscale Harris points - thousands of interest points<sup>σ<sup>4</sup></sup>



$$H = \sum_{(x,y) \in R} w(x,y) \nabla I(x,y) \nabla I(x,y)^T$$

$$f = \det(H) - \alpha^2 \text{trace}(H)$$



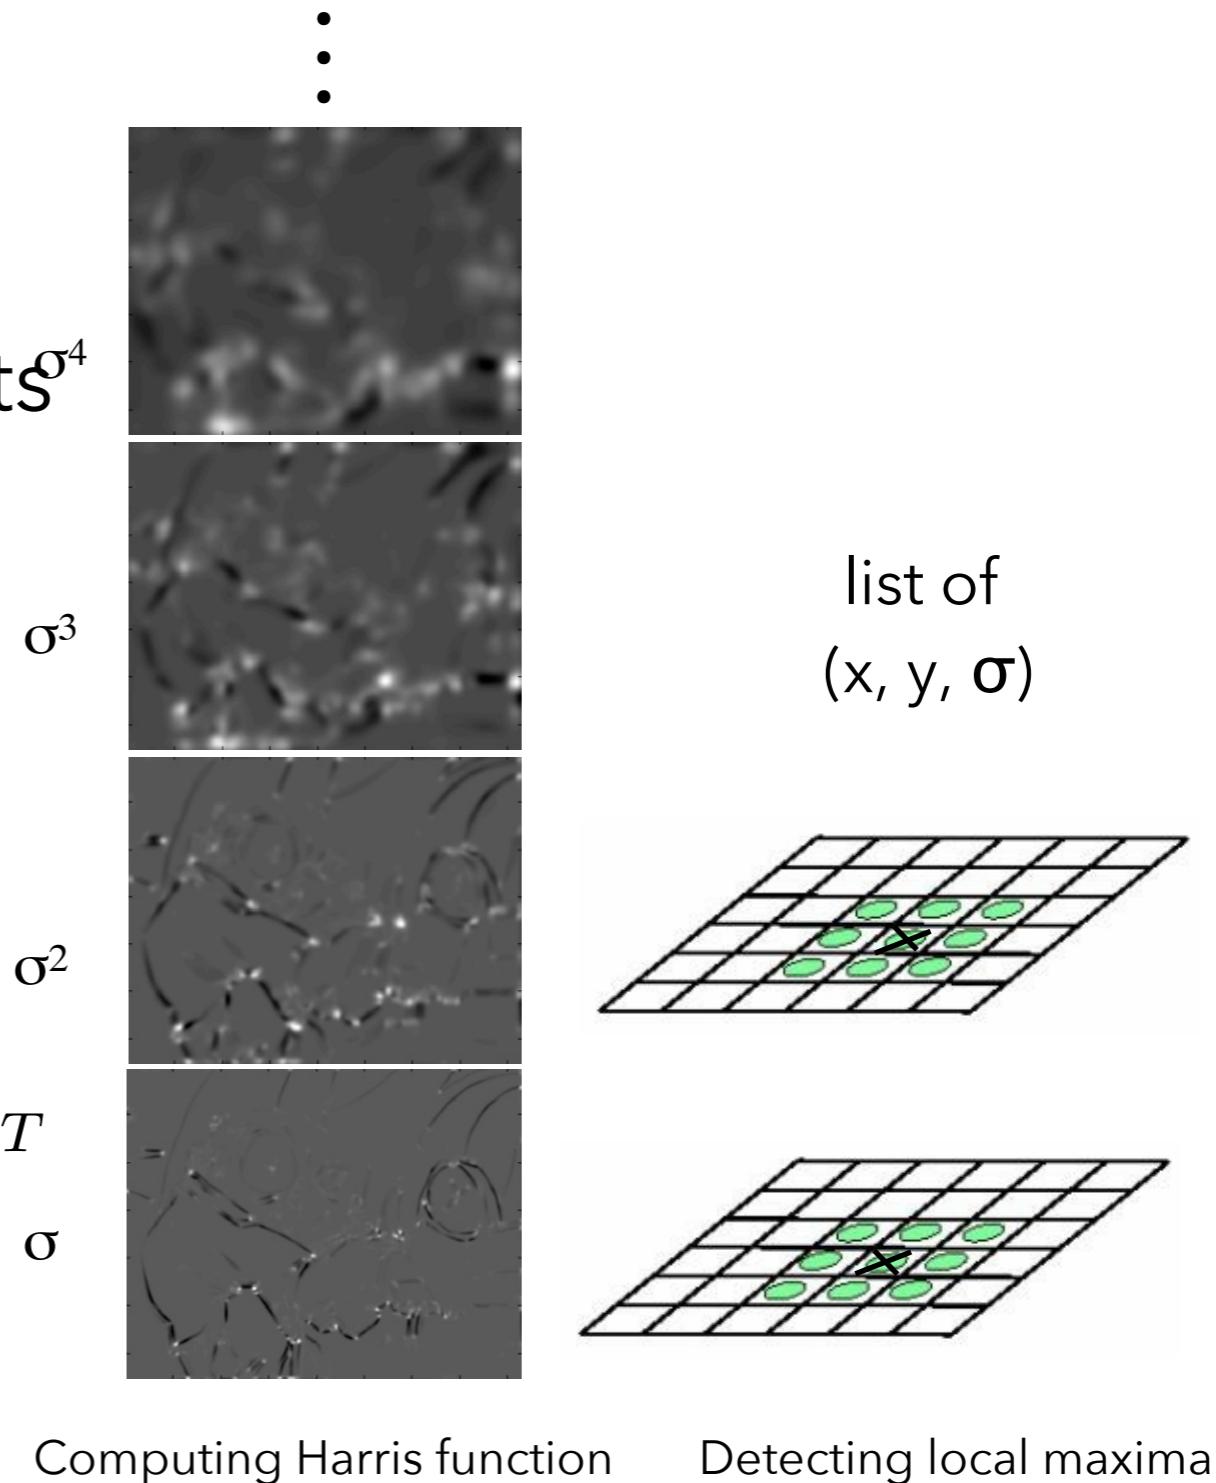
# Harris-Laplace (HarLap)

- ◆ Detecting multiscale Harris points - thousands of interest points<sup>4</sup>



$$H = \sum_{(x,y) \in R} w(x,y) \nabla I(x,y) \nabla I(x,y)^T$$

$$f = \det(H) - \alpha^2 \text{trace}(H)$$



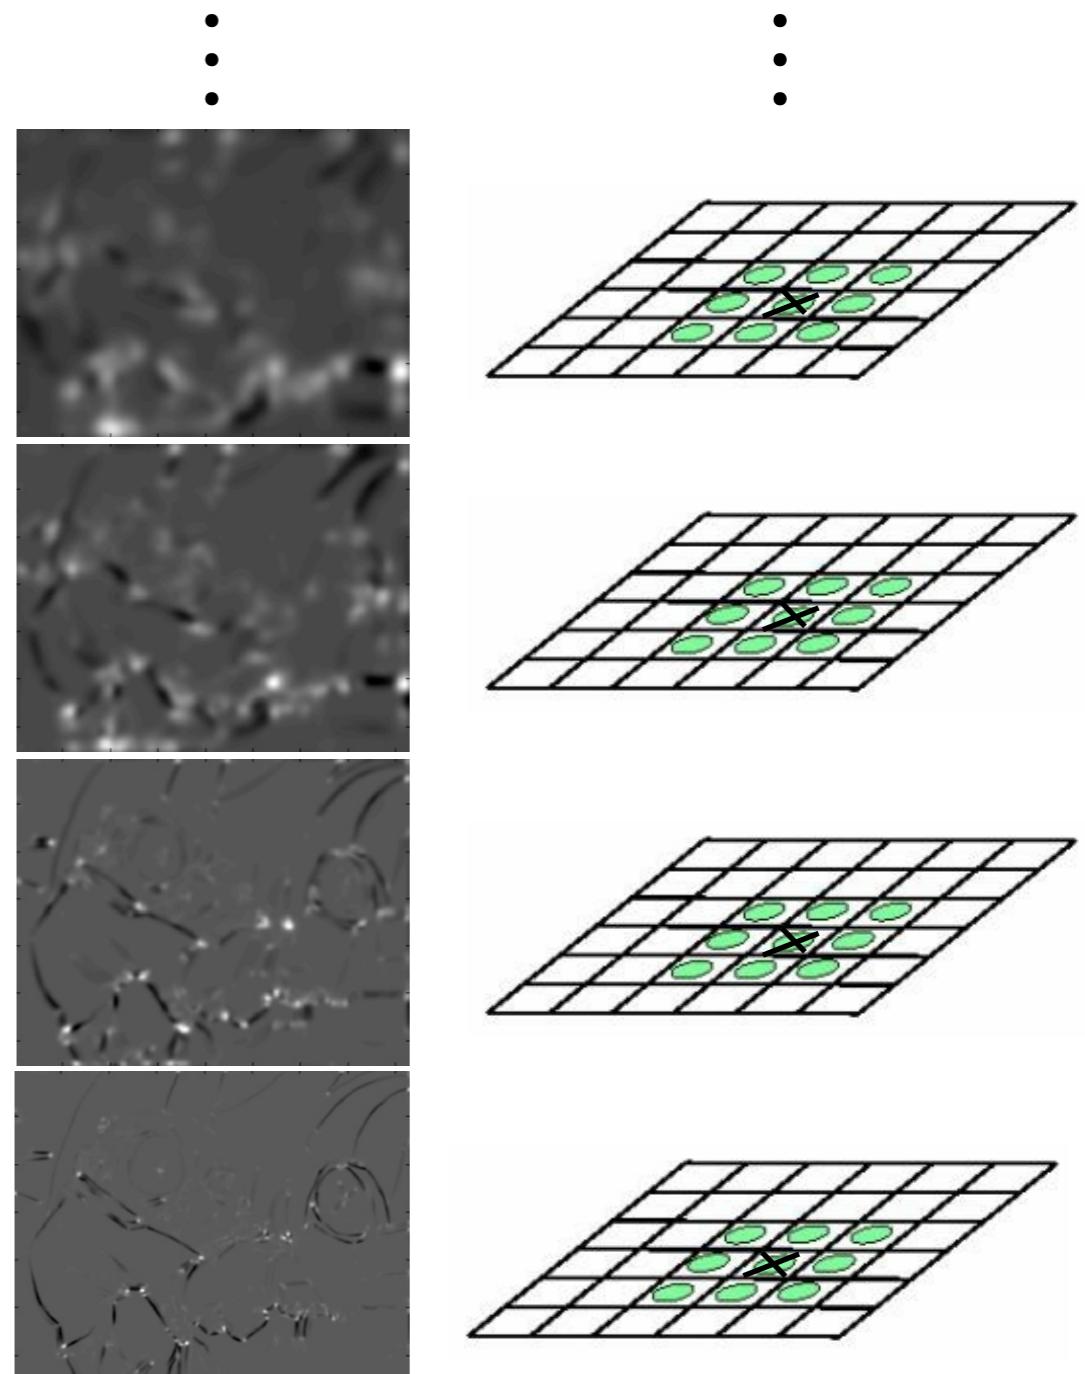
# Harris-Laplace (HarLap)



- ◆ Detecting multiscale Harris points - thousands of interest points<sup>σ<sup>4</sup></sup>



$$H = \sum_{(x,y) \in R} w(x,y) \nabla I(x,y) \nabla I(x,y)^T$$
$$f = \det(H) - \alpha^2 \text{trace}(H)$$

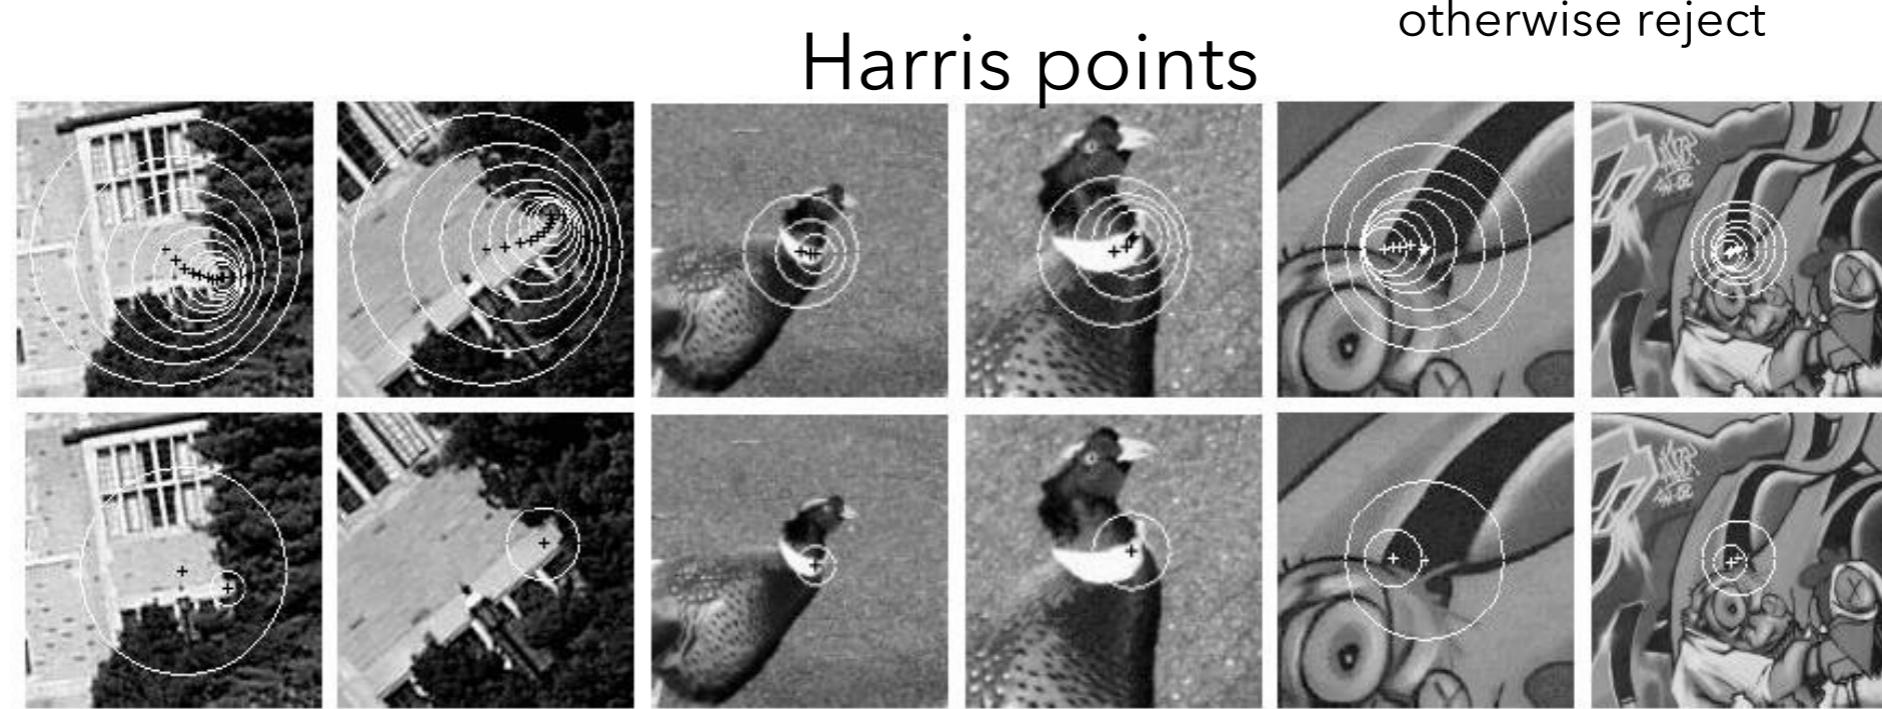


Computing Harris function

Detecting local maxima

# Harris-Laplace (HarLap)

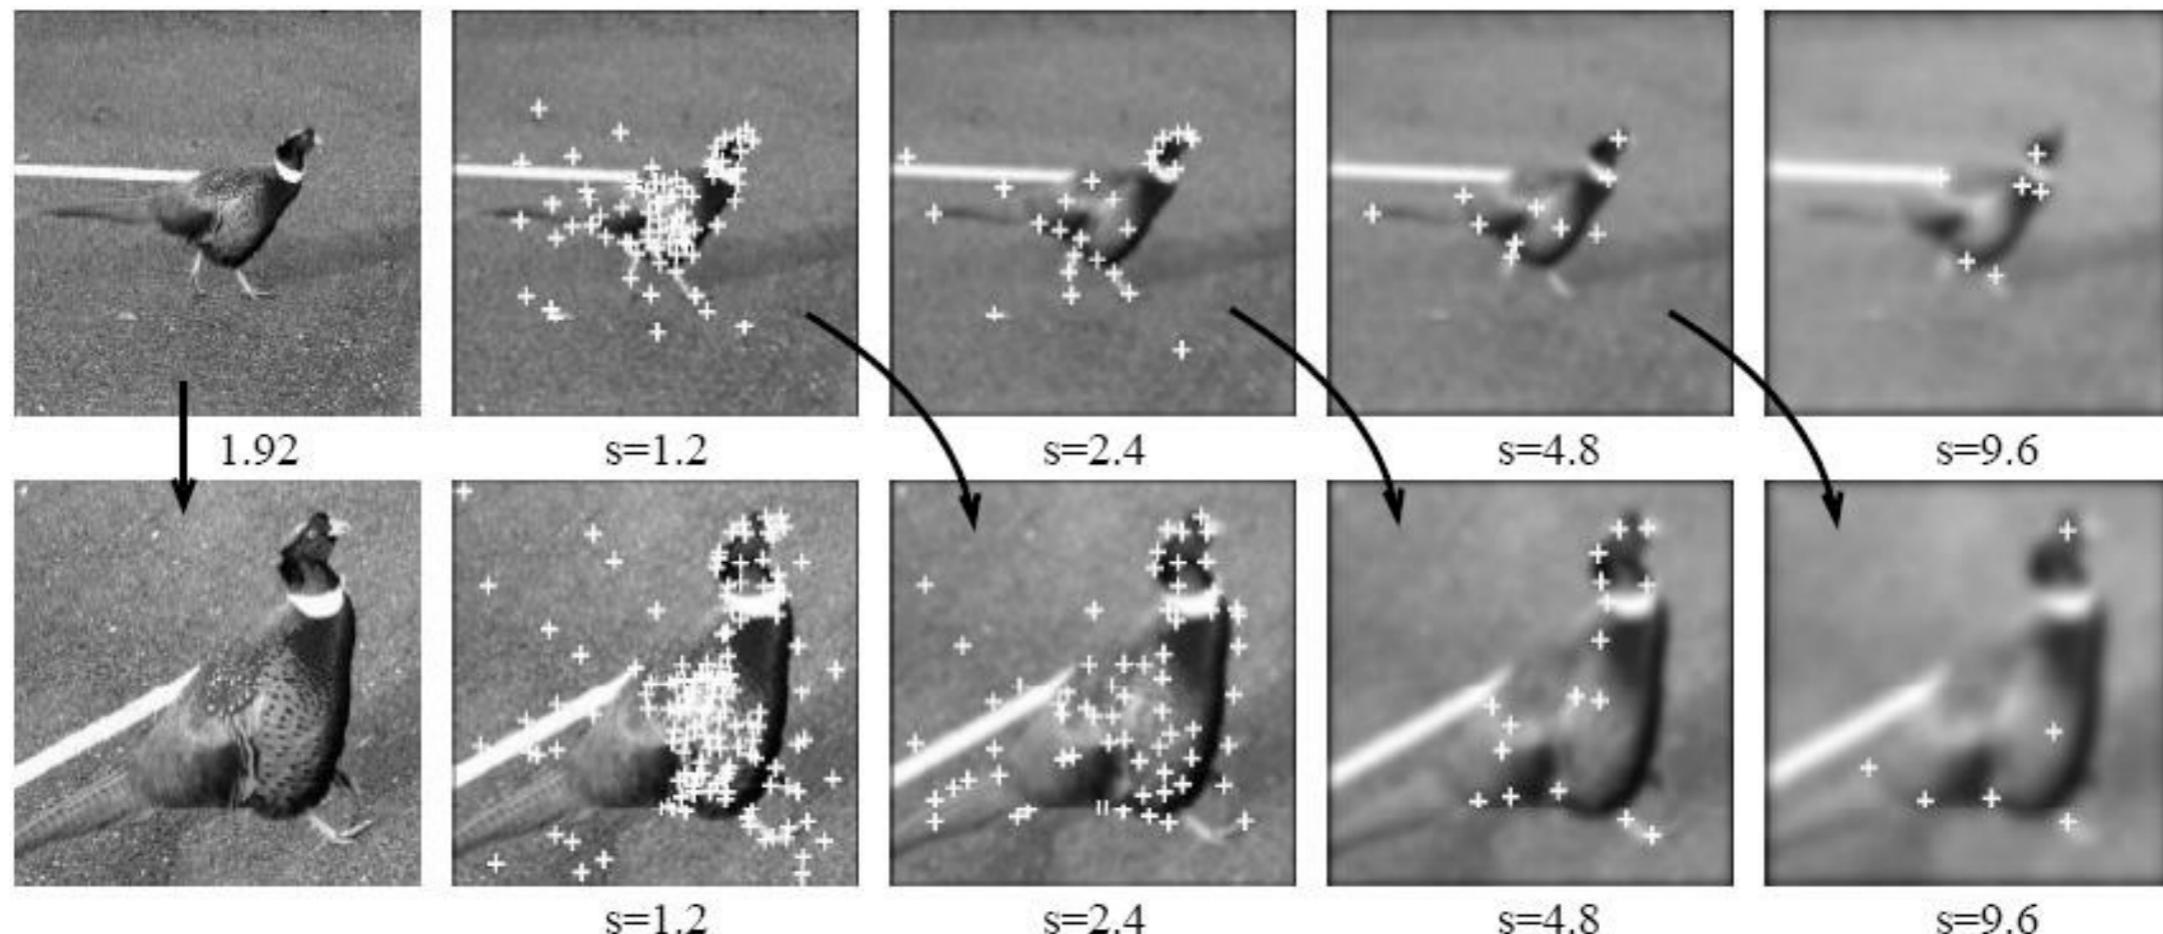
- ◆ Detecting multiscale Harris points
- ◆ Selecting points which maximize the Laplacian
  - ◆ Automatic scale selection
    - ◆ Given a point  $(x, y, \sigma_n)$
    - ◆ If  $(L_{xx}(\sigma_n) + L_{yy}(\sigma_n)) > (L_{xx}(\sigma_{n-1}) + L_{yy}(\sigma_{n-1}))$  and  $(L_{xx}(\sigma_n) + L_{yy}(\sigma_n)) > (L_{xx}(\sigma_{n+1}) + L_{yy}(\sigma_{n+1}))$  keep the point, otherwise reject



Harris-Laplace points

# Harris-Laplace (HarLap)

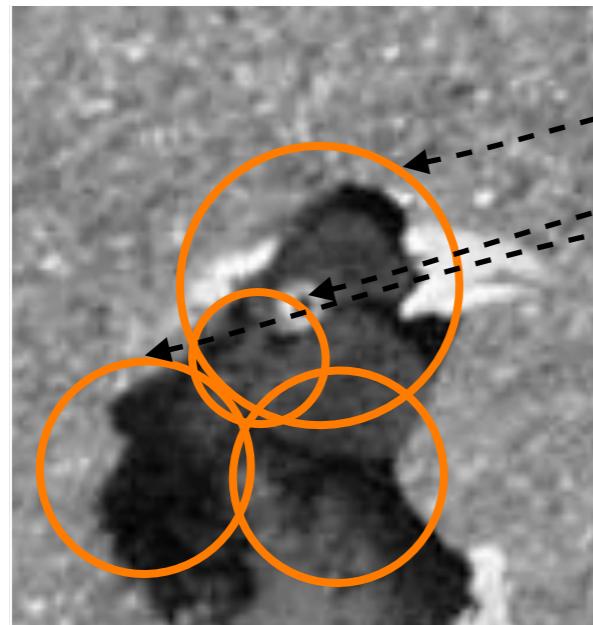
- ◆ Detecting multiscale Harris points
- ◆ Selecting Harris points which maximize the Laplacian
- ◆ Automatic scale selection



Harris-Laplace points

# Evaluation criteria

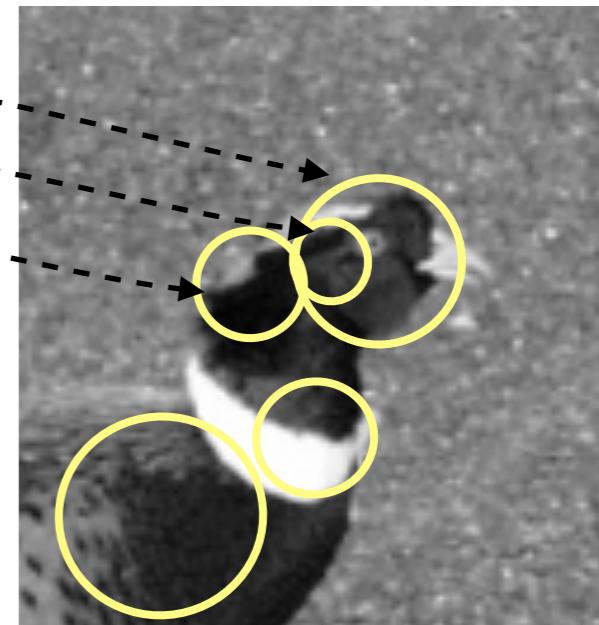
- ◆ Repeatability rate : percentage of corresponding points



#correspondences = 3

#detected = 4

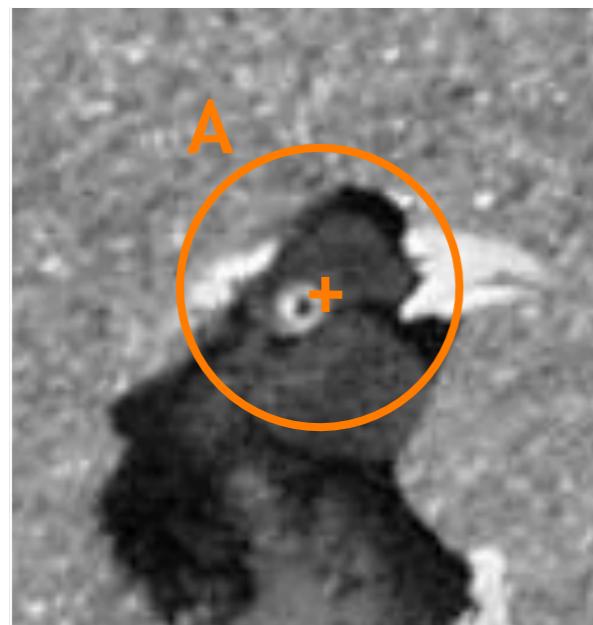
Repeatability=75%



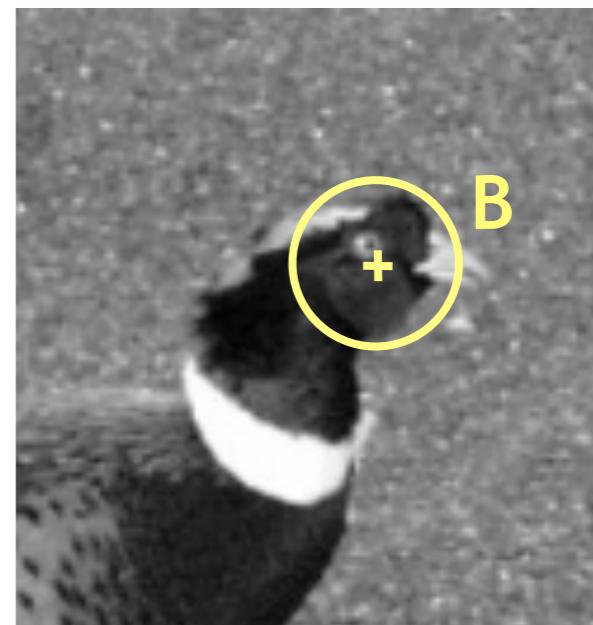
$$\text{repeatability} = \frac{\# \text{correspondences}}{\# \text{detected}} \cdot 100\%$$

# Evaluation criteria

- ◆ Repeatability rate : percentage of corresponding points

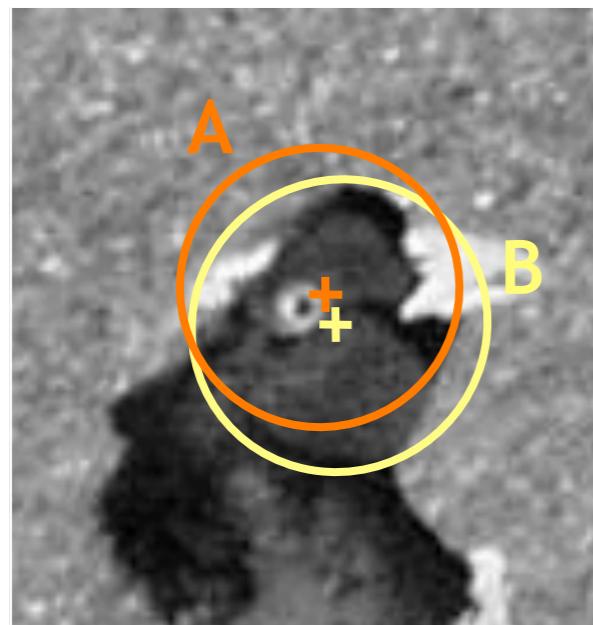


homography

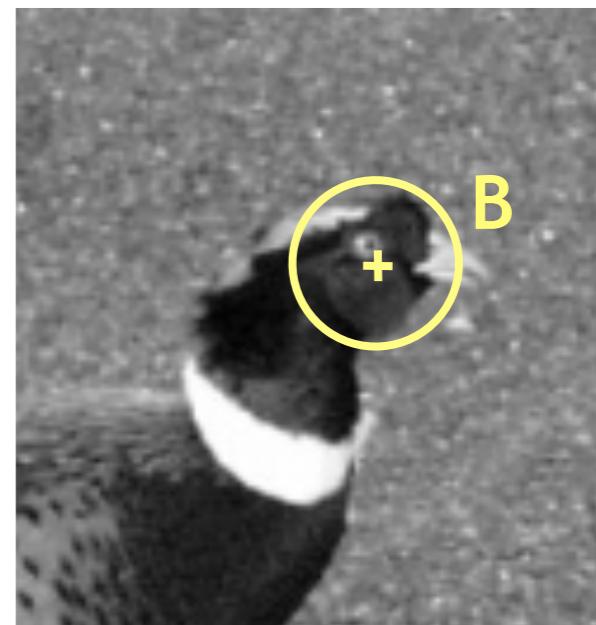


# Evaluation criteria

- ◆ Repeatability rate : percentage of corresponding points

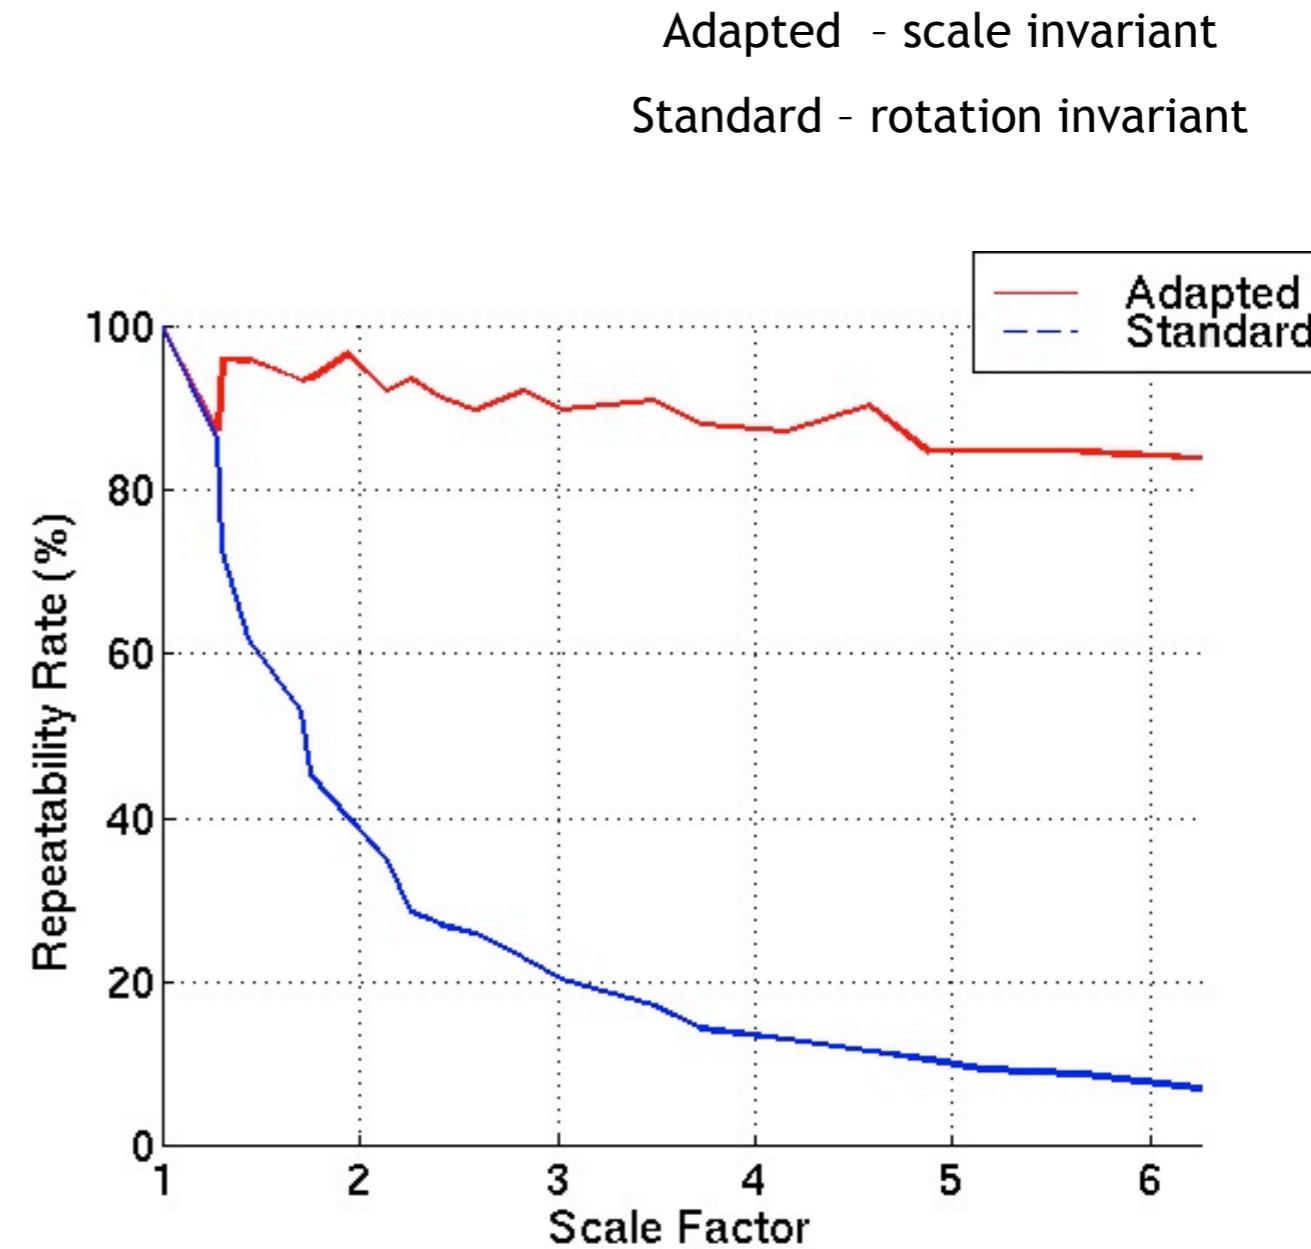
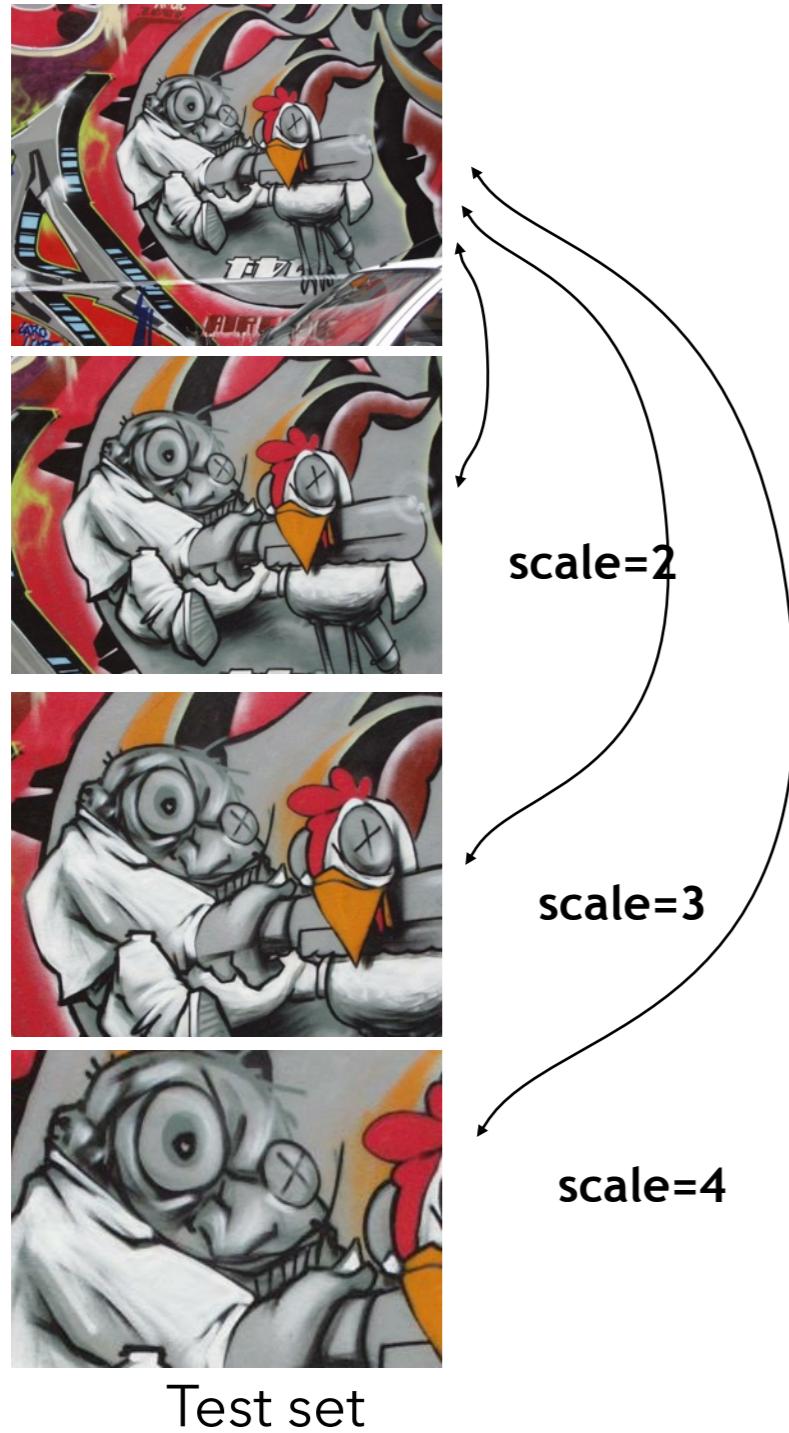


homography



Two points are corresponding if  $\frac{A \cap B}{A \cup B} > T$   
 $T=60\%$

# Evaluation criteria





# Overview Today

- ◆ 1 - Local Interest Point Detection
  - ◆ finding discriminant points (Harris, Hessian)
  - ◆ scale invariant interest point detection
- ◆ **2 - Local Descriptors (= Features)**
- ◆ 3 - Bag-of-Words Model (BoW)
  - ◆ BoW for Object Categorization

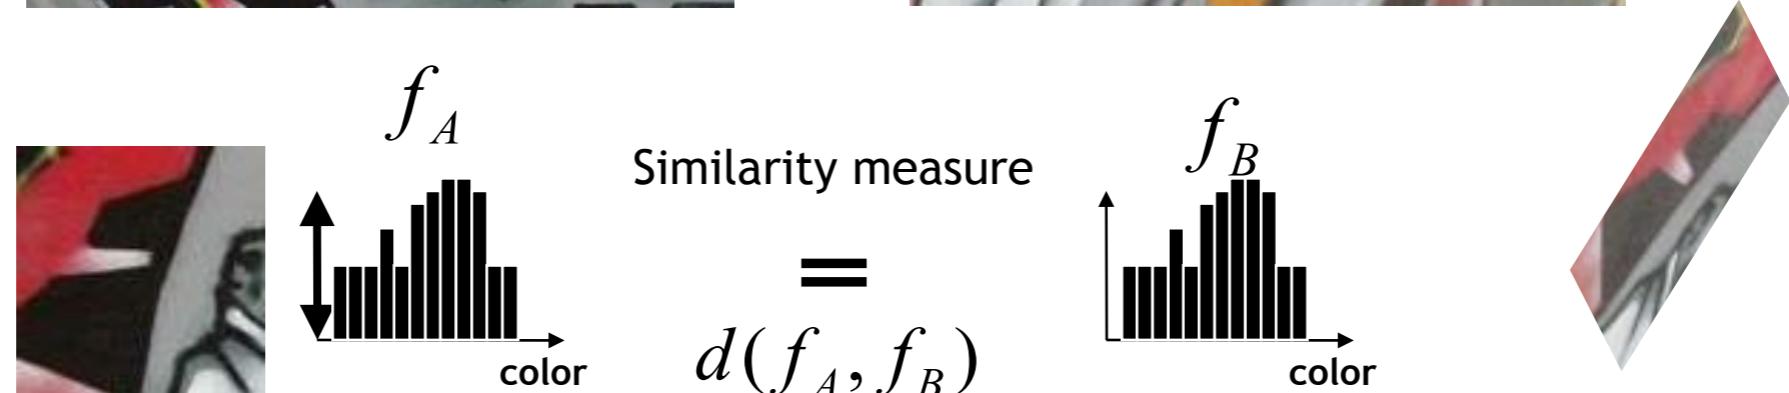
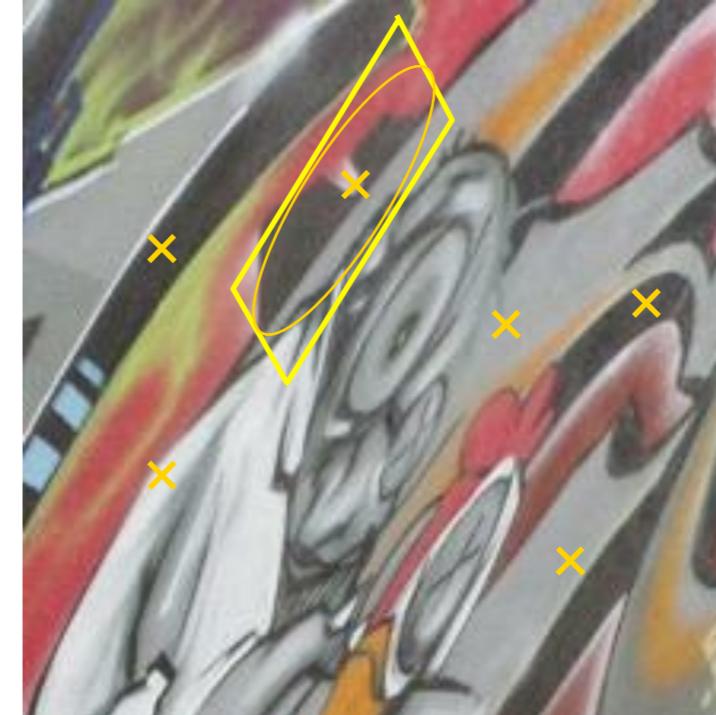
# 2 - Local Descriptors / Features



- ◆ Important properties of local descriptors
  - ◆ Distinctiveness, invariance, robustness, dimensionality, etc...
- ◆ Local descriptors
  - ◆ Image patches
  - ◆ SIFT
  - ◆ Shape context
- ◆ Evaluation criteria

# Local descriptors

- ◆ Detector finds location, scale and shape of interest regions
- ◆ Local descriptors are computed for interest regions

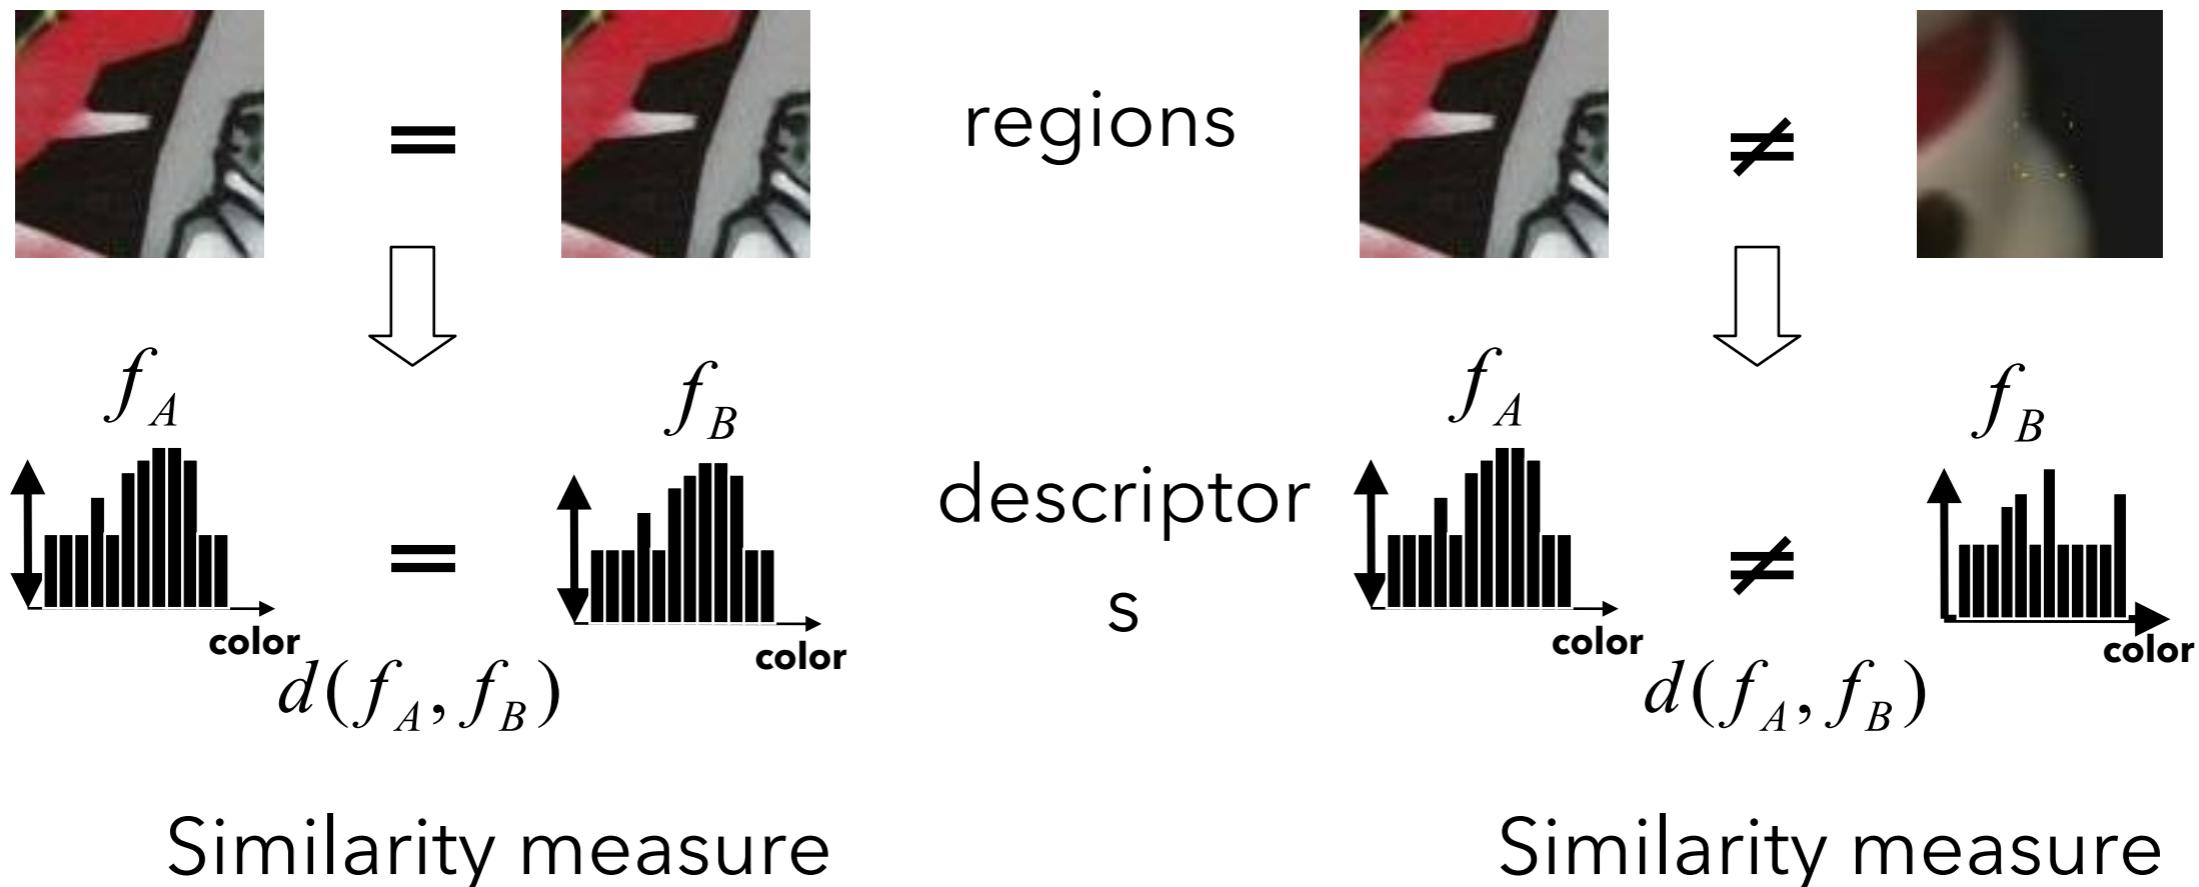


# Local descriptors

- ◆ Important properties of local descriptors

- ◆ Distinctiveness

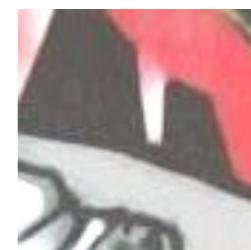
- ◆ Visually similar regions should have similar descriptors
- ◆ Different regions should have different descriptors



# Local descriptors

- ◆ Important properties of local descriptors
  - ◆ Distinctiveness
    - ◆ Visually similar regions should have similar descriptors
    - ◆ Different regions should have different descriptors
  - ◆ Invariance
    - ◆ Visually similar regions should have similar descriptors despite the transformation (geometric, photometric) i.e., rotation, brightness

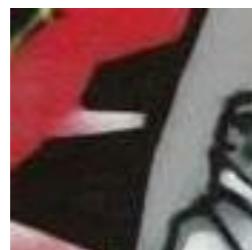
Two ways to obtain invariance



# Local descriptors

- ◆ Important properties of local descriptors
  - ◆ Distinctiveness
    - ◆ Visually similar regions should have similar descriptors
    - ◆ Different regions should have different descriptors
  - ◆ Invariance
    - ◆ Visually similar regions should have similar descriptors despite the transformation (geometric, photometric) i.e., rotation, brightness

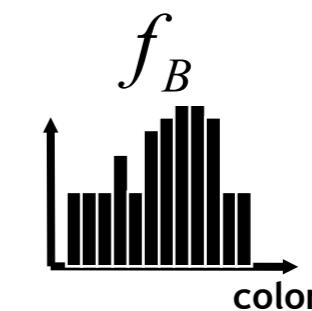
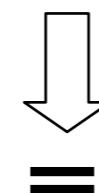
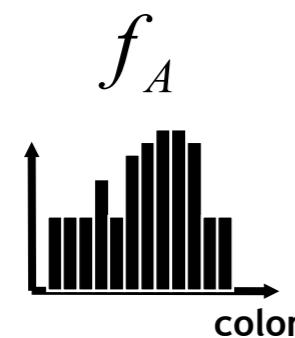
Two ways to obtain invariance



=



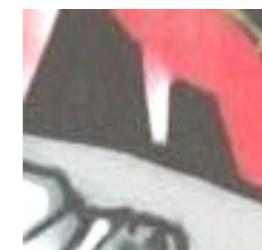
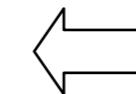
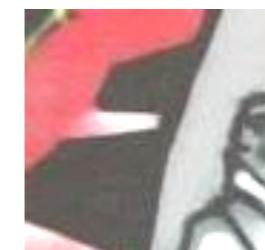
1. Computing invariant descriptor



# Local descriptors

- ◆ Important properties of local descriptors
  - ◆ Distinctiveness
    - ◆ Visually similar regions should have similar descriptors
    - ◆ Different regions should have different descriptors
  - ◆ Invariance
    - ◆ Visually similar regions should have similar descriptors despite the transformation (geometric, photometric) i.e., rotation, brightness

Two ways to obtain invariance

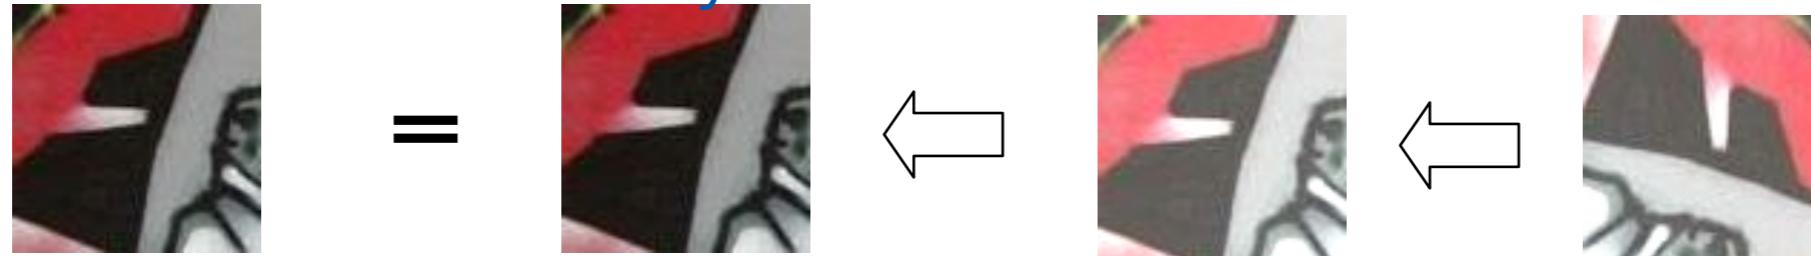


2. Geometric normalization

# Local descriptors

- ◆ Important properties of local descriptors
  - ◆ Distinctiveness
    - ◆ Visually similar regions should have similar descriptors
    - ◆ Different regions should have different descriptors
  - ◆ Invariance
    - ◆ Visually similar regions should have similar descriptors despite the transformation (geometric, photometric) i.e., rotation, brightness

Two ways to obtain invariance

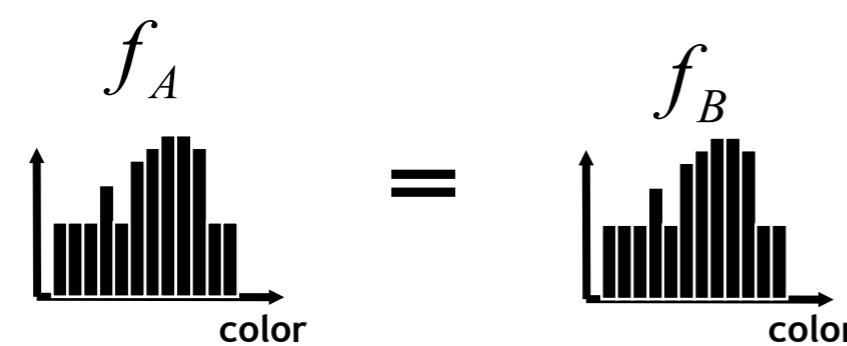
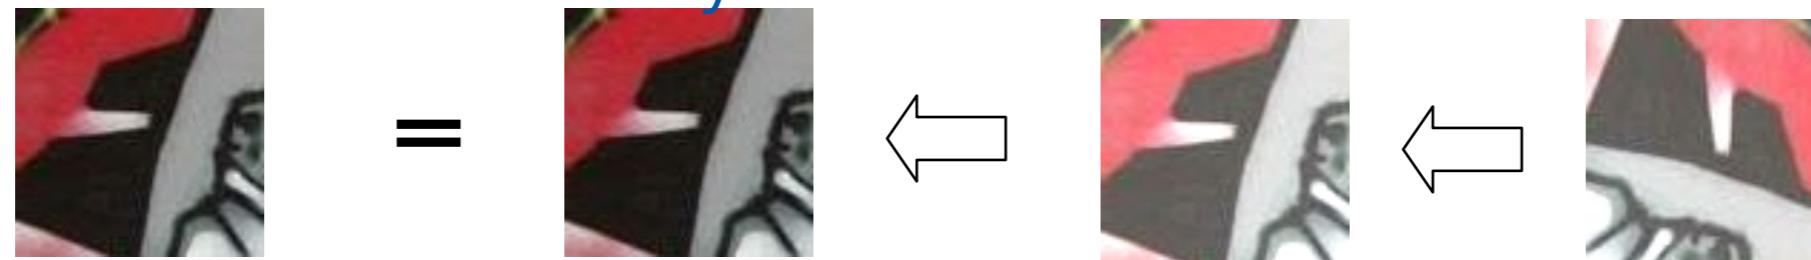


2. Geometric normalization  
+ photometric normalization

# Local descriptors

- ◆ Important properties of local descriptors
  - ◆ Distinctiveness
    - ◆ Visually similar regions should have similar descriptors
    - ◆ Different regions should have different descriptors
  - ◆ Invariance
    - ◆ Visually similar regions should have similar descriptors despite the transformation (geometric, photometric) i.e., rotation, brightness

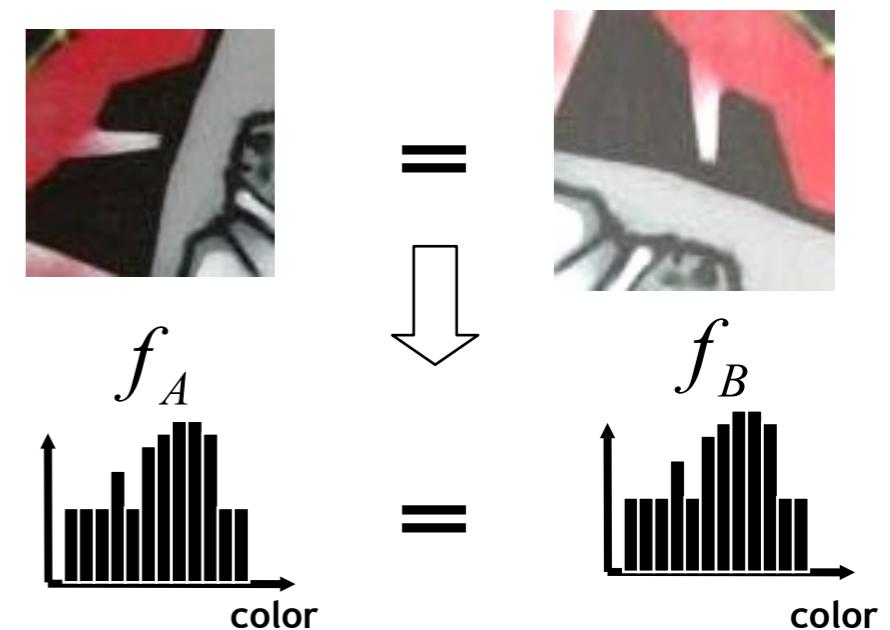
Two ways to obtain invariance



2. Geometric normalization  
+ photometric normalization  
+ computing descriptor

# Local descriptors

- ◆ Important properties of local descriptors
  - ◆ Distinctiveness
    - ◆ Visually similar regions should have similar descriptors
    - ◆ Different regions should have different descriptors
  - ◆ Invariance
    - ◆ Visually similar regions should have similar descriptors despite the transformation (geometric, photometric) i.e., rotation, brightness
  - ◆ Robustness
    - ◆ Visually similar regions should have similar descriptors despite noise (geometric, photometric)



# Local descriptors

- ◆ Important properties of local descriptors
  - ◆ Distinctiveness
    - ◆ Visually similar regions should have similar descriptors
    - ◆ Different regions should have different descriptors
  - ◆ Invariance
    - ◆ Visually similar regions should have similar descriptors despite the transformation (geometric, photometric) i.e., rotation, brightness
  - ◆ Robustness
    - ◆ Visually similar regions should have similar descriptors despite the noise (geometric, photometric)
  - ◆ Dimensionality
    - ◆ Descriptors should be low dimensional, i.e. small number of histogram bins.
      - ◆ Efficiency (large databases)
      - ◆ Generalization property

# 2 - Local Descriptors / Features



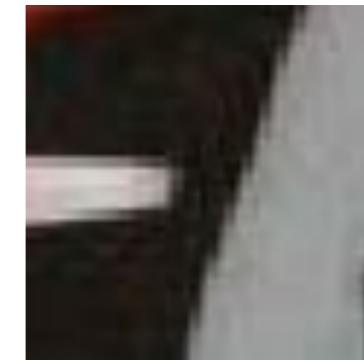
- ◆ Important properties of local descriptors
  - ◆ Distinctiveness, invariance, robustness, dimensionality, etc...
- ◆ Local descriptors
  - ◆ Image patches
  - ◆ SIFT
  - ◆ Shape context
- ◆ Evaluation criteria

# Local Descriptors

- ◆ Image patch (vector of pixels)
  - ◆ Invariant only when computed on normalized patches



$$f_1 = \begin{bmatrix} I_1(0,0) \\ I_1(0,1) \\ I_1(0,2) \\ \vdots \end{bmatrix}$$

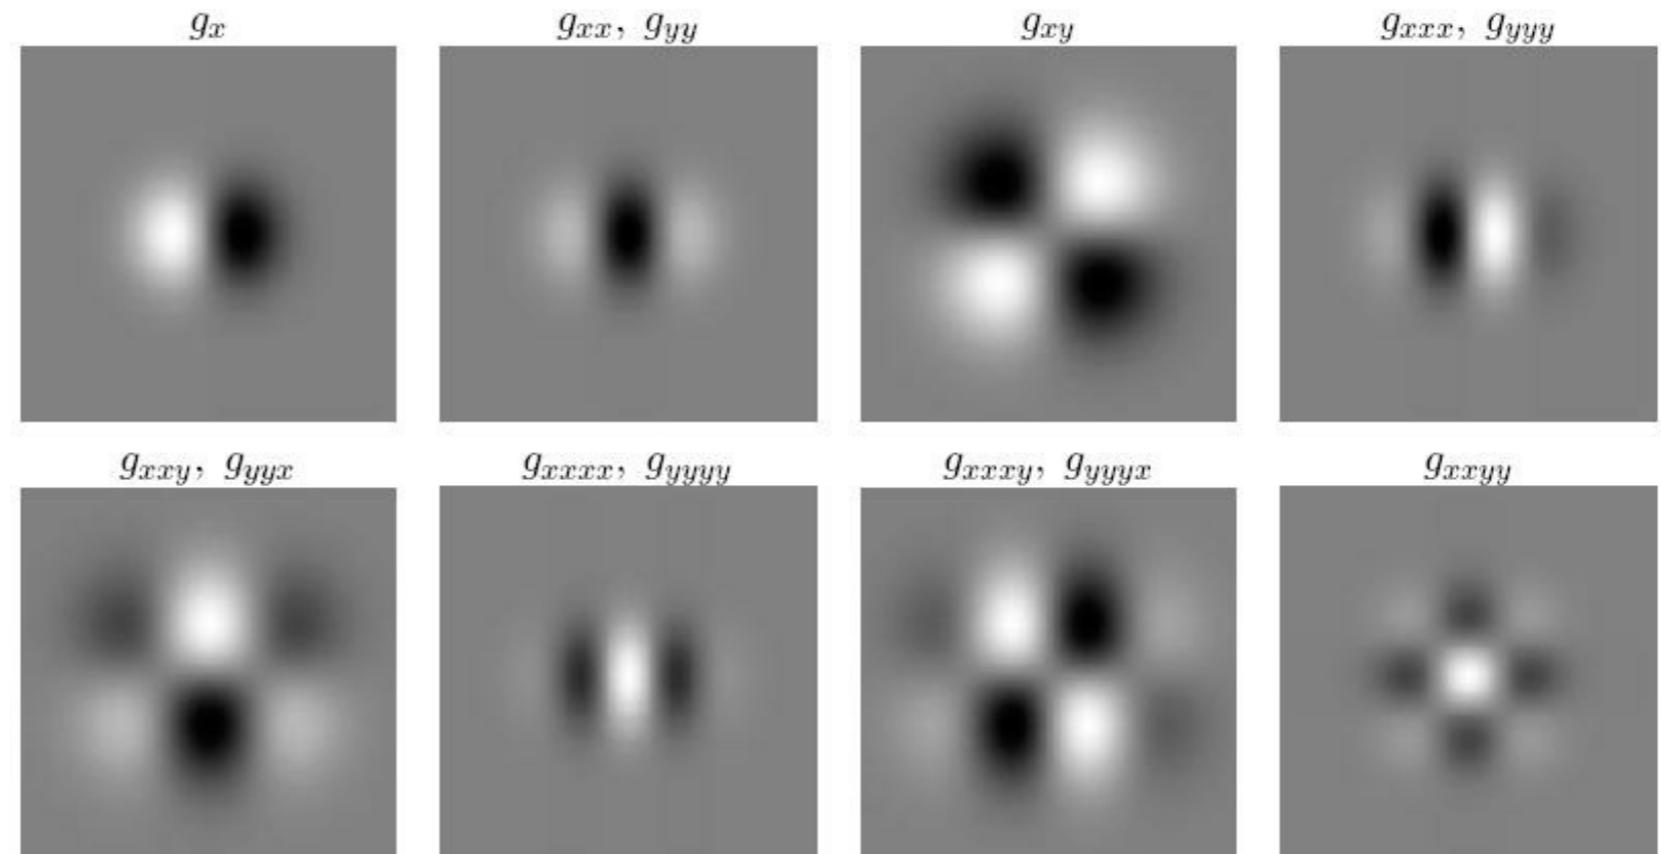


$$f_2 = \begin{bmatrix} I_2(0,0) \\ I_2(0,1) \\ I_2(0,2) \\ \vdots \end{bmatrix}$$

- ◆ Distinctive & easy to implement
- ◆ But high-dimensional and not very robust
- ◆ Match descriptors using:
  - ◆ Euclidean distance
  - ◆ Cross correlation

# Local Descriptors

- ◆ Bank of filter responses ("jet")
  - ◆ Invariant only when computed on normalized patch



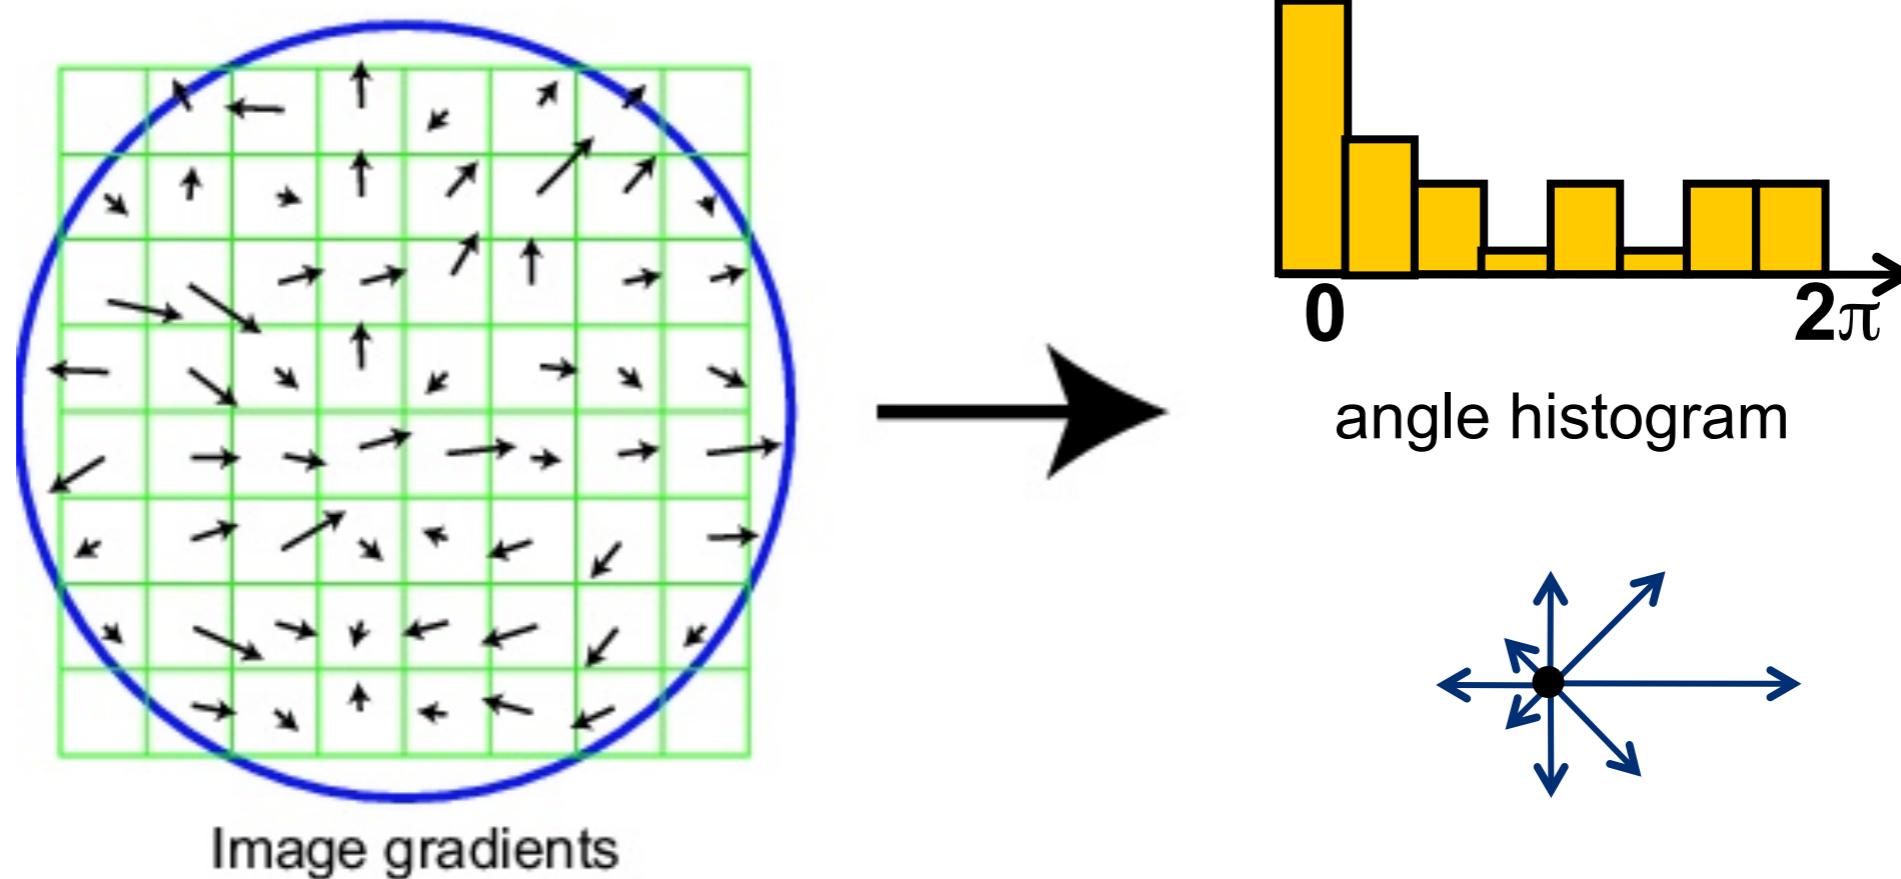
$$[g_x, g_y, g_{xx}, g_{xy}, g_{yy}, \dots]$$

Responses of Gaussian derivatives

# Scale Invariant Feature Transform

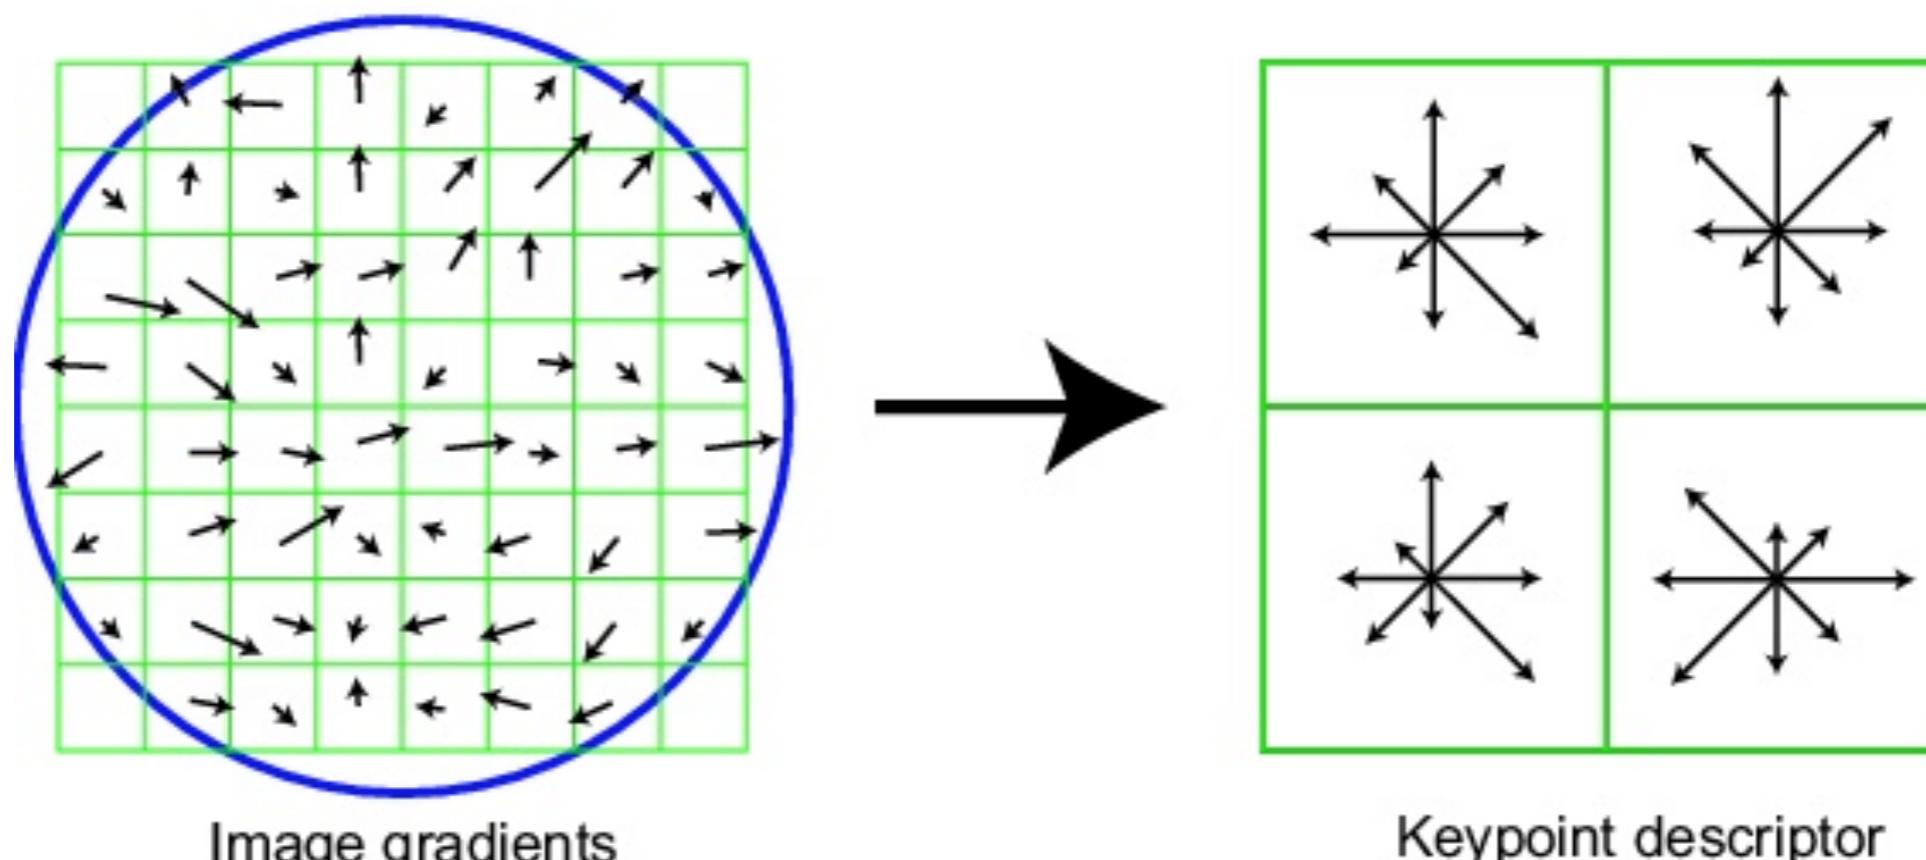
## ◆ Basic idea:

- ◆ Take 16x16 square window around detected feature
- ◆ Compute edge orientation (angle of the gradient - 90°) for each pixel
- ◆ Throw out weak edges (threshold gradient magnitude)
- ◆ Create histogram of surviving edge orientations



# SIFT Descriptor

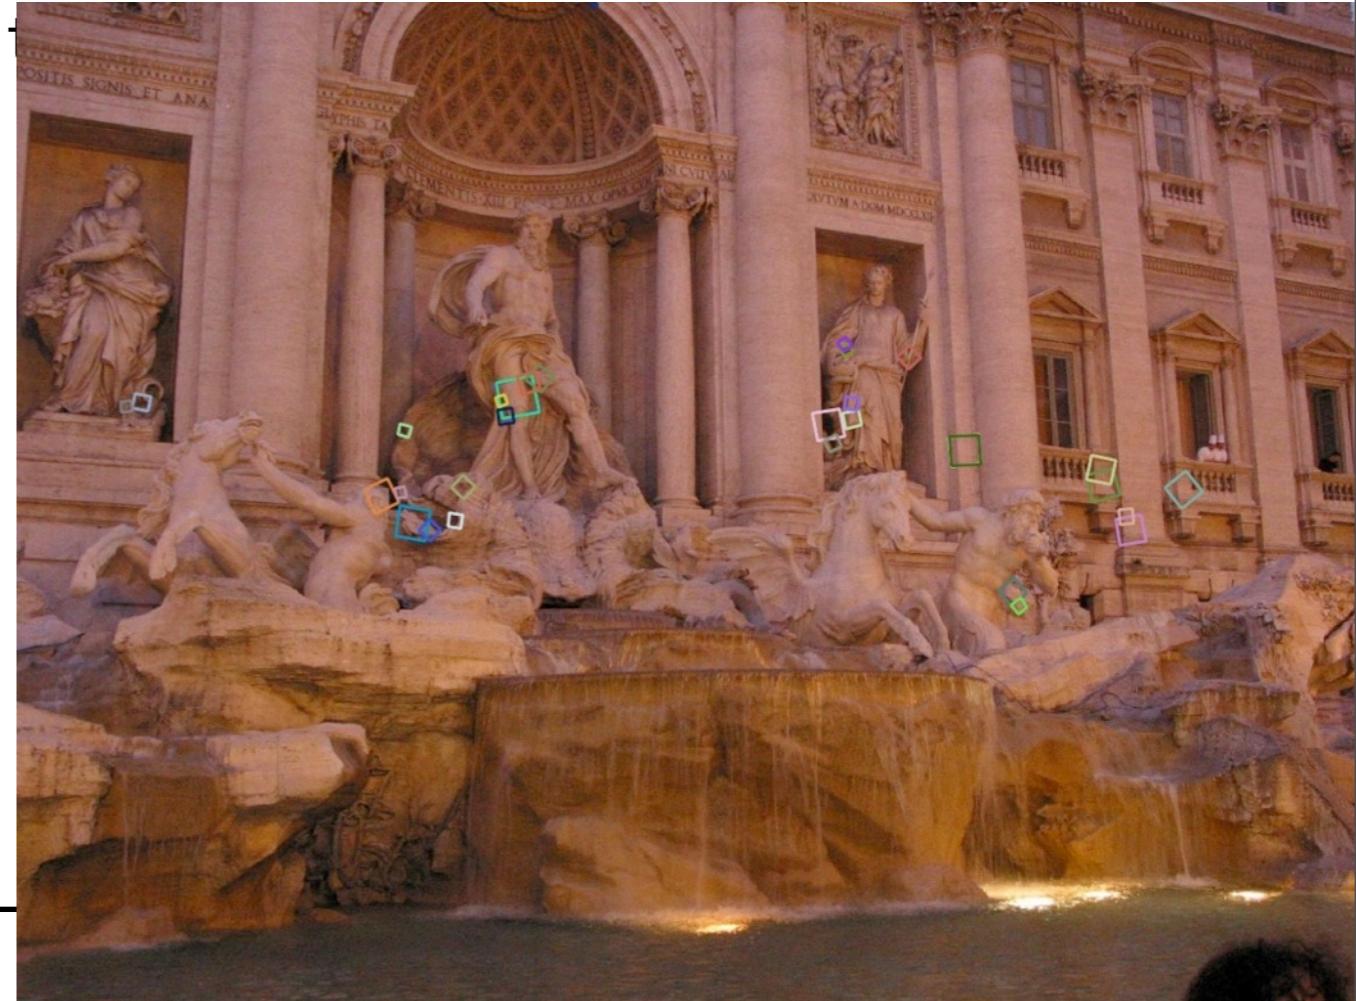
- ◆ Full version:
  - ◆ Divide the 16x16 window into a 4x4 grid of cells (2x2 case shown below)
  - ◆ Compute an orientation histogram for each cell
  - ◆  $16 \text{ cells} * 8 \text{ orientations} = 128 \text{ dimensional descriptor}$



[Lowe]

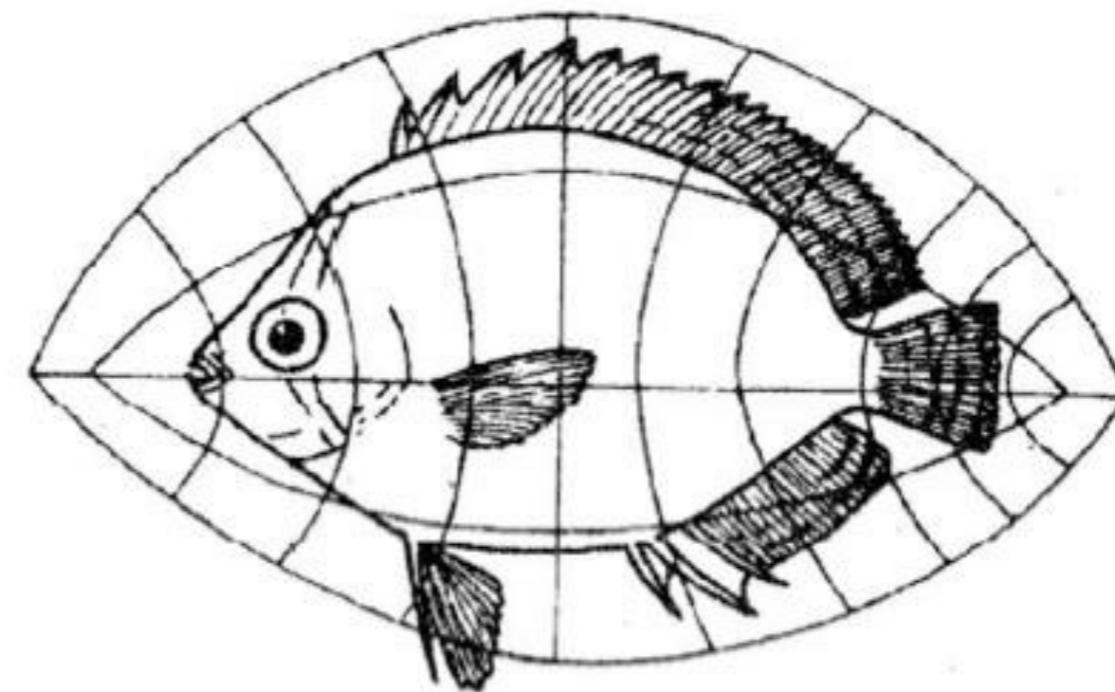
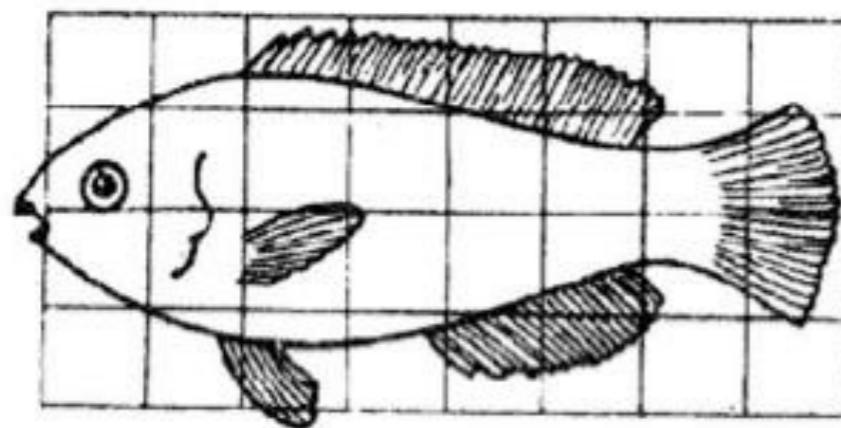
# Properties of SIFT

- ◆ Extraordinarily robust matching technique
  - ◆ Can handle changes in viewpoint
    - ◆ Up to about 60 degree out of plane rotation
  - ◆ Can handle significant changes in illumination
    - ◆ Sometimes even day vs. night (below)



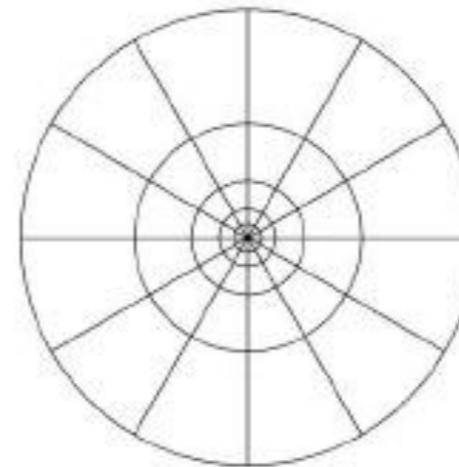
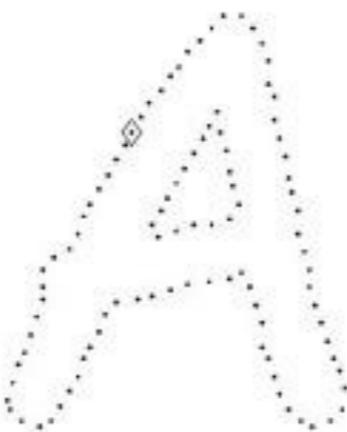
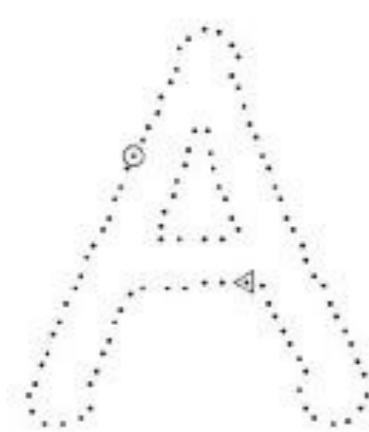
# Shape context

- ◆ Motivation

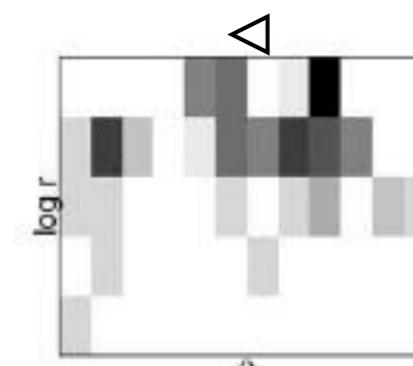
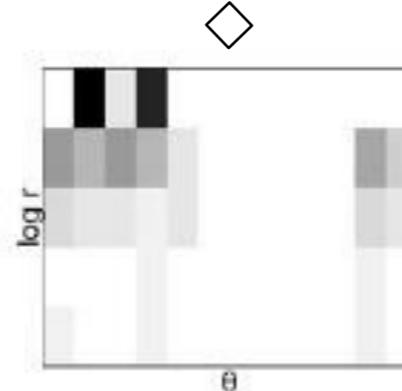
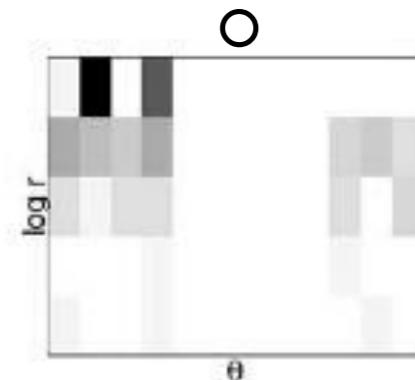


# Shape context

- ◆ Represent histogram of edge locations



Log polar  
coordinate  
system



Histogram of edge points in log-polar coordinate system

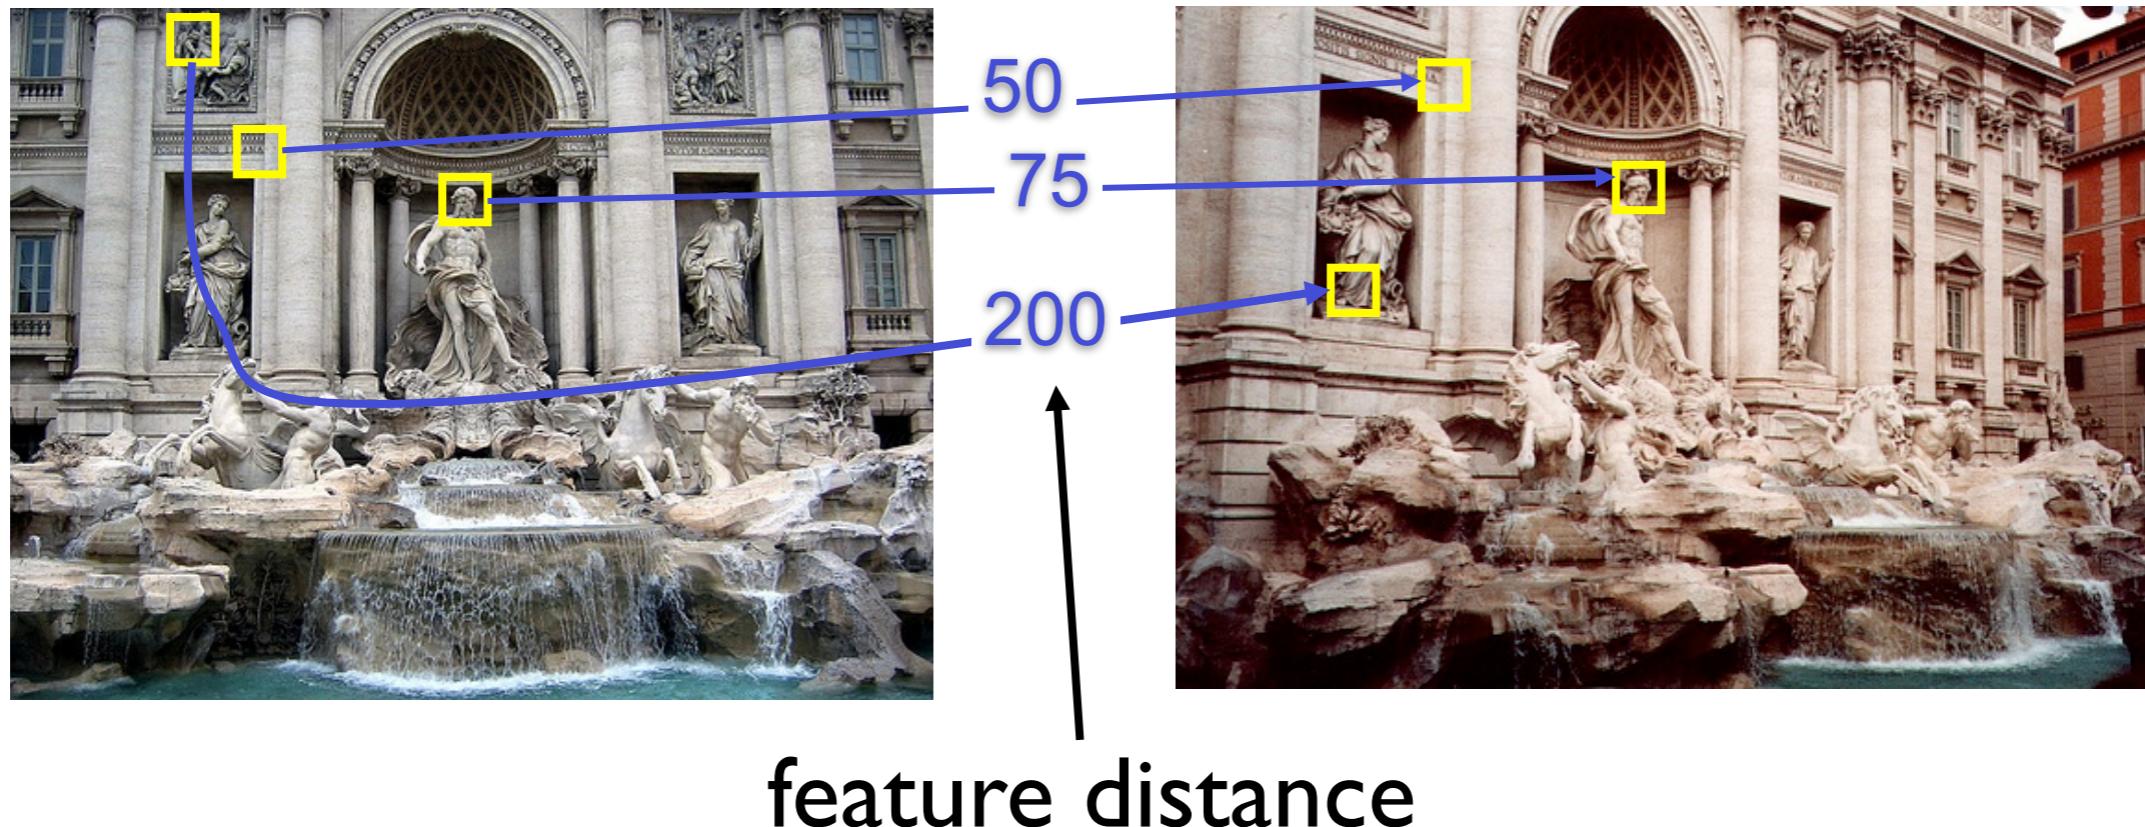
# 2 - Local Descriptors / Features



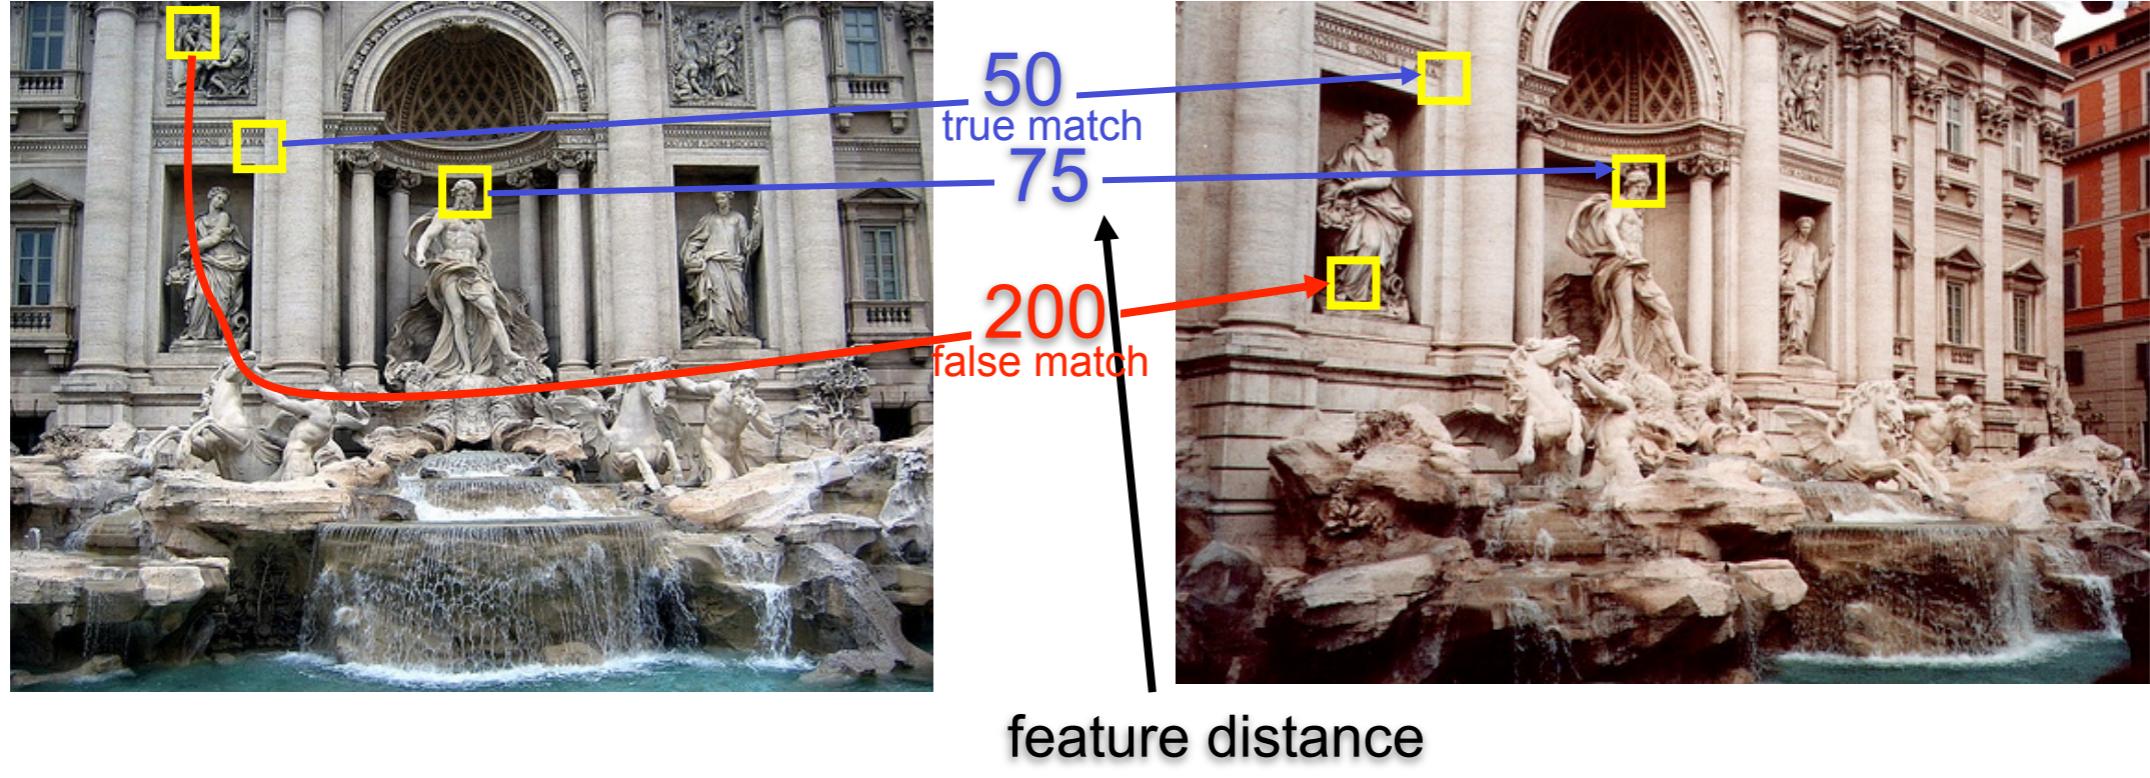
- ◆ Important properties of local descriptors
  - ◆ Distinctiveness, invariance, robustness, dimensionality, etc...
- ◆ Local descriptors
  - ◆ Image patches
  - ◆ SIFT
  - ◆ Shape context
- ◆ Evaluation criteria

# Evaluating the results

- ◆ How can we measure the performance of a feature matcher (descriptor + distance/classifier)?



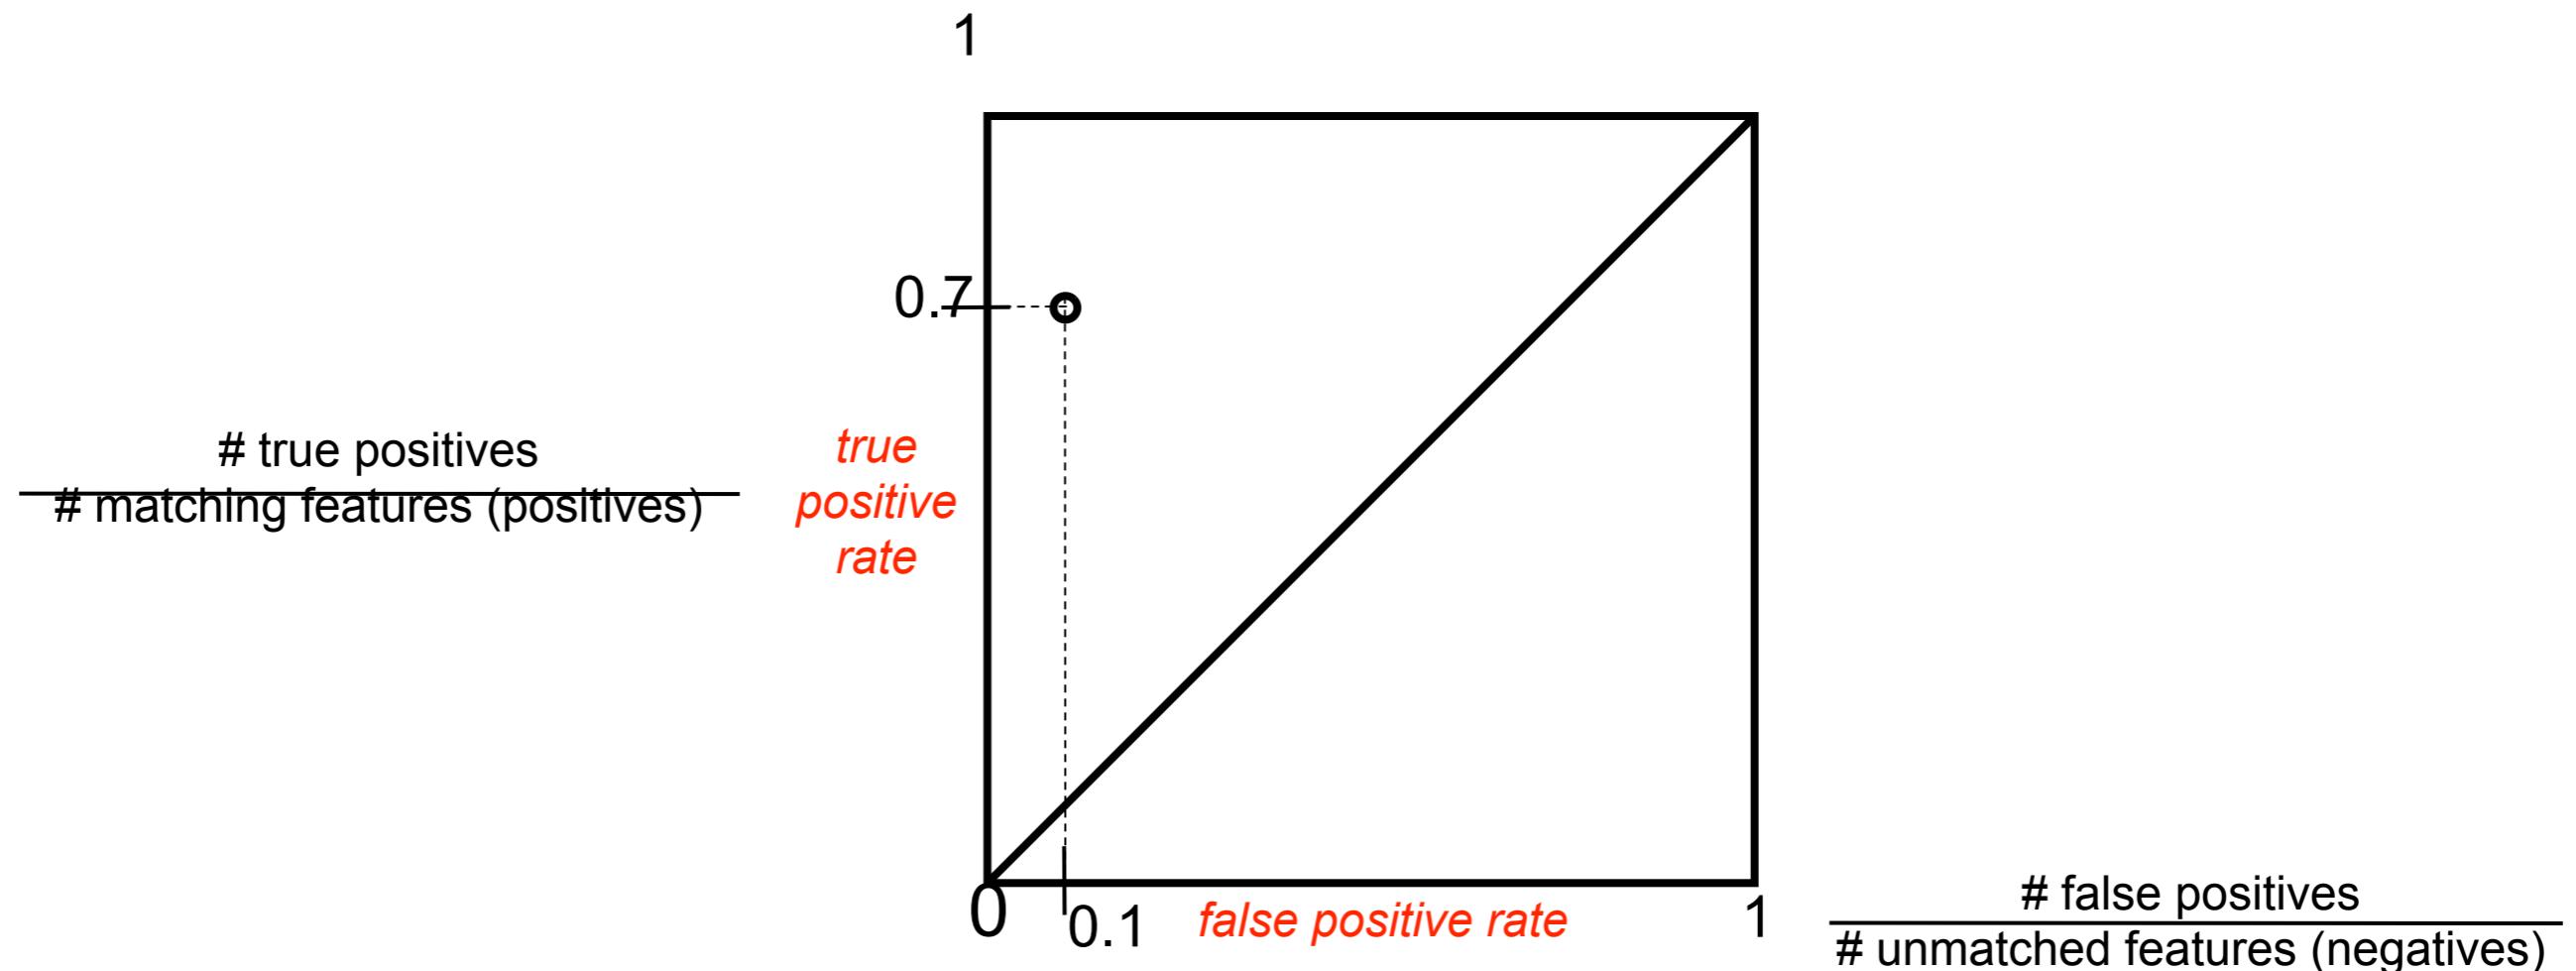
# True/false positives



- ◆ The distance threshold affects performance
  - ◆ True positives = # of detected matches that are correct
    - ◆ Suppose we want to maximize these—how to choose threshold?
  - ◆ False positives = # of detected matches that are incorrect
    - ◆ Suppose we want to minimize these—how to choose threshold?

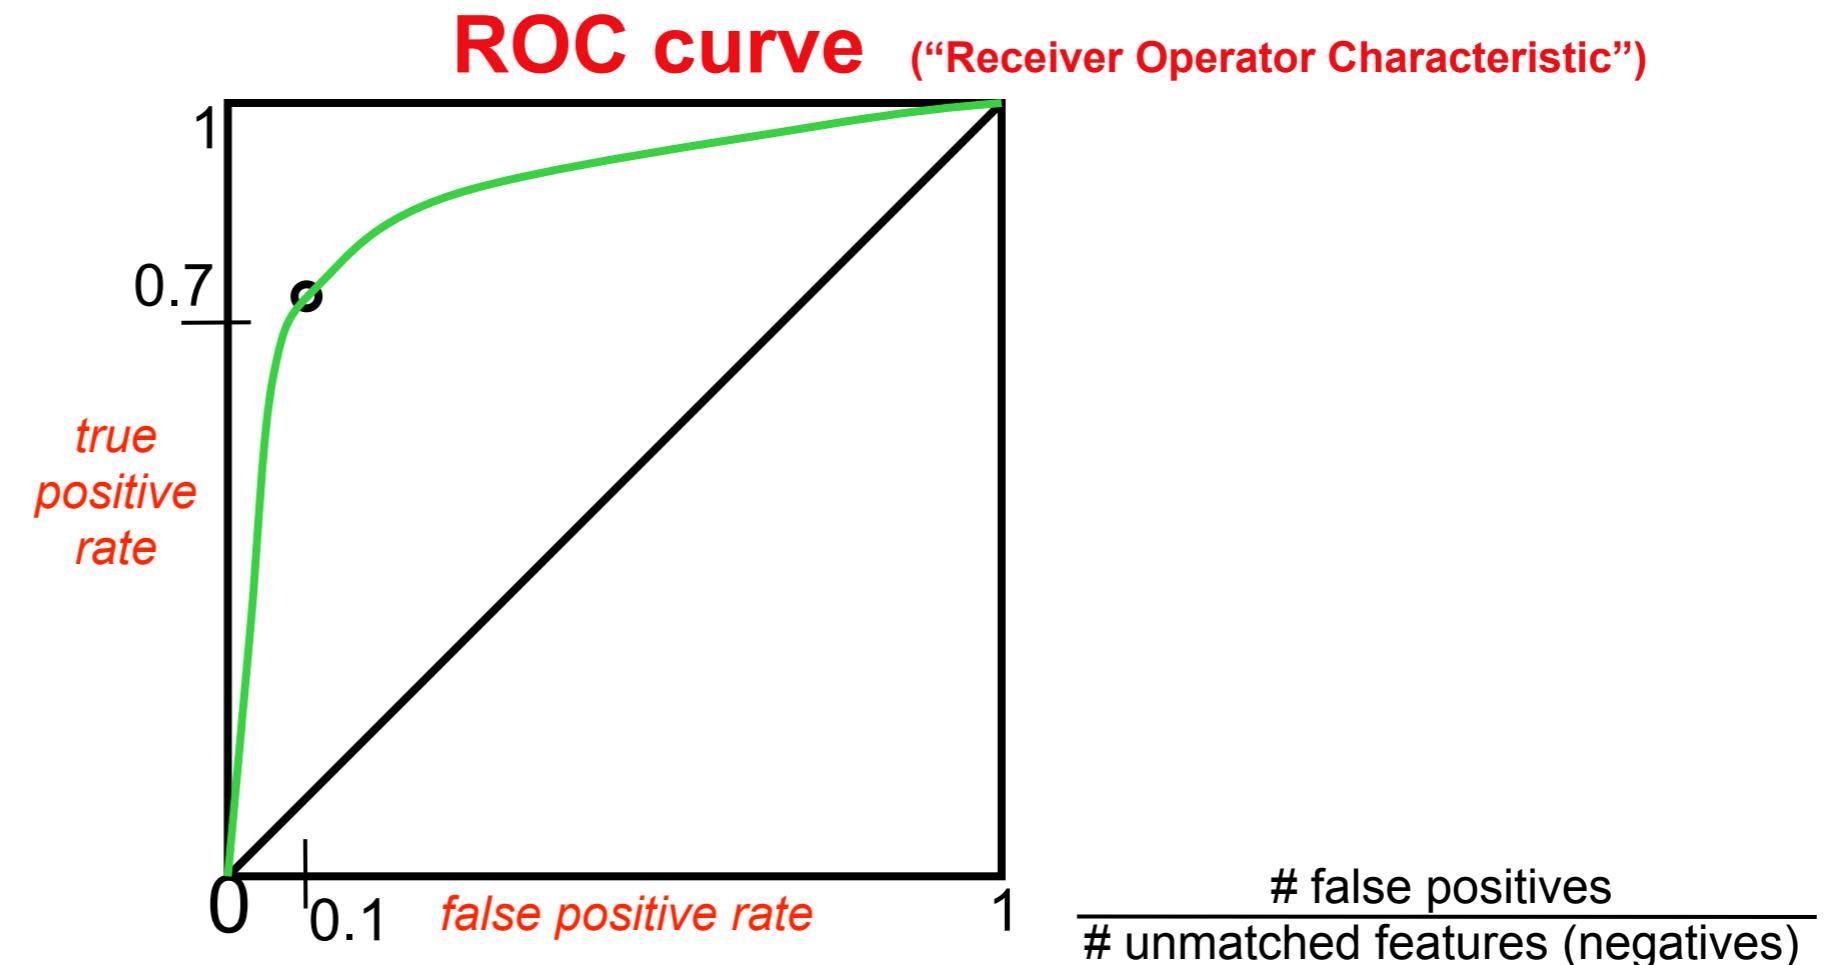
# Evaluating the results

- How can we measure the performance of a feature matcher (descriptor + distance/classifier)?



# Evaluating the results

- How can we measure the performance of a feature matcher (descriptor + distance/classifier)?

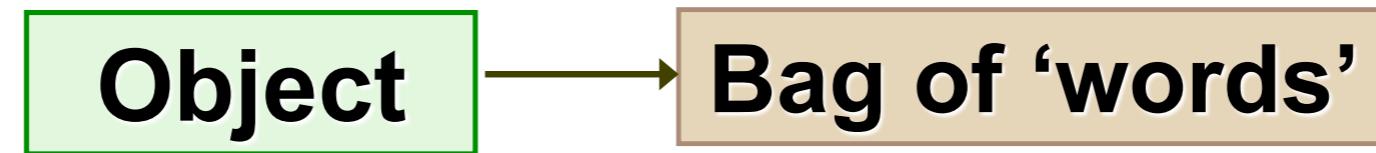


- ROC Curves
- Generated by counting # current/incorrect matches, for different thresholds
- Want to maximize area under the curve (AUC)

# Overview Today

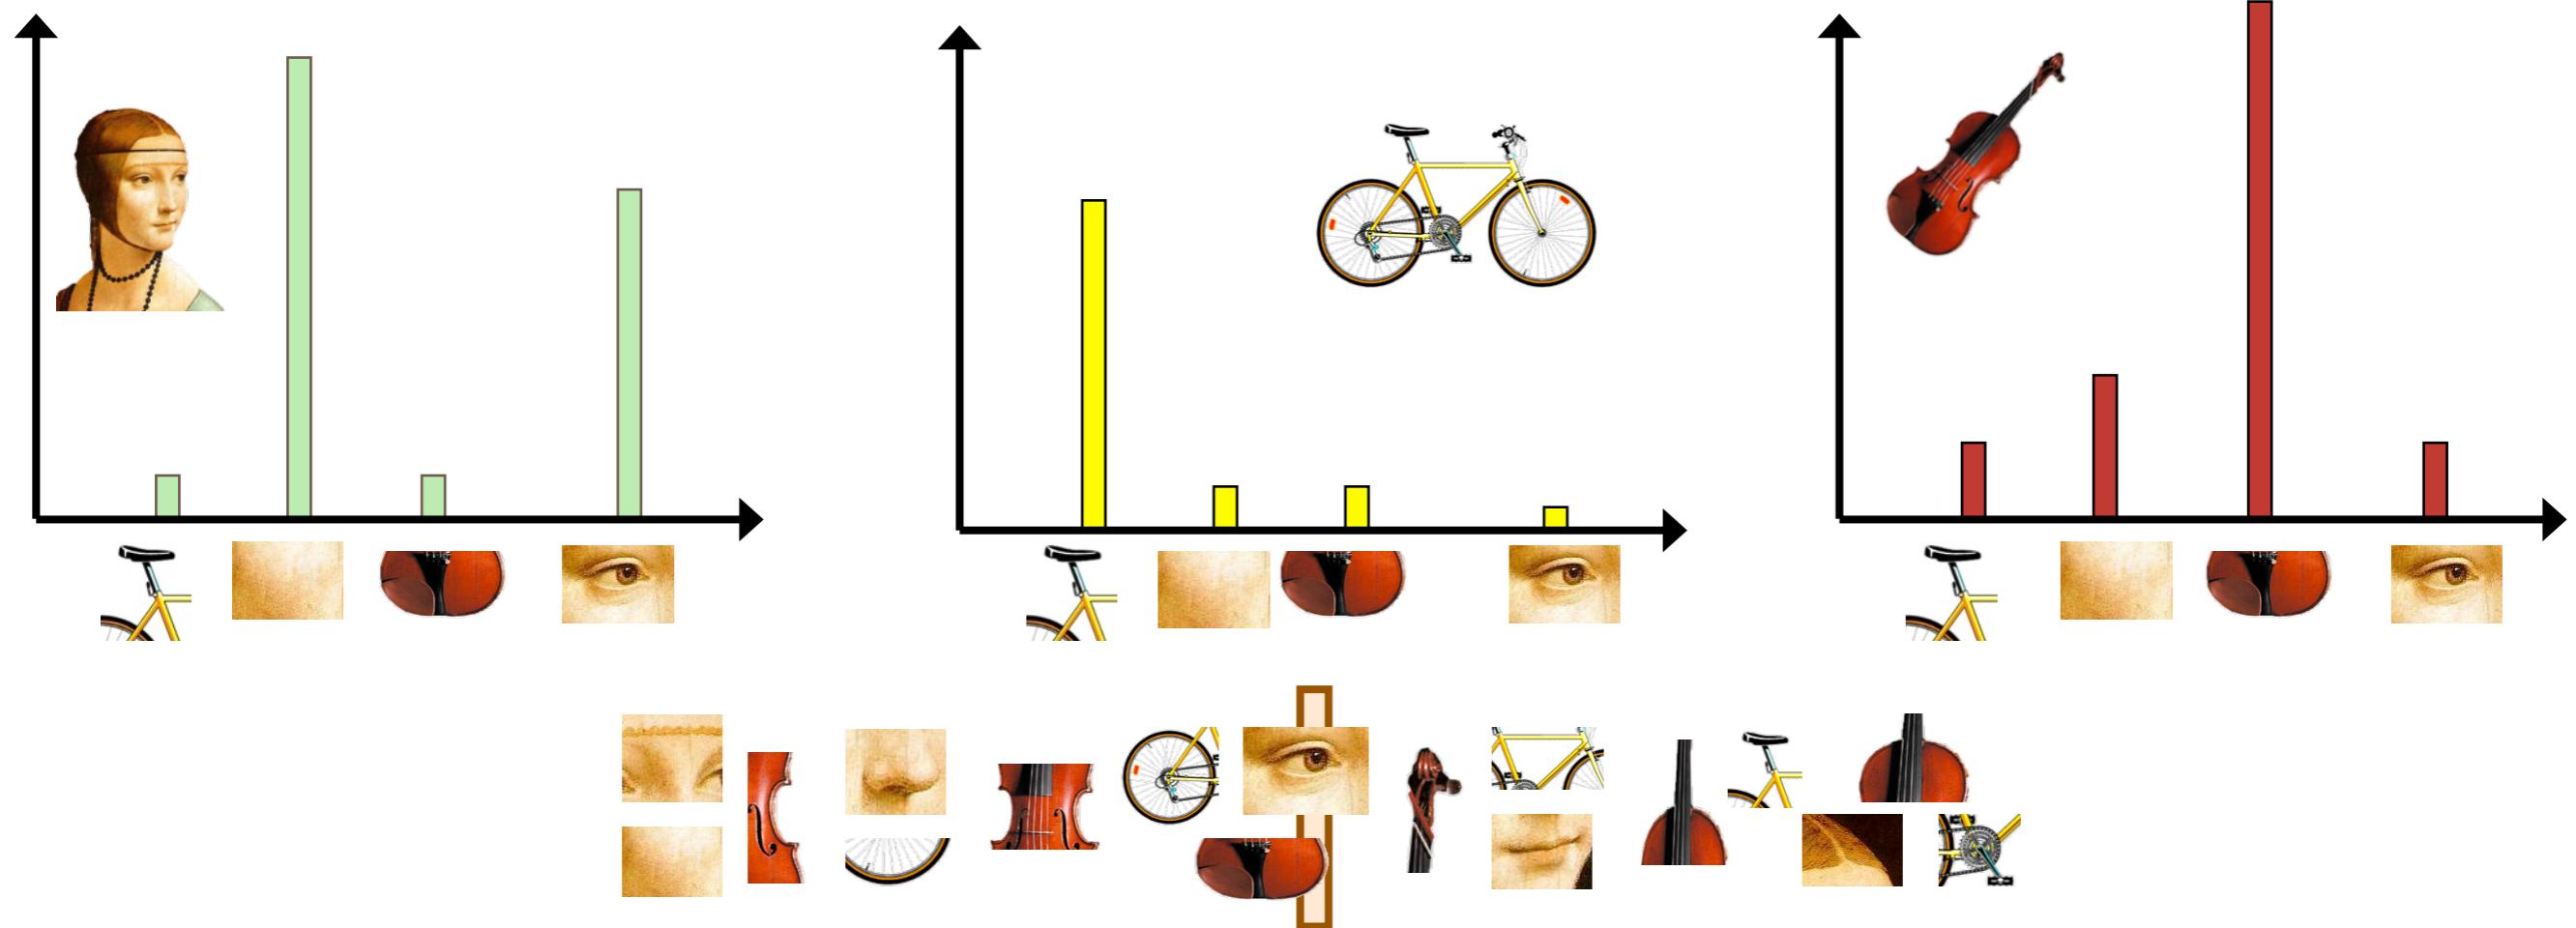
- ◆ 1 - Local Interest Point Detection
  - ◆ finding discriminant points (Harris, Hessian)
  - ◆ scale invariant interest point detection
- ◆ 2 - Local Descriptors (= Features)
- ◆ **3 - Bag-of-Words Model (BoW)**
  - ◆ **BoW for Object Categorization**

# 3 - Bag-of-Words Model (BoW) now for Object Categorization

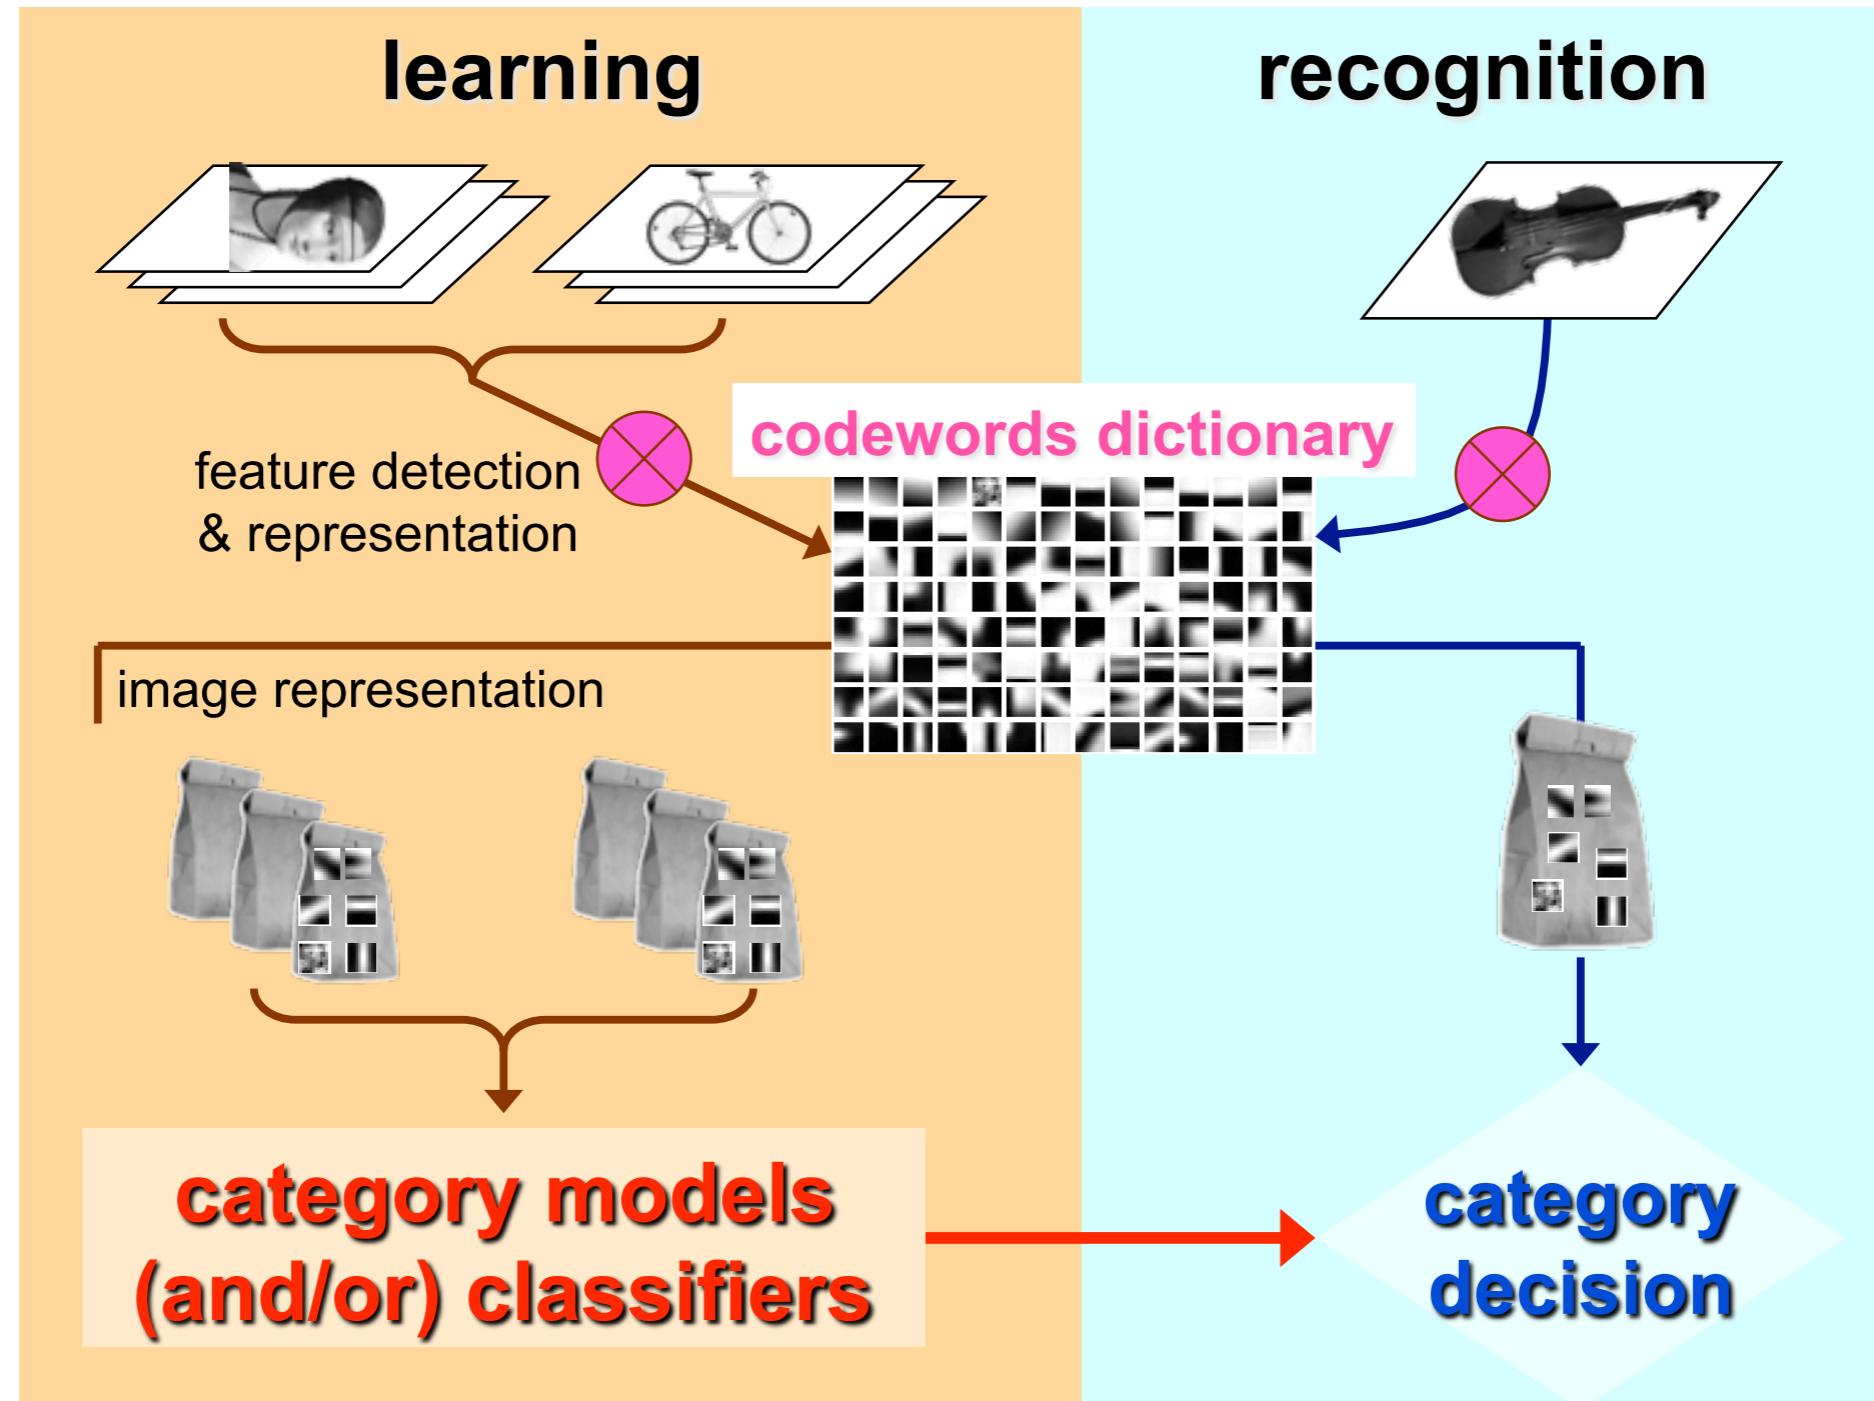


# Bag of Words

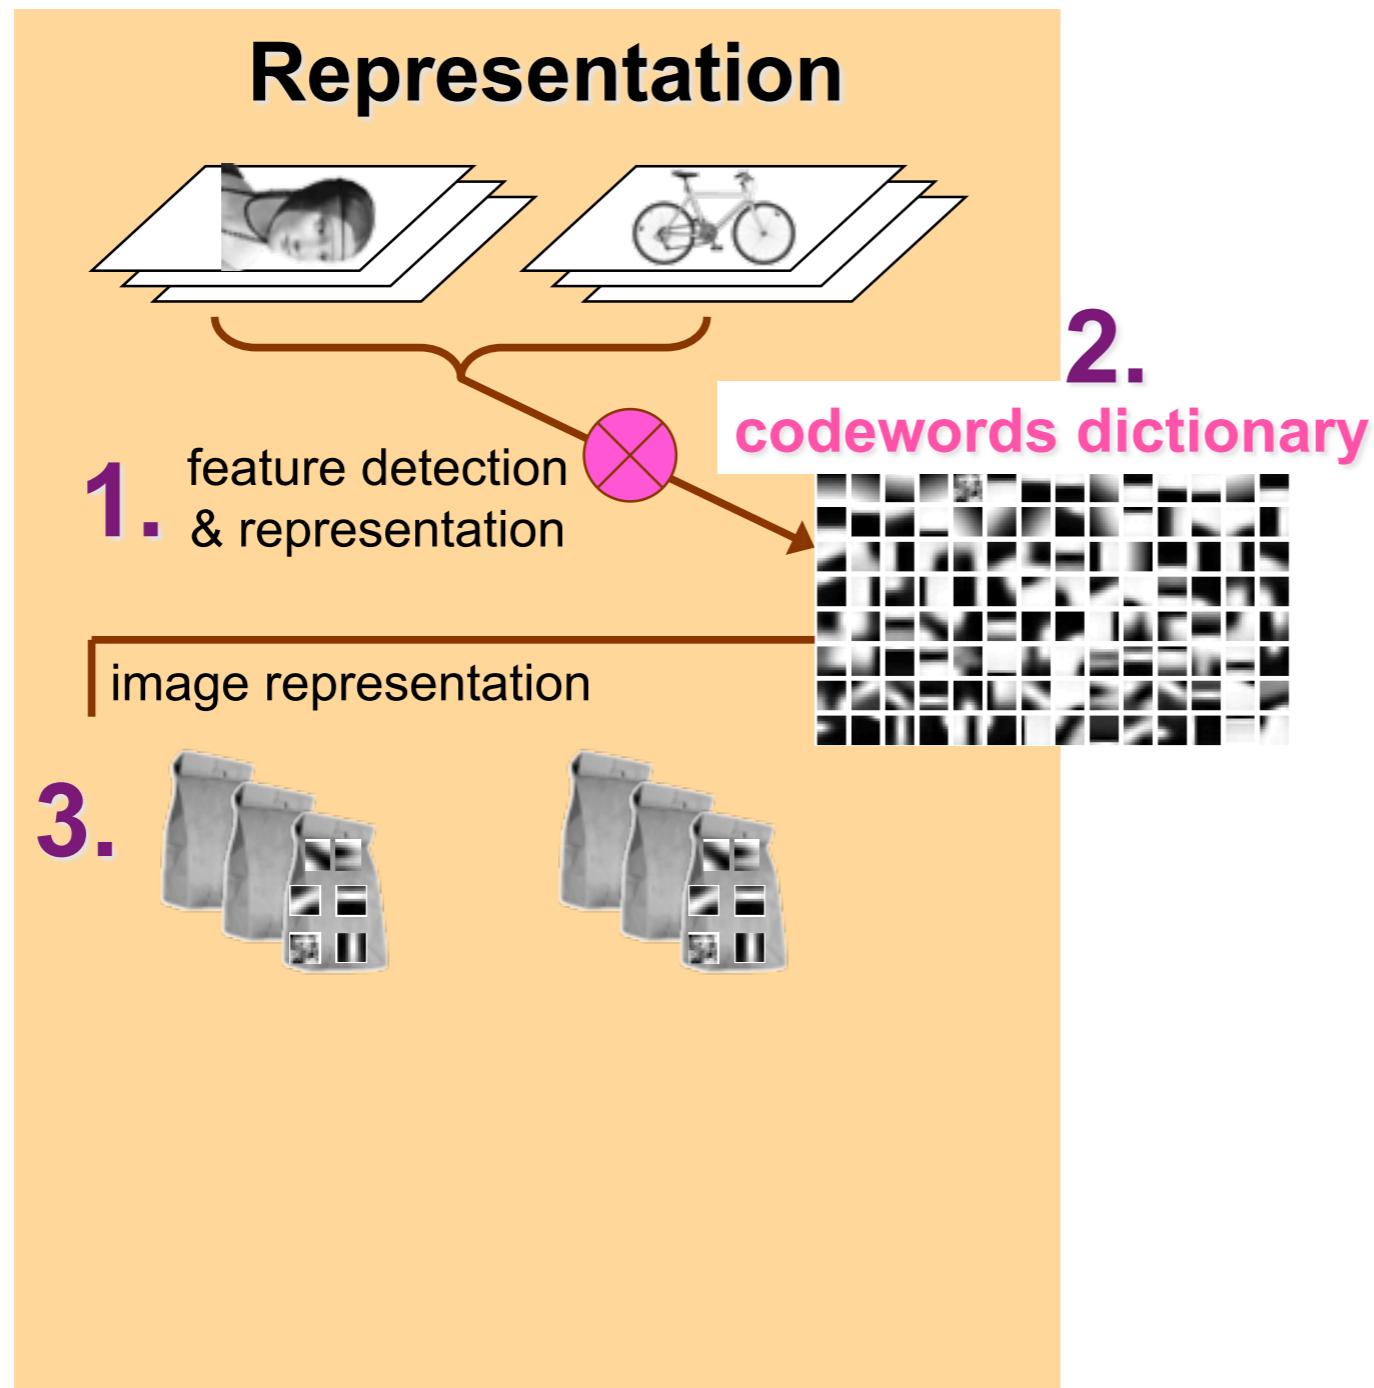
- ◆ Independent features
- ◆ Histogram representation



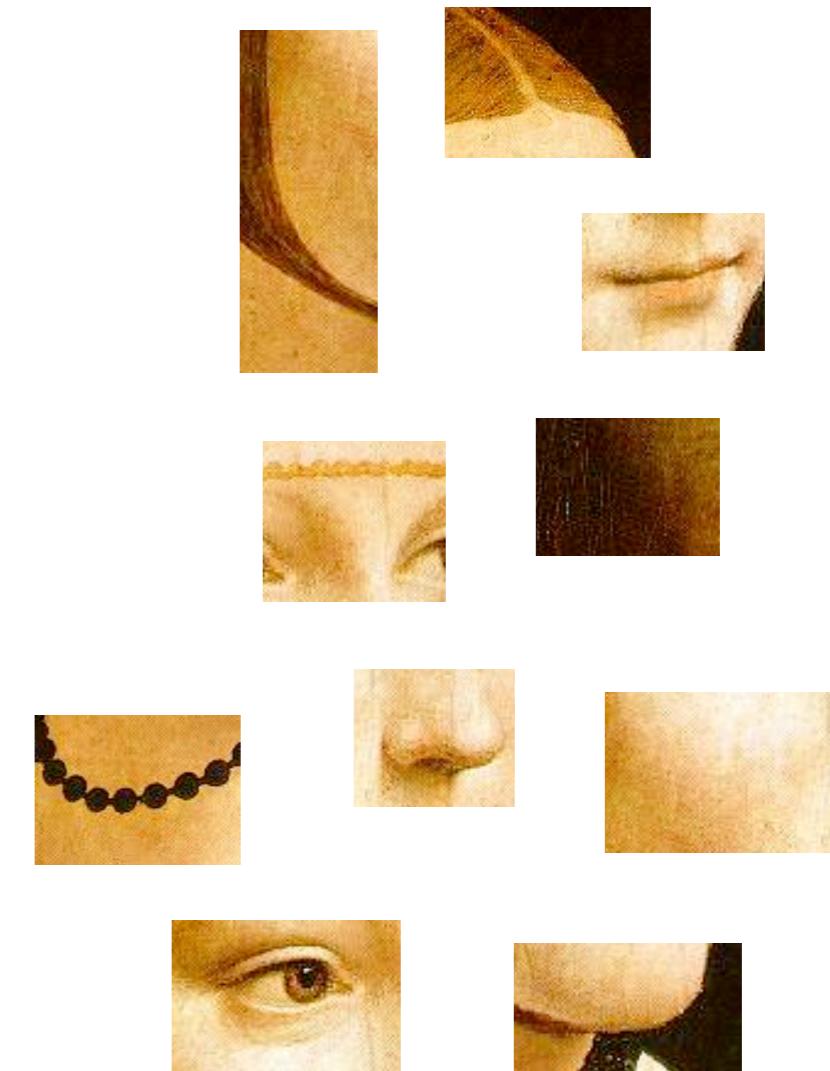
# Bag-of-Words Model: Overview



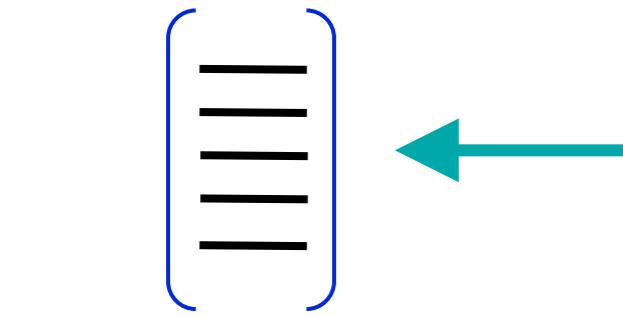
# Bag-of-Words Model: Object Representation & Learning



# BoW-1. Feature detection and representation

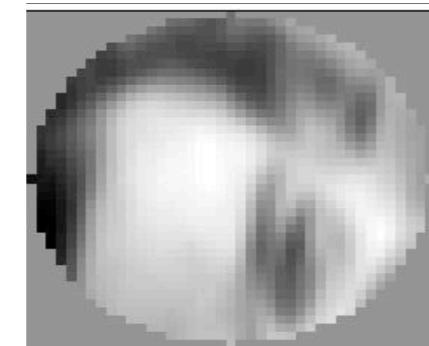


# BoW-1. Feature detection and representation

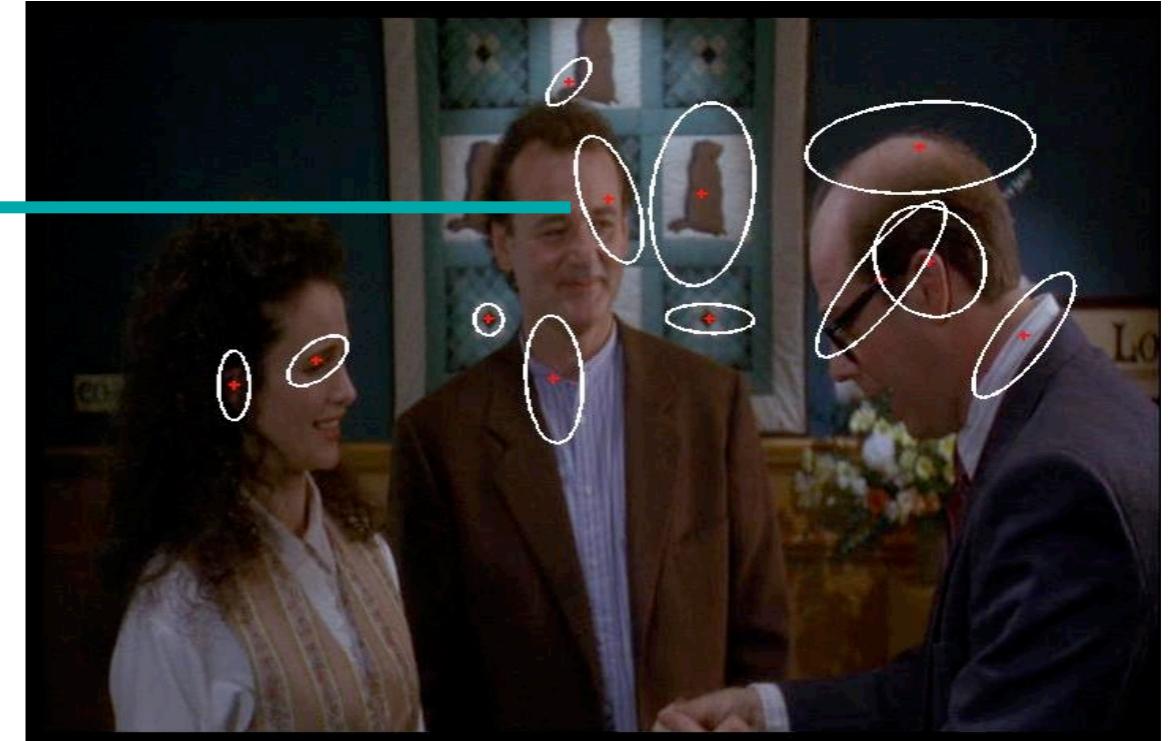


**Compute  
descriptor**

e.g. SIFT, shape  
context, etc.



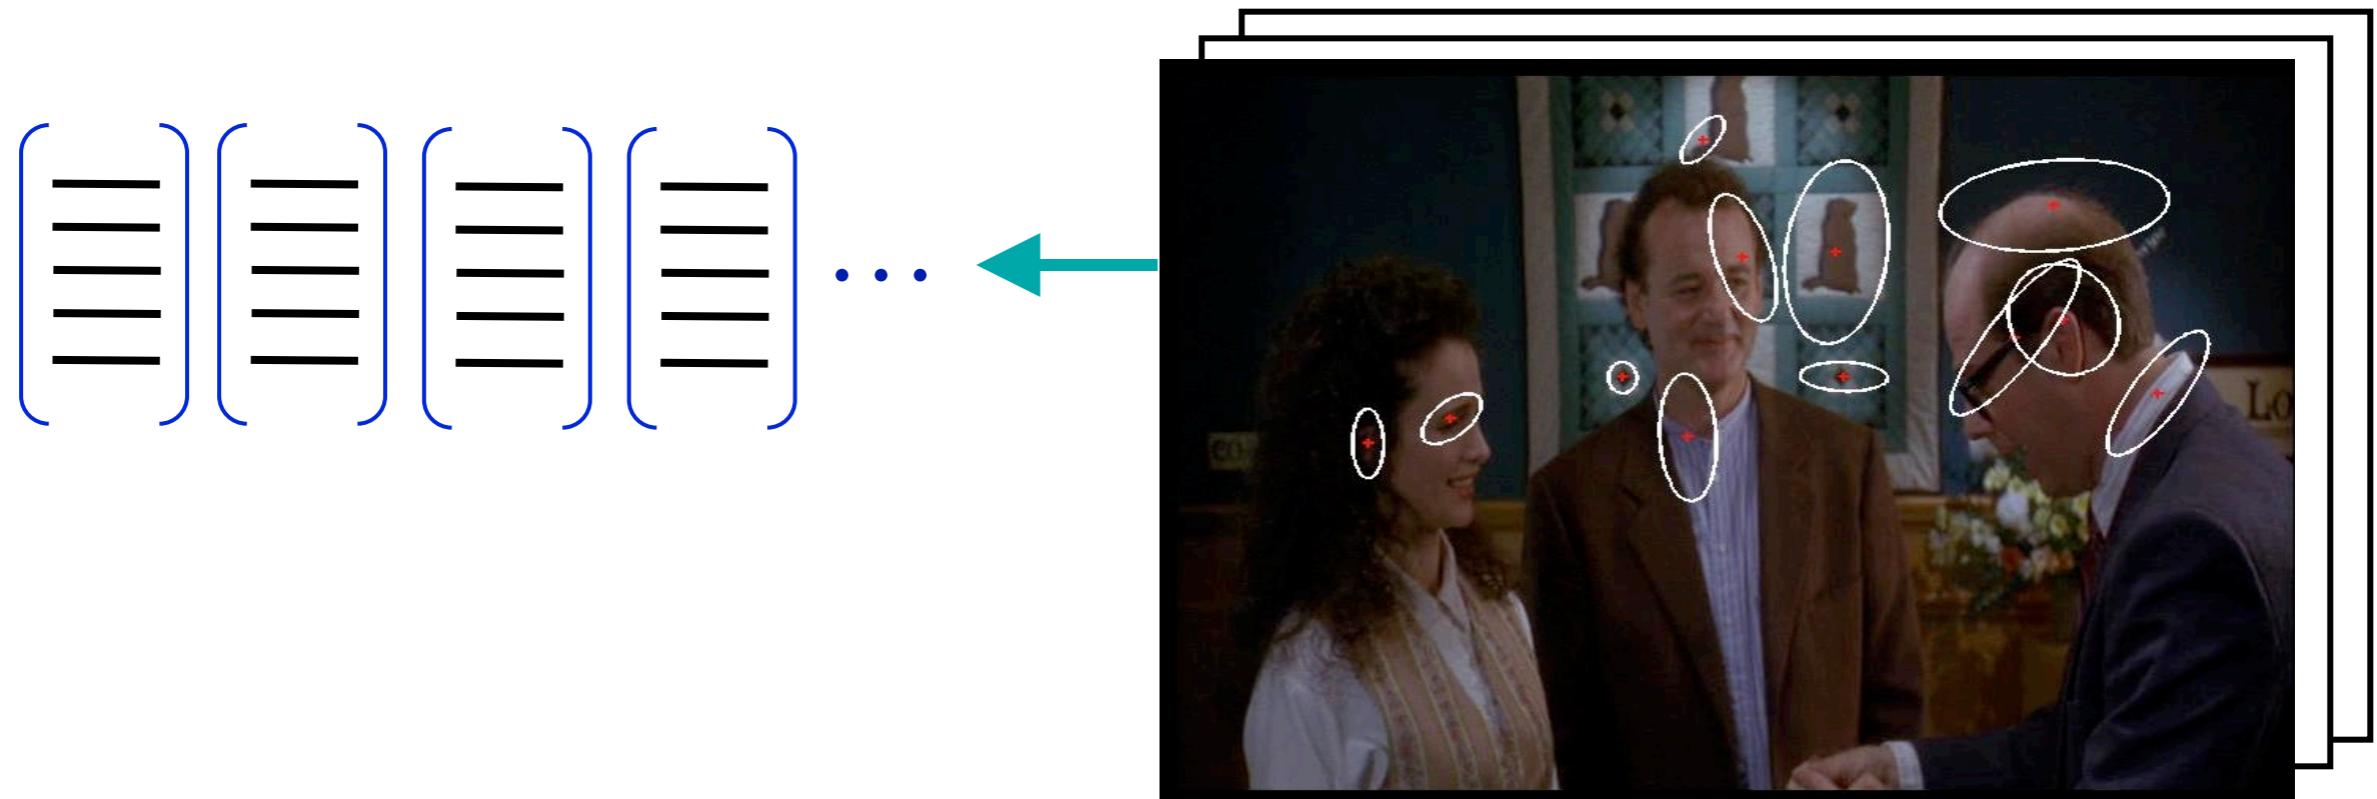
**Normalize  
patch**



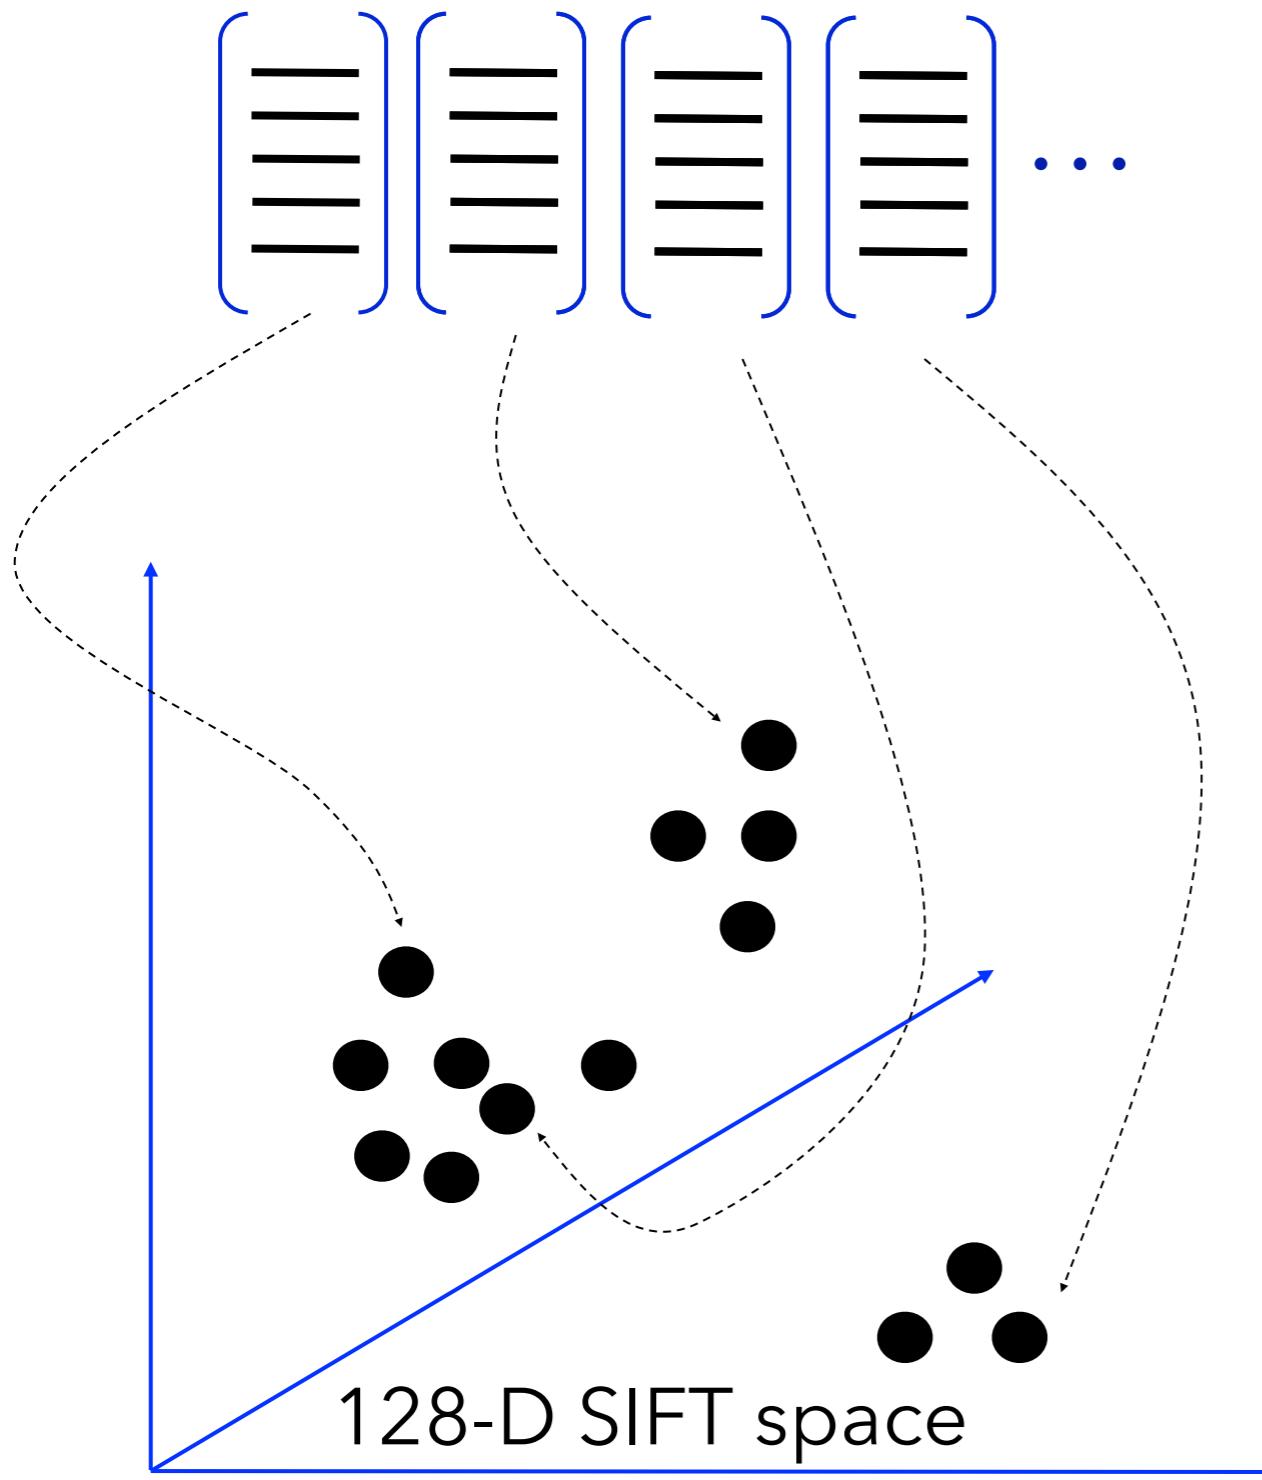
**Detect patches**

Local interest operator  
(e.g. Harris-Laplace)  
or regular grid

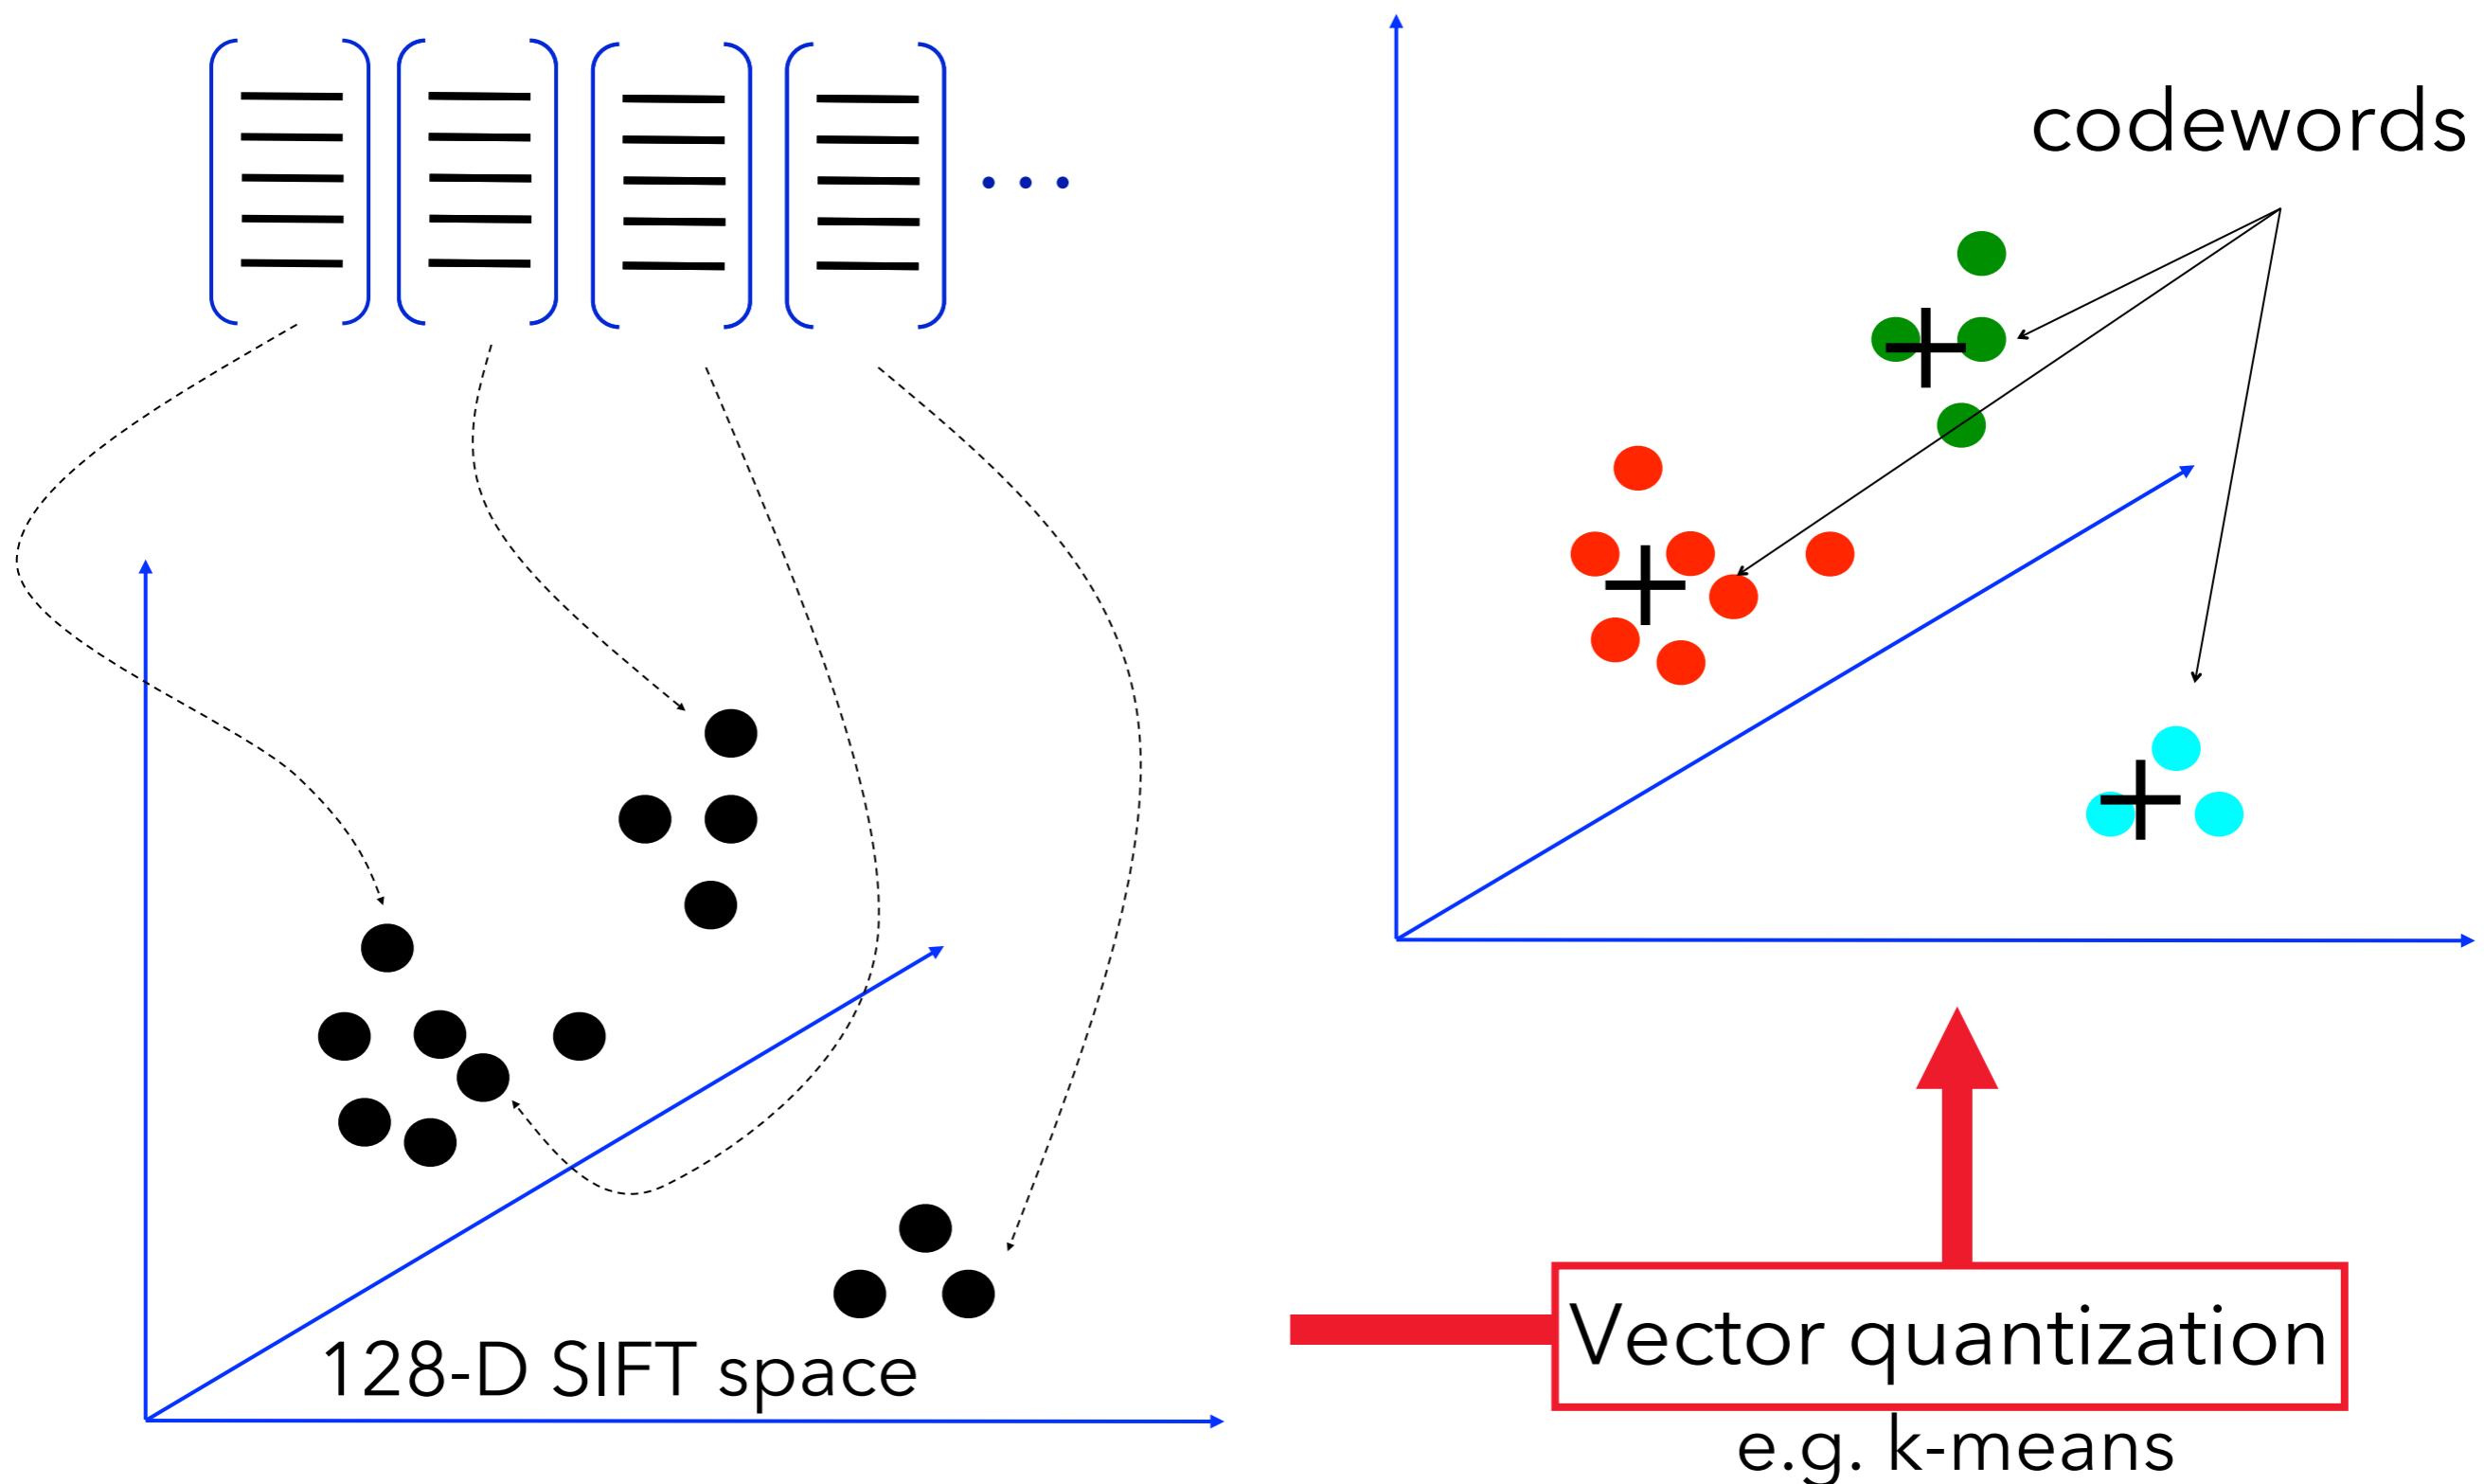
# BoW-1. Feature detection and representation



# BoW-2. Codeword dictionary formation



# BoW-2. Codeword dictionary formation



# Vector quantization

- ◆ Via clustering
- ◆ Simple, but good method: K-means clustering

Choose  $k$  data points to act as cluster centers

Until the cluster centers are unchanged

    Allocate each data point to cluster whose center is nearest

    Now ensure that every cluster has at least one data point; possible techniques for doing this include supplying empty clusters with a point chosen at random from points far from their cluster center.

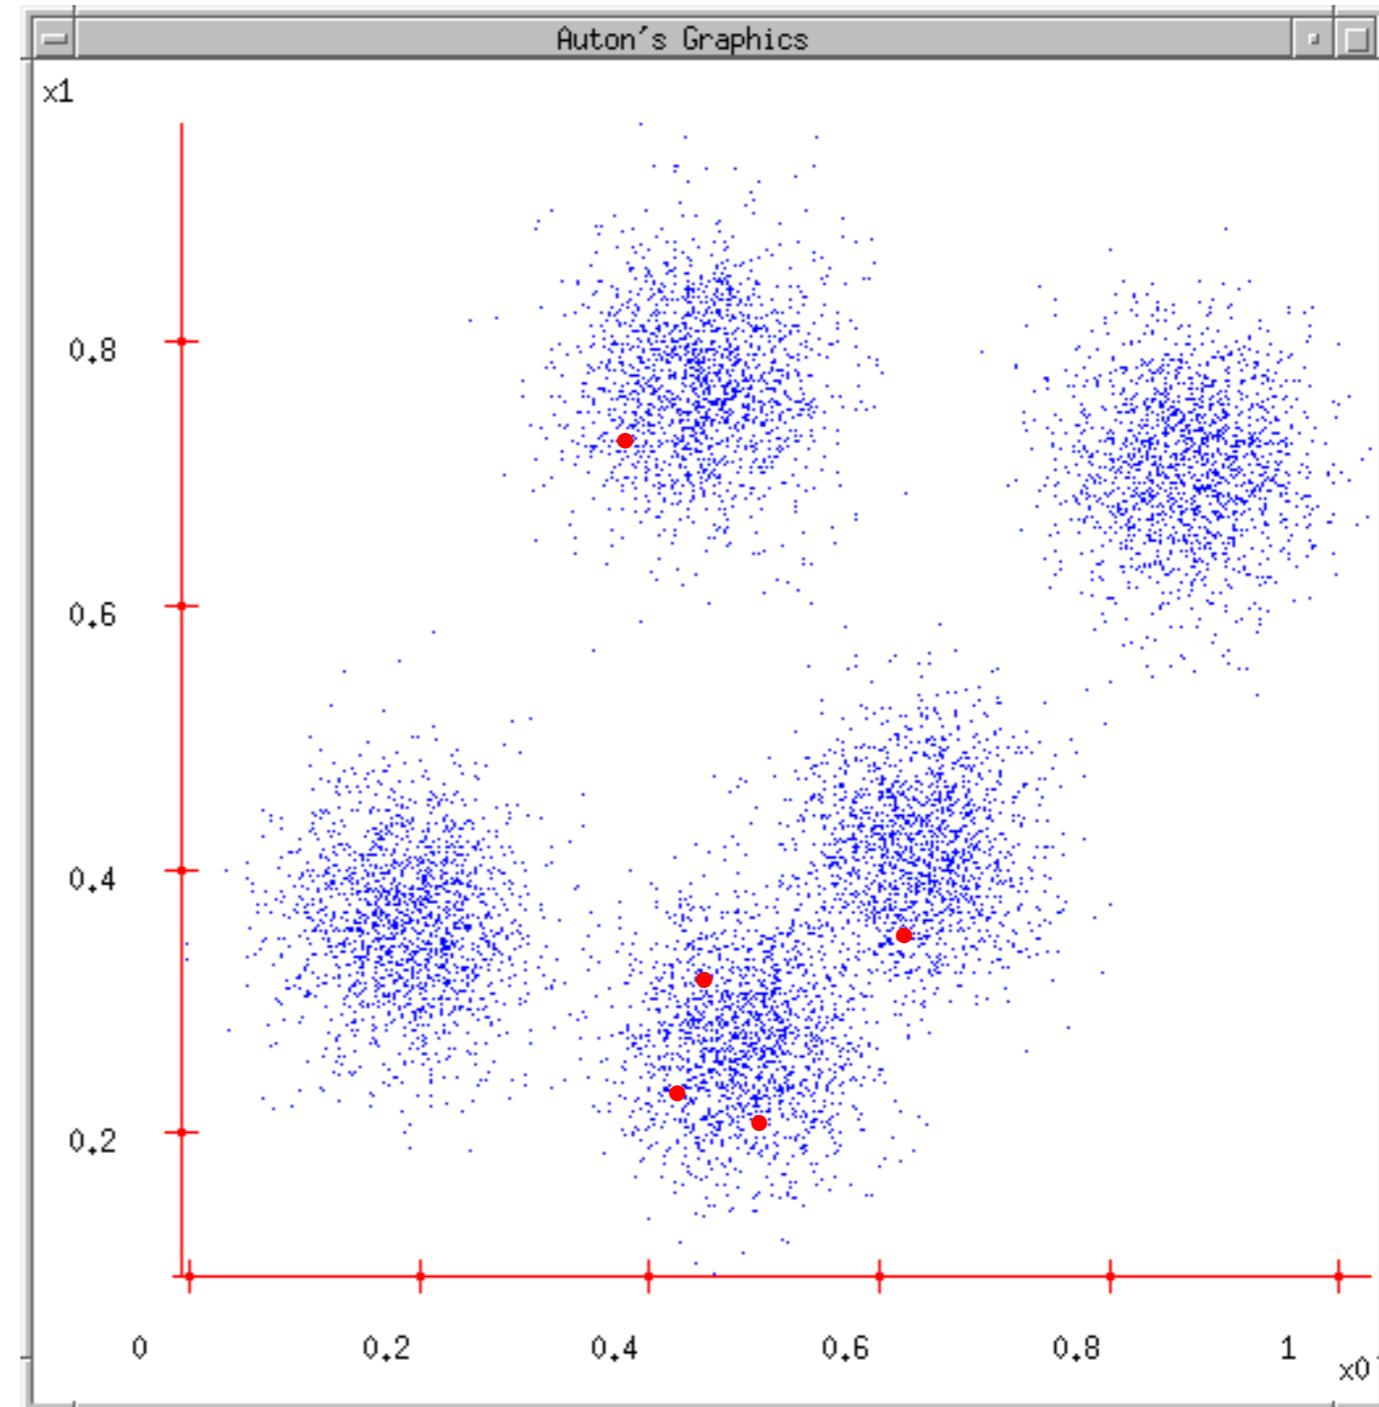
    Replace the cluster centers with the mean of the elements in their clusters.

end

[FP]

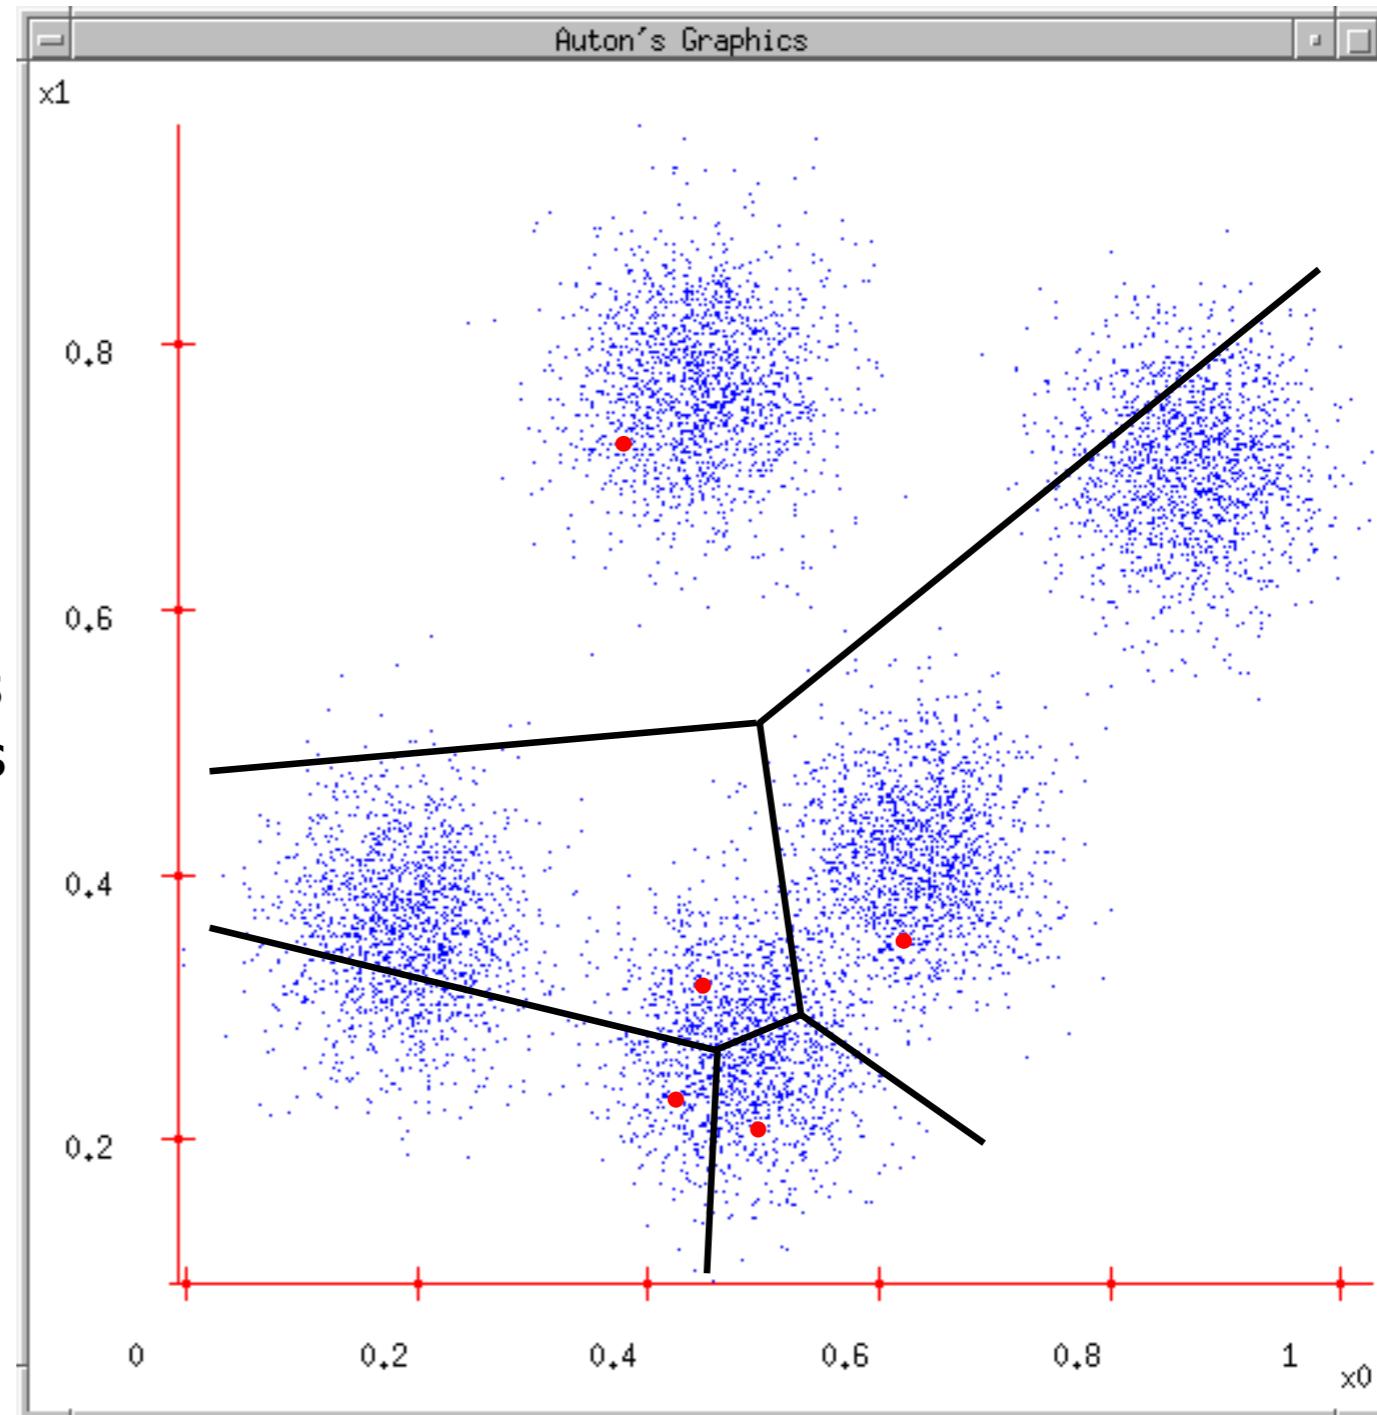
## K-means

1. Ask user how many clusters they'd like.  
*(e.g.  $k=5$ )*
2. Randomly guess  $k$  cluster Center locations



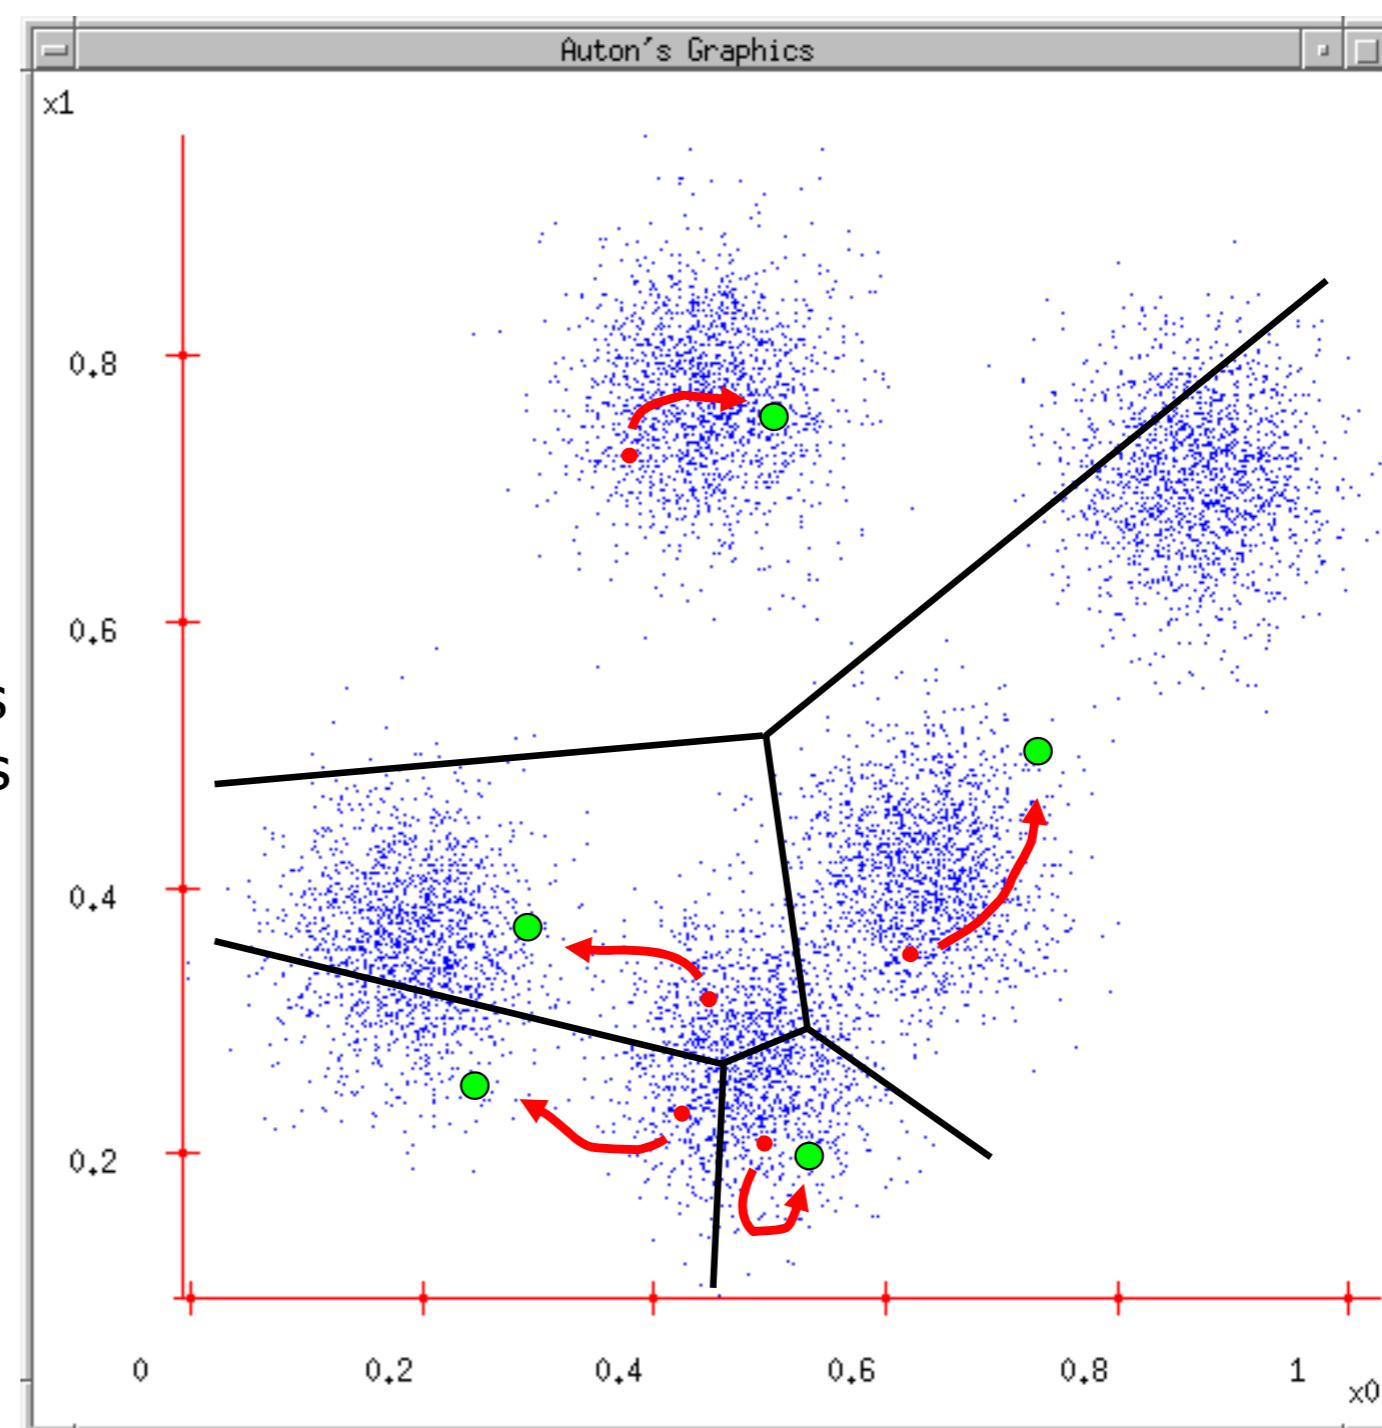
## K-means

1. Ask user how many clusters they'd like.  
*(e.g. k=5)*
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to. (Thus each Center "owns" a set of datapoints)



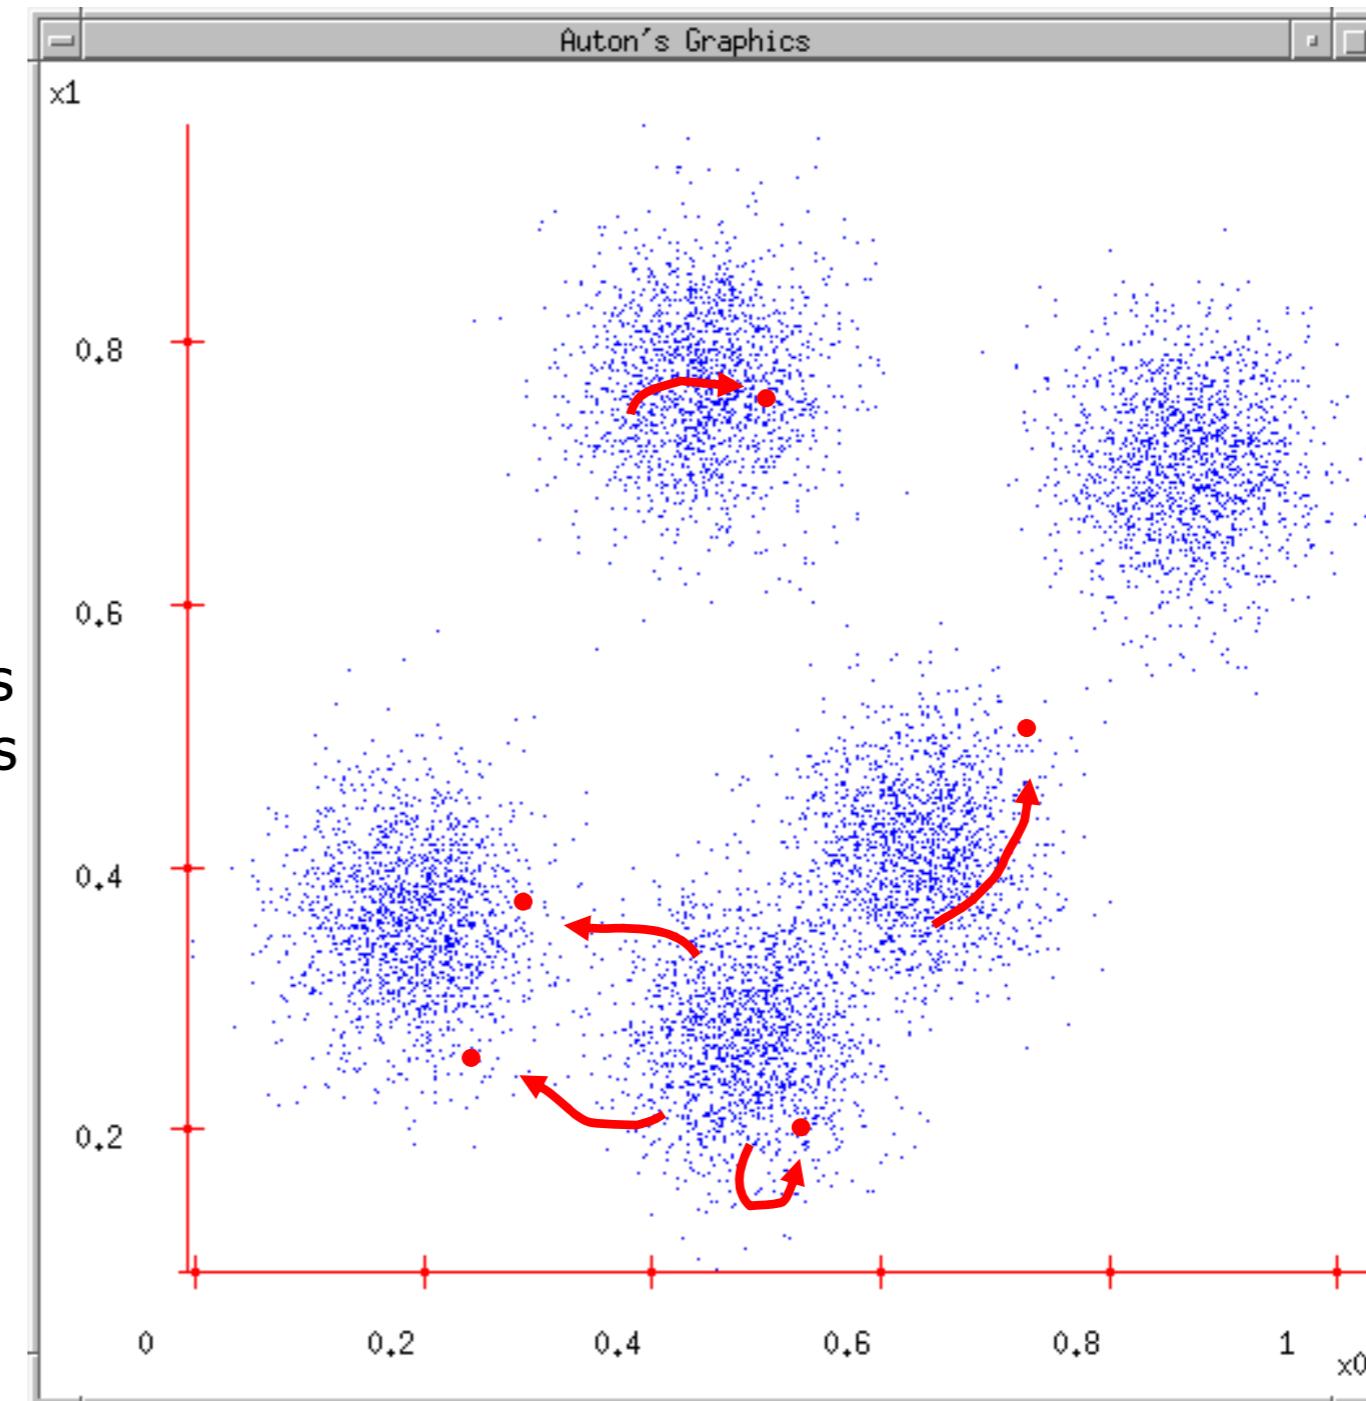
## K-means

1. Ask user how many clusters they'd like.  
*(e.g. k=5)*
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to.
4. Each Center finds the centroid of the points it owns

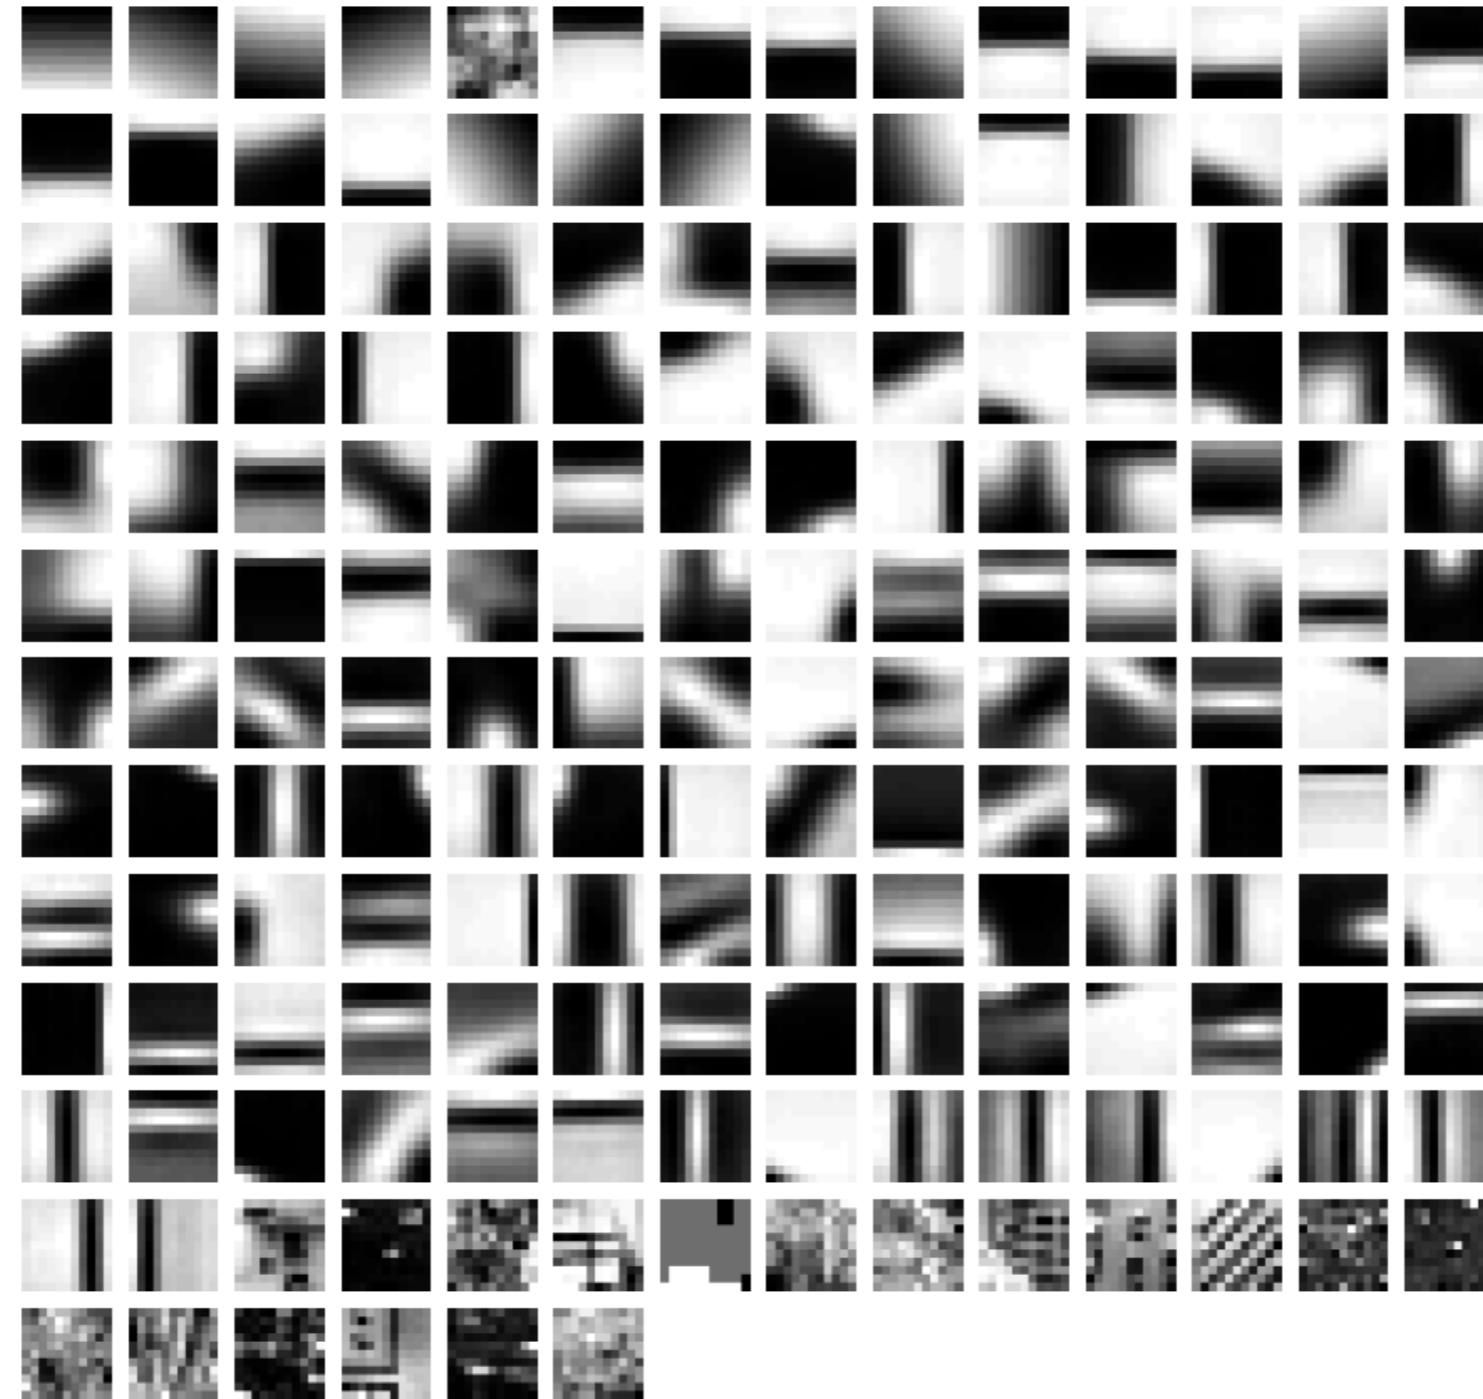


## K-means

1. Ask user how many clusters they'd like.  
*(e.g. k=5)*
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to.
4. Each Center finds the centroid of the points it owns...
5. ...and jumps there
6. ...Repeat until terminated!

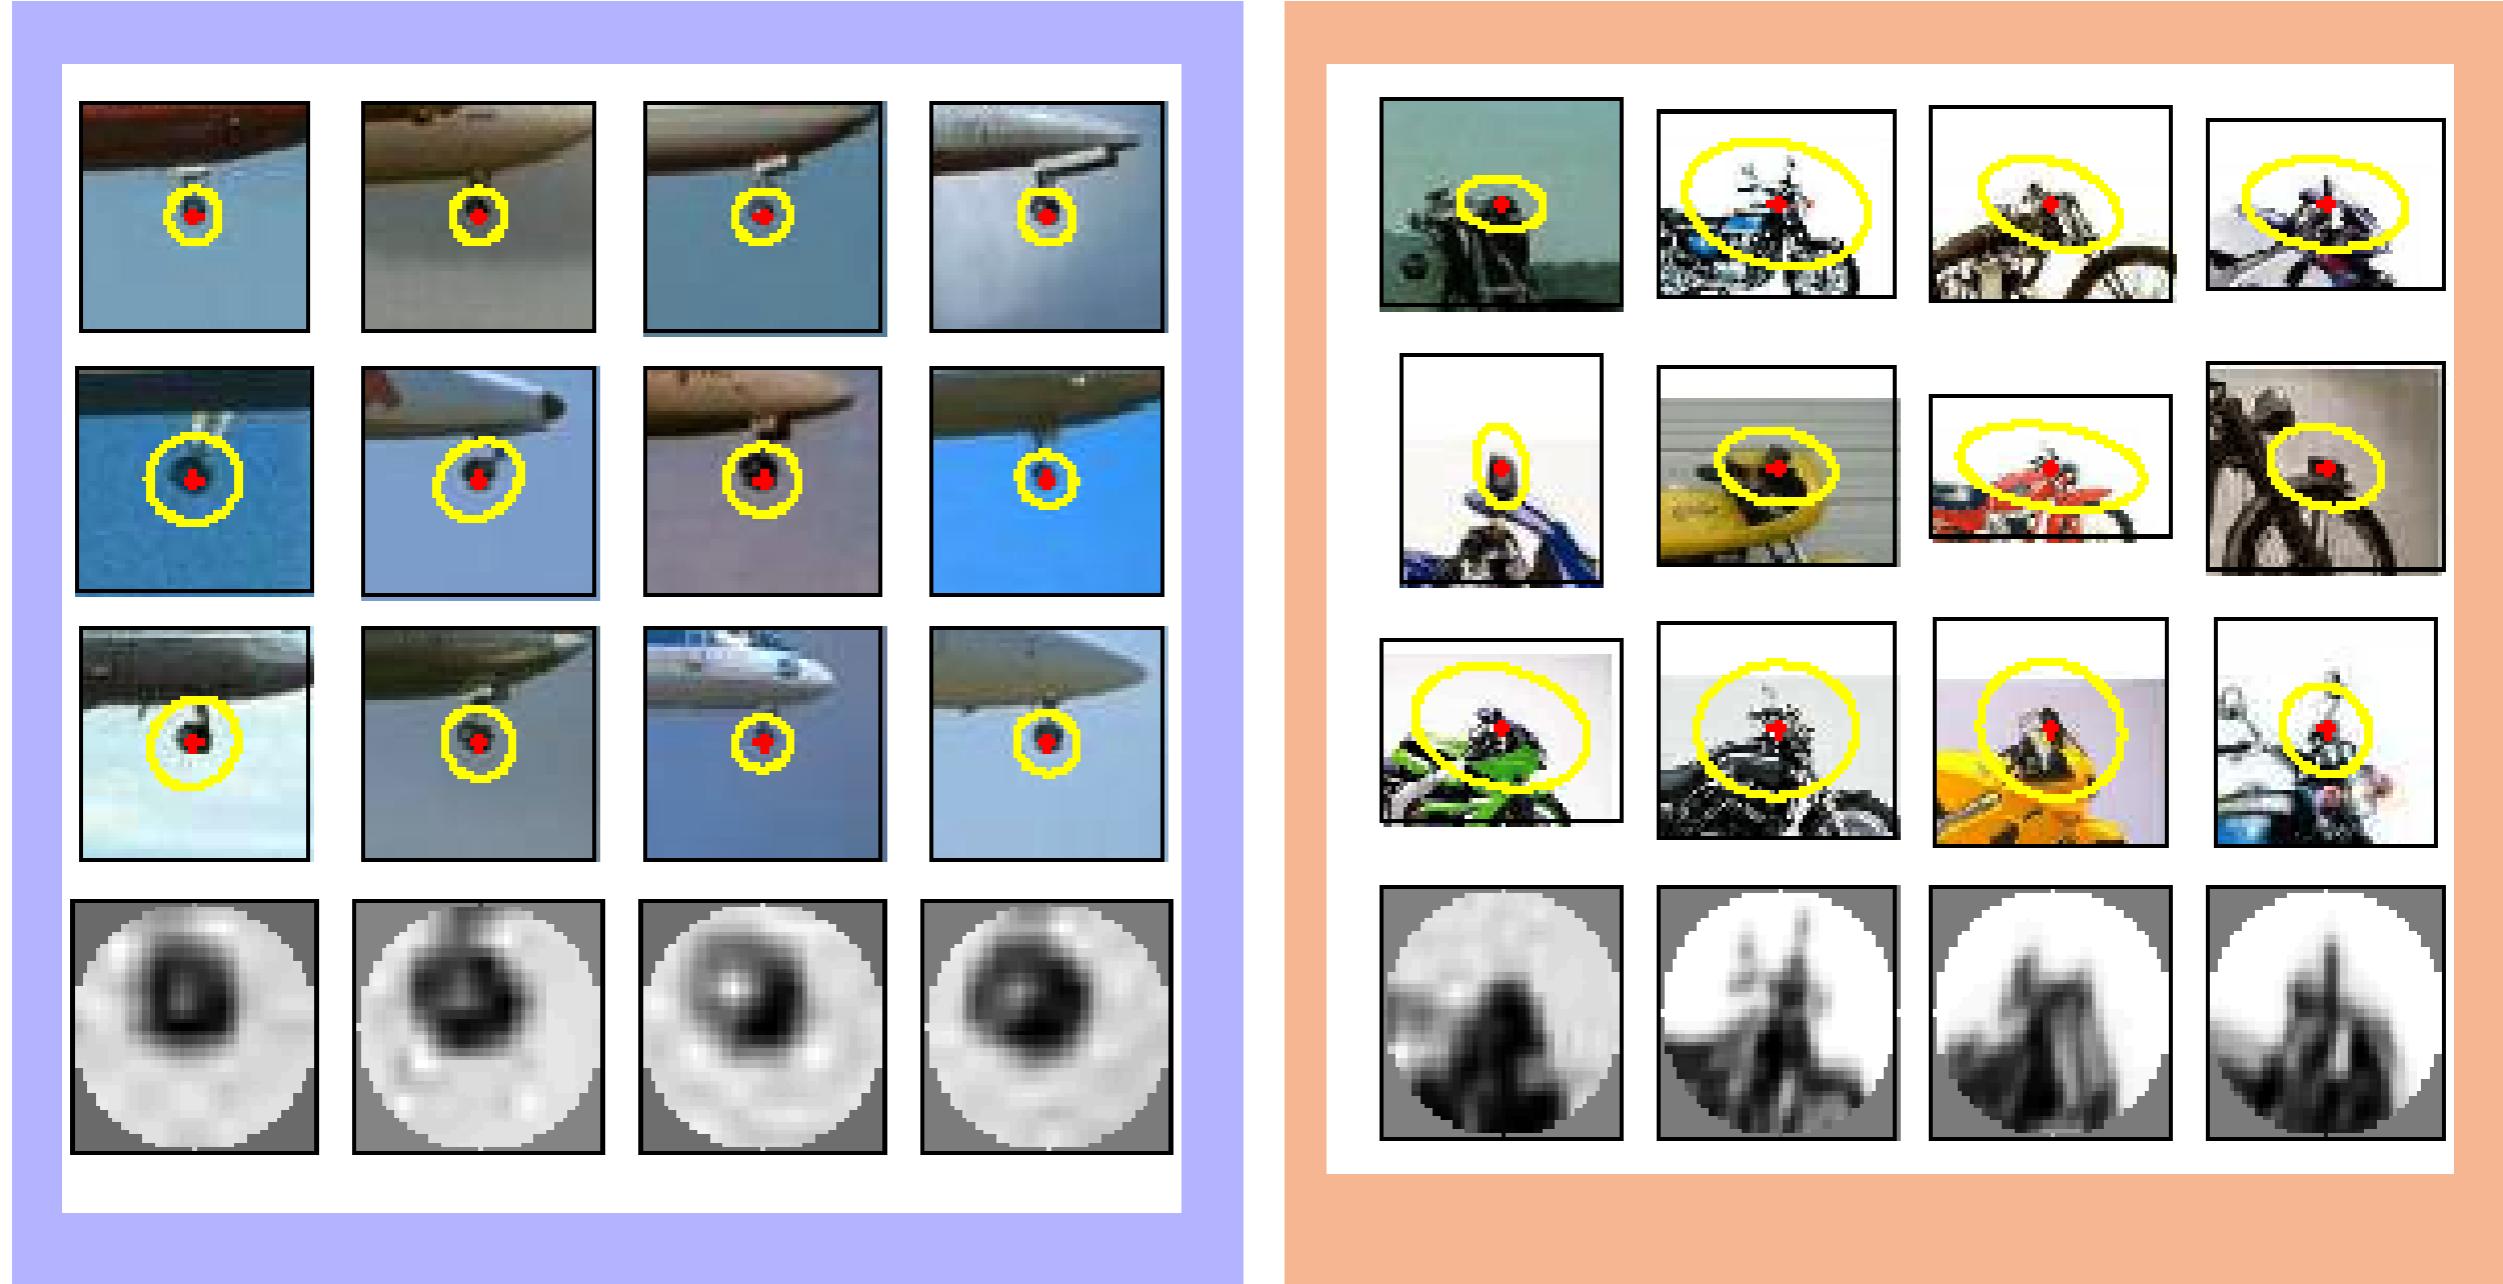


# BoW-2. Codeword dictionary formation



[Fei-Fei]

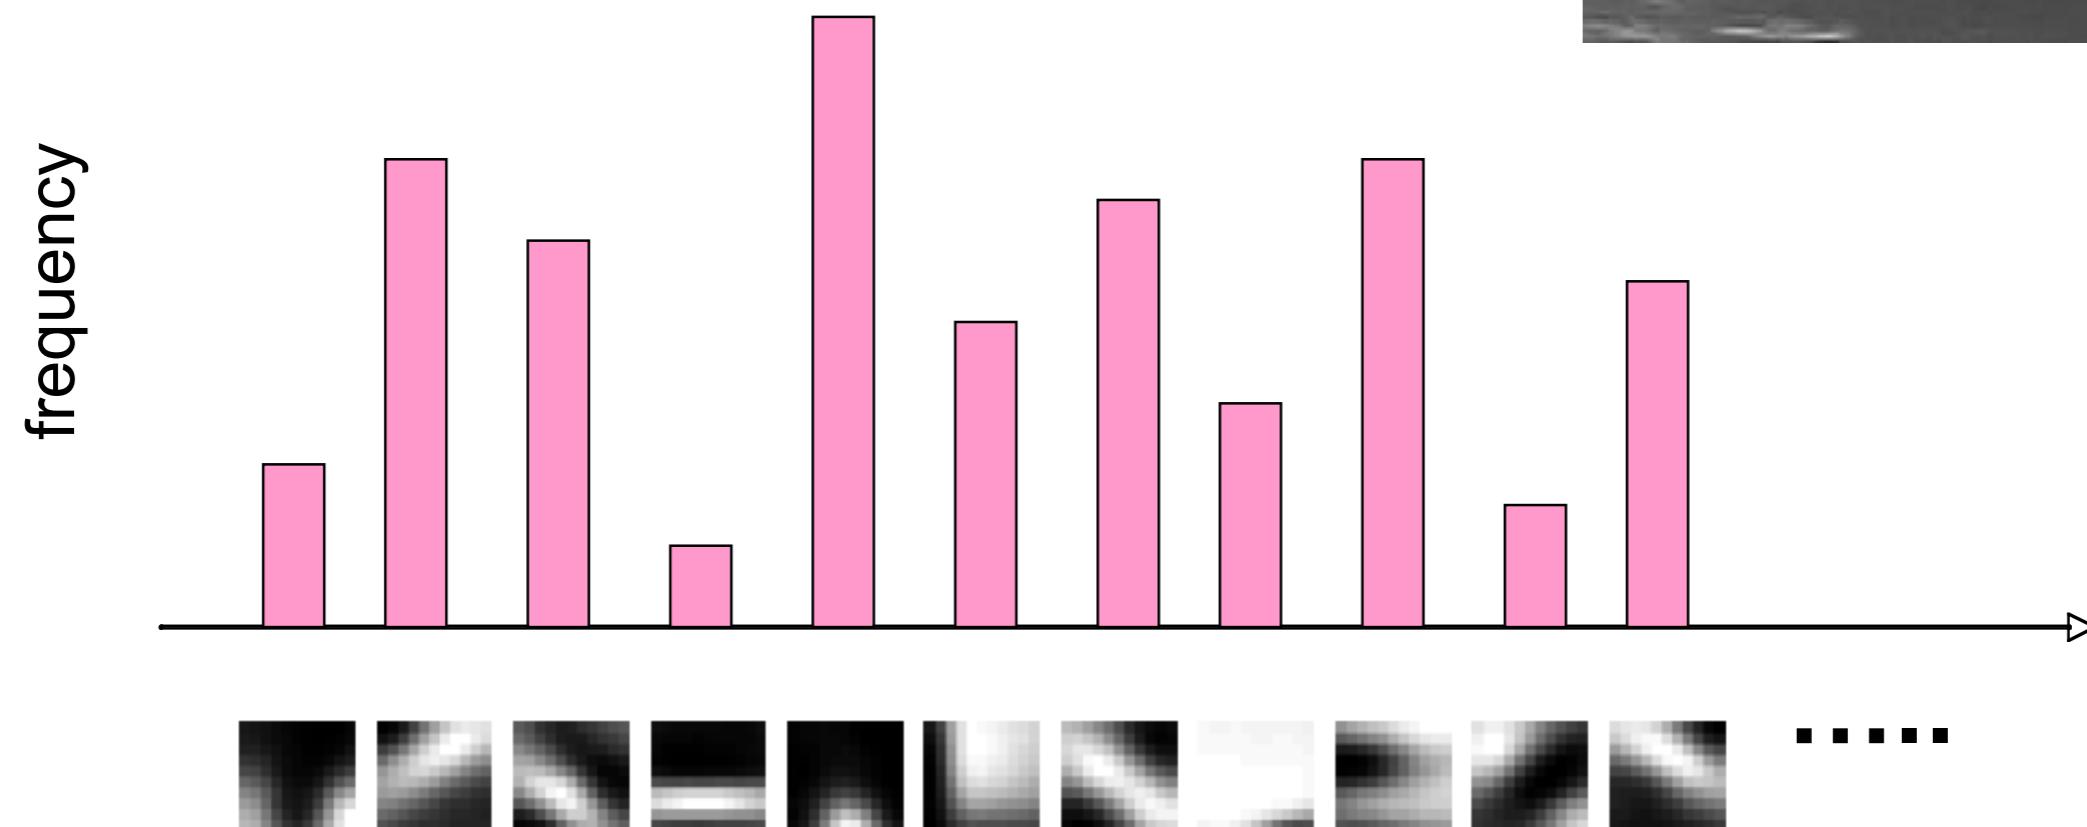
# Image patch examples of codewords



[Sivic]

# BoW-3. Image representation

- ◆ Histogram of features assigned to each cluster



# Next: Actual Recognition

