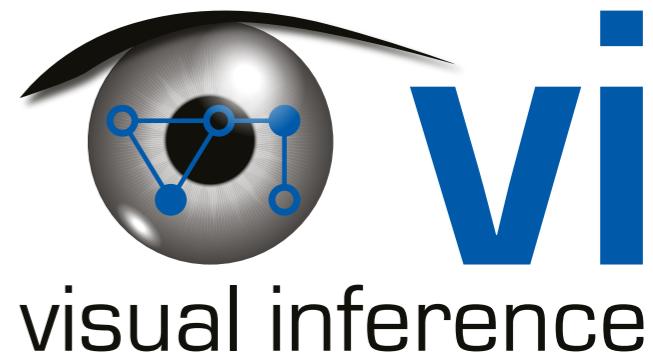


# Global Motion Models & Dense Optical Flow

21.05.2014



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT



# Announcements

- ◆ Exam:
  - ◆ Written exam on Wednesday, 6. Aug. 2014, 10:00-12:00
  - ◆ Room S2|02 C205
  - ◆ See TUCaN
  - ◆ Please make sure that you are registered!
- ◆ TU meet and move:
  - ◆ 4. June 2014
  - ◆ We will have our lecture and finish in time (11:30)

# What is Image Motion?



# Optical Flow

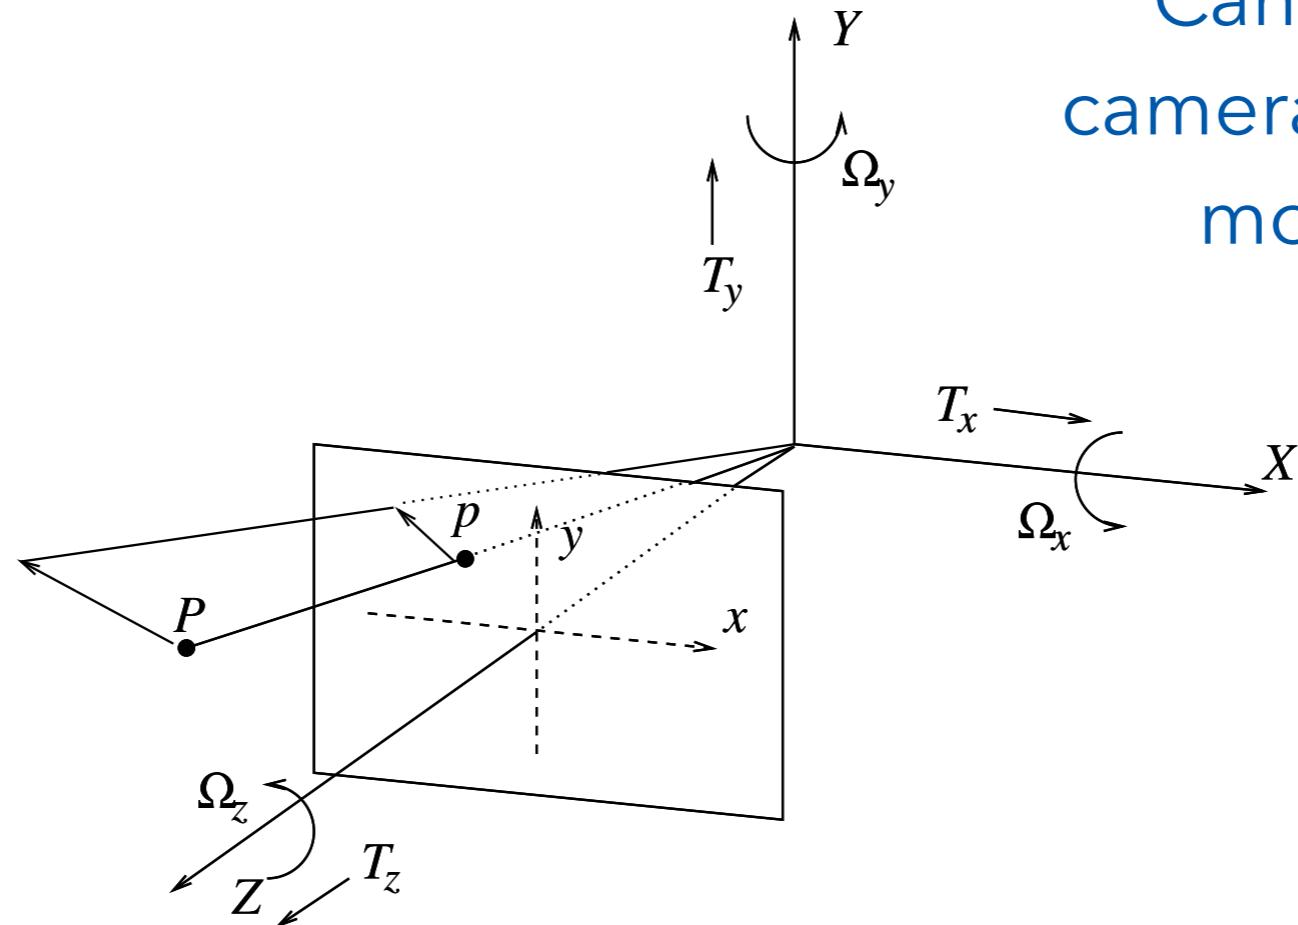
Note: Most slides borrowed from Michael Black.



J. J. Gibson, The Ecological Approach to Visual Perception

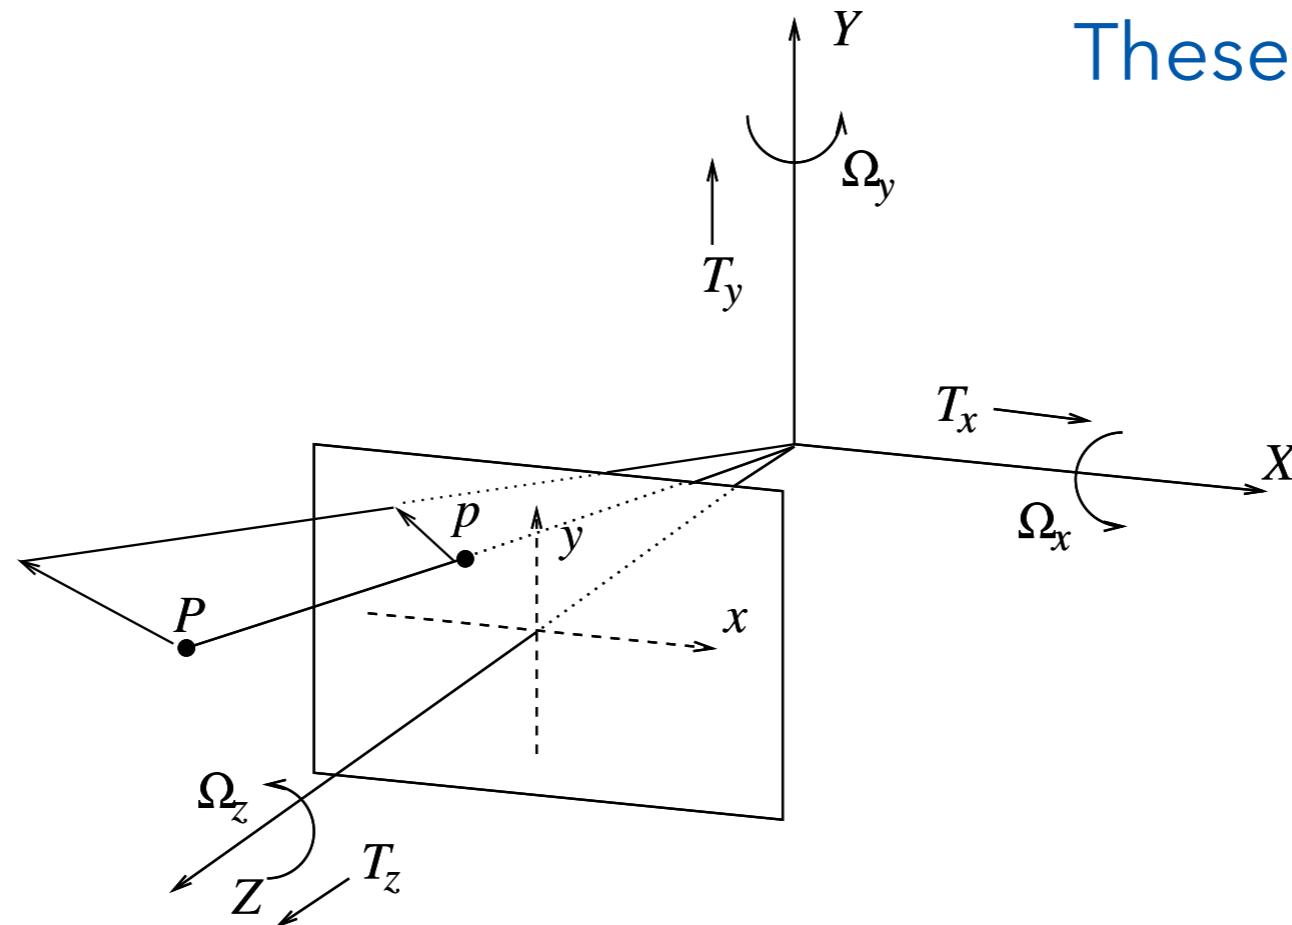
# Motion Field

Can be the result of  
camera motion or object  
motion (or both)!



**Motion field** = 2D motion field representing the projection of the 3D motion of points in the scene onto the image plane.

# Optical Flow



These are not the same!  
Why?

**Optical flow** = 2D velocity field describing the apparent motion in the images.

# Thought Experiment 1

Lambertian ball rotating  
in 3D

What does the 2D  
motion field look like?

What does the 2D  
optical flow field look  
like?

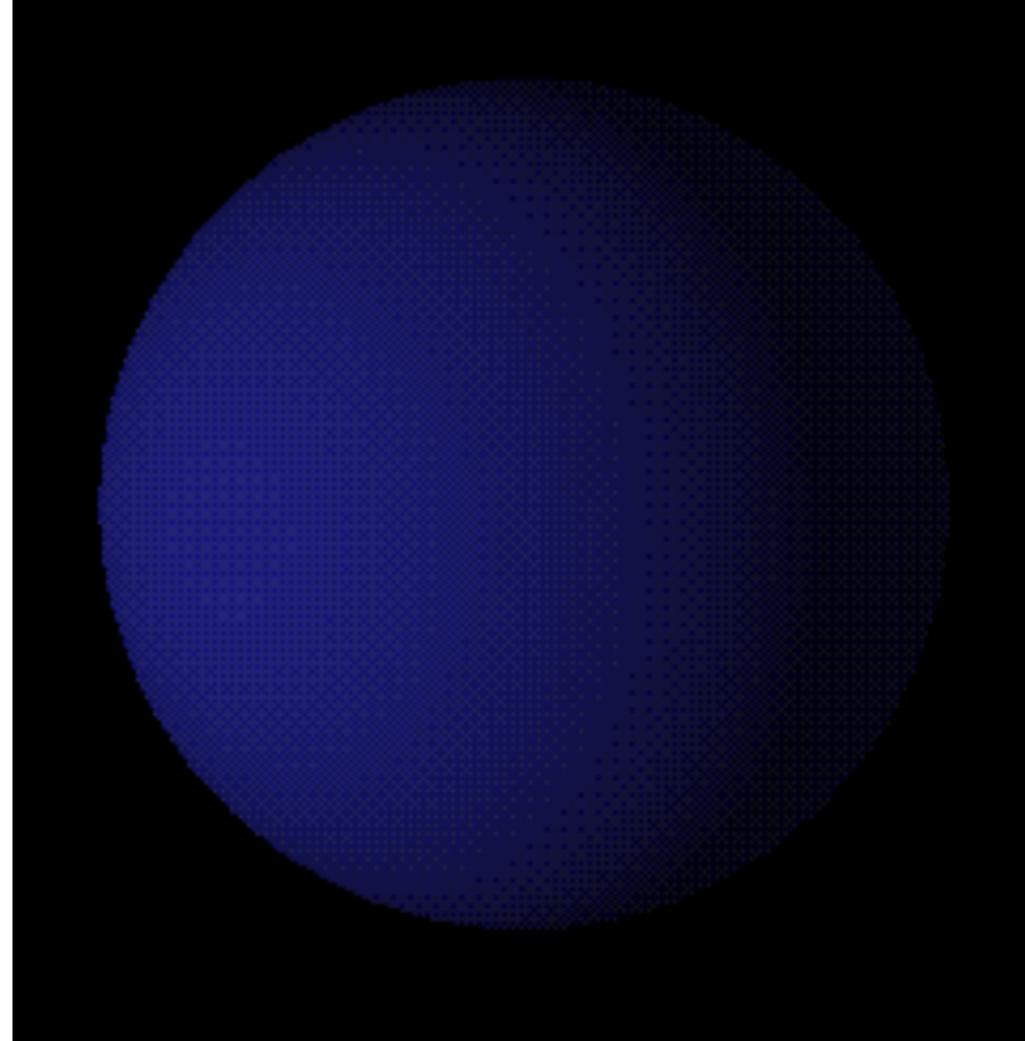


Image source: <http://www.evl.uic.edu/aej/488/lecture12.html>

# Thought Experiment 2

Stationary specular ball,  
moving light source.

What does the 2D  
motion field look like?

What does the 2D  
optical flow field look  
like?

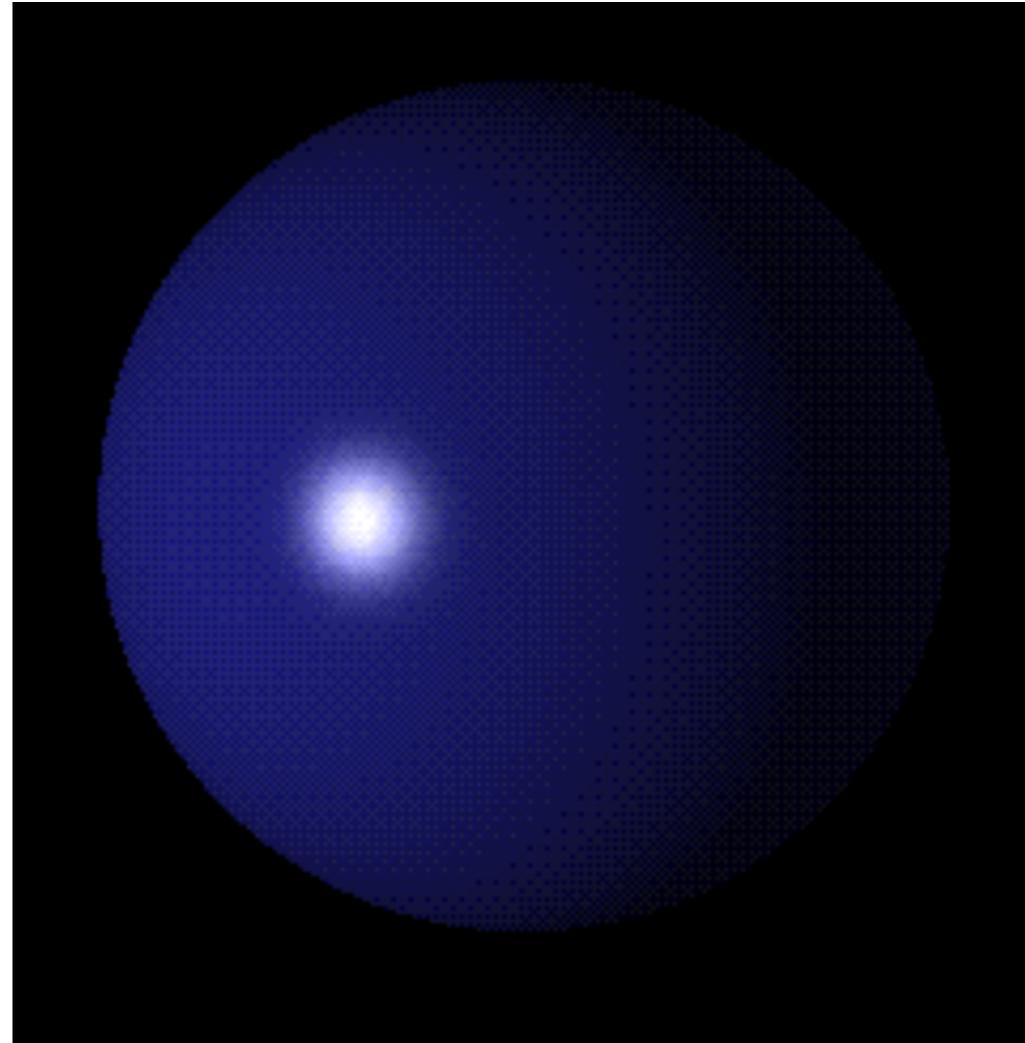
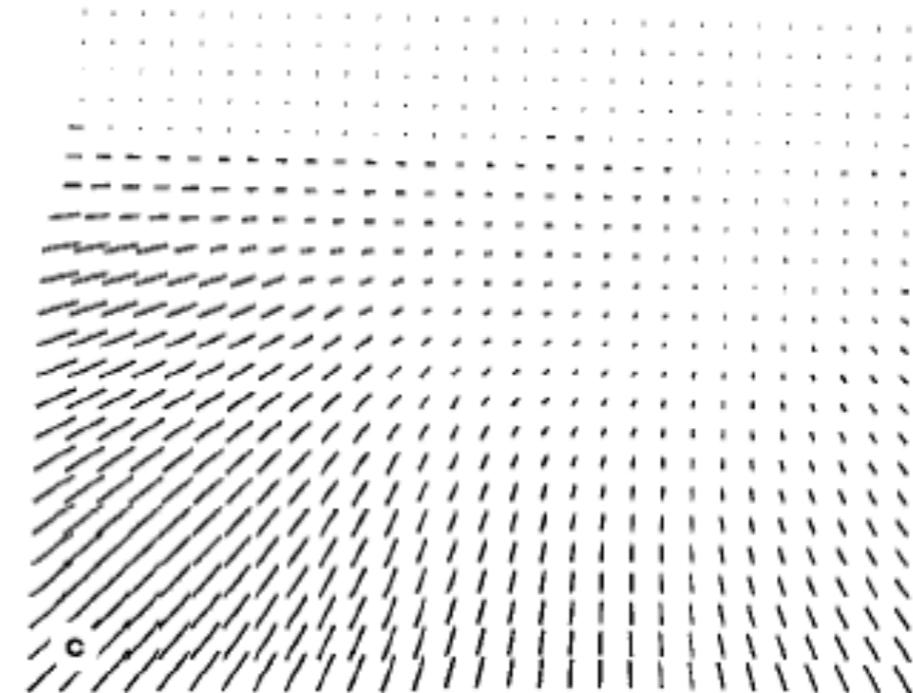
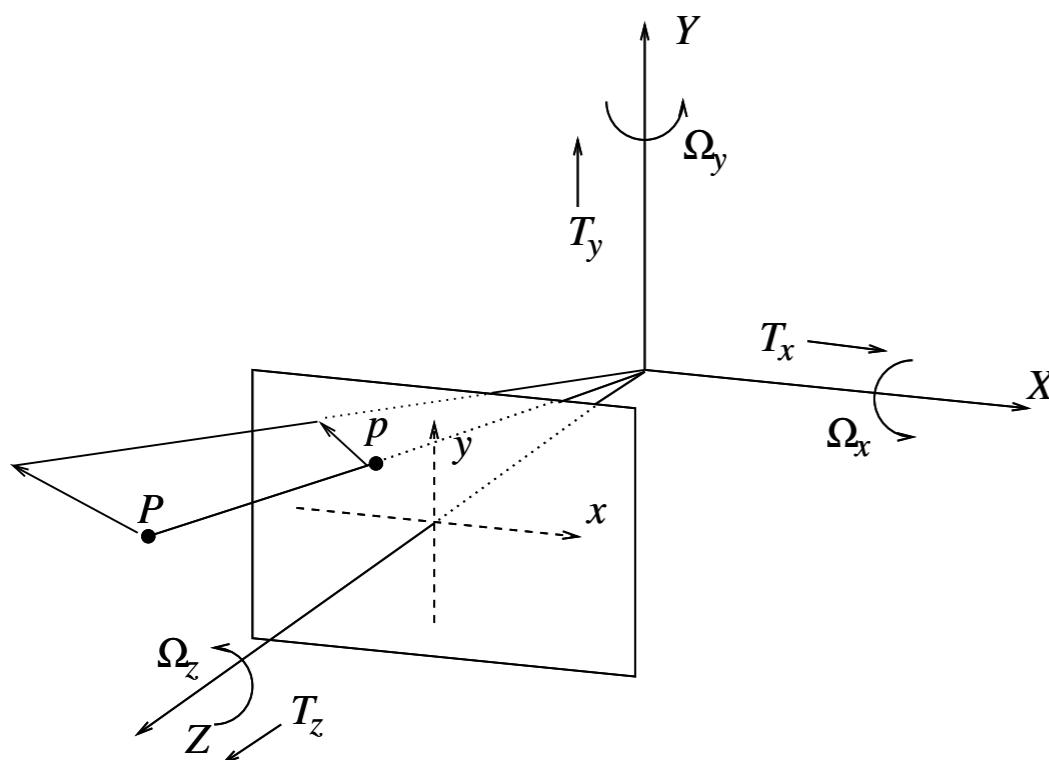


Image source: <http://www.evl.uic.edu/aej/488/lecture12.html>

# Optical Flow Field

Image irradiance at time  $t$   
and location  $\mathbf{x} = (x, y)$ :

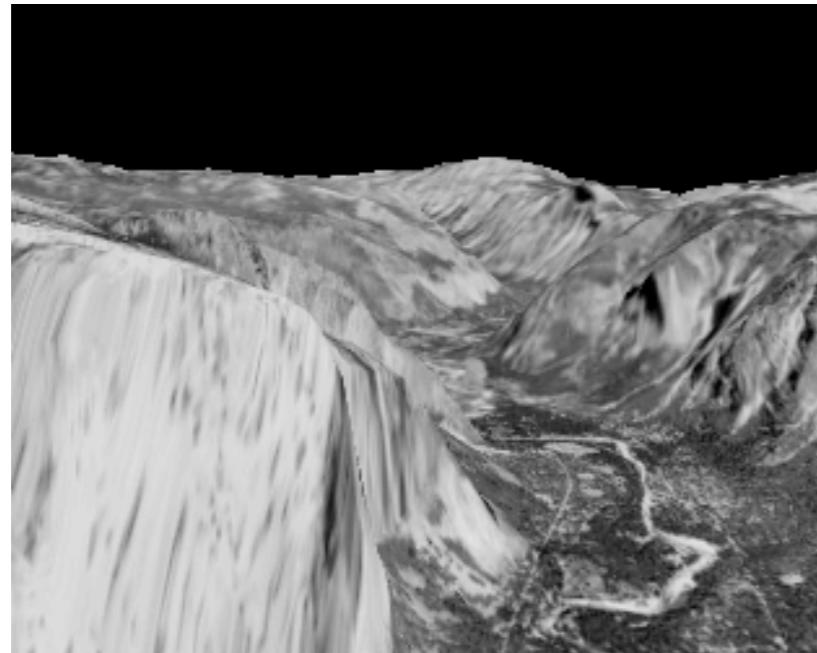
$$I(x, y, t)$$



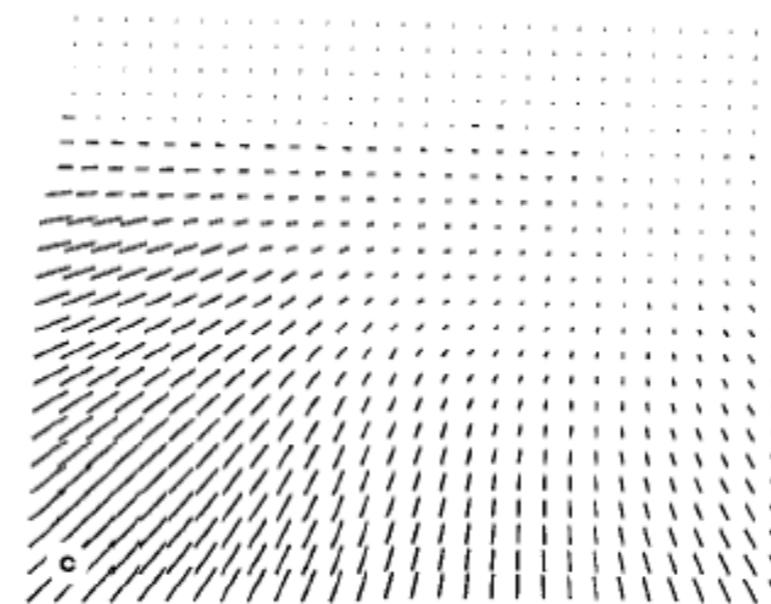
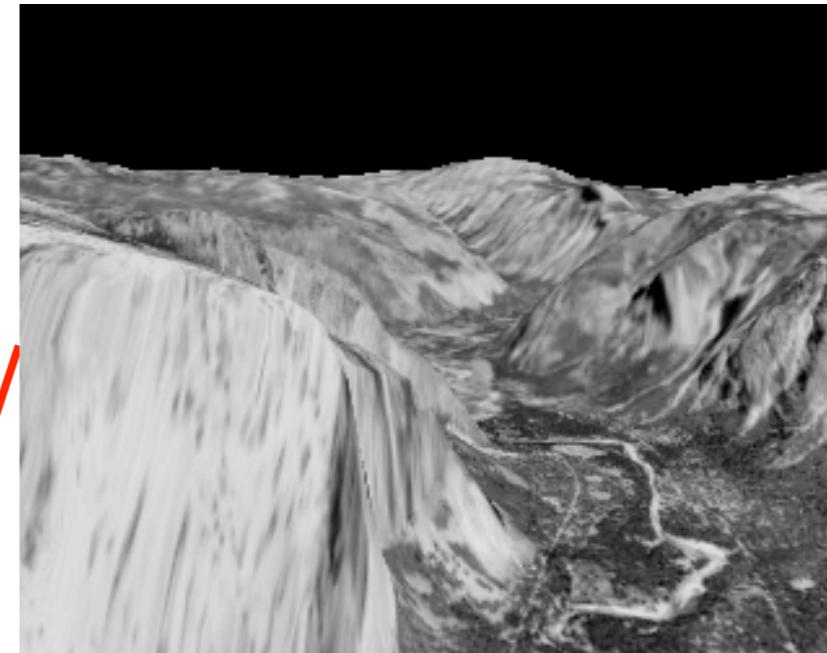
$u(x, y)$  Horizontal component  
 $v(x, y)$  Vertical component

# Optical Flow Estimation

time t



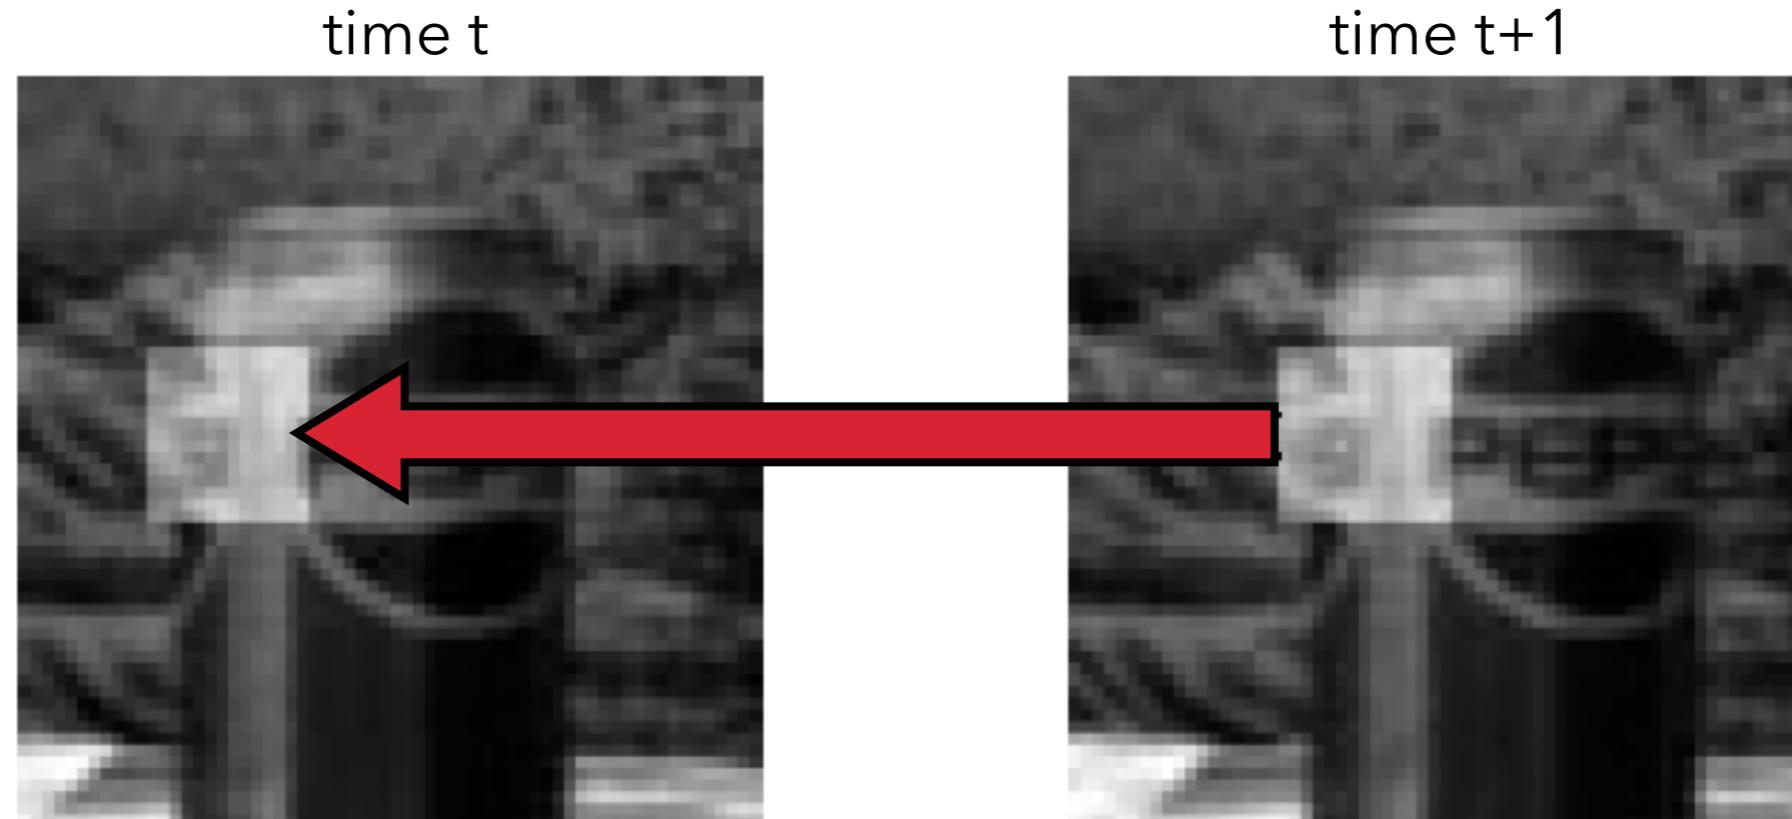
time t+1



# How do we compute flow?

- ◆ Reminder from CV1:
  - ◆ Make 3 important assumptions first.
  - ◆ Then exploit those computationally.

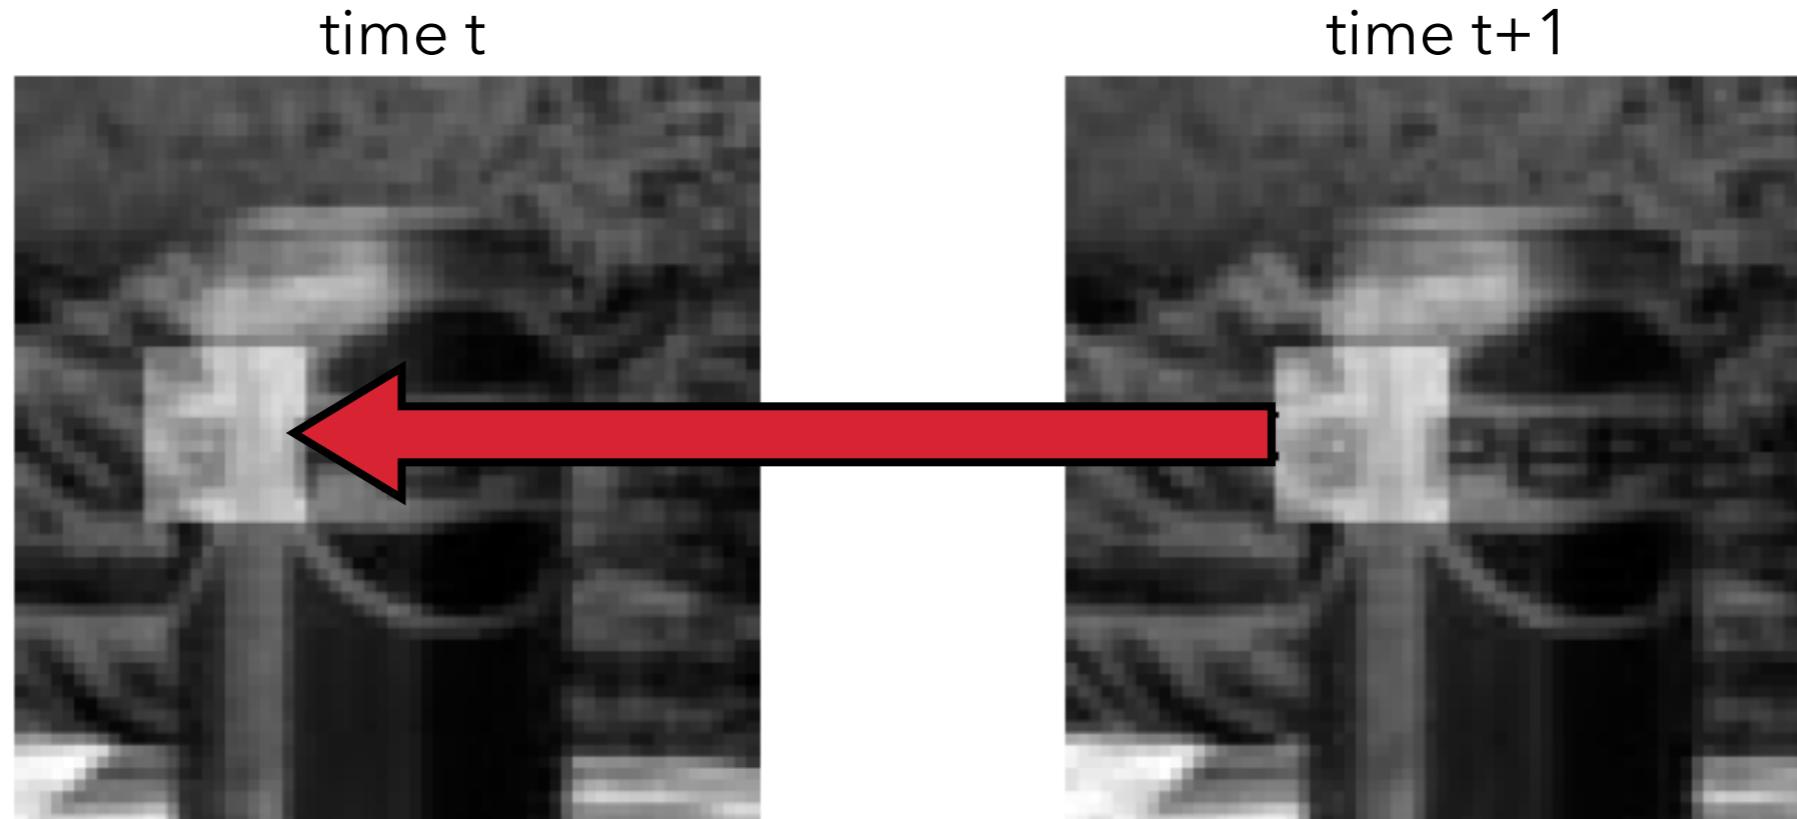
# Brightness Constancy



## Assumption 1:

Image measurements (e.g. brightness) in a small region remain the same although their location may change.

# Brightness Constancy



**Assumption 1:**

$$I(x + u(x, y), y + v(x, y), t + 1) = I(x, y, t)$$

shifted by horizontal flow      shifted by vertical flow

# Brightness Constancy

time t



time t+1



- This is a lot like what we saw in stereo!
- Instead of horizontal disparity, we have more general horizontal and vertical motion.



## Assumption 1:

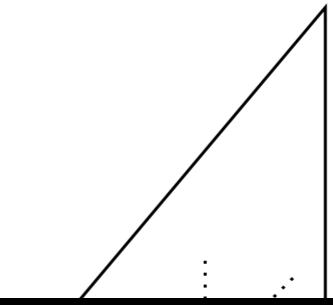
$$I(x + u(x, y), y + v(x, y), t + 1) = I(x, y, t)$$

shifted by                                  shifted by  
horizontal flow                              vertical flow

# Spatial Coherence



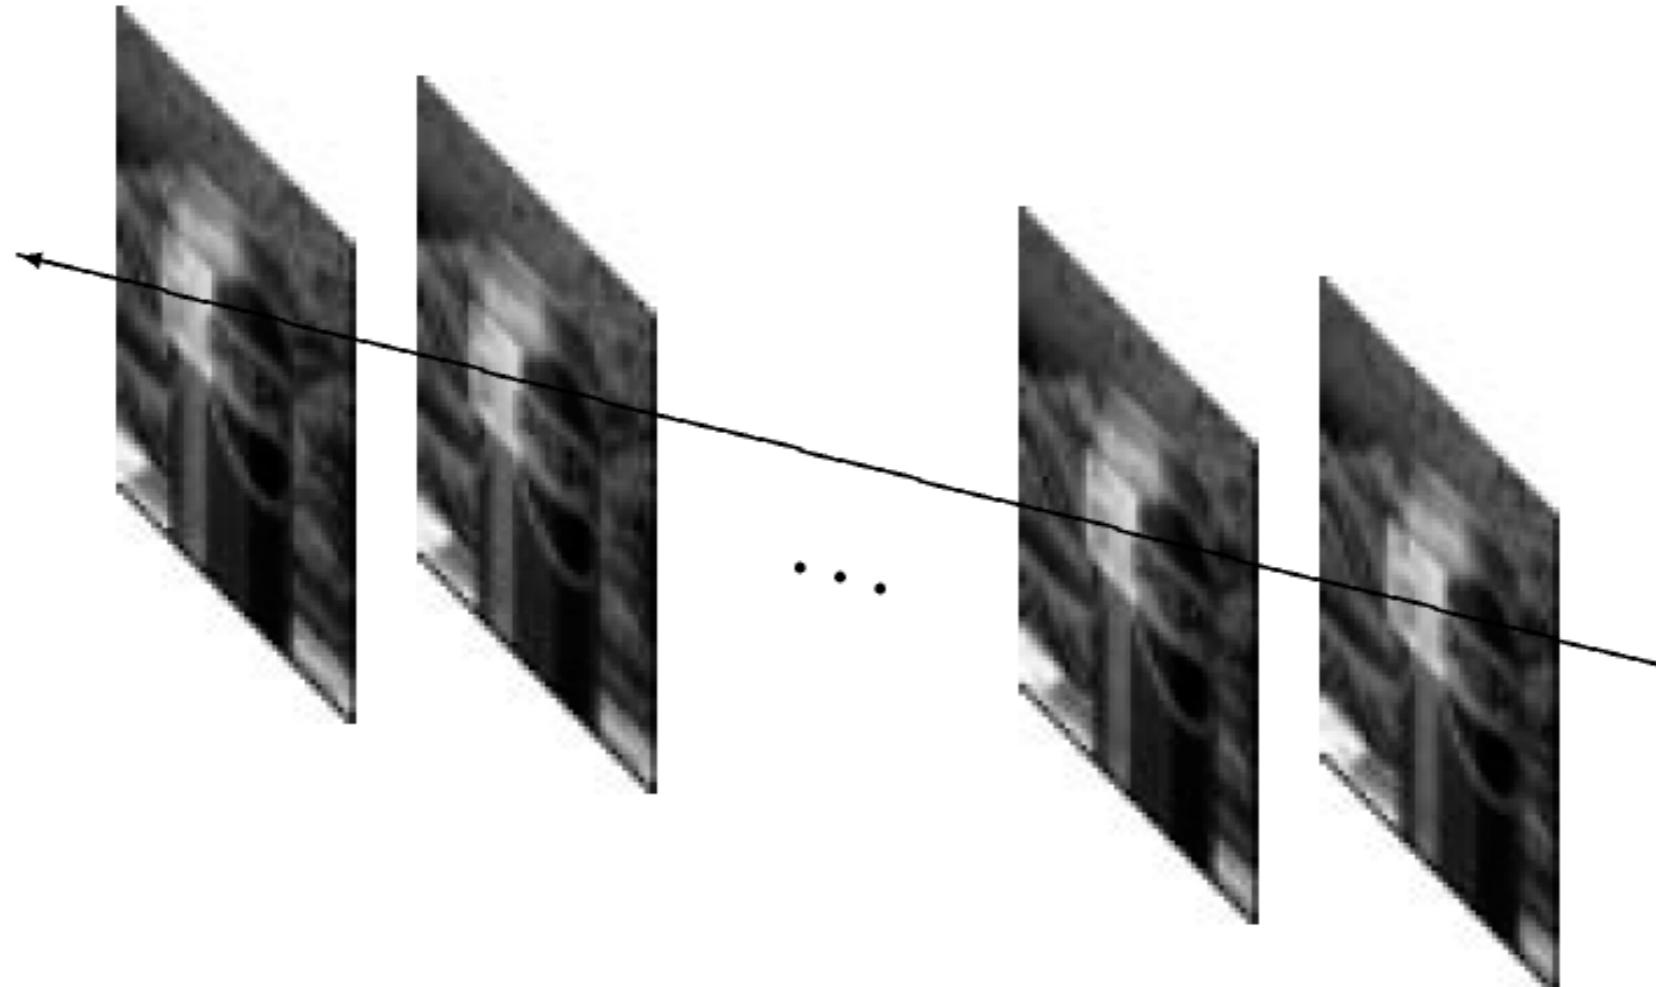
- This is a lot like what we saw in stereo!
- Corresponding notion of spatial regularity



## Assumption 2:

- Neighboring points in the scene typically belong to the same surface and hence typically have similar 3D motions.
- Since they also project to nearby points in the image, we expect spatial coherence in the image flow.

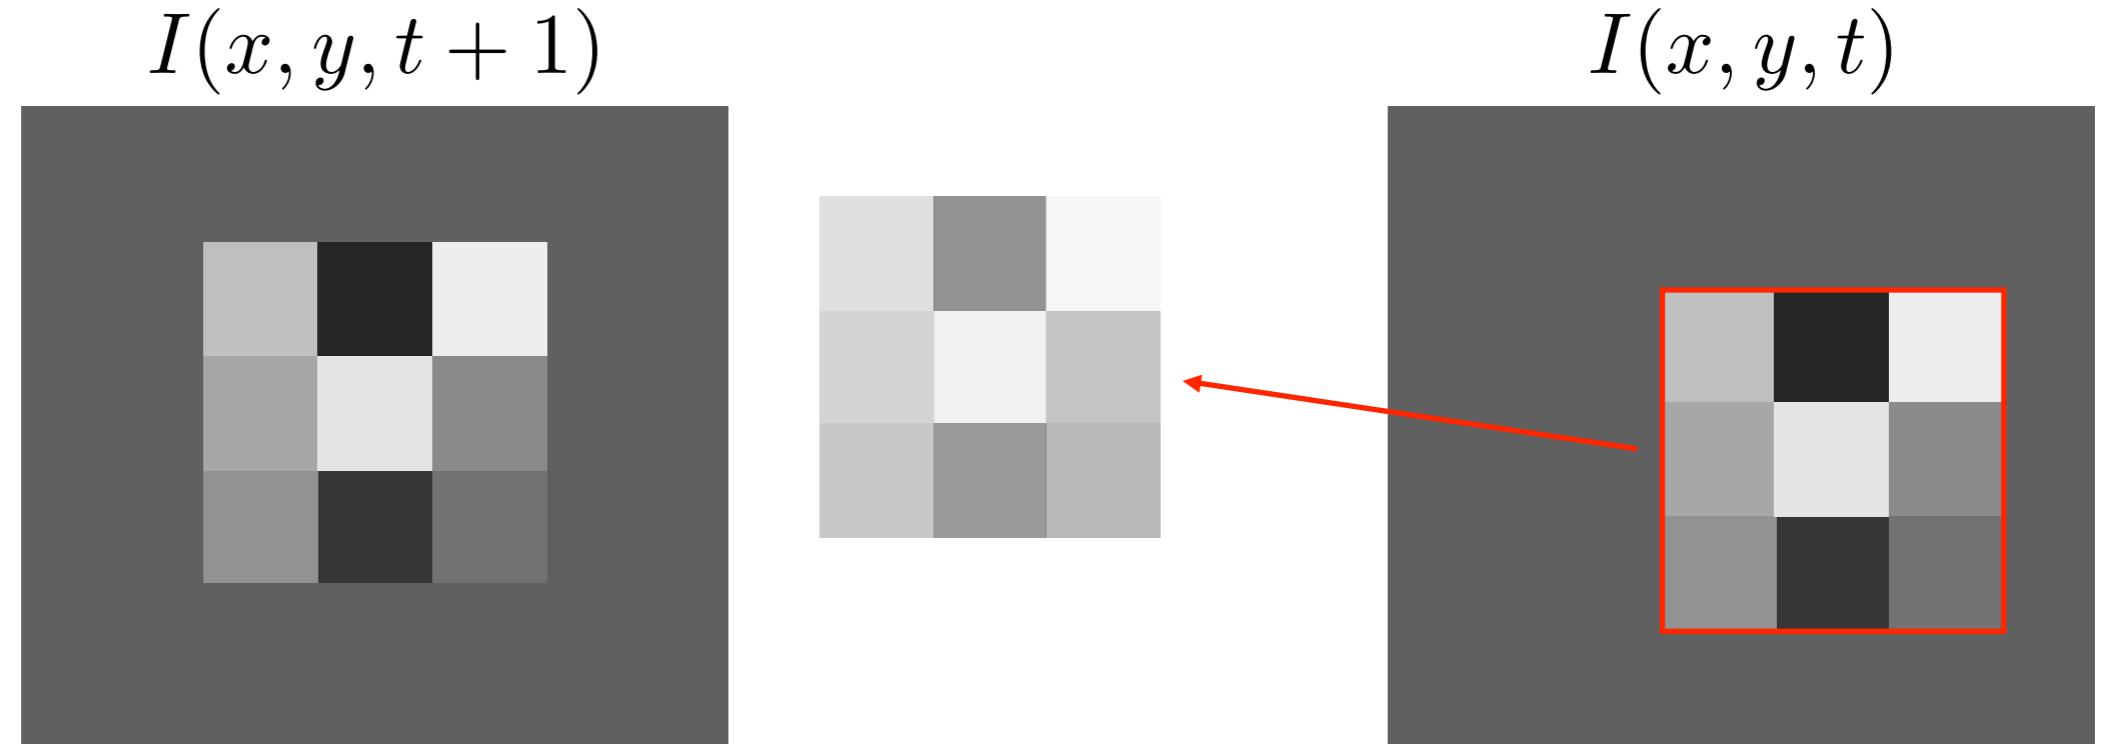
# Temporal Persistence



## Assumption 3:

The image motion of a surface patch changes gradually over time.

# Minimize Brightness Difference



$$E_{SSD}(u, v) = \sum_{(x,y) \in R} (I(x + u, y + v, t + 1) - I(x, y, t))^2$$

# Simple Flow Estimation Algorithm

$$E_{SSD}(u, v) = \sum_{(x,y) \in R} (I(x + u, y + v, t + 1) - I(x, y, t))^2$$

- ◆ Simple way of flow estimation:
  - ◆ Discretize the space of possible motions.
  - ◆ For each pixel, try all possible motions.
  - ◆ Take the one that minimizes the SSD.
- ◆ Problems with this approach:
  - ◆ Very inefficient.
  - ◆ The motions are discrete, but we may not want that.
  - ◆ So how can we optimize this nonlinear and non-convex function?

# Can we approximate this?

$$E_{SSD}(u, v) = \sum_{(x,y) \in R} (I(x + u, y + v, t + 1) - I(x, y, t))^2$$

Let us look at the **general case**:

$$I(x + \Delta_x, y + \Delta_y, t + \Delta_t)$$

**Taylor series approximation:**

$$I(x, y, t) + \Delta_x \frac{\partial}{\partial x} I(x, y, t) + \Delta_y \frac{\partial}{\partial y} I(x, y, t) + \Delta_t \frac{\partial}{\partial t} I(x, y, t) + \epsilon(\Delta_x^2, \Delta_y^2, \Delta_t^2)$$

Approximation error

# Can we approximate this?

$$E_{SSD}(u, v) = \sum_{(x,y) \in R} (I(x + u, y + v, t + 1) - I(x, y, t))^2$$

Assume small motion



Brightness varies smoothly

$$I(x, y, t) + u \frac{\partial}{\partial x} I(x, y, t) + v \frac{\partial}{\partial y} I(x, y, t) + \frac{\partial}{\partial t} I(x, y, t) + \epsilon(u, v, 1)$$

## SSD approximation:

$$\begin{aligned} E_{SSD}(u, v) &\approx \sum_{(x,y) \in R} (I(x, y, t) + u \frac{\partial}{\partial x} I(x, y, t) + v \frac{\partial}{\partial y} I(x, y, t) + \frac{\partial}{\partial t} I(x, y, t) - I(x, y, t))^2 \\ &= \sum_{(x,y) \in R} (u \frac{\partial}{\partial x} I(x, y, t) + v \frac{\partial}{\partial y} I(x, y, t) + \frac{\partial}{\partial t} I(x, y, t))^2 \\ &= \sum_{(x,y) \in R} (u \cdot I_x(x, y, t) + v \cdot I_y(x, y, t) + I_t(x, y, t))^2 \end{aligned}$$

# Optical Flow Constraint Equation

$$E_{SSD}(u, v) \approx \sum_{(x,y) \in R} (u \cdot I_x(x, y, t) + v \cdot I_y(x, y, t) + I_t(x, y, t))^2$$

- ◆ This approximated SSD objective is a **convex function** of the motion  $u$  and  $v$ .
- ◆ It becomes much easier to optimize.
  - ◆ We will see that shortly.
- ◆ But, this only holds for **small motions!**

# Optical Flow Constraint Equation

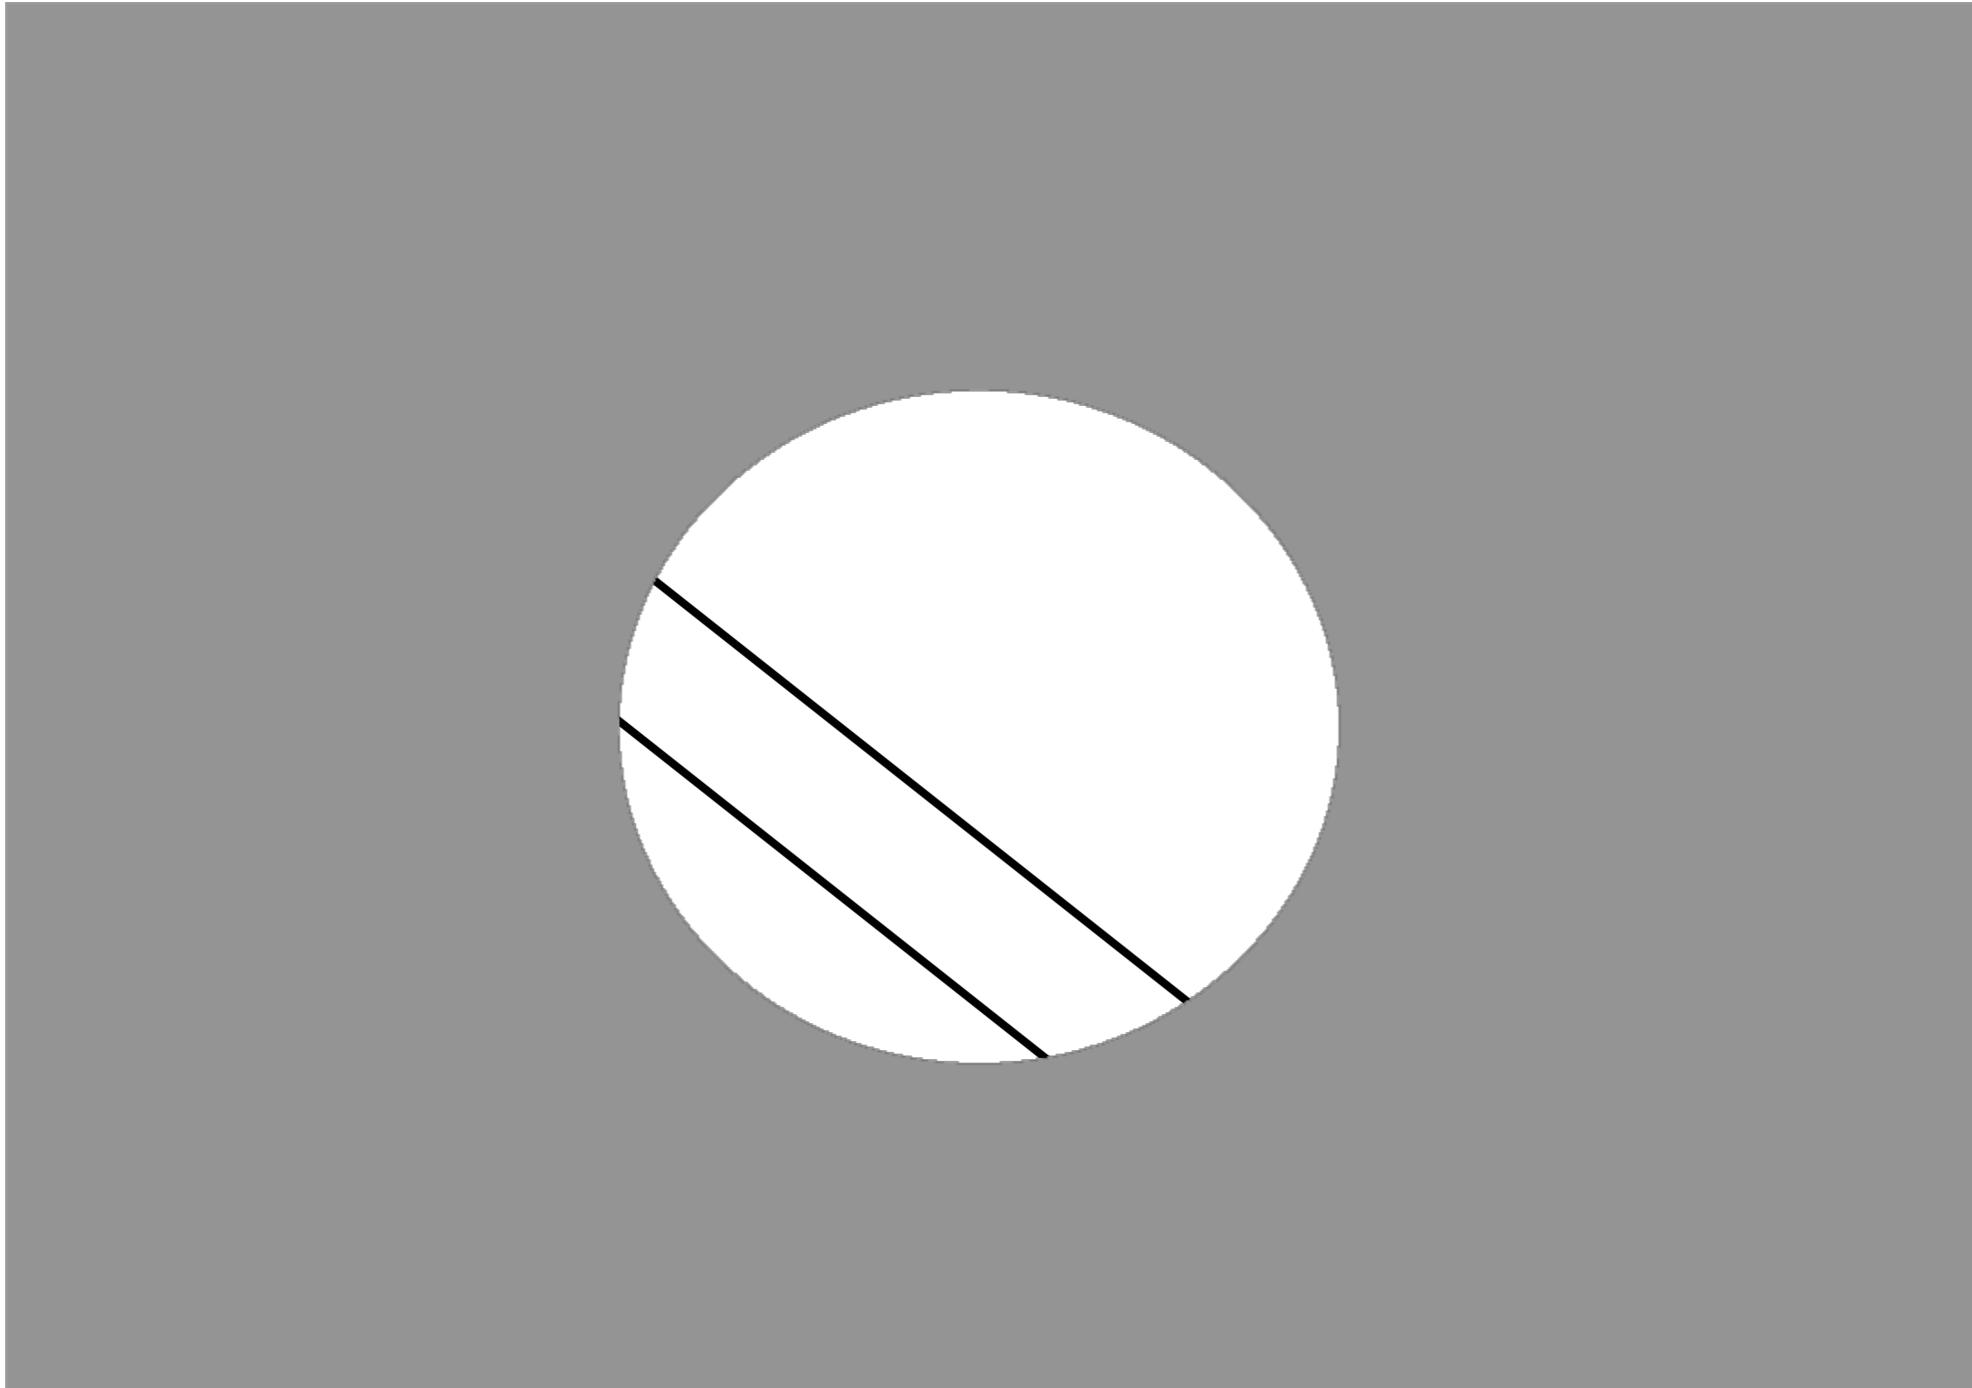
$$E_{SSD}(u, v) \approx \sum_{(x,y) \in R} (u \cdot I_x(x, y, t) + v \cdot I_y(x, y, t) + I_t(x, y, t))^2$$

- ◆ By minimizing this Taylor series approximation to the SSD, we are trying to enforce the so-called **optical flow constraint equation** (OFCE):

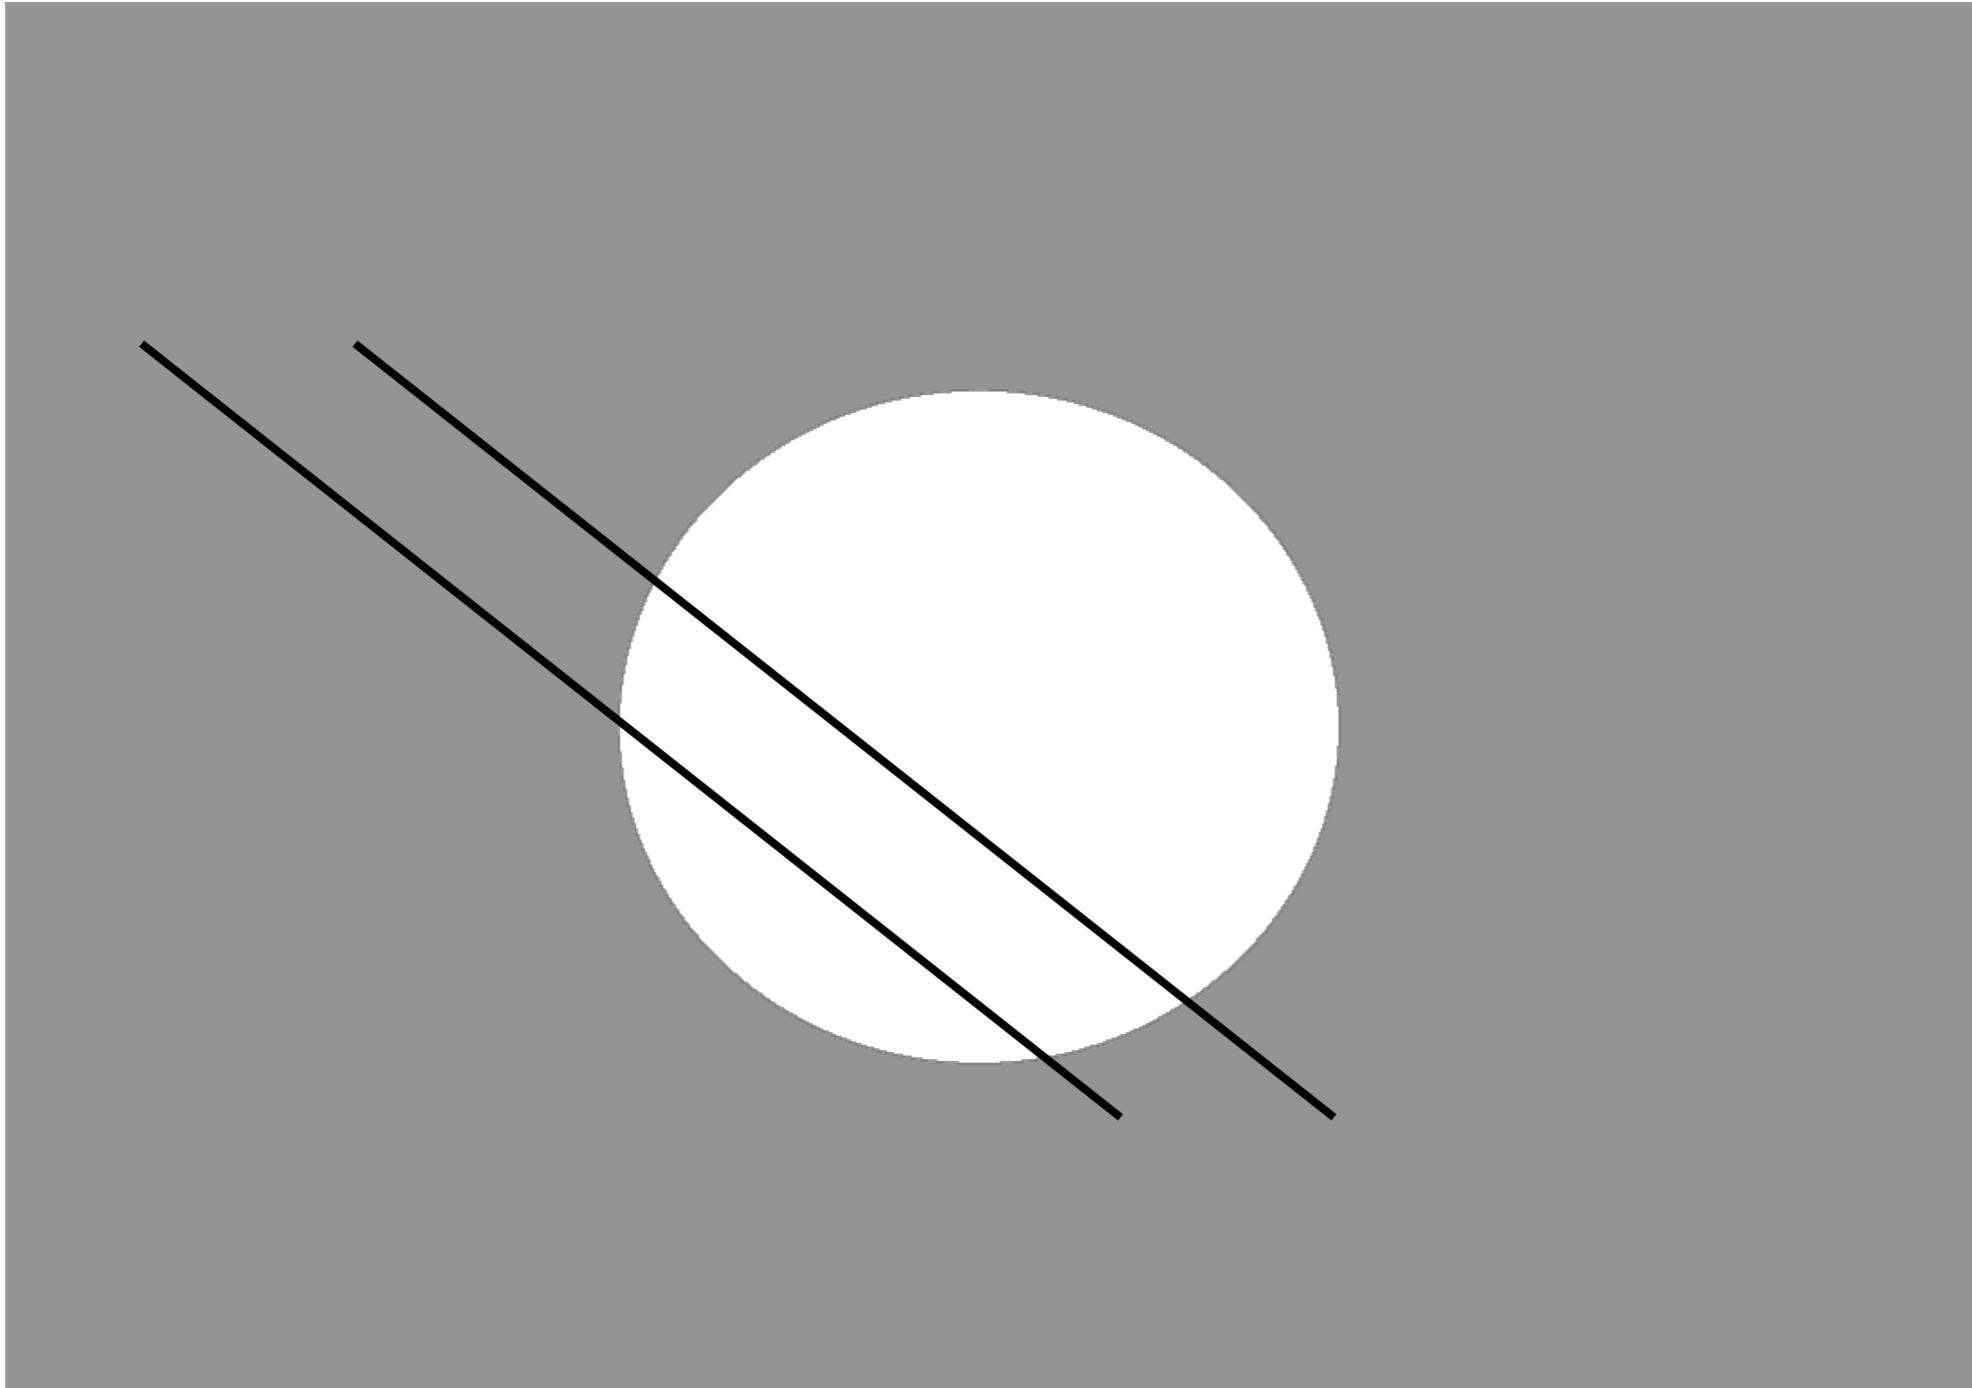
$$u \cdot I_x + v \cdot I_y + I_t = 0$$

- ◆ This is also called the linearized brightness constancy constraint.

# Aperture Problem

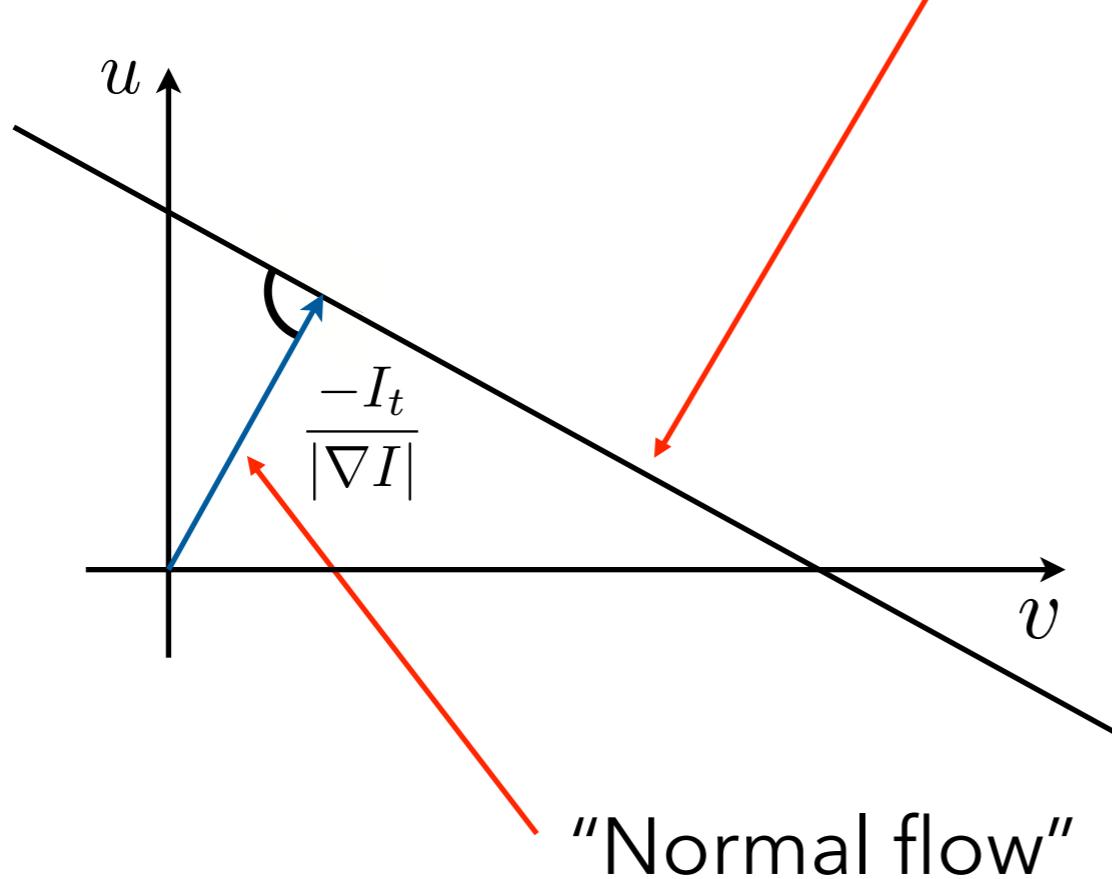


# Aperture Problem



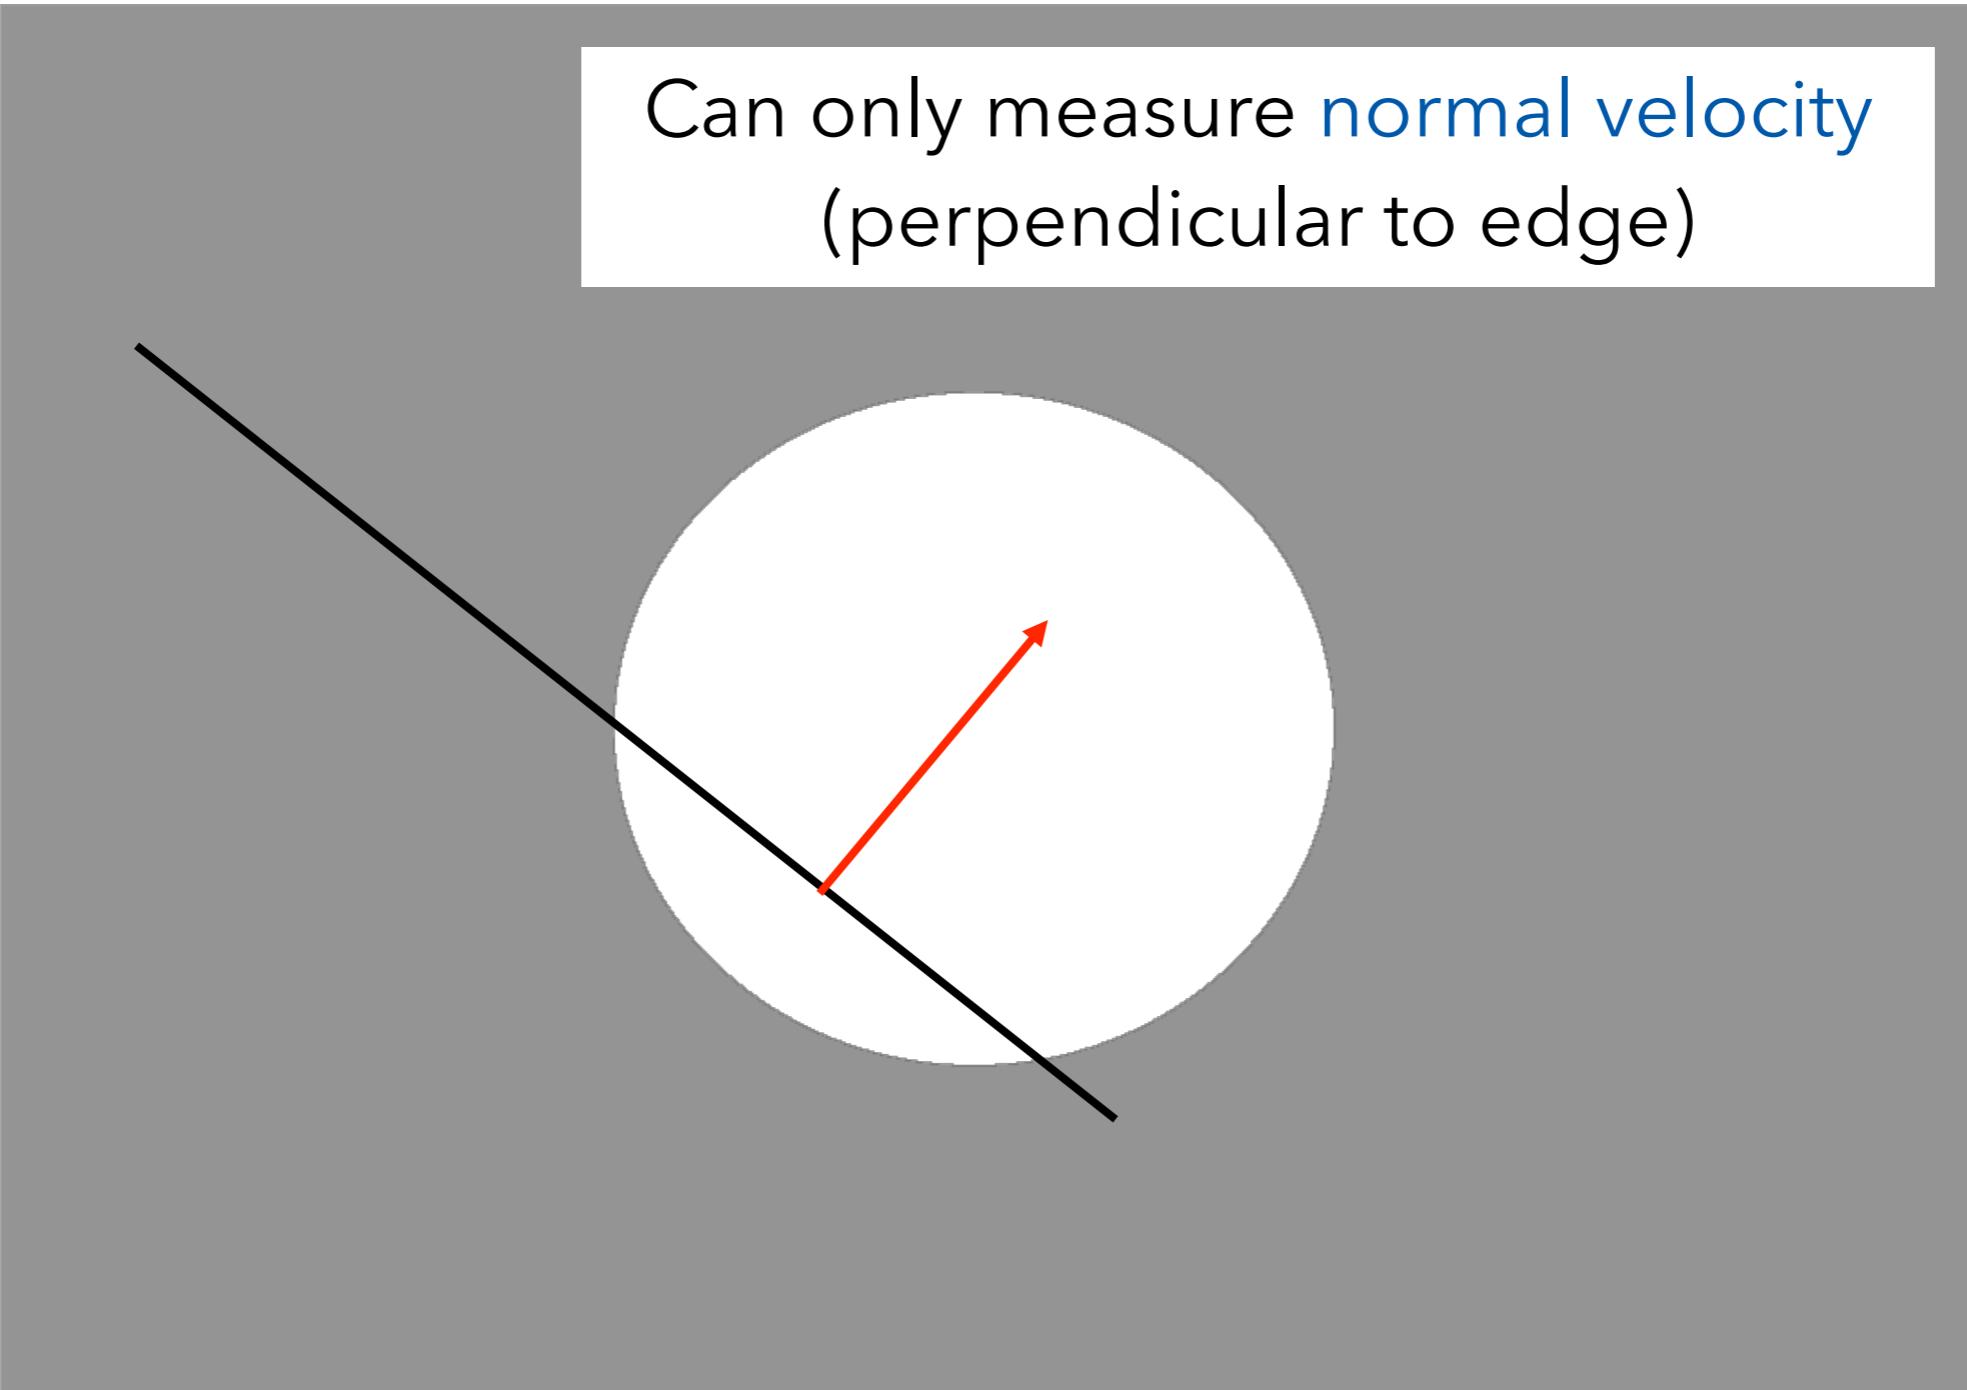
If we take a single image pixel, we get a constraint line:

$$u \cdot I_x + v \cdot I_y + I_t = 0$$

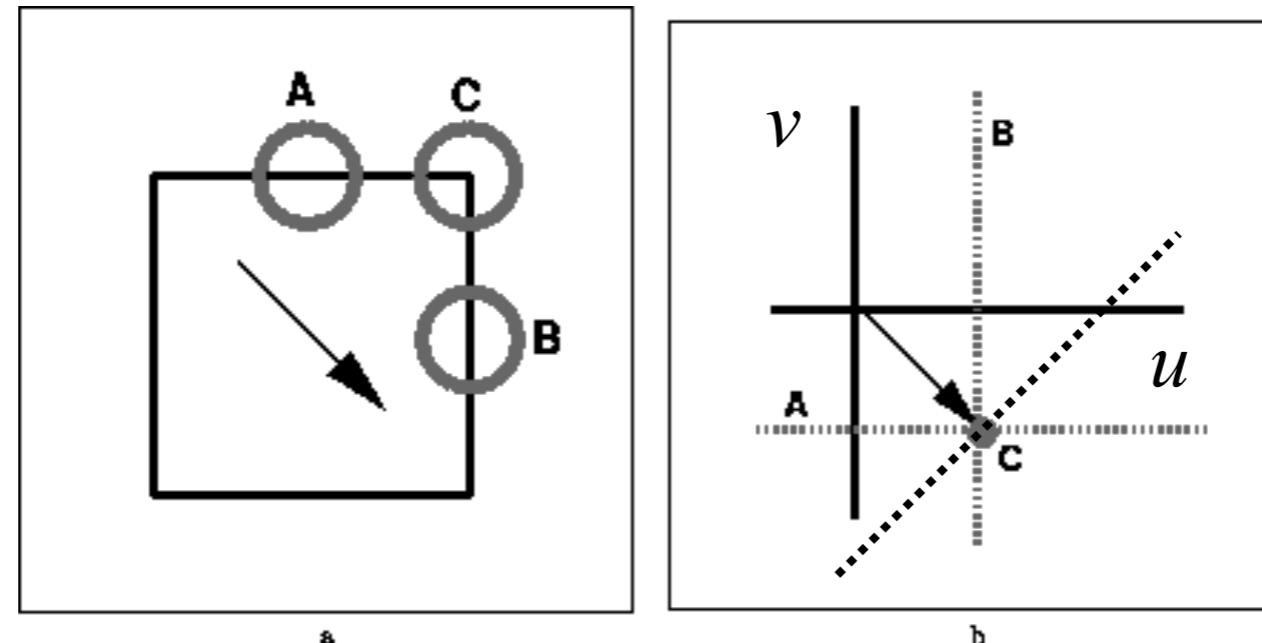


# Aperture Problem

Can only measure **normal velocity**  
(perpendicular to edge)



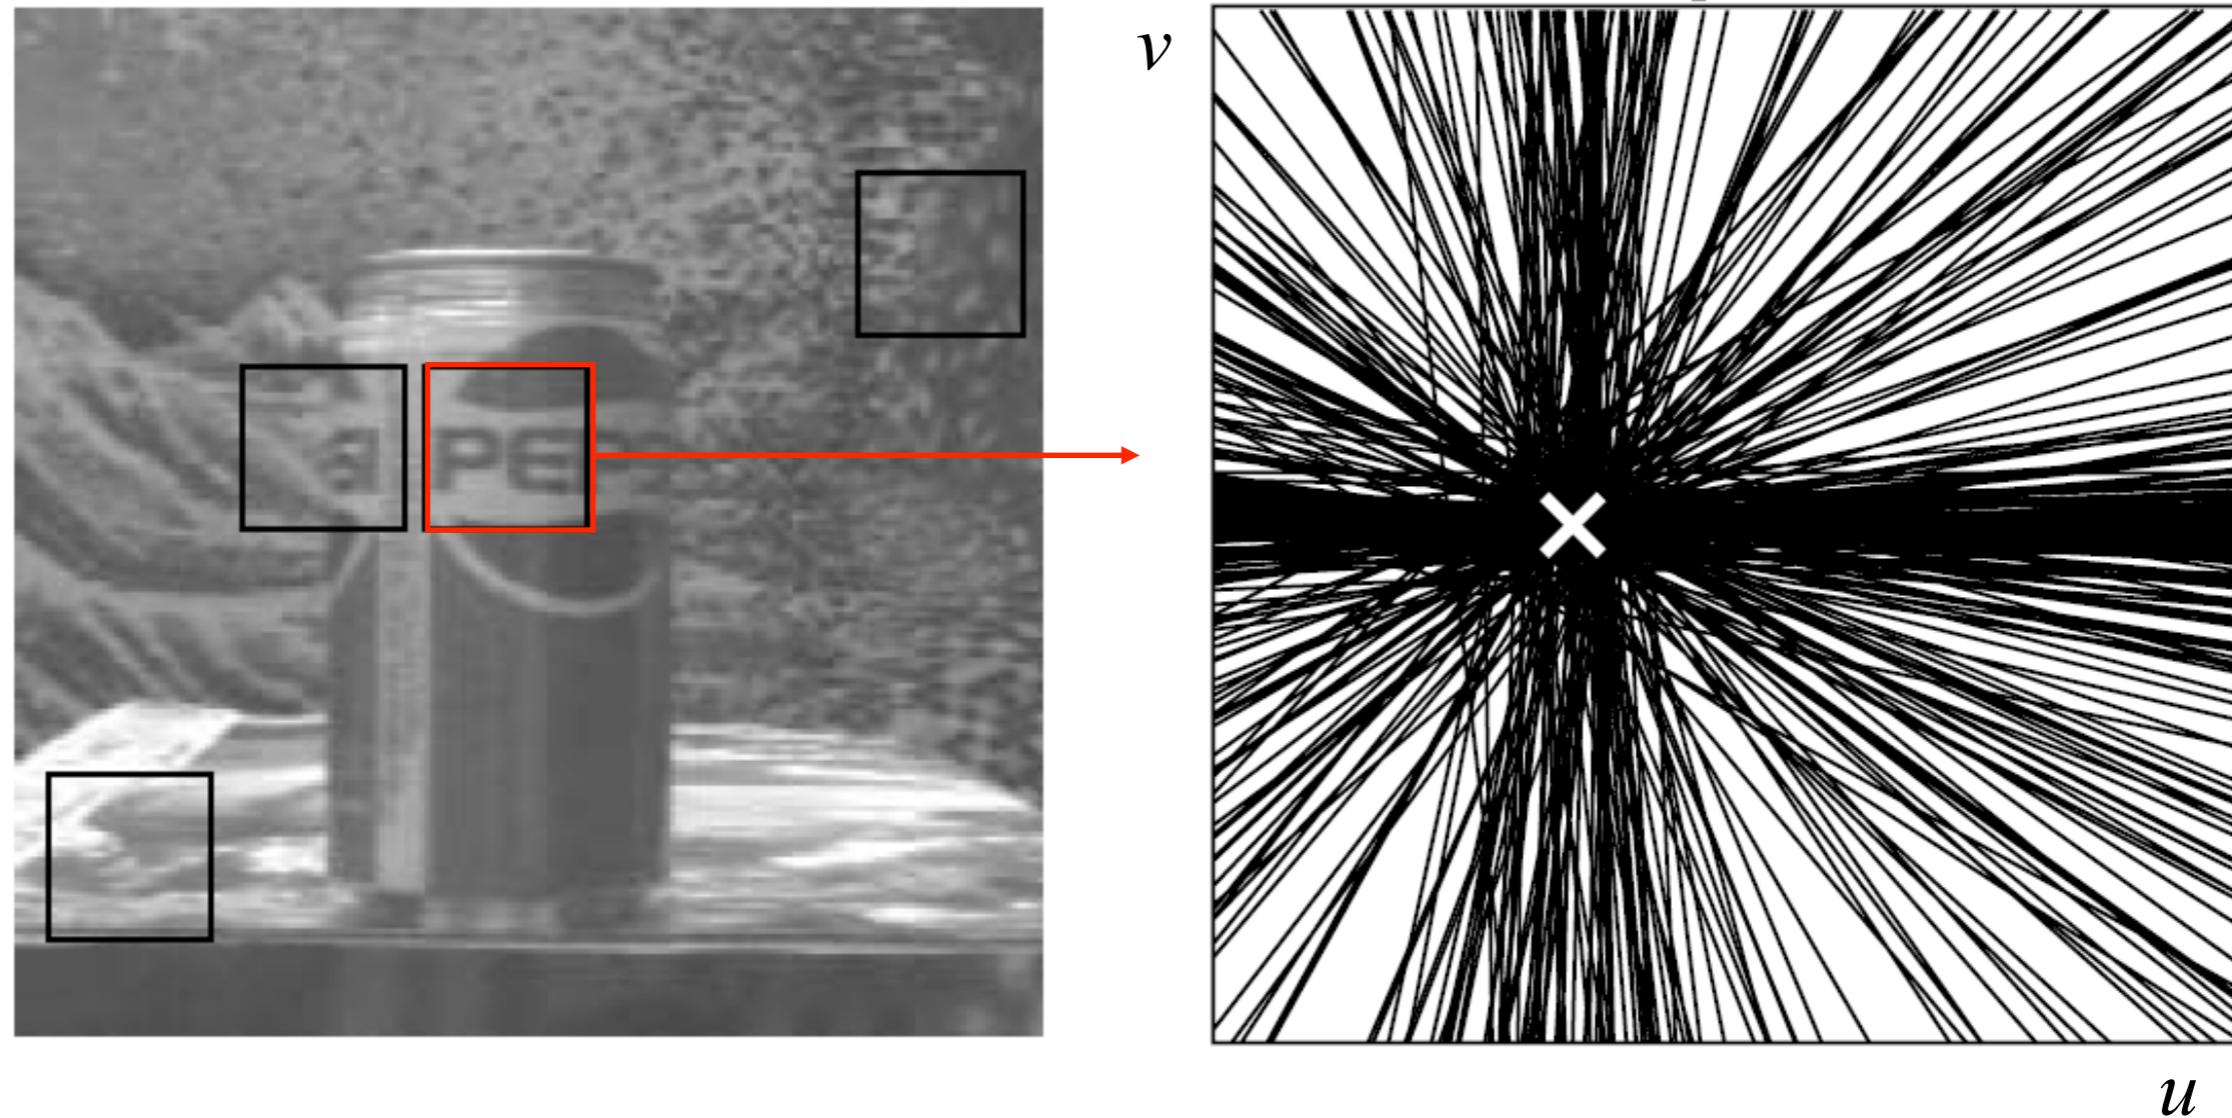
# Multiple Constraints



[Yair Weiss]

Combine multiple constraints to get an estimate of the velocity (not only the normal component).

# Multiple Constraints



# Area-Based Flow

- ◆ How do we combine multiple constraints?
  - ◆ Remember our assumptions.
  - ◆ We will assume **spatial smoothness** of the flow.
- ◆ More specifically:
  - ◆ Assume that the flow is **constant** in a region.

$$E_{SSD}(u, v) \approx \sum_{(x,y) \in R} (u \cdot I_x(x, y, t) + v \cdot I_y(x, y, t) + I_t(x, y, t))^2$$

This is what we have been doing (sliding window).

- ◆ But how do we solve for the motion?

# Optimization

$$E_{SSD}(u, v) \approx \sum_{(x,y) \in R} (u \cdot I_x(x, y, t) + v \cdot I_y(x, y, t) + I_t(x, y, t))^2$$

Differentiate with respect to  $u$  and  $v$  and set this to zero.

$$\frac{\partial}{\partial u} E_{SSD}(u, v) \approx 2 \sum_{(x,y) \in R} (u \cdot I_x(x, y, t) + v \cdot I_y(x, y, t) + I_t(x, y, t)) I_x(x, y, t) = 0$$

$$\frac{\partial}{\partial v} E_{SSD}(u, v) \approx 2 \sum_{(x,y) \in R} (u \cdot I_x(x, y, t) + v \cdot I_y(x, y, t) + I_t(x, y, t)) I_y(x, y, t) = 0$$

# Optimization

$$\frac{\partial E_{SSD}}{\partial u} \approx 2 \sum_R (u \cdot I_x + v \cdot I_y + I_t) I_x = 0$$

$$\frac{\partial E_{SSD}}{\partial v} \approx 2 \sum_R (u \cdot I_x + v \cdot I_y + I_t) I_y = 0$$

Rearrange the terms and drop the constant:

$$\left[ \sum_R I_x^2 \right] u + \left[ \sum_R I_x I_y \right] v = - \left[ \sum_R I_x I_t \right]$$

$$\left[ \sum_R I_x I_y \right] u + \left[ \sum_R I_y^2 \right] v = - \left[ \sum_R I_y I_t \right]$$

# Optimization

$$\left[ \sum_R I_x^2 \right] u + \left[ \sum_R I_x I_y \right] v = - \left[ \sum_R I_x I_t \right]$$

$$\left[ \sum_R I_x I_y \right] u + \left[ \sum_R I_y^2 \right] v = - \left[ \sum_R I_y I_t \right]$$

System of 2 equations in 2 unknowns:

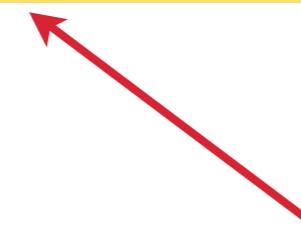
$$\begin{bmatrix} \sum_R I_x^2 & \sum_R I_x I_y \\ \sum_R I_x I_y & \sum_R I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} -\sum_R I_x I_t \\ -\sum_R I_y I_t \end{bmatrix}$$



Symmetric positive definite (invertible)

# Structure tensor

$$\begin{bmatrix} \sum_R I_x^2 & \sum_R I_x I_y \\ \sum_R I_x I_y & \sum_R I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} -\sum_R I_x I_t \\ -\sum_R I_y I_t \end{bmatrix}$$



**Structure tensor**

- ◆ The structure tensor is an important tool in vision.
  - ◆ The eigenvectors give the principal directions of the local image variation.
  - ◆ The eigenvalues indicate their strength.
- ◆ You have seen this in CV I when talking about feature detectors!

# Optimization

$$\begin{bmatrix} \sum_R I_x^2 & \sum_R I_x I_y \\ \sum_R I_x I_y & \sum_R I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} -\sum_R I_x I_t \\ -\sum_R I_y I_t \end{bmatrix}$$

- ◆ Rewrite using the abbreviations from before:

$$\left( \sum_R \nabla I \nabla I^T \right) \mathbf{u} = - \sum_R I_t \nabla I$$

- ◆ Invert structure tensor to obtain flow:

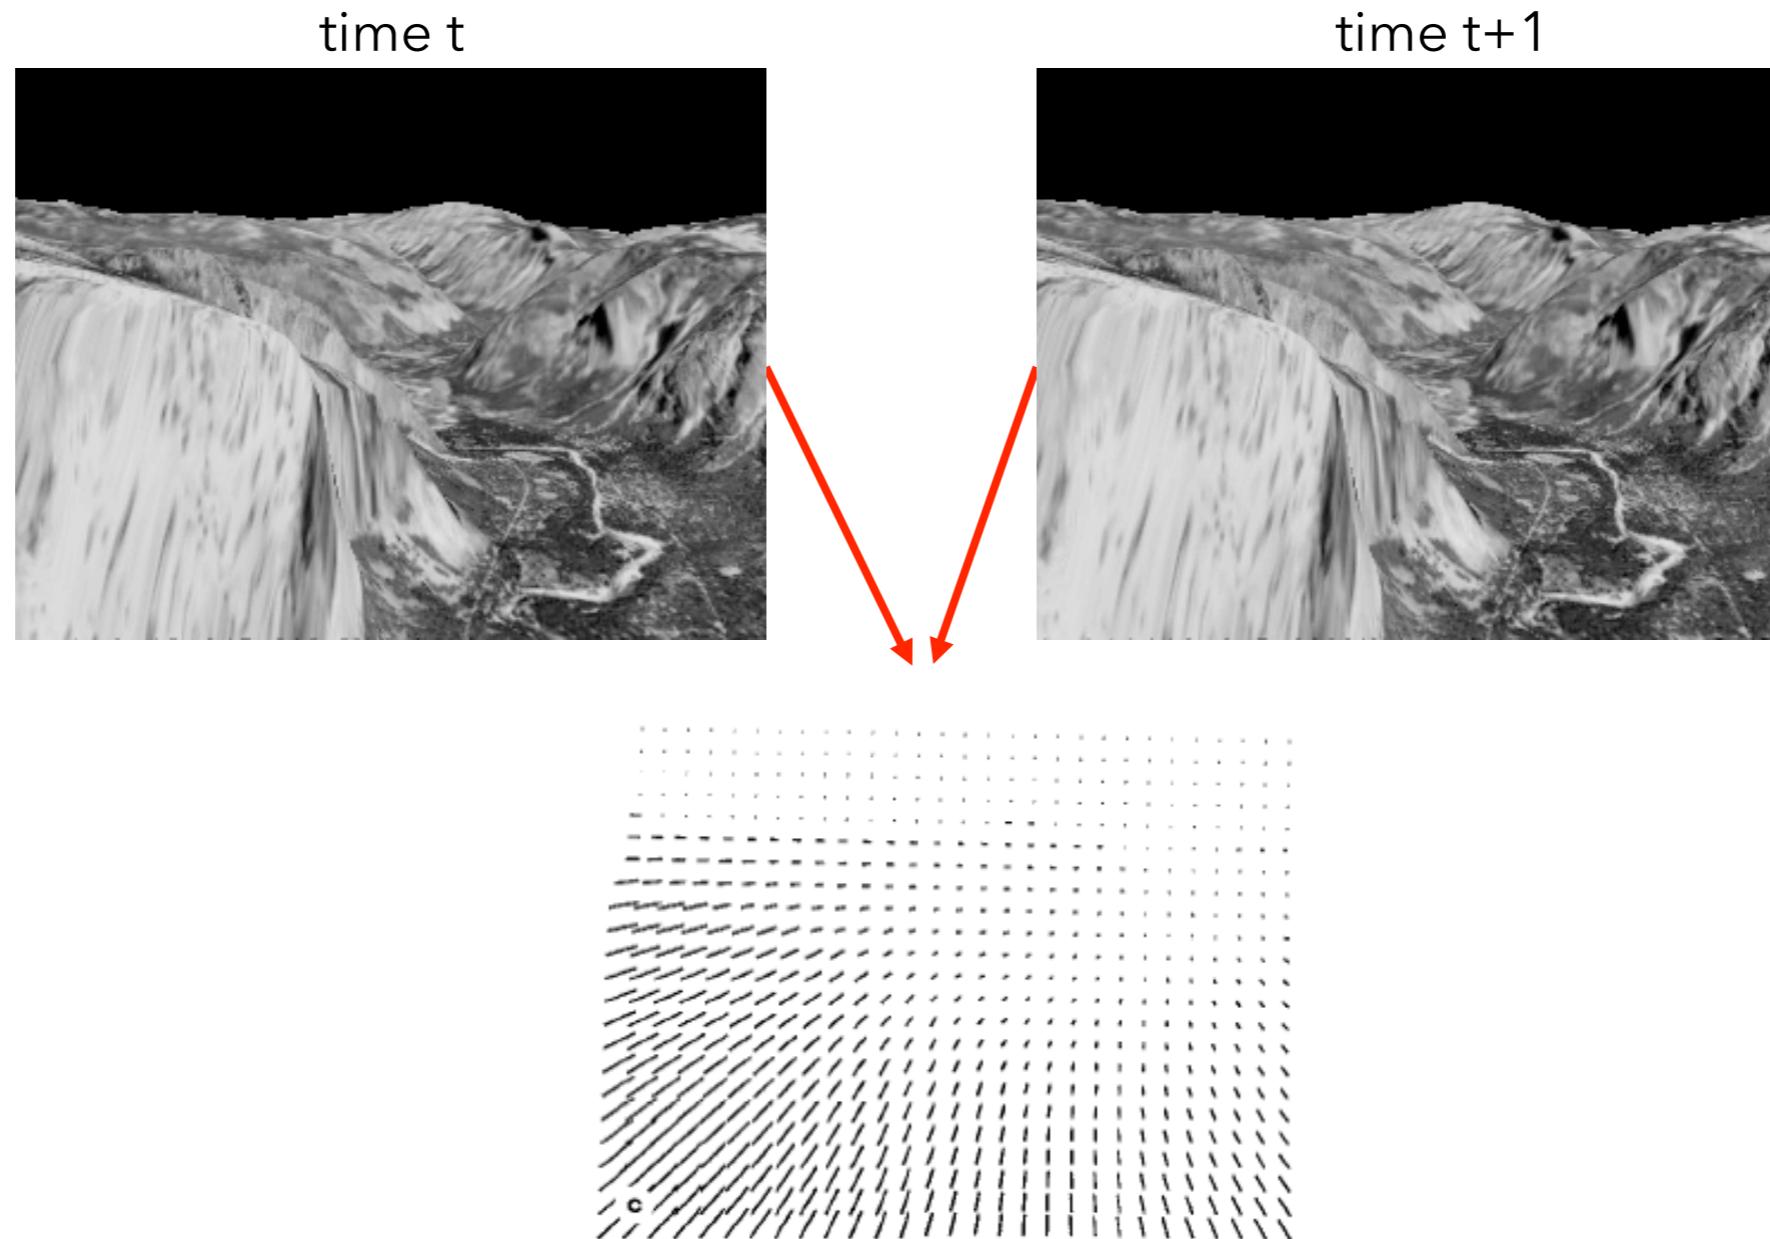
$$\mathbf{u} = - \left( \sum_R \nabla I \nabla I^T \right)^{-1} \left( \sum_R I_t \nabla I \right)$$

# Solving for $\mathbf{u}$

$$\mathbf{u} = - \left( \sum_R \nabla I \nabla I^T \right)^{-1} \left( \sum_R I_t \nabla I \right)$$

- ◆ This is a classical flow technique:  
**B. D. Lucas and T. Kanade**. An iterative image registration technique with an application to stereo vision. IJCAI, pp. 674–679, 1981.

# Optical Flow Estimation



Our goal was to estimate a “dense” flow field with a motion vector for every pixel.

# What have we achieved so far?



- ◆ We talked about the [Lucas-Kanade method](#) and estimated the flow using the following expression:

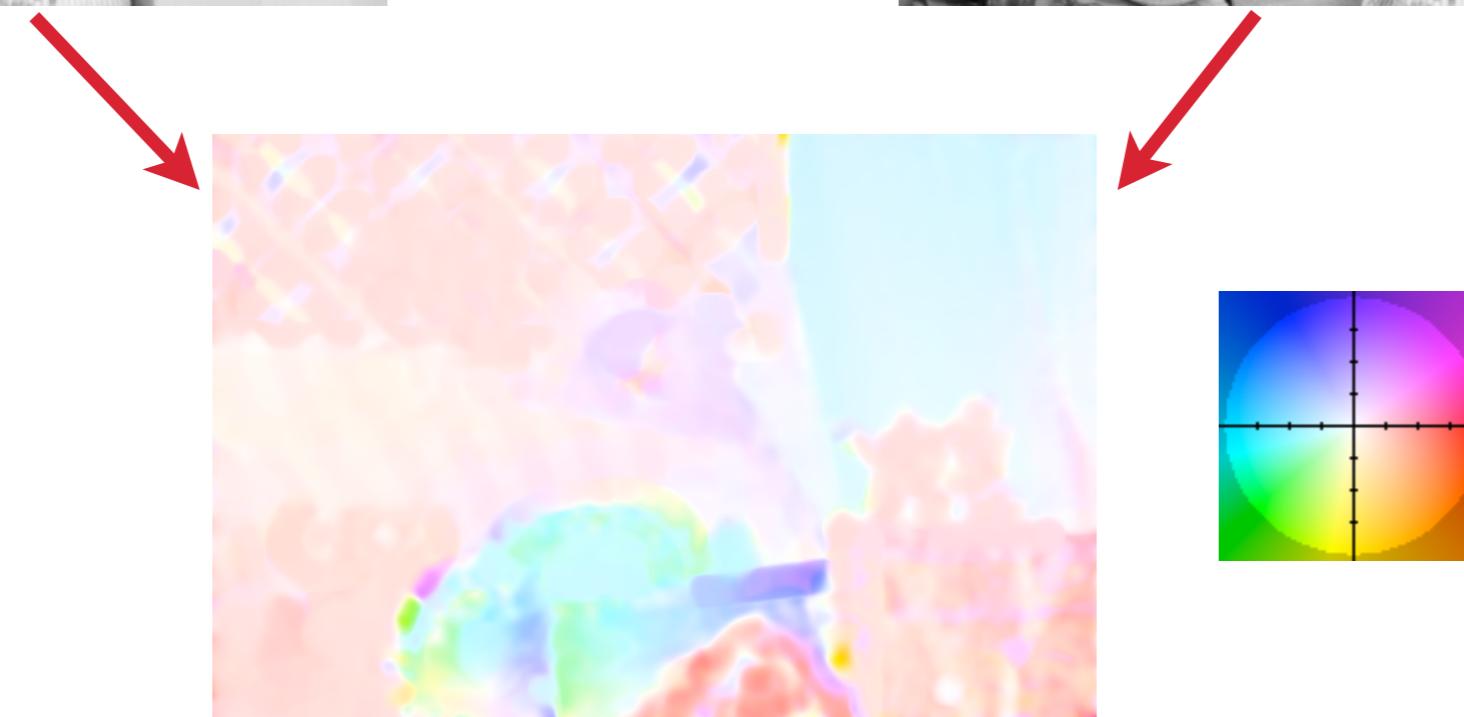
$$\mathbf{u} = - \left( \sum_R \nabla I \nabla I^T \right)^{-1} \left( \sum_R I_t \nabla I \right)$$

- ◆ Initially we motivated this using a small window  $R$ .
- ◆ In the end we took the region to mean the entire image and computed a single motion vector for the image.
- ◆ This allows us to do [image registration](#).
- ◆ We refined this based on an affine flow model.

- ◆ What if we now want to estimate **dense flow**?
  - ◆ We can just take the region  $R$  to be a small region around every pixel and compute a flow vector for every pixel.
- ◆ But we face a crucial problem:
  - ◆ The LK method only works for **small motions**.
- ◆ Two workarounds (use both):
  - ◆ Iterative estimation.
  - ◆ Coarse-to-fine estimation.

# Iterative Estimation

- ◆ Take image pair and compute LK flow (21x21 window):





Frame 1



Frame 2



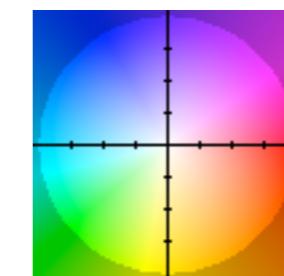
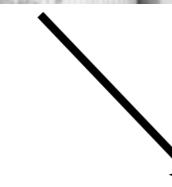
Frame 2 (warped)



Frame 1

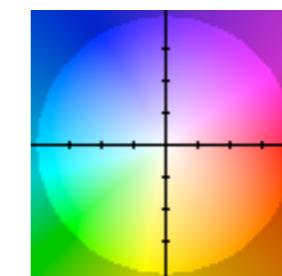
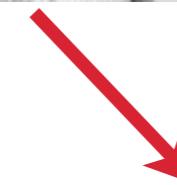
# Iterative Estimation

- ◆ Backward warp the second image toward the first:



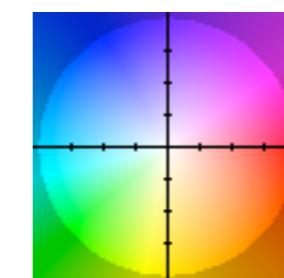
# Iterative Estimation

- ◆ Estimate incremental flow from warped image pair:



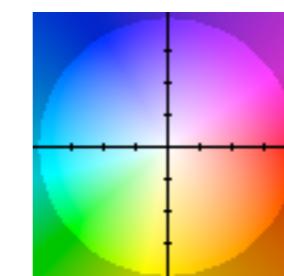
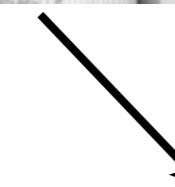
# Iterative Estimation

- ◆ Add incremental flow to previous estimate:



# Iterative Estimation

- ◆ Warp again and so on... until convergence.



# Coarse-to-fine Estimation

- ◆ Of coarse, we also apply the usual trick and use a Gaussian pyramid:
  - ◆ Initialize with the flow from a coarser level.
  - ◆ If we do this on the previous image pair, we get this:

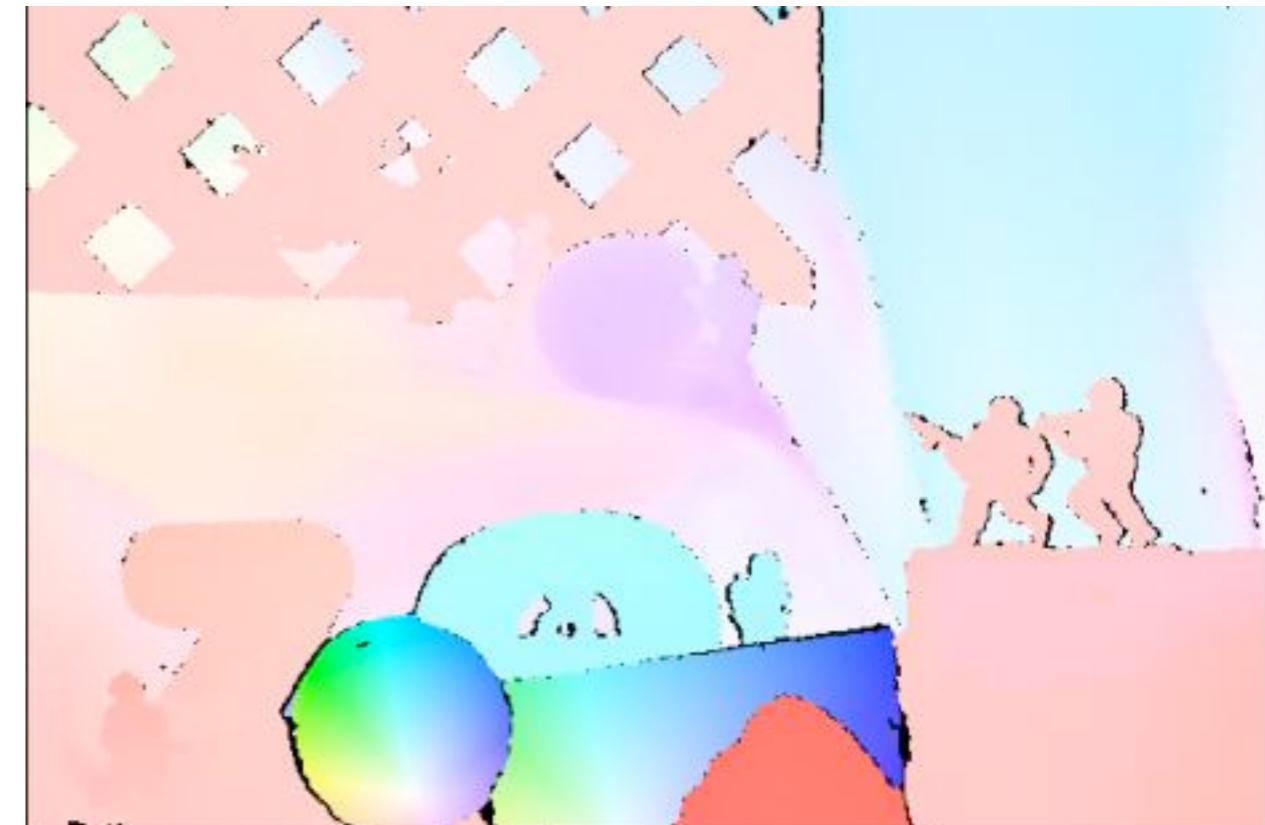


# Is that a good result?

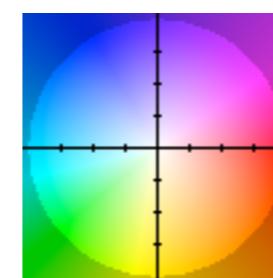
- ◆ Let's see:
- ◆ Maybe not...



Pyramid LK method

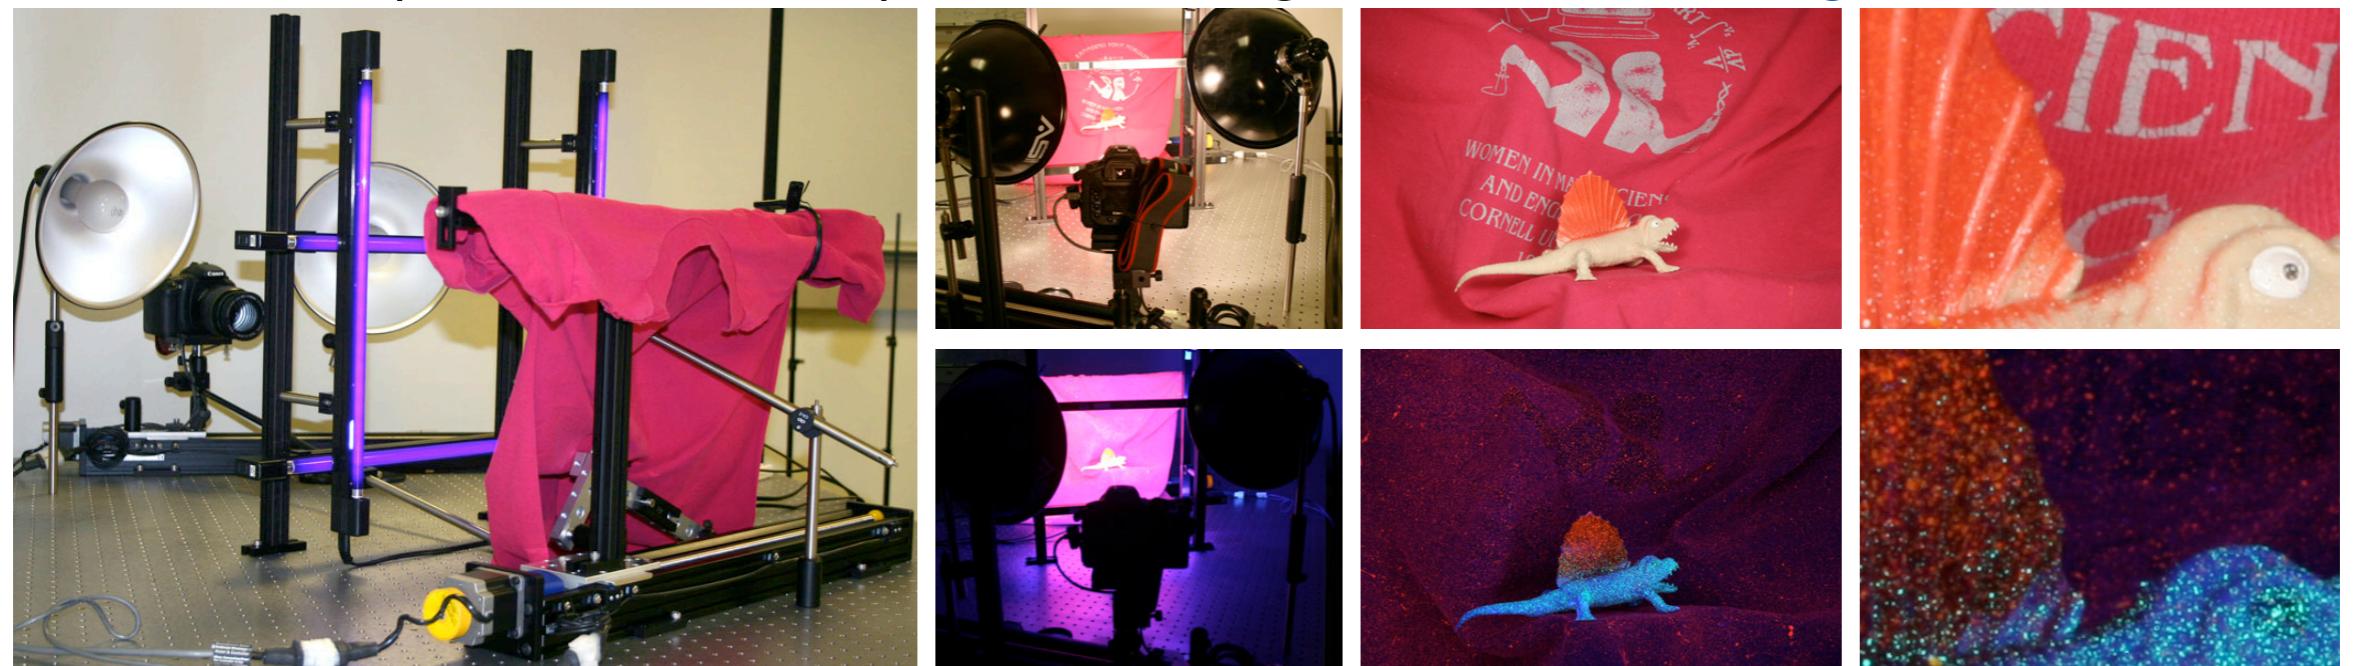


Ground truth



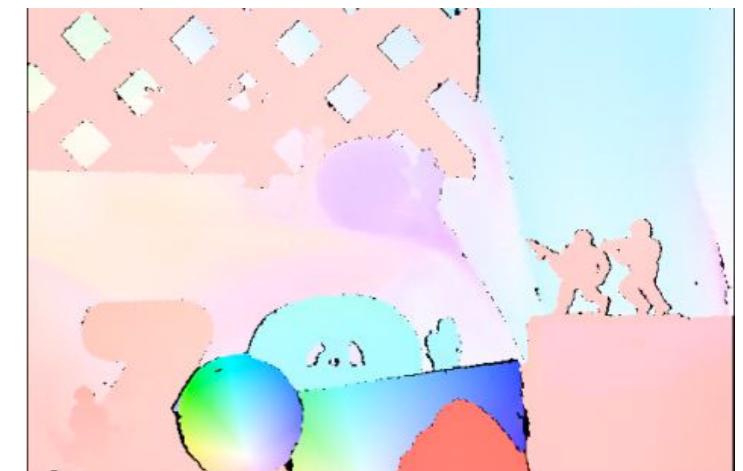
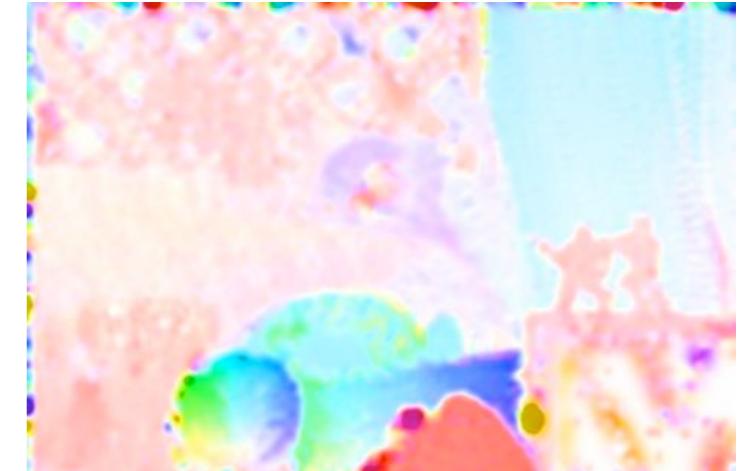
# Aside: How to get the ground truth?

- ◆ Until recently, there were only very few image sequences with ground truth flow available.
  - ◆ All of them were **synthetic** and furthermore very simple.
- ◆ Middlebury optical flow benchmark [Baker et al., 07]:
  - ◆ Contains a number of complex synthetic and real sequences with ground truth.
  - ◆ Real sequences are captured using normal + **UV light**:



# What is the problem?

- ◆ The window is **too big**!
  - ◆ All the **discontinuities** in the flow are **smoothed** over.
  - ◆ But discontinuities do exist, e.g., at motion boundaries.
- ◆ But: The window is also **too small**!
  - ◆ In some areas, the flow is really bad, because there is **not enough image information** in the window.
  - ◆ This happens in areas with little texture.
- ◆ LK is a **local** optical flow method.



# First Option: Revisiting the assumptions

- ◆ Before we move on to something more complex, revisit the assumptions:
  - ◆ brightness constancy: SSD error as objective
  - ◆ spatial coherence: motion within a patch is constant!
- ◆ Constant motion: Only translational motion within the image plane is allowed

# Translational Model



What's wrong with the translational assumption (i.e. constant motion within a region R)?

How can we generalize it?

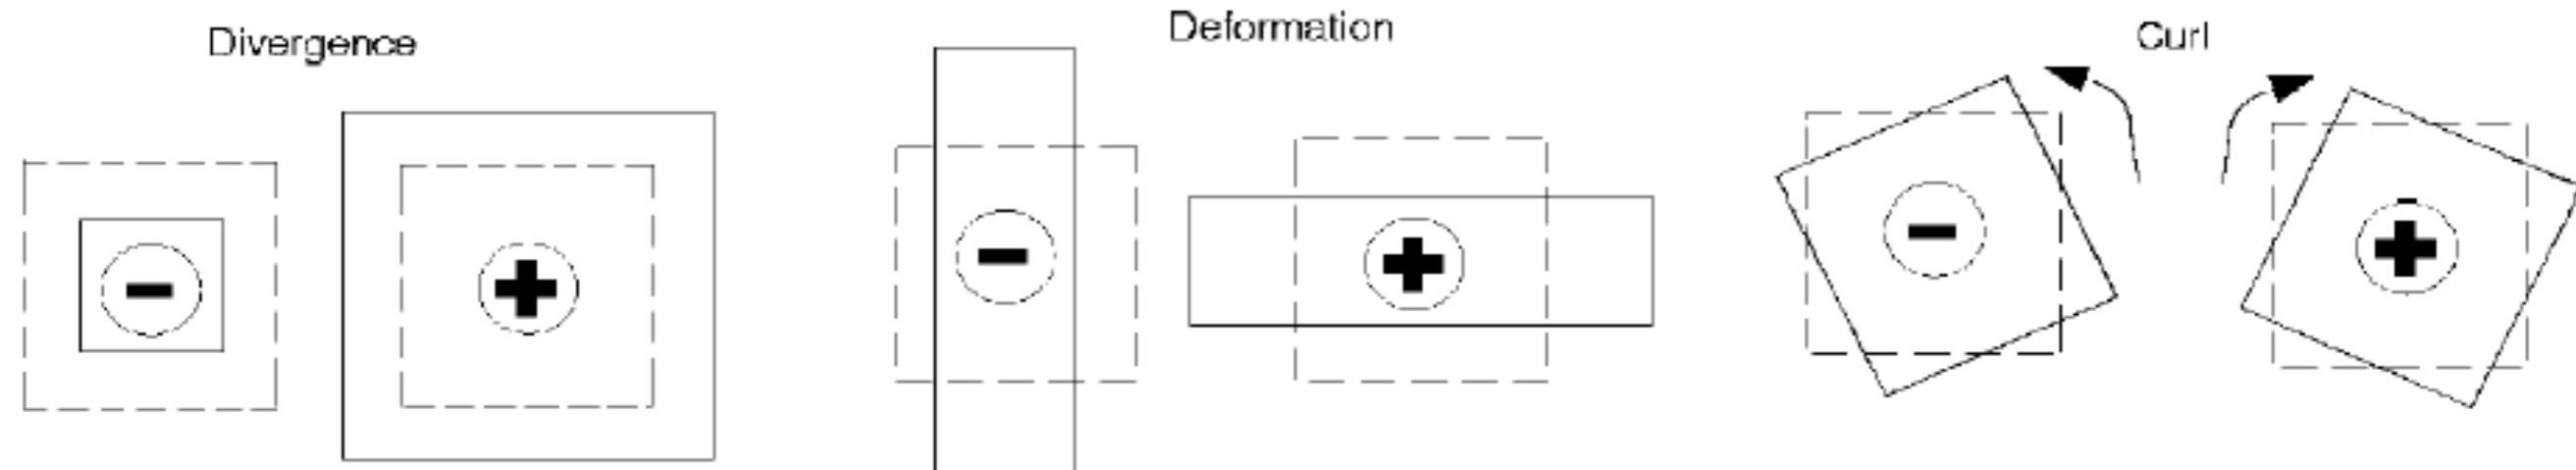
- ◆ Assume a parametric motion model within the region
  - ◆ Old idea
  - ◆ Recent renewed interest: efficient to implement

# Affine Flow

- ◆ Parameterize the flow within the region:

$$\mathbf{u}(\mathbf{x}; \mathbf{a}) = \begin{bmatrix} u(\mathbf{x}; \mathbf{a}) \\ v(\mathbf{x}; \mathbf{a}) \end{bmatrix} = \begin{bmatrix} a_1 + a_2x + a_3y \\ a_4 + a_5x + a_6y \end{bmatrix}$$

$$E(\mathbf{a}) = \sum_{\mathbf{x} \in R} (\nabla I(\mathbf{x})^T \mathbf{u}(\mathbf{x}; \mathbf{a}) + I_t(\mathbf{x}))^2$$



# Linear Basis

This can be thought of as a **basis representation** of the flow within the region.

$\mathbf{u}(\mathbf{x}; \mathbf{c})$

$$\begin{matrix} \nearrow \searrow \nwarrow \swarrow \\ \downarrow \uparrow \downarrow \uparrow \\ \leftarrow \rightarrow \leftarrow \rightarrow \\ \uparrow \downarrow \uparrow \downarrow \\ \searrow \nearrow \swarrow \nwarrow \end{matrix} = c_1 * \begin{matrix} \nearrow \searrow \nwarrow \swarrow \\ \downarrow \uparrow \downarrow \uparrow \\ \leftarrow \rightarrow \leftarrow \rightarrow \\ \uparrow \downarrow \uparrow \downarrow \\ \searrow \nearrow \swarrow \nwarrow \end{matrix} + c_2 * \begin{matrix} \nearrow \searrow \nwarrow \swarrow \\ \downarrow \uparrow \downarrow \uparrow \\ \leftarrow \rightarrow \leftarrow \rightarrow \\ \uparrow \downarrow \uparrow \downarrow \\ \searrow \nearrow \swarrow \nwarrow \end{matrix} + c_3 * \begin{matrix} \nearrow \searrow \nwarrow \swarrow \\ \downarrow \uparrow \downarrow \uparrow \\ \leftarrow \rightarrow \leftarrow \rightarrow \\ \uparrow \downarrow \uparrow \downarrow \\ \searrow \nearrow \swarrow \nwarrow \end{matrix} + c_4 * \begin{matrix} \nearrow \searrow \nwarrow \swarrow \\ \downarrow \uparrow \downarrow \uparrow \\ \leftarrow \rightarrow \leftarrow \rightarrow \\ \uparrow \downarrow \uparrow \downarrow \\ \searrow \nearrow \swarrow \nwarrow \end{matrix} + c_5 * \begin{matrix} \nearrow \searrow \nwarrow \swarrow \\ \downarrow \uparrow \downarrow \uparrow \\ \leftarrow \rightarrow \leftarrow \rightarrow \\ \uparrow \downarrow \uparrow \downarrow \\ \searrow \nearrow \swarrow \nwarrow \end{matrix} + c_6 * \begin{matrix} \nearrow \searrow \nwarrow \swarrow \\ \downarrow \uparrow \downarrow \uparrow \\ \leftarrow \rightarrow \leftarrow \rightarrow \\ \uparrow \downarrow \uparrow \downarrow \\ \searrow \nearrow \swarrow \nwarrow \end{matrix}$$

$$\mathbf{u}(\mathbf{x}; \mathbf{c}) = \sum_{j=1}^n c_j \mathbf{B}_j(\mathbf{x})$$

# Optimization

$$E(A) = \sum_{(x,y) \in R} (I_x a_1 + I_x a_2 x + I_x a_3 y + I_y a_4 + I_y a_5 x + I_y a_6 y + I_t)^2$$

Differentiate wrt the  $a_i$  and set equal to zero:

$$\begin{bmatrix} \sum I_x^2 & \sum I_x^2 x & \sum I_x^2 y & \sum I_x I_y & \sum I_x I_y x & \sum I_x I_y y \\ \sum I_x^2 x & \sum I_x^2 x^2 & \sum I_x^2 xy & \sum I_x I_y x & \sum I_x I_y x^2 & \sum I_x I_y xy \\ & & & \vdots & & \end{bmatrix} \begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \\ a_5 \\ a_6 \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_x I_t x \\ \sum I_x I_t y \\ \sum I_y I_t \\ \sum I_y I_t x \\ \sum I_y I_t y \end{bmatrix}$$

# Affine Flow Estimation

- ◆ As before, we solve for the parameters (now only more) by inverting the matrix.
- ◆ We can use the iterative image registration (CV1)
  - ◆ help overcome small motion assumption
- ◆ Need warping in its full generality:

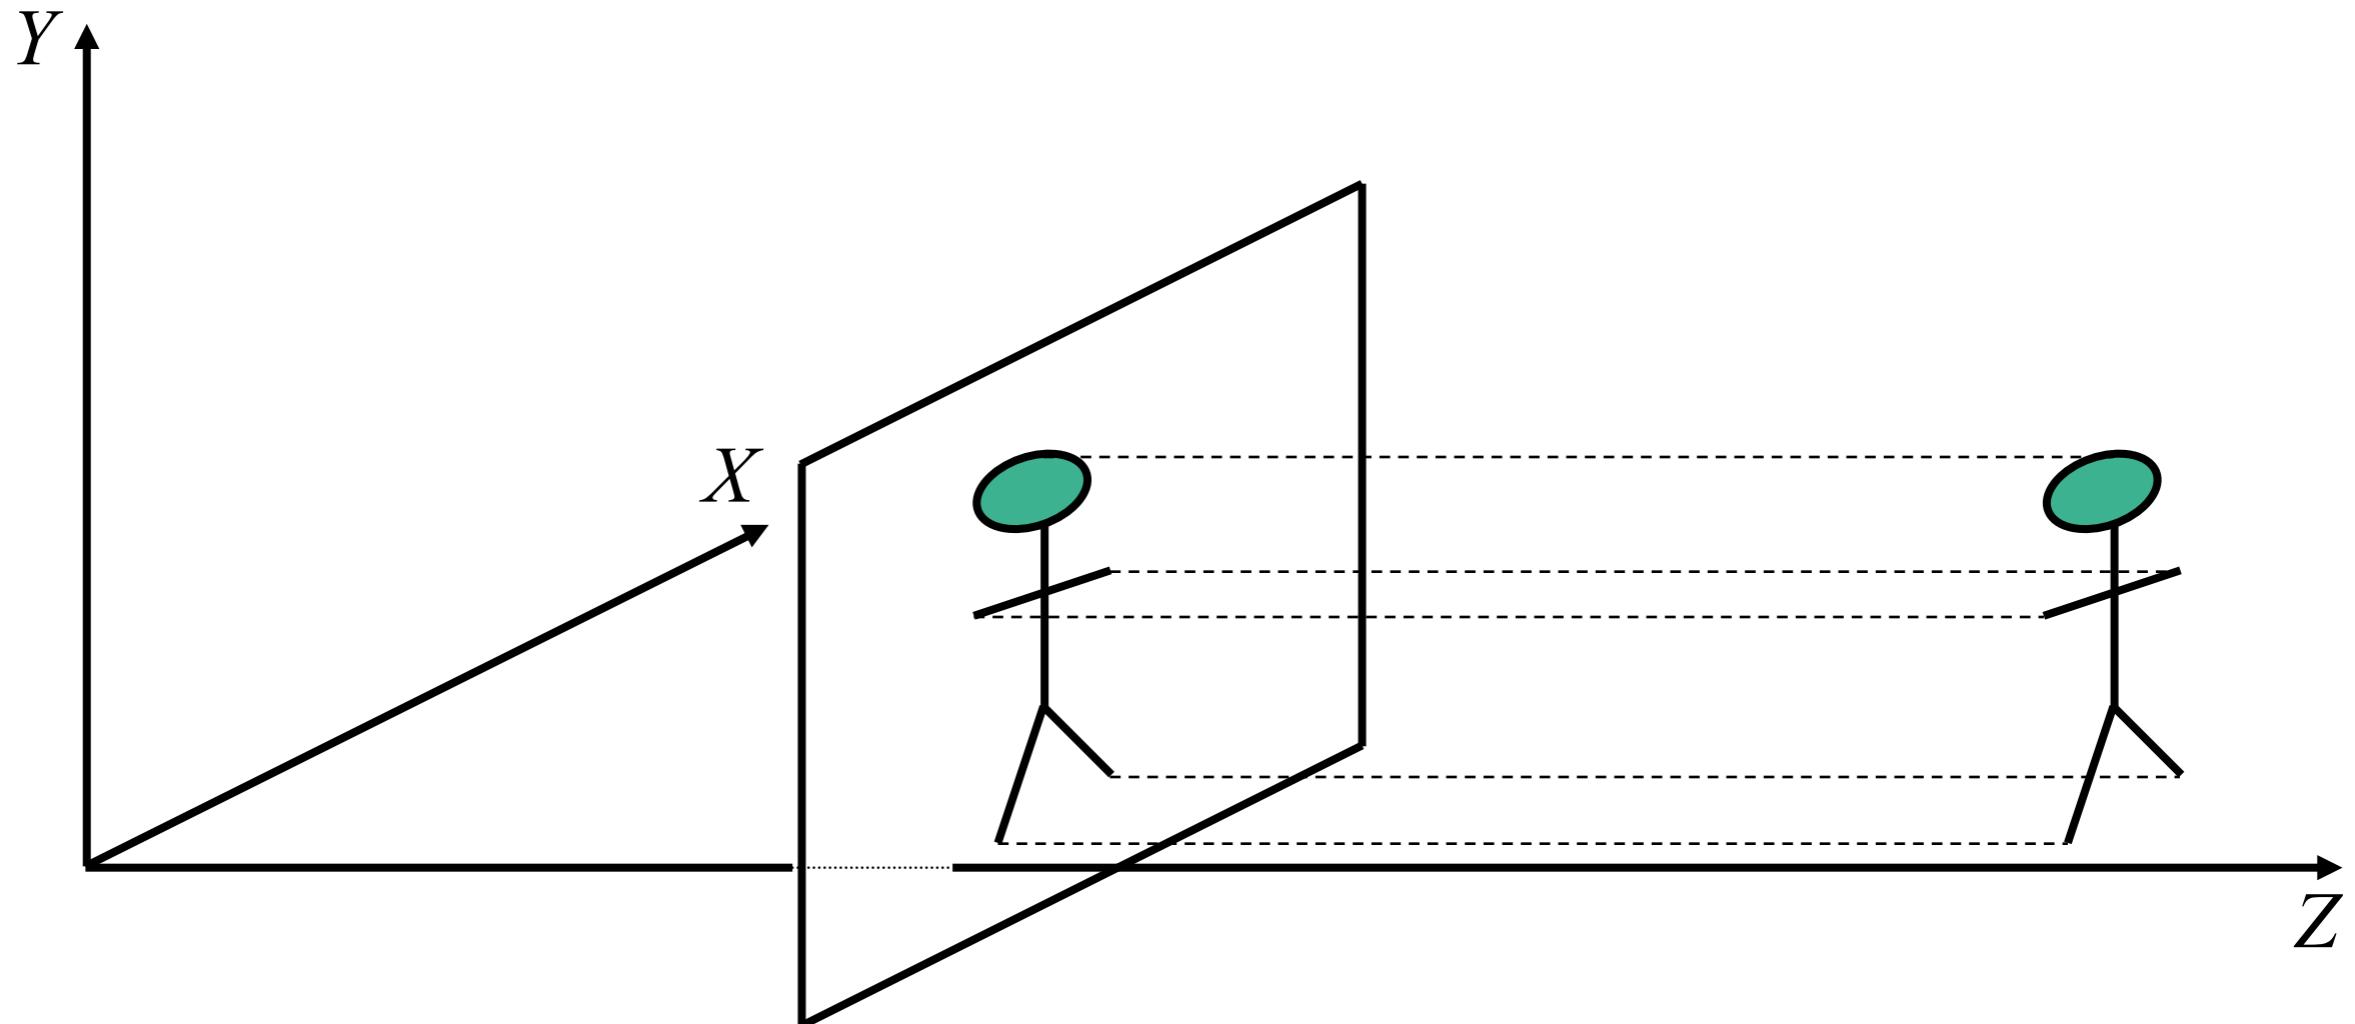


affine

# Why Affine?

- ◆ Where does this affine approximation come from?
- ◆ All our models are approximations to the world. What are the assumptions in the affine approximation?
- ◆ For this we need some geometry.

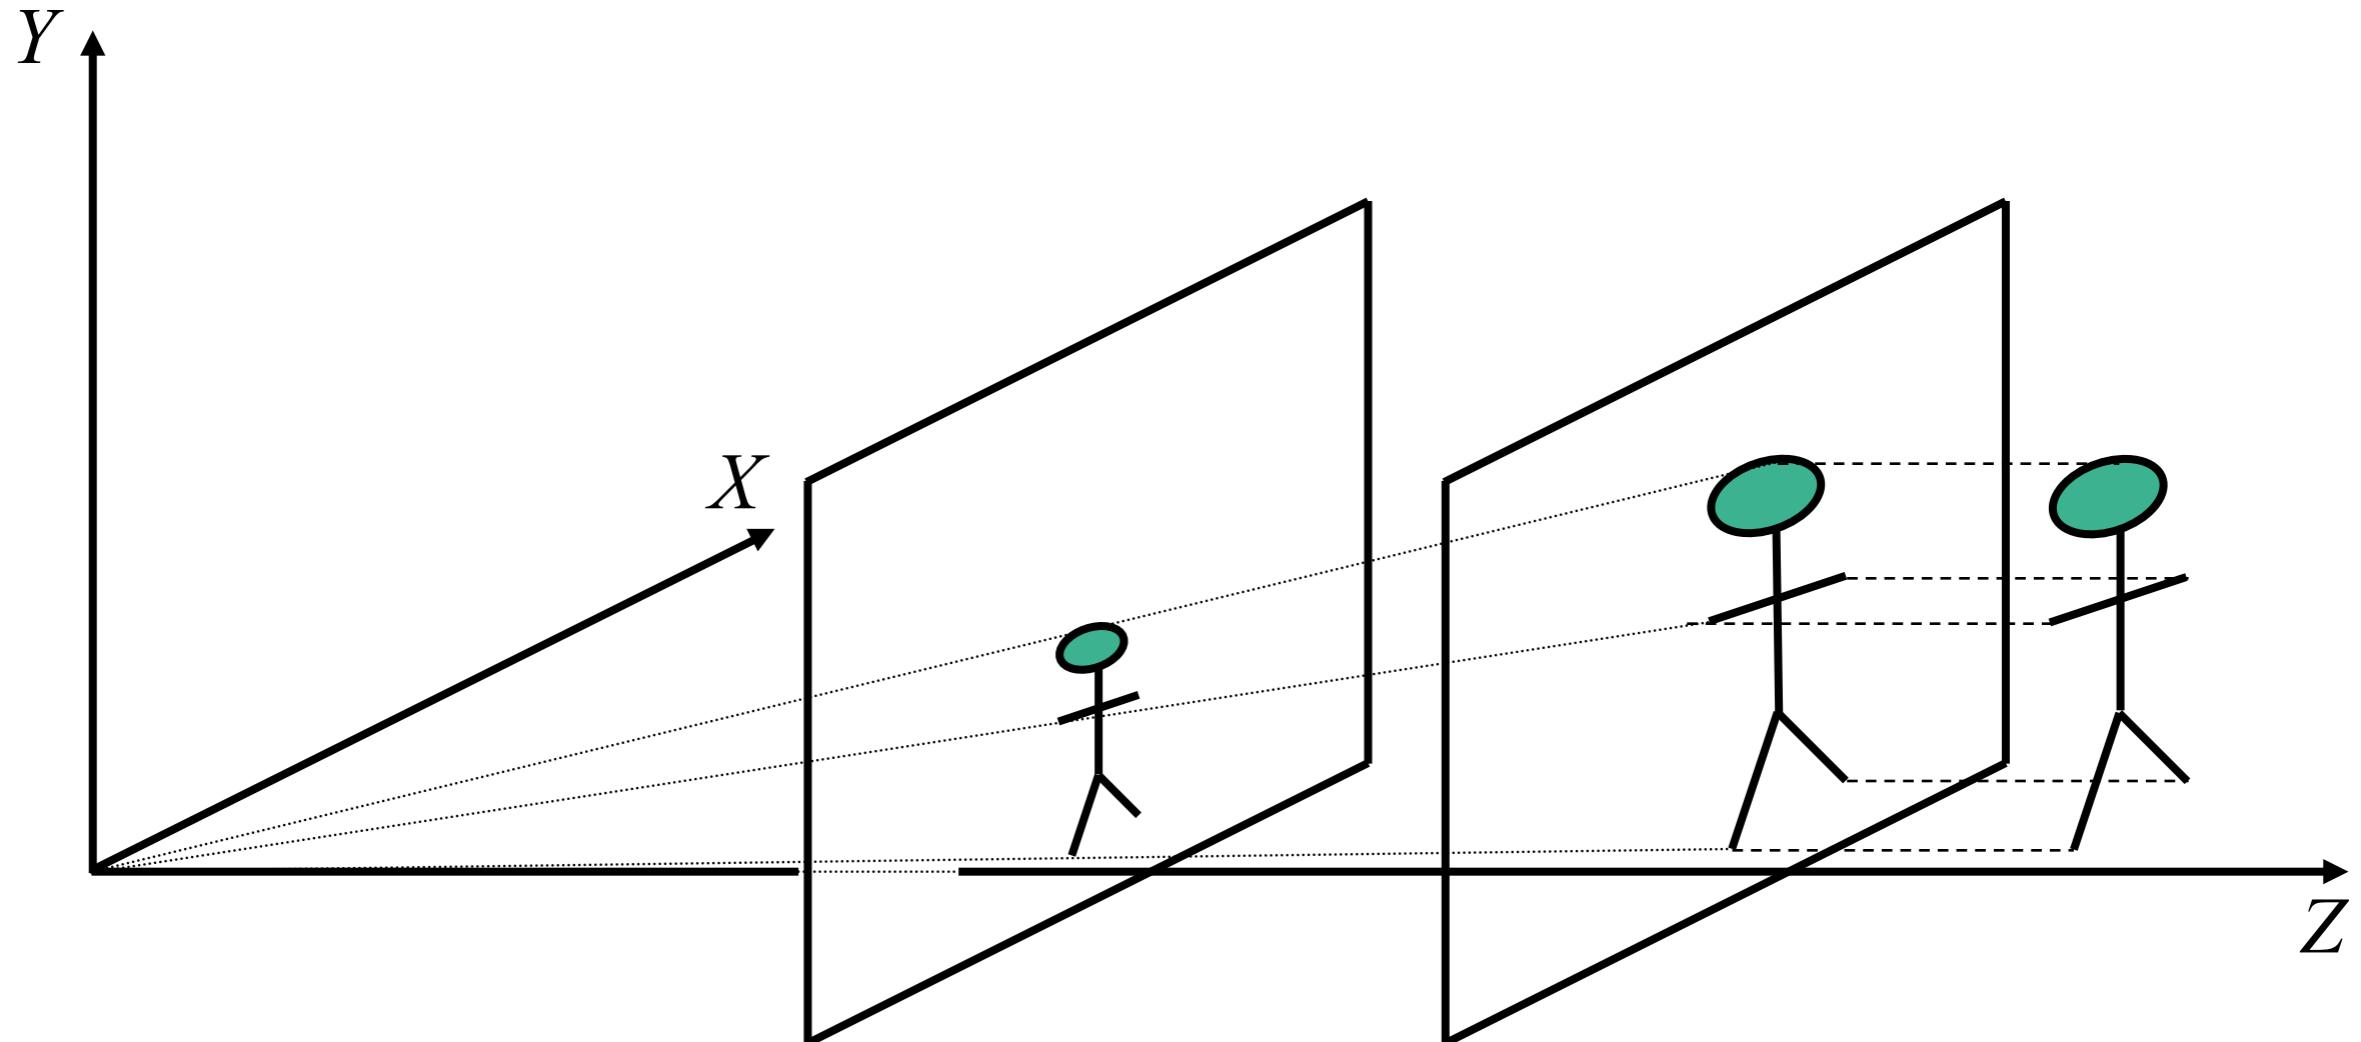
# Recall: Orthographic Projection



Assume the world is infinitely far away.

$$(x, y, z) \rightarrow (x', y')$$

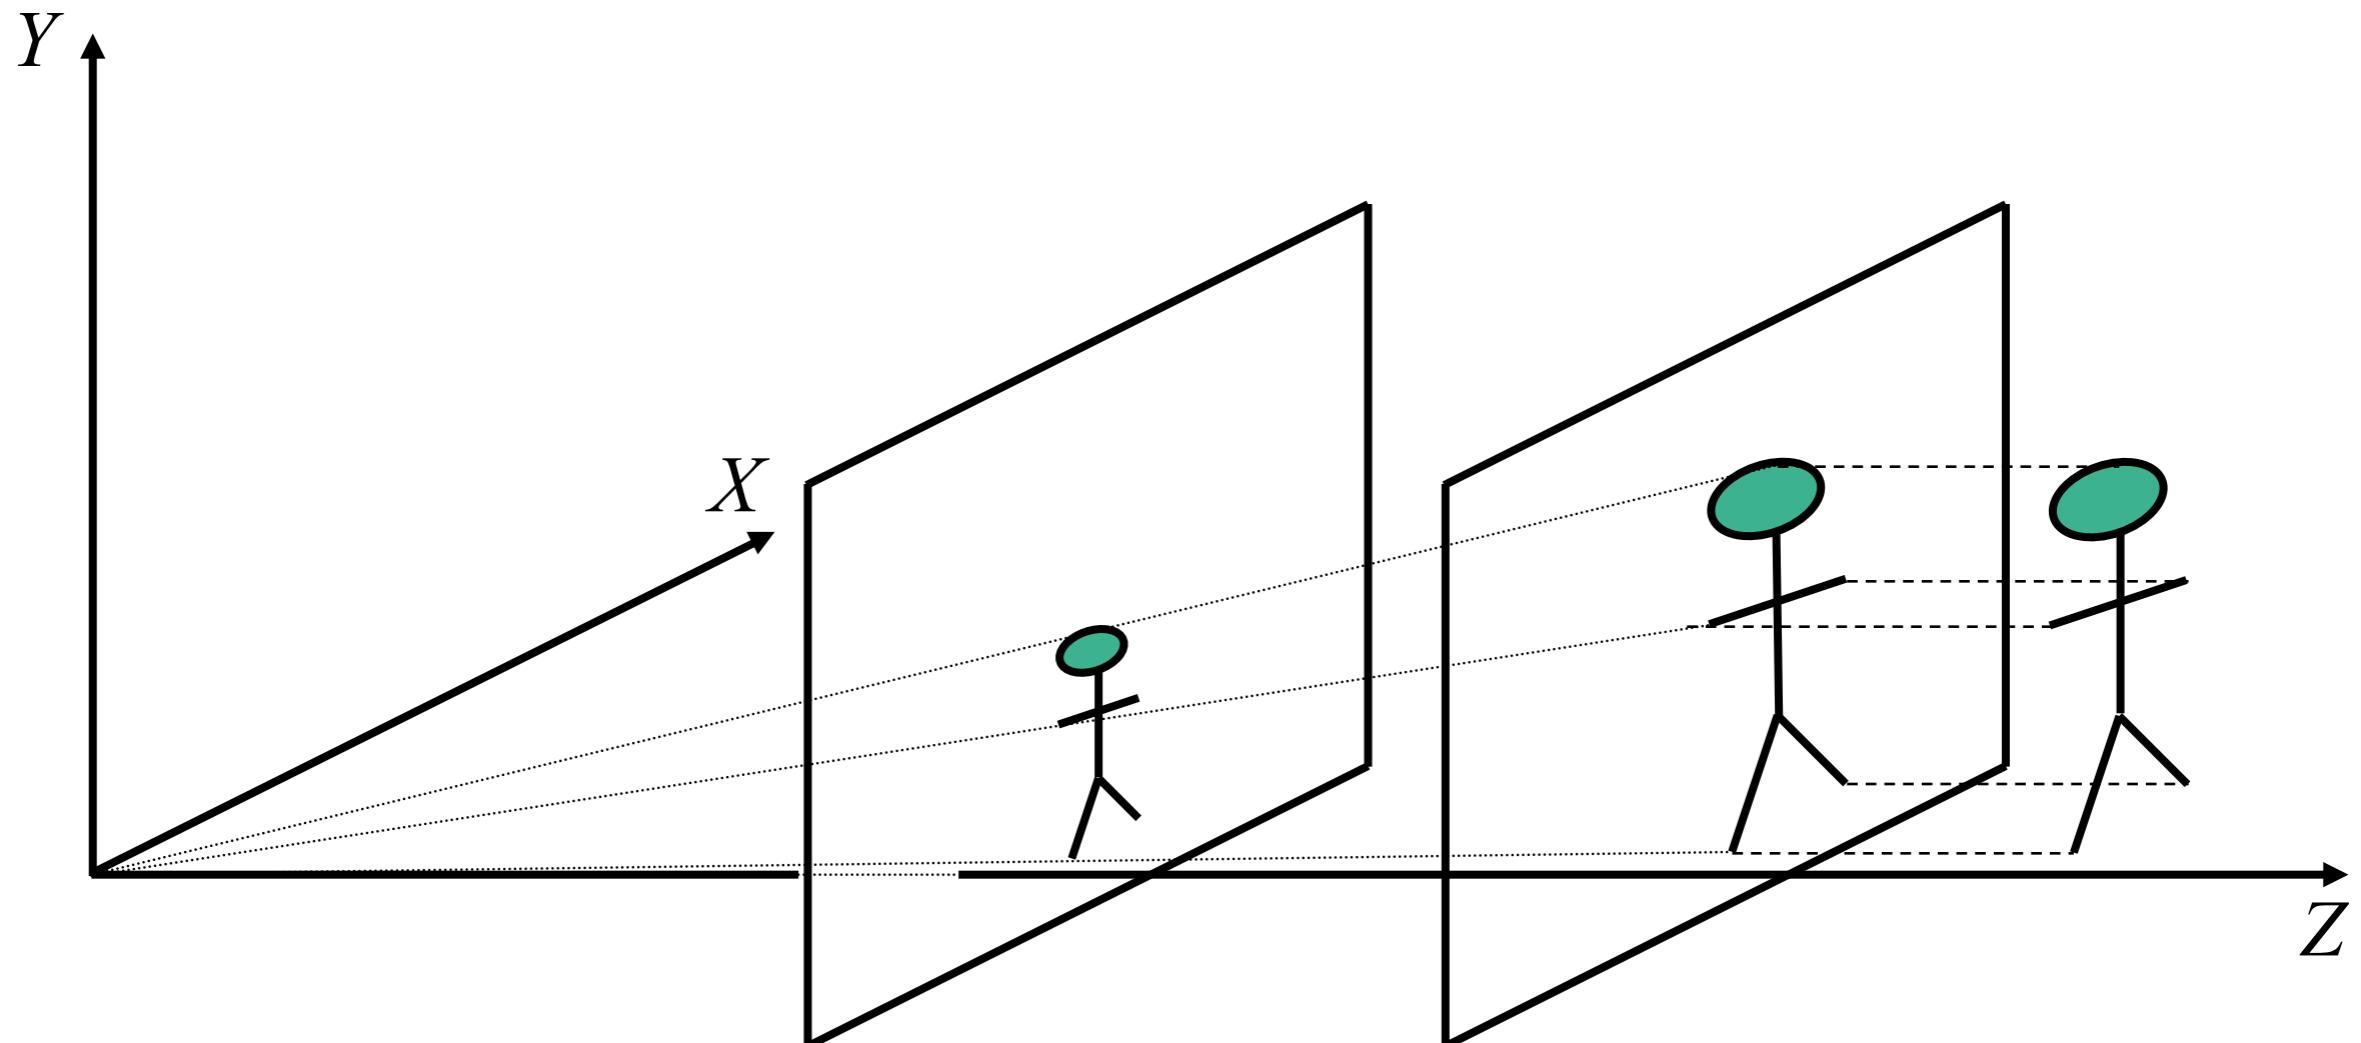
# Weak Perspective (scaled orthography)



Assume variation in depth is small relative to the distance from the camera.

Approximate scene as a fronto-parallel plane

# Weak Perspective (scaled orthography)



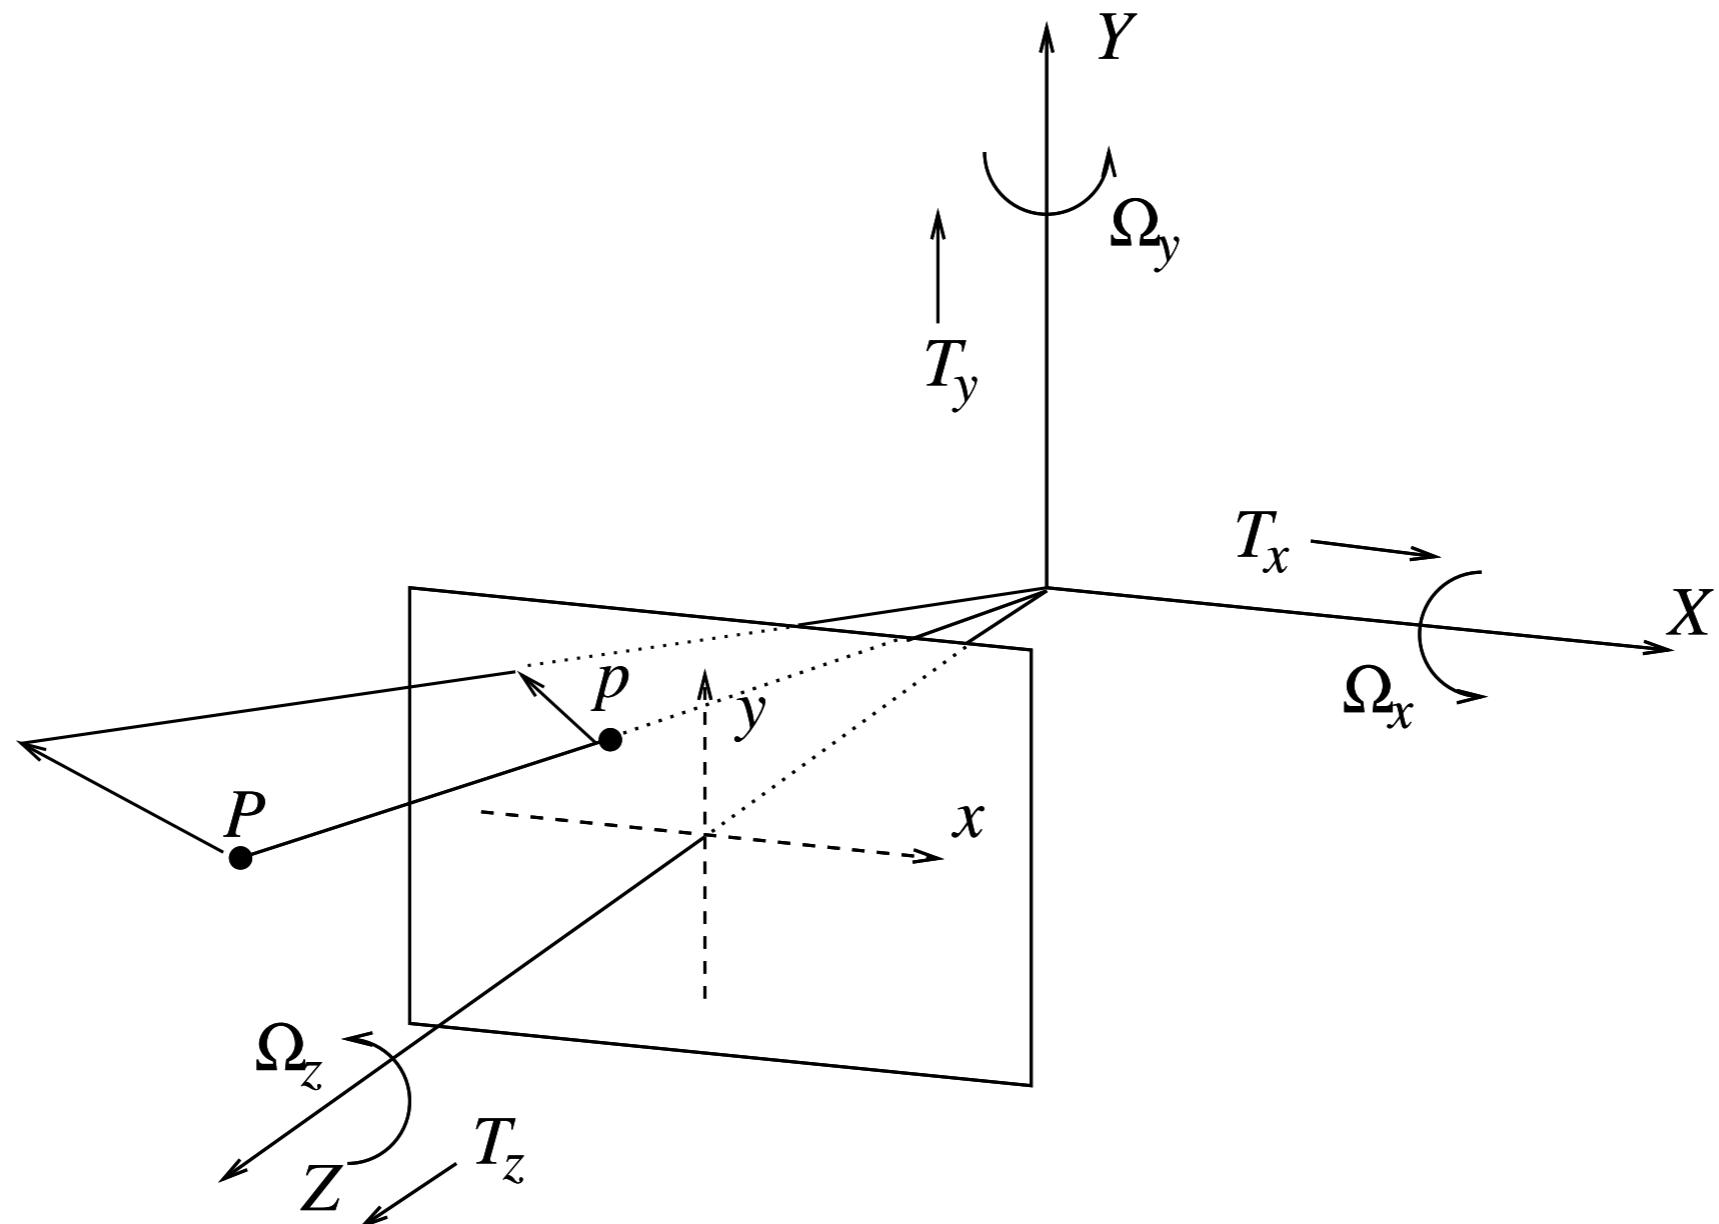
$$(x', y') = (sx, sy)$$

# Claim

- ◆ For **small motions**, affine flow approximates the motion of a **plane** viewed under **orthographic projection**.

$$\mathbf{u}(\mathbf{x}; \mathbf{a}) = \begin{bmatrix} u(\mathbf{x}; \mathbf{a}) \\ v(\mathbf{x}; \mathbf{a}) \end{bmatrix} = \begin{bmatrix} a_1 + a_2x + a_3y \\ a_4 + a_5x + a_6y \end{bmatrix}$$

# Recall: 3D motion



# Motion Models



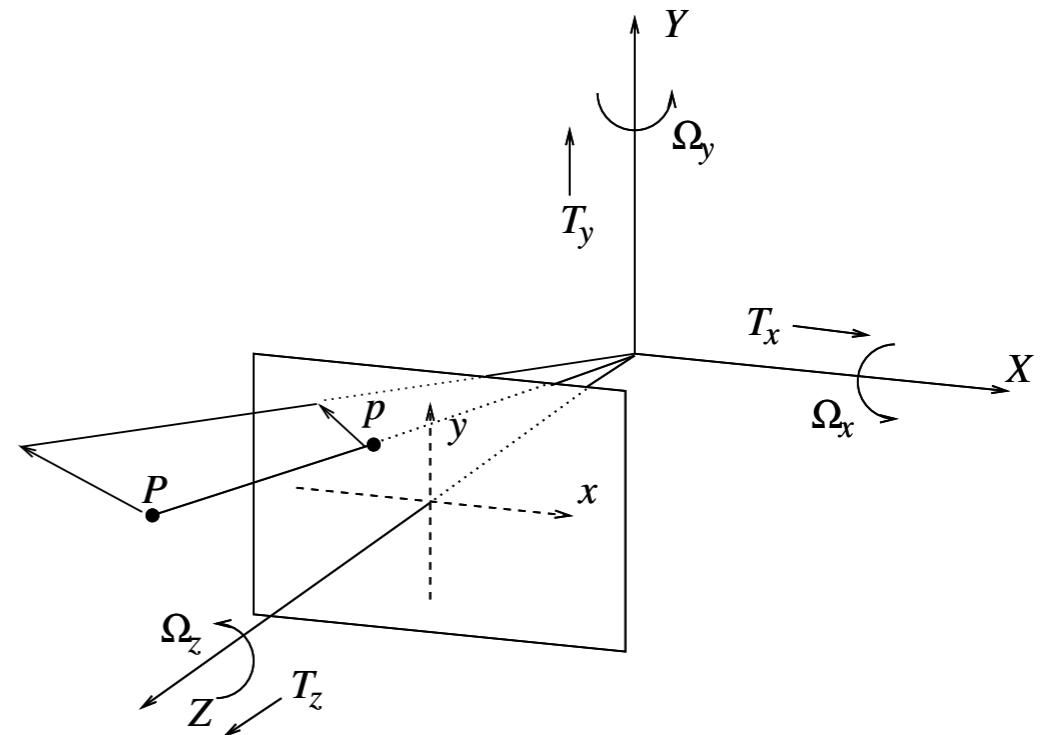
## 3D Rigid Motion

rotation

translation

$$\begin{bmatrix} X' \\ Y' \\ Z' \end{bmatrix} = \mathbf{R} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} + \mathbf{T} = \mathbf{R}_Z^{\Omega_z} \mathbf{R}_Y^{\Omega_y} \mathbf{R}_X^{\Omega_x} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} + \mathbf{T}$$

Euler angles

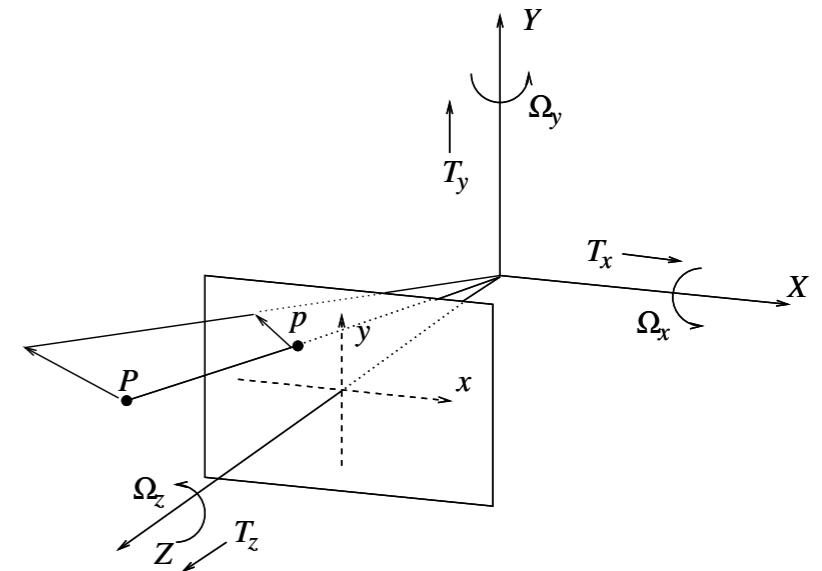


# Rotation



Review: FP Ch 2.1, Euclidean Geometry

$$\mathbf{R}_Z^{\Omega_Z} = \begin{bmatrix} \cos \Omega_Z & -\sin \Omega_Z & 0 \\ \sin \Omega_Z & \cos \Omega_Z & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



$$\mathbf{R}_Y^{\Omega_Y} = \begin{bmatrix} \cos \Omega_Y & 0 & \sin \Omega_Y \\ 0 & 1 & 0 \\ -\sin \Omega_Y & 0 & \cos \Omega_Y \end{bmatrix}$$

$$\mathbf{R}_X^{\Omega_X} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos \Omega_X & -\sin \Omega_X \\ 0 & \sin \Omega_X & \cos \Omega_X \end{bmatrix}$$

$$\begin{bmatrix} X' \\ Y' \\ Z' \end{bmatrix} = \begin{bmatrix} \cos \Omega_Y \cos \Omega_Z & \sin \Omega_X \sin \Omega_Y \cos \Omega_Z - \cos \Omega_X \sin \Omega_Z & \cos \Omega_X \sin \Omega_Y \cos \Omega_Z + \sin \Omega_X \sin \Omega_Z \\ \cos \Omega_Y \sin \Omega_Z & \sin \Omega_X \sin \Omega_Y \sin \Omega_Z + \cos \Omega_X \cos \Omega_Z & \cos \Omega_X \sin \Omega_Y \sin \Omega_Z - \sin \Omega_X \cos \Omega_Z \\ -\sin \Omega_Y & \sin \Omega_X \cos \Omega_Y & \cos \Omega_X \cos \Omega_Y \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} + \begin{bmatrix} T_X \\ T_Y \\ T_Z \end{bmatrix}$$

# Assumption: Small Motion

$$\begin{bmatrix} X' \\ Y' \\ Z' \end{bmatrix} = \begin{bmatrix} \cos \Omega_Y \cos \Omega_Z & \sin \Omega_X \sin \Omega_Y \cos \Omega_Z - \cos \Omega_X \sin \Omega_Z & \cos \Omega_X \sin \Omega_Y \cos \Omega_Z + \sin \Omega_X \sin \Omega_Z \\ \cos \Omega_Y \sin \Omega_Z & \sin \Omega_X \sin \Omega_Y \sin \Omega_Z + \cos \Omega_X \cos \Omega_Z & \cos \Omega_X \sin \Omega_Y \sin \Omega_Z - \sin \Omega_X \cos \Omega_Z \\ -\sin \Omega_Y & \sin \Omega_X \cos \Omega_Y & \cos \Omega_X \cos \Omega_Y \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} + \begin{bmatrix} T_X \\ T_Y \\ T_Z \end{bmatrix}$$

Assume that rotation is  
small



$$\cos \theta \approx 1 \quad \sin \theta \approx \theta$$

$$\begin{bmatrix} X' \\ Y' \\ Z' \end{bmatrix} = \begin{bmatrix} 1 & -\Omega_Z & \Omega_Y \\ \Omega_Z & 1 & -\Omega_X \\ -\Omega_Y & \Omega_X & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} + \begin{bmatrix} T_X \\ T_Y \\ T_Z \end{bmatrix}$$

# 3D Motion

$$\begin{bmatrix} X' \\ Y' \\ Z' \end{bmatrix} = \left( \begin{bmatrix} 0 & -\Omega_Z & \Omega_Y \\ \Omega_Z & 0 & -\Omega_X \\ -\Omega_Y & \Omega_X & 0 \end{bmatrix} + \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \right) \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} + \begin{bmatrix} T_X \\ T_Y \\ T_Z \end{bmatrix}$$

$$\begin{bmatrix} V_X \\ V_Y \\ V_Z \end{bmatrix} = \begin{bmatrix} X' - X \\ Y' - Y \\ Z' - Z \end{bmatrix} = \begin{bmatrix} 0 & -\Omega_Z & \Omega_Y \\ \Omega_Z & 0 & -\Omega_X \\ -\Omega_Y & \Omega_X & 0 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} + \begin{bmatrix} T_X \\ T_Y \\ T_Z \end{bmatrix}$$

$$V_X = -\Omega_Z Y + \Omega_Y Z + T_X$$

$$V_Y = \Omega_Z X - \Omega_X Z + T_Y$$

$$V_Z = -\Omega_Y X + \Omega_X Y + T_Z$$

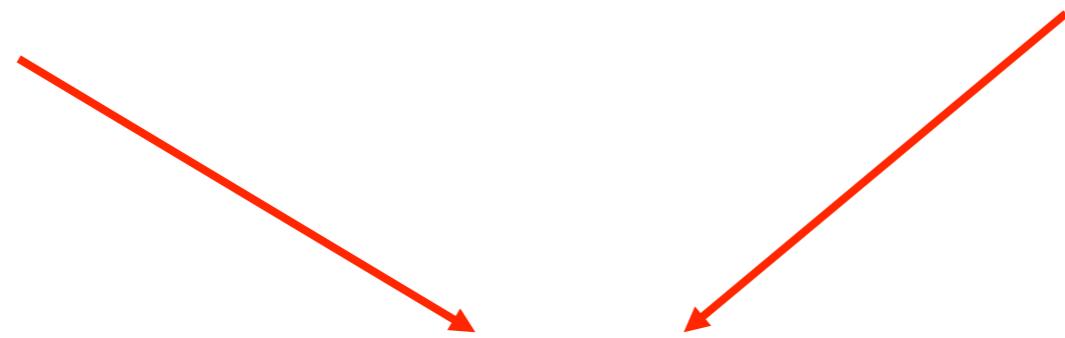
# Assumption:

## Planar World + Orthographic Projection



$$\begin{aligned}x &= X \\y &= Y \\Z &= a + bX + cY\end{aligned}$$

$$\begin{aligned}u &= V_X = -\Omega_Z y + \Omega_Y Z + T_X \\v &= V_Y = \Omega_Z x - \Omega_X Z + T_Y\end{aligned}$$



$$\begin{aligned}u &= V_X = -\Omega_Z y + \Omega_Y(a + bx + cy) + T_X \\v &= V_Y = \Omega_Z x - \Omega_X(a + bx + cy) + T_Y\end{aligned}$$

# Assumption: Planar World + Orthographic Projection

$$u = -\Omega_{ZY}y + \Omega_Y(a + bx + cy) + T_X$$

$$v = \Omega_Zx - \Omega_X(a + bx + cy) + T_Y$$

$$u = \Omega_Ybx + (\Omega_Yc - \Omega_Z)y + (\Omega_Ya + T_X)$$

$$v = (\Omega_Z - \Omega_Xb)x - \Omega_Xcy + (T_Y - \Omega_Xa)$$

$$\begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} a_1 + a_2x + a_3y \\ a_4 + a_5x + a_6y \end{bmatrix}$$



Substitute:

$$\begin{aligned} a_1 &= \Omega_Ya + T_X \\ a_2 &= \Omega_Yb \\ a_3 &= \Omega_Yc - \Omega_Z \\ a_4 &= T_Y - \Omega_Xa \\ a_5 &= \Omega_Z - \Omega_Xb \\ a_6 &= -\Omega_Xc \end{aligned}$$

# Affine Flow

- ◆ Small motion assumption
  - ◆ e.g. at [video frame rate](#)
- ◆ Planar surface
  - ◆ look at only a [small region](#) of the scene
- ◆ Orthographic projection
  - ◆ [surface distant](#) from camera
  - ◆ long focal length

$$\begin{bmatrix} u \\ y \end{bmatrix} = \begin{bmatrix} a_1 + a_2x + a_3y \\ a_4 + a_5x + a_6y \end{bmatrix}$$

# Second Option: Global regularization

- ◆ Remember: We were faced with challenges similar to those in LK in stereo.
  - ◆ Window-based approaches turned out to be problematic.
  - ◆ In stereo, we introduced prior knowledge which supplied global regularity on the disparity map.
- ◆ In optical flow, we can do the same:
  - ◆ We impose **global regularity** of the flow field, and by that impose **prior knowledge**.

- ◆ The classical global optical flow technique was developed by **Horn & Schunck**:
  - ◆ B. K. P. Horn and B. G. Schunck. Determining optical flow. Artificial Intelligence, 17(1-3):185-203, 1981.
- ◆ It formulates optical flow as a problem of energy minimization.
  - ◆ Ideally they would like to minimize the following energy:

$$E(u, v) = \int \left( I(x + u(x, y), y + v(x, y), t + 1) - I(x, y, t) \right)^2 + \lambda \cdot (||\nabla u(x, y)||^2 + ||\nabla v(x, y)||^2) \, dx \, dy$$

# Where does this energy come from?

Brightness difference between corresponding pixels

Quadratic penalty for brightness changes

$$E(u, v) = \int \left( I(x + u(x, y), y + v(x, y), t + 1) - I(x, y, t) \right)^2 +$$

$$\lambda \cdot \left( \|\nabla u(x, y)\|^2 + \|\nabla v(x, y)\|^2 \right) dx dy$$

Regularization parameter

Gradient magnitude of horz. flow

Penalizes changes in flow

Same for vertical flow

# Linearizing the Brightness Constancy Assumption



$$E(u, v) = \int (I(x + u(x, y), y + v(x, y), t + 1) - I(x, y, t))^2 + \lambda \cdot (||\nabla u(x, y)||^2 + ||\nabla v(x, y)||^2) \, dx \, dy$$

- ◆ If we tried to minimize this directly to find the flow, we would run into problems, because the energy is **non-convex** and has many local optima.
  - ◆ We've had this problem before.
  - ◆ We **linearized the brightness constancy assumption**.
  - ◆ Will do the same here:

$$E(u, v) = \int (I_x(x, y, t)u(x, y) + I_y(x, y, t)v(x, y) + I_t(x, y, t))^2 + \lambda \cdot (||\nabla u(x, y)||^2 + ||\nabla v(x, y)||^2) \, dx \, dy$$

# Actual H&S Energy

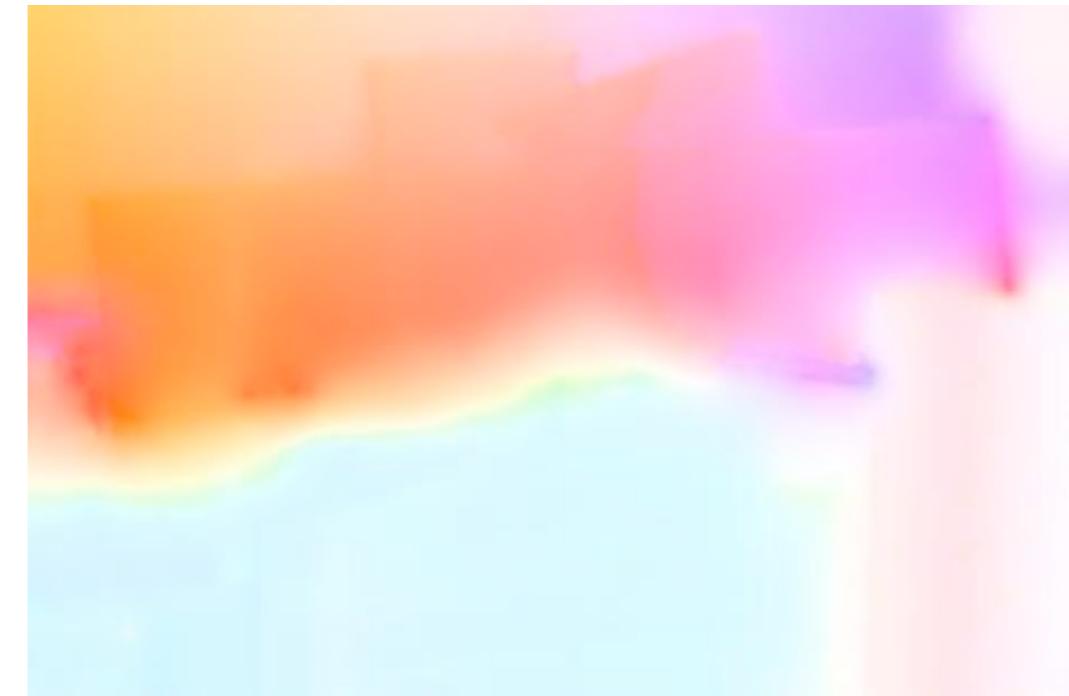
$$E(u, v) = \int (I_x(x, y, t)u(x, y) + I_y(x, y, t)v(x, y) + I_t(x, y, t))^2 + \lambda \cdot (||\nabla u(x, y)||^2 + ||\nabla v(x, y)||^2) \, dx \, dy$$

- ◆ This energy imposes a quadratic penalty on the **optical flow constraint** (OFC).
- ◆ It is **convex** and thus has a **unique optimum**.
- ◆ We estimate flow by discretizing it spatially and doing gradient descent.
- ◆ Old problem:
  - ◆ The linearization only works for small motions.
  - ◆ Hence we estimate the flow using **coarse-to-fine refinement with warping**.

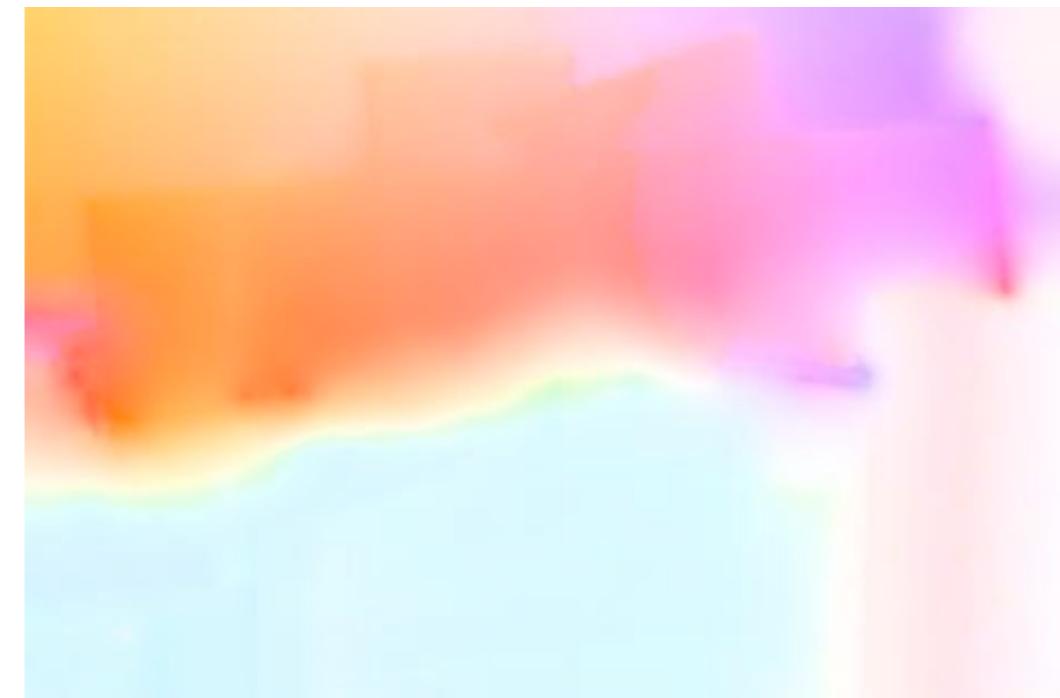
# Some H&S results



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT



# Some H&S results



# Problem with H&S

- ◆ These results are quite a bit better than what we can get with the LK method, but one problem remains.
- ◆ The flow is **very smooth**:
  - ◆ In particular, the method smoothes over the discontinuities that are in the flow.
- ◆ Why?
  - ◆ We use a quadratic penalty to penalize changes in the flow.
  - ◆ We have seen that this does not allow for discontinuities.
  - ◆ It penalizes large changes too much.

# Probabilistic Formulation

- ◆ We will do what we have already done several times:
  - ◆ We model the problem of optical flow estimation as one of probabilistic inference.
  - ◆ This allows us to motivate the formulation from observations on the data.
- ◆ Posterior for optical flow estimation:

$$p(\mathbf{u}, \mathbf{v} | I^0, I^1) \propto p(I^1 | \mathbf{u}, \mathbf{v}, I^0) \cdot p(\mathbf{u}, \mathbf{v} | I^0)$$


  
 Observation model  
 embodies brightness constancy

"Prior" on the flow field

# Some More Results



Method by [Black & Anandan, 96]

# Some More Results



Method by [Black & Anandan, 96]

# Some More Results



Method by [Sun et al. 2010]

# Some More Results



Method by [Sun et al. 2010]